Niching Methods for Genetic Algorithms

Samir W. Mahfoud
3665 E. Bay Dr. #204-429
Largo, FL 34641

Illinois Genetic Algorithms Laboratory (IlliGAL)
Department of General Engineering
University of Illinois at Urbana-Champaign
117 Transportation Building
104 South Mathews Avenue
Urbana, IL 61801

NICHING METHODS FOR GENETIC ALGORITHMS

BY

SAMIR W. MAHFOUD

B.S., Murray State University, 1985
M.S., University of Wisconsin-Madison, 1987

THESIS

Submitted in partial fulfillment of the requirements
for the degree of Doctor of Philosophy in Computer Science
in the Graduate College of the
University of Illinois at Urbana-Champaign, 1995

Urbana, Illinois

Abstract

Niching methods extend genetic algorithms to domains that require the location and maintenance of multiple solutions. Such domains include classification and machine learning, multimodal function optimization, multiobjective function optimization, and simulation of complex and adaptive systems.

This study presents a comprehensive treatment of niching methods and the related topic of population diversity. Its purpose is to analyze existing niching methods and to design improved niching methods. To achieve this purpose, it first develops a general framework for the modelling of niching methods, and then applies this framework to construct models of individual niching methods, specifically *crowding* and *sharing* methods.

Using a constructed model of crowding, this study determines why crowding methods over the last two decades have not made effective niching methods. A series of tests and design modifications results in the development of a highly effective form of crowding, called *deterministic crowding*. Further analysis of deterministic crowding focuses upon the distribution of population elements among niches, that arises from the combination of crossover and replacement selection. Interactions among niches are isolated and explained. The concept of *crossover hillclimbing* is introduced.

Using constructed models of fitness sharing, this study derives lower bounds on the population size required to maintain, with probability γ, a fixed number of desired niches. It also derives expressions for the expected time to disappearance of a desired niche, and relates disappearance time to population size. Models are presented of sharing under selection, and sharing under both selection and crossover. Some models assume that all niches are equivalent with respect to fitness. Others allow niches to differ with respect to fitness.

Focusing on the differences between *parallel* and *sequential* niching methods, this study compares and further examines four niching methods — crowding, sharing, sequential niching, and parallel hillclimbing. The four niching methods undergo rigorous testing on optimization and classification problems of increasing difficulty; a new niching-based technique is introduced that extends genetic algorithms to classification problems.

Acknowledgments

Many thanks go to the following people:

- My advisor, Dr. David Goldberg, for numerous valuable discussions concerning this research; for motivation, guidance, and financial support in the study of genetic algorithms; and for his willingness to take on a new student back in 1990.

- The other members of my committee — Dr. Larry Rendell, for his assistance with machine-learning issues and projects; Dr. Gul Agha, for meeting with this (former) student on several occasions; Dr. Geneva Belford, for her willingness to explore the strange new research area of genetic algorithms; and Dr. Jay Mittenthal, for his meticulous proofreading and thought-provoking biological analogies.

- Dr. Arthur Baskin, for giving me the initial assistantship that finalized my decision to study at the University of Illinois.

- Ganesh Mani, for pointing me in the direction of my two most recent jobs, and for the free lodging over the years.

- Dr. Michael Faiman and the secretarial staff of the Computer Science and General Engineering Departments, for their regular and very professional assistance.

- Everyone at the Illinois Genetic Algorithms Laboratory (IlliGAL), for rich discussions of a variety of topics; for their assistance and moral support over several years; and especially for keeping quiet when I needed to work those last few weeks and for showing up at my final defense (regardless of the cookies and brownies).

- Lynnea Magnuson, for service above and beyond the call of duty.

Table of Contents

List of Tables

List of Figures

Chapter 1

Introduction

Genetic algorithms (GAs) are search methods for solving complex problems. GAs have successfully been applied to optimization of functions, schedules, and arbitrary physical and programming structures. GAs have also shown promising initial results in classification and machine learning, as well as in simulation of biological, ecological, economic, financial, and social systems.

Genetic algorithms are based upon principles from biological genetics and operate analogous to evolution. However, while natural evolutionary processes maintain a variety of species, each occupying a separate ecological niche, traditional GAs rapidly push an artificial population toward convergence. That is, all individuals in the population soon become nearly identical. Even when multiple solutions to a problem exist, a traditional GA locates only one of them.

Niching methods allow genetic algorithms to maintain a population of diverse individuals. GAs that incorporate niching methods are capable of locating multiple, optimal solutions within a single population. Effective niching techniques are critical to the success of GAs in classification and machine learning, multimodal function optimization, multiobjective function optimization, and simulation of complex and adaptive systems.

This is a study of niching methods. The purpose of this study is to analyze existing techniques and to design improved techniques for the formation and maintenance of stable subpopulations, or *niches*, in GAs. To achieve these two objectives, we first develop a framework for modelling those GAs which incorporate niching methods (*niching GAs*). We next use this framework to construct models of individual niching GAs. Through the models, we isolate, ex-

plain, and predict (i.e., analyze) significant behavioral characteristics of niching GAs. Finally, we propose algorithmic changes that have the potential to elicit desired behavior. Modified algorithms sometimes themselves become the subject of further analysis and redesign.

The cycle of modelling, analysis, and design is intended to yield several practical results. One such result is bounds for setting the control parameters of various niching GAs. The most important control parameter is typically population size; if the user creates a sufficiently large population, a GA will likely produce desirable results. Minor control parameters include crossover probability, mutation probability, and number of generations. Another practical result we strive to attain is the ability to assess, both quantitatively and qualitatively, the impact of various design alterations on the niching GAs we examine. The ultimate aim of this research is to make significant strides toward the development of niching GAs which are both powerful and practical to use.

One can not undertake a comprehensive study of niching in genetic algorithms without encountering the broader topic of diversity, a frequently recurring issue in GA research. An original motivation for developing niching methods was, in fact, to promote diversity in the traditional GA. Diversity can serve two purposes in GAs. The first purpose is to delay convergence in order to increase exploration, so that a better, single solution can be located. (Convergence to undesirable solutions or to nonglobal optima has been dubbed *premature convergence* in the GA literature.) The second purpose is to locate multiple, final solutions. We will examine techniques for delaying convergence, but will concentrate upon methods that locate multiple, final solutions, since these are the stronger niching mechanisms. Note that the two purposes are not mutually exclusive — techniques which locate multiple, final solutions should also be highly effective at forestalling convergence. However, techniques which are designed solely to delay convergence are typically not useful for locating multiple, final solutions.

We will visit the subject of diversity towards the beginning of this thesis. Specifically, we will examine prior techniques for promoting diversity in a population, define general notions and rigorous measures of population diversity, and place prior definitions of diversity within the scope of our framework. This study will thus take a preliminary step towards a comprehensive theory of diversity for genetic algorithms.

As mentioned previously, GAs are applicable to a variety of problems in general areas such as optimization, machine learning, and simulation. However, much of the GA research to date

has been conducted in the domain of static function optimization, under the assumption that results achieved in optimizing static functions will extend to other problem domains. This assumption has held up relatively well in past research, although some domains have required more specialization of the GA than others. The success of GAs in complex and dynamic domains very likely hinges upon their success in the arena of static function optimization, a simpler and better understood domain.

We will examine the niching GA primarily from the standpoint of multimodal function optimization. That is, given a function with multiple optima, the goal of our niching GAs will be to locate several of the best optima. Like prior GA researchers, we will operate with the understanding that niching methods which optimize multimodal functions are also applicable to other domains that require the formation or maintenance of multiple solutions. In fact, we will extend our niching methods to the solution of classification problems. We will also mention successful application of a method developed in this thesis, to delaying the GA's convergence, in order to locate a better, single solution.

Our overriding goal in undertaking this research, as stated earlier, is to design improved niching methods. Given this goal, several questions arise that we will strive to answer in this study. These include the following:

- What potential niching mechanisms already exist in some form? Is it difficult to isolate them from the GAs in which they are embedded?

- Do logical categories of niching mechanisms exist? If so, what type of behavior can be expected from niching methods within each category?

- What kind of general-purpose modelling framework can we devise for niching methods?

- Given a modelling framework, what models can we construct?

- How indicative of actual behavior are the models we construct?

- What does a given model tell us about setting control parameters for its corresponding niching method?

- What hints does a given model provide for improving the design of its corresponding niching method? What is the expected impact of a particular design alteration?

- Given a model, what conclusions can we draw about its corresponding niching method?

- Given a model, what experiments can we conduct to increase our understanding of its corresponding niching method?

- How good is a given niching method?

In terms of traditional scientific methodology, we will be testing the following underlying hypotheses in this study —

- that undiscovered niching methods exist that have general applicability and that are highly effective.

- that it is possible to devise a general and useful framework for modelling niching GAs.

- that it is possible to construct explanatory and predictive models from the framework.

- that constructed models will help us understand the various behaviors of niching methods.

- that through modelling, we will be able to bound control parameters for niching GAs, and to otherwise improve the designs of niching GAs.

- that the modelling, analysis, and design process will lead to the discovery of better niching methods.

In addition, many intermediate hypotheses will present themselves over the course of this research, and will need to be tested before proceeding further.

We must emphasize that in the design of GAs, flexibility is essential, and results match preconceived hypotheses only on occasion. Since the space of all possible GA designs is highly nonlinear, incremental design improvements often lead to dead ends. In the literature, GAs are currently undergoing a process of testing and redesign, mostly through ad hoc experimentation. "New and improved" GAs are unveiled seemingly every day, with the driving motivation often being novelty; in many cases, previously introduced GAs are simpler, and have identical or better performance. At the opposite extreme is research into precise mathematical modelling of GAs. Inevitably, exact models of GAs are more complex and analytically unwieldy than the GAs themselves. Precise mathematical modelling has not yet resulted in improved designs for GAs.

4

The middle ground between the two extremes of approaches — ad hoc experimentation and precise mathematical modelling — has, with a few notable exceptions, remained vacant. Much of the difficulty lies in the fact that GAs are complex systems, and successfully modifying the design of a complex system is an involved, multifaceted undertaking. What is often required is an approach that is a combination of intuition, coarse analysis, and thoughtful experimentation. Goldberg (1993a) connects the design of effective GAs to the invention of powered aircraft by the Wright Brothers in 1903. "The invention of broadly efficient genetic algorithms is a design challenge as stiff as the most difficult that have been faced this century." In place of ad hoc design changes ("hacking"), rigorous mathematics, or "systematic application of scientific method" — paths to flight that inventors had pursued for over a century — he recommends decomposing the GA-design task into "quasi-separable subproblems", much as the Wright Brothers did when inventing the airplane. We consider this approach in greater detail in Section 2.4.

The remainder of this thesis is organized as follows. Chapter 2 is an overview of genetic algorithms. It starts with a review of the traditional GA, proceeds to explain basic GA theory, and finishes with a coarse-grained overview of current GA research. The reader should consult the research overview in Section 2.4 to obtain a better grasp of the role this thesis plays in the broader spectrum of GA research.

Chapter 3 discusses diversity. It starts by motivating the need for diversity in GAs; it presents simple optimization problems on which the traditional GA fails because of a lack of diversity. It next explores three potential culprits in the loss of population diversity: selection noise, selection pressure, and operator disruption. The chapter then covers, in depth, previous diversification methods. Afterwards, it presents a formal framework for the study of diversity in GAs, and looks at prior measures of diversity within this framework. Finally, Chapter 3 defines the notion of useful diversity, and partially distinguishes between a diversity mechanism and a niching method.

Chapter 4 introduces niching methods. It begins by defining a niching method and answering common questions about niching methods. It next reviews, in depth, previous research into niching mechanisms, and describes niching methods along two dimensions of behavior; it presents several general categories of niching methods, some of which are still largely unexplored. Chapter 4 then extends our formal diversity framework to multimodal function opti-

mization, our chosen domain of application. Finally, it chooses the specific niching methods to be examined in the remainder of the thesis.

Chapter 5 presents the framework for modelling niching methods. It builds upon the framework for diversity defined in Chapter 3 and the specialization of that framework given in Chapter 4. It starts by reviewing previous methods of modelling GAs. It next presents the modelling framework of this thesis, including definitions, abstractions, and simplifying assumptions. Finally, it describes the test functions used in subsequent chapters, and presents a new methodology for applying niching methods to classification problems.

Chapter 6 analyzes crowding, using two performance criteria to determine why the original algorithm is not an effective niching method. Through a series of tests and modifications, the chapter designs a form of crowding that is highly effective. The resulting algorithm is called *deterministic crowding*. Chapter 6 concludes by examining bitwise diversity as a potential performance criterion, and by describing both current and future research extensions.

Chapter 7 further analyzes deterministic crowding, concentrating upon the distribution of population elements among niches, that arises from the combination of crossover and replacement selection. Interactions among niches are isolated and explained. The chapter finishes by introducing the concept of crossover hillclimbing, an abstraction that can be used to characterize fitness landscapes.

Chapter 8 models fitness sharing in the absence of crossover. It first analyzes and illustrates sharing's distributional properties. It next examines the probability that sharing loses important solutions, and derives closed-form expressions for the expected time to loss. The chapter then constructs models of fitness sharing that handle cases of varying generality. The models are used to derive lower bounds on required population size. Derived properties are verified and illustrated through several experiments.

Chapter 9 extends Chapter 8's models of sharing to include crossover. It discusses how the methodology for incorporating crossover can be copied in order to incorporate mutation and other sources of noise. It also mentions the possible relaxation of modelling assumptions, and the construction of extended models.

Chapter 10 tests the parallel niching methods modelled in previous chapters, on sets of problems of increasing difficulty. It compares their performances to those of two basic algorithms — sequential niching and parallel hillclimbing. The chapter especially highlights the differences

between parallel and sequential niching methods, and illustrates the strengths and weaknesses of all methods considered. The terms, parallel and sequential, refer not to physical processing elements, but to conceptual parallelism within an algorithm.

Chapter 11, the final chapter, provides a summary of this thesis, and highlights its contributions to GA research. It discusses paths for future research, and draws overall conclusions from the research presented in this thesis.

Chapter 2

Genetic Algorithms

Genetic algorithms (Goldberg, 1989c; Holland, 1975, 1992) are general purpose, parallel search procedures that are based upon genetic and evolutionary principles. A genetic algorithm works by repeatedly modifying a population of artificial structures through the application of genetic operators. GAs are typically black-box methods that use fitness information exclusively; they do not require gradient information or other internal knowledge of the problem.

This chapter first reviews the traditional genetic algorithm, also called the *simple GA*, in the context of function optimization. It next presents basic GA theory, followed by a coarse overview of the current state of GA research. Finally, it shows how this thesis fits into the broad field of GA research.

2.1 Components

The goal in optimization is to find the best possible solution or solutions to a problem, with respect to one or more criteria. In order to use a genetic algorithm, one must first choose a suitable structure for representing those solutions. In the terminology of state-space search, an instance of this data structure represents a state or point in the search space of all possible solutions.

A genetic algorithm's data structure consists of one or more *chromosomes* (usually one). A chromosome is typically a string of bits, so the term *string* is often used instead. GAs, however, are not restricted to bit-string representations. Other possibilities include vectors of real numbers (L. Davis, 1991b; Eshelman & Schaffer, 1993; Goldberg, 1991a, 1991b), and high-

level computer programs (Koza, 1992). Although variable-length structures are appropriate for many problems, fixed-length structures are the norm. In this study, we restrict our attention to structures that are single strings of l bits.

Each chromosome (string) is a concatenation of a number of subcomponents called *genes*. Genes occur at various positions or *loci* of the chromosome, and take on values called *alleles*. In bit-string representations, a gene is a bit, a locus is its position in the string, and an allele is its value (0 or 1). The biological term *genotype* refers to the overall genetic makeup of an individual, and corresponds to a structure in the GA. The term *phenotype* refers to the outward characteristics of an individual, and corresponds to a decoded structure in the GA.

An extremely simple, but illustrative example of a genetic optimization problem is maximization of the following function of two variables:

$$f(x_1, x_2) = x_1 + x_2 \quad , \tag{2.1}$$

where $0 \leq x_1 \leq 1$ and $0 \leq x_2 \leq 1$. A common encoding technique for real variables is to transform them into binary integer strings of sufficient length to provide a desired degree of precision. Assuming 8-bit encodings are sufficient for both x_1 and x_2, these encoded variables are decoded through normalization of the corresponding binary integer — division of the integer by $2^8 - 1$. For example, 00000000 represents 0/127 or 0, while 11111111 represents 127/127 or 1. The data structure to be optimized is a 16-bit string, representing the concatenation of the encodings for x_1 and x_2. The variable x_1 resides in the leftmost 8 bit-positions, while x_2 resides in the rightmost. An individual's genotype is a 16-bit string, while its phenotype is an instance of the tuple, $\langle x_1, x_2 \rangle$. The genotype is a point in the 16-dimensional Hamming space that the GA searches. The phenotype is a point in the two-dimensional space of decoded variables.

To optimize a structure using a GA, one must be able to assign some measure of quality to each structure in the search space. The *fitness function* is responsible for this task. In function maximization, the objective function often acts as a fitness function; such is the case for example function (2.1), in which the goal is to find the value of $\langle x_1, x_2 \rangle$ that maximizes f. GAs perform maximization by default; for minimization problems, objective function values can be negated, then translated into positive territory to yield fitnesses.

Natural fitness functions also exist in domains of application other than mathematical function optimization. For combinatorial optimization problems such as the travelling salesman problem (Goldberg, 1989c; Homaifar, Guan, & Liepins, 1993), tour length, once negated and properly translated, makes a good choice. For weight optimization in neural networks (Whitley & Hanson, 1989; Whitley, Starkweather, & Bogart, 1990), the sum of squared errors on a set of training examples can serve as a fitness function (after negation and translation).

2.2   Mechanics

The mechanics of a simple genetic algorithm are as follows. The simple GA randomly generates an initial population of n structures. The GA proceeds for a fixed number of *generations* or until it satisfies some stopping criterion. During each generation, the GA performs *fitness-proportionate selection*, followed by *single-point crossover*, followed by *mutation*. First, fitness-proportionate selection assigns each individual structure i in the population a probability of selection $p_s(i)$, according to the ratio of i's fitness to overall population fitness:

$$p_s(i) \; = \; \frac{f(i)}{\sum_{j=1}^{n} f(j)} \; . \tag{2.2}$$

Then it selects (with replacement) a total of n individuals for further genetic processing, according to the distribution defined by the $p_s(i)$. The simplest variety of fitness-proportionate selection, *roulette-wheel selection* (Goldberg, 1989c), chooses individuals through n simulated spins of a roulette wheel. The roulette wheel contains one slot for each population element. The size of each slot is directly proportional to its respective $p_s(i)$. Note that population members with higher fitnesses are likely to be selected more often than those with lower fitnesses.

After selection, the n selected individuals undergo crossover (also called *recombination*) with fixed probability p_c per string. The n selected strings are paired at random, yielding $n/2$ couples. For each couple, crossover may or may not occur. With probability $1 - p_c$, crossover does not occur, and both individuals proceed to the mutation stage. Otherwise, the couple produces two offspring via crossover, and only the offspring continue to the mutation stage.

Single-point crossover works as follows. First, one of $l - 1$ potential crossing sites is chosen at random. (Crossing sites are between a string's neighboring bits.) Two parent strings are each sliced at the selected crossing site, into two segments apiece. Appropriate segments from

different parents are then concatenated to yield two offspring. For example, suppose one parent consists of 16 zeros, while the other parent consists of 16 ones. Suppose further that site number 12 is randomly selected from the 15 possible crossing sites. The parents and their offspring are shown below.

Parents		Offspring
0000000000000000	$\longrightarrow$	0000000000001111
1111111111111111	crossover	1111111111110000

After the crossover stage has finished, the mutation stage begins. For every string that advances to the mutation stage, each of its bits is flipped with probability p_m. The population resulting from the mutation stage then overwrites the old population (the one prior to selection), completing one generation. Subsequent generations follow the same cycle of selection, crossover, and mutation.

Many alternatives exist to the selection, crossover, and mutation operators presented above. We briefly mention a few that are relevant to this thesis. First of all, *tournament selection* (Brindle, 1981; Goldberg & Deb, 1991) is one alternative to roulette-wheel selection that is sensitive to relative rather than absolute fitnesses. Tournament selection holds n tournaments to choose n individuals. Each tournament consists of sampling k elements from the population, and choosing the fittest one. The most common variation, *binary tournament selection*, uses $k = 2$.

Elitist selection methods (De Jong, 1975) ensure that the best element or elements of the population survive from generation to generation. The most basic elitist strategy copies the best element from the current population to the next population, if that element has not been transferred through the normal process of selection, crossover, and mutation. Any standard selection method can be made elitist.

Two-point crossover (Cavicchio, 1970; Goldberg, 1989c) and *uniform crossover* (Syswerda, 1989) are alternatives to single-point crossover. In two-point crossover, two crossing sites are selected at random, and parent chromosomes exchange the segment that lies between the two crossing sites. In uniform crossover, each bit of the first parent is transferred to the first offspring with fixed probability; otherwise the bit is transferred to the second offspring. The rest of the loci on both offspring are filled in using the bits of the second parent.

2.3 Theory

Newcomers to genetic algorithms often develop the impression that GAs simply improve upon populations of strings. This section examines the inner workings of a simple GA, including the mechanisms by which the GA arrives at its final solutions. The section reviews basic GA theory, including schemata, building blocks, implicit parallelism, the schema theorem, and the building block hypothesis.

Schemata

While a GA on the surface processes strings, it implicitly processes schemata, which represent similarities between strings (Goldberg, 1989c; Holland, 1992). A GA can not, as a practical matter, visit every point in the search space. It can, however, sample a sufficient number of hyperplanes in highly fit regions of the search space. Each such hyperplane corresponds to a set of highly fit, similar substrings.

A *schema* is a string of total length l (the same overall length as the population's strings), taken from the alphabet $\{0, 1, *\}$, where '$*$' is a wild-card or "don't care" character. Each schema represents the set of all binary strings of length l, whose corresponding bit-positions contain bits identical to those '0' and '1' bits of the schema. For example, the schema, $10**1$, represents the set of five-bit strings, $\{10001, 10011, 10101, 10111\}$. Schemata are also called *similarity subsets* because they represent subsets of strings with similarities at certain, fixed bit-positions. Two properties of schemata are their *order* and *defining length*. Order is the number of fixed bit-positions (non-wild-cards) in a schema. Defining length is the distance between a schema's outermost, fixed bit-positions. For example, the above schema is of order 3, written $o(10**1) = 3$, and has a defining length of 4, written $\delta(10**1) = 4$. Each string in the population is an element of 2^l schemata.

Building blocks

Building blocks are low-order, short defining-length, highly fit schemata (Goldberg, 1989c), where the fitness of a schema is defined as the average fitness of the elements it contains. Building blocks represent similarities (between strings) that are significant to the GA's solution of a particular problem.

Selection chooses strings with higher fitnesses for further processing. Hence strings that are members of highly fit schemata are selected a greater number of times. Crossover infrequently disrupts schemata with shorter defining lengths, and mutation infrequently disrupts lower order schemata. Therefore, highly fit, short defining-length, low-order schemata, otherwise known as building blocks, are likely to proliferate from generation to generation. From this fact comes the claim that GAs process building blocks, also known as *useful schemata*. Holland (1992) estimates that while a GA processes n strings each generation, it processes on the order of n^3 useful schemata. He calls this phenomenon *implicit parallelism*. For the solution of real-world problems, the presence of implicit parallelism means that a larger population has the potential to locate a solution in polynomially faster time than a smaller population.

Schema theorem

The simple GA, prior to significant convergence, allocates an exponentially increasing number of trials to useful schemata or building blocks. This is illustrated by the re-derivation of the following theorem. Let $m(H, t)$ be the number of instances of schema H present in the population at generation t. We calculate the expected number of instances of H at the next generation, or $m(H, t + 1)$, in terms of $m(H, t)$. Recall that the simple GA assigns a string a probability of selection directly proportional to fitness. It follows from Equation 2.2 that H can expect to be selected $m(H, t) \cdot (f(H)/\bar{f})$ times, where $\bar{f}$ is average population fitness and $f(H)$ is the average fitness of those strings in the population that are elements of H.

The probability that single-point crossover disrupts a schema is precisely the probability that the crossover point falls within the schema's defining positions (those outermost, fixed bit-positions used to calculate the defining length). The probability that H survives crossover is greater than or equal to the term, $1 - p_c \times \frac{\delta(H)}{l-1}$. This survival probability is an inequality, because a disrupted schema might regain its composition if it crosses with a similar schema. The probability that H survives mutation is $(1 - p_m)^{o(H)}$, which can be approximated as $1 - o(H)p_m$ for small p_m and small $o(H)$. The product of the expected number of selections and the survival probabilities (with the smallest multiplicative term omitted) yields what is known as the *schema theorem*:

$$m(H, t + 1) \geq m(H, t) \cdot \frac{f(H)}{\bar{f}} \cdot \left(1 - p_c \frac{\delta(H)}{l-1} - o(H)p_m\right) \ . \qquad (2.3)$$

The schema theorem states that building blocks grow exponentially over time, while below-average schemata decay at a similar rate. Holland (1992) connects the competing schemata of GAs to the multiarmed bandits of statistical decision theory. He states that an optimal solution to the bandit problem is to allocate greater than exponential numbers of trials to the observed best arm. The connection is that GAs perform sampling in a near optimal fashion.

Goldberg (1983, 1989c), in his detailed presentations of the schema theorem, puts forth the *building block hypothesis*, which states that "building blocks combine to form better strings". That is, the recombination and exponential growth of building blocks leads to the formation of better building blocks, which in turn leads to highly fit final solutions.

While the schema theorem is often predictive of schema growth, it is overly simplistic in some regards for describing a GA's behavior. First of all, $f(H)$ and $\bar{f}$ do not stay constant from generation to generation. Fitnesses of population elements shift significantly after the first few generations. Secondly, the schema theorem accounts for schema losses but not for schema gains. Schemata are often constructed by crossover and mutation. In addition, as a GA progresses, population elements look more and more alike, so that schemata which are disrupted by crossover tend to be regained immediately. Finally, the schema theorem is a theory of expectation and hence does not take variance into account; in many interesting problems, schema fitness variance can be quite high, making the detection of schemata that contain the global optimum a noisy process (Goldberg & Rudnick, 1991; Rudnick & Goldberg, 1991). Significant schema fitness variance can lead to suboptimal or "premature" convergence.

Despite its simplicity, the schema theorem describes several important aspects of a GA's behavior. Higher mutation probabilities increasingly disrupt higher order schemata, while higher crossover probabilities increasingly disrupt higher defining-length schemata. When selection is factored in, a population converges at a rate proportional to the ratio of the best individual's fitness to average population fitness; this ratio is one measure of *selection pressure* (Bäck, 1994). Increasing either p_c or p_m, or decreasing the selection pressure, leads to increased sampling or exploration of the search space, but does not allow as much exploitation of good schemata that the GA locates. Decreasing either p_c or p_m, or increasing the selection pressure, leads to increased use or exploitation of the better schemata, but does not allow as much exploration for good schemata. The GA must maintain an often delicate balance between what are commonly known as *exploration* and *exploitation*.

Some researchers have criticized the typically fast convergence of the GA, stating that the simultaneous testing of enormous quantities of overlapping schemata requires more sampling and slower, more controlled convergence. While higher population sizes increase schema sampling (Goldberg, Deb, & Clark, 1992; Mahfoud & Goldberg, 1995), a methodology is still needed for controlling convergence in the simple GA.

2.4 Research

The most recent International Conference on Genetic Algorithms (Forrest, 1993) showcased the work of over 200 researchers, in both the theory and application of GAs. Theoretical research covered the modelling and analysis of GAs using Markov chains and other statistical methods; the search for optimal control-parameter settings; the design of problem representations and genetic operators; the construction and solution of difficult problems; the design of mechanisms for niching and for the maintenance of population diversity; the parallel implementation of GAs; the design of hybrid GAs that incorporate ideas borrowed from neural networks, simulated annealing, fuzzy logic, hillclimbing, and tabu search; and the comparison of algorithms. Applied research covered problems in classification, combinatorial optimization, design, function optimization, information retrieval, machine learning, noise-tolerant problem solving, scheduling, search, simulation, structural optimization, and tracking nonstationary environments. Current applications of GAs range from evolving drawings (E. Baker, 1993; Caldwell & Johnston, 1991; L. Davis, 1994), to composing music (Horner & Goldberg, 1991; Laine & Kuuskankare, 1994; Mc Intyre, 1994), to investing money (Bauer, 1994; Mahfoud & Mani, in press).

A line of GA research influential to this thesis is the methodology for analysis and design promoted by Goldberg and colleagues, based upon the Wright Brothers analogy mentioned in Chapter 1 (Goldberg, 1993a, 1993b, 1993c). In the remainder of the current chapter, we review this methodology, mention the possibility of fast-convergence proofs for GAs, and place this thesis into the broad spectrum of GA research.

The Wright Brothers decomposed the design of the powered aircraft into three subproblems — "lateral stability", "lift and drag", and "propulsion" (Goldberg, 1993a, 1993c). They tackled each subproblem separately, and later pieced together the results. Their approach was highly successful, despite the fact that the subproblems were not entirely separable. Goldberg

recommends designing other complex systems, such as GAs, by following the Wright Brothers'
approach. Even though a complex system's subcomponents may not be completely separable, he recommends finding an approximate and intuitive decomposition into "quasi-separable subproblems". After decomposition, he recommends examining each subproblem in isolation, using whatever tools are most effective. Tools include intuition, experimentation, and "rough analysis" of models "that isolate one or two main effects" (Goldberg, 1993b). The final step is to reassemble the full problem, using experimentation to fine-tune the design to compensate for "unforeseen interactions".

Goldberg, Deb, and Clark (1992) have put forth a decomposition for designing a general-purpose GA in which the goal is to locate a single, global solution. They decompose the design of a "selecto-recombinative" GA in terms of building-block processing. (They do not incorporate mutation in the design; this removal of mutation can be viewed as an immediate decomposition of the simple GA.) Specifically, the authors present several subproblems, involving building blocks, whose solution should lead to the design of a better GA.

The first subproblem is to bound the difficulty of problems that the GA must solve. Much work has been done to characterize the functions that are most difficult for a GA to optimize (Bethke, 1980; Brindle, 1981; Das & Whitley, 1991; Davidor, 1991a; Deb & Goldberg, 1993, 1994; Forrest & Mitchell, 1993; Goldberg, 1987, 1989a, 1989b, 1991a, 1991b, 1992; Goldberg, Korb, & Deb, 1989; Grefenstette, 1993; Kargupta, Deb, & Goldberg, 1992; Kinnear, 1994; Liepins & Vose, 1990, 1991; Mason, 1991; Whitley, 1991). It is generally recognized that GA-hard optimization problems possess one or more of the following properties (Goldberg, 1993a; Horn & Goldberg, in press): multiple optima; isolation of desired optima; *misleading* optima (extraneous optima that lead away from the desired optima); noise within schema partitions; and noise between schema partitions. (A *schema partition* of the search space, given a set of bit-positions, is the set of all schemata that have those bit-positions fixed.) Multimodality, by itself, can be ignored as a source of hardness for the GA, if one wishes to locate only a single optimum (Brindle, 1981; Goldberg, Deb, & Horn, 1992; Horn & Goldberg, in press). The problem of noise can be handled through appropriate population sizing (Goldberg, Deb, & Clark, 1992). That leaves isolation and misleadingness as the remaining sources of hardness.

Problems which have both misleading optima and isolated, desirable optima should be the most difficult for the GA to solve. *Deceptive* problems (Goldberg, 1987) are widely studied

examples. Classic deceptive problems have a global optimum and another local optimum, called the *deceptive optimum.* The global optimum has a small basin of attraction, while the deceptive optimum has a large basin of attraction. To make matters worse, the deceptive optimum is similar in fitness to the global optimum.

Formally, deceptive problems are defined in terms of schema partitions. A schema partition is deceptive if the schema containing the deceptive optimum is fitter than all other schemata in the partition. A problem is order-x deceptive if all partitions containing schemata of less than order-x are deceptive.

Deceptive problems are useful in the analysis, testing, and design of GAs because they are of bounded difficulty (order-x). GAs that solve deceptive problems should also be able to solve other problems of up to the same level of difficulty. Furthermore, analysis can determine the order of deception of any objective function (Goldberg, 1989b; Homaifar, Qi, & Fost, 1991). However, one must be careful that a problem which has a certain order of deception does not have a higher level of difficulty along another dimension of problem hardness.

The second subproblem of Goldberg, Deb, and Clark's decomposition is to ensure correct decision-making in the presence of schema noise. They solve this subproblem through adequate population sizing. The main idea is that the population must be large enough to generate an amount of signal sufficient to overcome a level of estimated noise (Goldberg, Deb, & Clark, 1992; Goldberg & Rudnick, 1991; Kargupta & Goldberg, 1994; Rudnick & Goldberg, 1991). Smith (1993) presents a method for adaptively estimating schema noise.

The third subproblem is to construct desired solutions via the beneficial mixing or exchange of building blocks. Two recent studies (Goldberg, Deb, & Thierens, 1993; Thierens & Goldberg, 1993) attack this subproblem. The authors first size their populations using Goldberg, Deb, and Clark's (1992) population-sizing formulas. They then relate mixing success and mixing failure to crossover probability and selection pressure.

An issue related to mixing is *linkage.* Recall that schemata with high defining lengths are likely to be disrupted by single-point crossover. If the GA is to solve problems in which the fixed bit-positions of useful schemata are highly separated, a mechanism must be introduced to bring those fixed bit-positions closer together. Holland (1992) suggests the *inversion* of string segments as a reordering operator. Although inversion is intuitively appealing, it must be applied sparingly in order to function properly (Holland, 1992; Goldberg & Bridges, 1990).

Goldberg, Korb, and Deb (1989) conclude that inversion is impractical for the GA, because its required time frame (to produce good results) is much greater than the time frames for selection and crossover. The *messy GA* (Deb, 1991; Goldberg, Deb, & Korb, 1990; Goldberg, Korb, & Deb, 1989; Merkle & Lamont, 1993) contains a promising approach to reordering that has solved order-5 deceptive problems of up to 150 bits, in subquadratic time with respect to the number of bits (Goldberg, Deb, Kargupta, & Harik, 1993).

A fourth subproblem is to provide the GA an "adequate supply of building blocks". One way to provide these building blocks is to introduce diversity throughout the GA run. Another way is to ensure that the required building blocks are present upon initialization, either through sufficiently high population-sizing (Goldberg, Deb, & Clark, 1992), partially enumerative initialization (Goldberg, Korb, & Deb, 1989), or probabilistically complete initialization (Goldberg, Deb, Kargupta, & Harik, 1993).

A final subproblem from Goldberg, Deb, and Clark's decomposition is to promote building block growth via selection. Goldberg and Deb (1991) find that convergence time for most selection methods is proportional to $\log n$, where n is population size. Selection pressure determines the base of the logarithm. Therefore, by adjusting selection pressure, the user can control the pace of selection and, consequently, the rate of building block growth.

With the above decomposition in mind, we briefly examine the tantalizing possibility of proving fast, global convergence for the simple GA. We emphasize *fast*, global convergence, because prior convergence proofs have relied upon either complete enumeration of the search space or infinite time for mutation to locate the global optimum. One of the more interesting results is the extension of simulated annealing's asymptotic convergence proofs to variations of the GA (Mahfoud & Goldberg, 1995). Unfortunately, simulated annealing's proofs require either an infinite number of iterations or an exponential number of coolings (Aarts & Korst, 1989; Romeo & Sangiovanni-Vincentelli, 1991). A proof of fast, global convergence for the GA would assume an upper bound on problem difficulty, and would be asymptotic in nature. An *asymptotic* proof ensures, with a fixed confidence, the location of a final solution within a certain qualitative distance of the global optimum; the right parameter settings can make both confidence and relative quality, arbitrarily high.

The existence of a fast-global-convergence proof is highly likely. One can model a GA with selection and crossover, as an absorbing Markov chain (Lial & Miller, 1989; Mahfoud, 1993),

with all possible populations as states. This Markov chain has 2^l absorbing states, where l is the string length. Each absorbing state represents a population that has converged to all of one string. By definition, there exists a nonzero probability of ultimate transition from a variety of starting states to each absorbing state. It follows that there exists a nonzero probability of ultimate transition from a variety of starting states to an absorbing state containing all globally optimal strings. We call such an absorbing state a *globally optimal state.* A proof would relate parameters such as population size, probability of crossover, and selection pressure, so that the probability of moving from starting states with nearly uniformly distributed alleles, to within some qualitative distance of a globally optimal state, reached a sufficiently high level. As population size increases, this probability of near-global convergence should also increase, under a variety of settings for the other control parameters.

Despite the likely existence of a fast-global-convergence proof, it is undoubtedly still too early in the evolution of GA research to focus attention upon one. Many small steps remain to be taken in GA modelling, analysis, and design, that will lead to a better understanding of how GAs work. The decomposition approach based on the Wright-Brothers analogy has already taken steps toward increasing our understanding.

Like the approach of Goldberg and colleagues, this thesis follows a methodology of problem decomposition, this time applied to the modelling, analysis, and design of niching GAs. Since the goal in niching is to find multiple solutions, multimodality becomes the major source of problem hardness. However, isolation and misleadingness are also significant sources of difficulty. We will separate out and then ignore, for the most part, the issue of noise, since prior results can be used directly (Goldberg, Deb, & Clark, 1992; Goldberg & Rudnick, 1991; Kargupta & Goldberg, 1994; Rudnick & Goldberg, 1991). We will make use of test functions from the literature that are both massively multimodal and deceptive (Deb, Horn, & Goldberg, 1993; Goldberg, Deb, & Horn, 1992; Horn & Goldberg, in press).

As stated in the previous chapter, the ultimate goal of this research is to make progress in the design of general-purpose niching GAs. Therefore, this thesis builds upon prior research in population diversity and niching methods. This thesis also takes prior research on the modelling of GAs, in new, more productive directions. In subsequent chapters, we review prior methods for promoting population diversity, for niching, and for modelling GAs.

Chapter 3

Diversity

In prior GA literature, niching methods have been the by-products of a quest for techniques to promote population diversity. The simple GA's selection mechanism replicates higher fitness solutions and discards lower fitness solutions, leading to convergence of the population. However, given alternative solutions of identical fitness, the population will still converge. The simple GA loses solutions and subsolutions due to three effects: *selection pressure, selection noise,* and *operator disruption.* Selection pressure is the result of the expected value of the selection process: lower fitness solutions are expected to disappear from a finite population. Selection noise results from the variance of the selection process: in a finite population, random choices among identically fit, competing subsolutions add noise to the expected count for each individual, eventually forcing good solutions from the population. Operator disruption results from the application of crossover and mutation, which are capable of directly destroying good solutions.

Techniques for diversifying a population typically reduce either selection pressure, selection noise, or operator disruption (or some combination). This chapter examines prior techniques for promoting population diversity that do not qualify as niching methods. The boundary between diversification techniques and bona fide niching methods is explained at the end of this chapter and at the beginning of Chapter 4. Chapter 4 reviews prior niching methods.

The remainder of the present chapter commences by motivating the need for diversity in GAs. Next, it constructs simple optimization problems on which the traditional GA fails because of a lack of diversity. It explores the variance and the expectation of the selection

scheme, as well as the disruption due to other operators, as potential culprits. Afterward, the chapter examines previous diversification methods and the approaches they take to promote diversity. It then presents a formal framework for the study of diversity in GAs, and looks at how prior measures of diversity fit within this framework. Finally, this chapter explores the notion of useful diversity, and motivates Chapter 4's definition of a niching method.

3.1 Motivation

Consider the eight-bit, example function below, a small "royal road" function (Forrest & Mitchell, 1993):

$$E1(x) \; = \; \begin{cases} 20 \; , & \text{if } x = 11111111 \; ; \quad \text{otherwise} \; , \\ 10 \; , & \text{if } x \in {*}{*}{*}{*}1111 \; ; \quad \text{otherwise} \; , \\ 10 \; , & \text{if } x \in 1111{*}{*}{*}{*} \; ; \quad \text{otherwise} \; , \\ 1 \; . \end{cases} \tag{3.1}$$

The function $E1$ has a single global optimum that occurs when all bits are set to '1'. It also has two stepping stones to the global optimum — plateaus of intermediate fitness — each containing 16 solutions. All other solutions have the lowest fitness of 1.

We run the simple GA on $E1$, using roulette-wheel selection (RWS), population size $n = 16$, and typical settings for crossover probability ($p_c = .9$) and mutation probability ($p_m = .01$). The GA runs until it converges; we define the convergence point as the generation in which average population fitness is less than or equal to average population fitness from four generations prior. Results are shown in Table 3.1.

The population locates the global optimum after only one generation, and substantially converges after only three generations. The initial population (Generation 0) consists of 16 unique elements. The final population, however, contains only 7 different elements, and many of these elements differ in only one or two bits. The diversity that remains is due primarily to the random bit-flipping of mutation. (We examine mutation as a potential mechanism for diversity, later in this chapter.)

The simple GA with $n = 16$ successfully locates the global optimum of $E1$. We attribute the GA's success to its having an initial population large enough to contain building blocks

Table 3.1: The simple GA with $n = 16$, $p_c = .9$, and $p_m = .01$ runs on function *E1*. The generation number is g and the average population fitness is $\bar{f}$.

g	$\bar{f}$	The Population
0	2.13	10100000 11110010 11000100 11011100 10000100 10101100 11010000 01111111
		00000101 10001100 10001101 00100101 01000011 01010011 00000001 00100110
1	4.44	00000000 11010101 01110101 00000010 11110000 10000110 11110011 01111100
		11010010 11110000 01111110 10101101 01111111 01111011 00001110 11111111
2	9.13	01110000 11101111 11110000 11111111 11110011 11111111 01111100 11110011
		11111111 00001110 00000010 01110101 11111111 11110000 01111111 01111100
3	10.75	11111011 11011111 11111111 11110011 00000011 11110010 11110111 11111011
		11011111 11111111 11110111 11111000 11111111 11101111 11111101 01110111
4	10.63	11011111 11111111 11111001 11111011 11011111 11110111 11011111 11011111
		11101111 11110010 11101111 11111000 11111011 11111011 11011111 11110111
5	8.31	11110111 11101101 11111011 11011111 11110111 11101111 11011111 11111011
		11011111 11111011 11111010 11111011 11110111 10111111 11100111 11100111
6	10.06	11011011 11110111 11110111 11111011 11111011 11011111 11111111 11111011
		11110111 11101111 11101111 10111111 11101111 11111011 01101111 11101111
7	10.19	11101111 11111011 11101111 01101111 11110111 11111011 11101011 11111111
		11111111 11101011 11101111 11111111 11101111 11101111 11101101 01101111

sufficient to form the global optimum. Note that the two stepping stones to the global optimum are represented by one element apiece in the initial population.

In many real-world problems, time or memory limitations restrict the size of the population, so that it does not contain the required building blocks upon initialization. Consider what happens when the GA runs again on *E1*, using the smaller population size $n = 8$; all other parameter settings, including the convergence criterion, remain the same. Table 3.2 shows that the GA locates the global optimum at Generation 2, but loses it the next generation, and never recovers it. (The run ends after six generations.) The GA also manages to lose all representatives of one of the plateaus. After the final generation, five distinct elements remain. The same experiment with $n = 4$ fails to locate the global optimum at any point during the run, and never encounters either plateau.

What the smaller population GA apparently needs and what many researchers have suggested is a mechanism to either maintain or reintroduce diversity. The maintenance of diversity has most often been the focus. The reason for emphasizing maintenance is illustrated by run-

Table 3.2: The simple GA with $n = 8$, $p_c = .9$, and $p_m = .01$ runs on function *E1*. The generation number is g and the average population fitness is $\bar{f}$.

g	$\bar{f}$	The Population
0	3.25	10100000 11110010 11000100 11011100 10000100 10101100 11010000 01111111
1	5.50	01110100 11001111 01110111 01111010 11110010 11110010 11110000 11010000
2	10.13	11111111 11000010 11110010 11110010 11110010 11110110 11110010 11110010
3	8.88	11110010 11100110 11110010 11110110 11110010 11110010 11110011 11111110
4	8.88	11111110 11110010 11110010 11110010 11110011 11100110 11110010 11110110
5	7.75	11110111 11010010 11110011 11110010 11100110 11110010 11110011 11110010
6	7.75	11010011 11110011 11110010 11110111 11010011 11110110 11110011 11110010

ning the GA on example function *E2*, shown in Figure 3.1. An eight-bit function of unitation, *E2* contains two global maxima, one at 00000000 and one at 11111111; an equal number of points in the search space lead to each global maximum. *Unitation* functions are defined over the number of '1' bits contained in a string. For example, if x is an eight-bit string, and $count(x)$ is the number of ones that x contains, $f(count(x))$ is a function of unitation.

We run the previous GA (with RWS, $p_c = .9$, and $p_m = .01$) on *E2*, at various population sizes. Populations of size $n = 2$, $n = 4$, and $n = 8$ fail to locate either optimum at any point during the run. A population of size $n = 16$ finds one global optimum early in the run, but loses it by the end of the run. A run with $n = 32$ is illustrated in Table 3.3.

Note that the initial population is lucky enough to contain one of the global optima, 11111111. However, by Generation 4, this global optimum disappears, never to be seen again, and the global optimum from the opposite pole, 00000000, appears in the population. By the end of the run, most population elements converge to 00000000, and all nonglobal population elements are either one or two bits away from 00000000.

One might expect a larger population to locate and maintain both global optima. However, $n = 64$, like $n = 32$, converges about the global optimum, 00000000, after only 20 generations; $n = 128$ does too, after 26 generations, despite the fact that the initial population contains both global optima. After 25 generations, $n = 256$ also converges about 00000000; so does $n = 512$, this time after 29 generations. A population of size $n = 1024$ converges, after 29 generations, about the opposite global optimum, 11111111. Even $n = 2048$ converges, after only 33 generations, about 11111111, performing $69,632$ function evaluations in the process.

Figure 3.1: Function *E2* is displayed over unitation space.

Figure 3.2: Function *E3* is displayed over unitation space.

Table 3.3: The simple GA with $n = 32$, $p_c = .9$, and $p_m = .01$ runs on function *E2*. The generation number is g and the average population fitness is $\bar{f}$.

g	$\bar{f}$	The Population
0	2.25	10100000 11110010 11000100 11011100 10000100 10101100 11010000 01111111
		00000101 10001100 10001101 00100101 01000011 01010011 00000001 00100110
		11010011 01000001 00010101 01100111 10110011 10100101 11000001 11111111
		00010101 10100000 11000100 10111000 11010010 10100011 11011011 11011111
4	3.03	00000101 01111111 00011111 01100001 10011111 10101011 00010000 11000001
		01100101 01111001 10000001 00000000 00000001 00010001 01111101 01111101
		11111011 00010111 01100001 11011101 11000101 00000000 11011111 01111101
		10111111 01000000 10000000 00011111 00000000 01000000 10000011 00010100
8	3.50	11010000 00001111 00000101 00010100 10111111 01000000 00111000 10101110
		00111000 00000011 10000000 10001000 00111111 00000000 00000000 10000000
		10101000 10000000 00011111 11000000 00000000 00000000 01100000 00101000
		11000000 10000000 00000000 00000000 00000100 10000000 00000000 00000000
12	4.16	10000000 00001000 10000000 01000000 10000000 00001000 00110000 00000000
		10000000 00000000 10001000 10000000 00110000 10100000 00000000 10000100
		00000000 00000000 10000000 10001000 00000000 10000000 00000000 00000000
		00000000 00000000 10100000 00000000 10000000 11000000 00000000 00000000
16	4.69	00000000 10000000 00000000 00000000 00000000 00000000 00000000 00000000
		00000000 00000000 00000000 00000000 00000000 10000000 00000000 10000000
		10000000 10000000 00000000 00000000 00000000 00100000 10000000 00000000
		00000000 00000000 01000100 00000000 00000000 10000000 00000000 00000000
18	4.59	00000000 10000000 01000000 00000100 01000000 00000000 00000000 01000000
		00000000 00000000 01000100 00000000 00000000 00000000 00000000 00000010
		00000100 01010000 00000000 00000000 00000000 00000000 00000000 00000000
		00000000 01000000 00000000 00000000 00000010 00000000 00000000 00000000

This number of function evaluations is 272 times the size of the search space! Apparently, the population, regardless of size, is unable to maintain both global optima. We propose that this behavior, as well as the earlier loss of diversity on *E1*, can be attributed mainly to the high variance of roulette-wheel selection, with operator disruption possibly playing a minor part. We examine the first part of this conjecture in Section 3.2, and the second part in Section 3.5.

Let us consider one more example, eight-bit unitation function *E3*, shown in Figure 3.2. *E3* differs from *E2* in that the 11111111 optimum is higher and has a steeper path leading up to it. Function *E3* is similar to Ackley's (1987) "two-max" function. *E3* has one global

Table 3.4: The simple GA with $n = 8$, $p_c = .9$, and $p_m = .01$ runs on function $E3$. The generation number is g and the average population fitness is $\bar{f}$.

g	$\bar{f}$	The Population
0	3.00	10100000 11110010 11000100 11011100 10000100 10101100 11010000 01111111
1	2.38	01111100 11011111 00100111 01111000 11000100 11011100 11110000 10101100
2	3.25	11011100 01111111 11000101 11011110 00100111 00100011 11011110 11000100
3	3.63	00101110 11000011 11000101 00100011 01111110 11011111 11011111 01111110
4	4.25	01111110 11011111 11000011 11000011 00111111 11000011 11011111 11011111
5	6.25	00111011 11111111 00011111 11111111 11011111 11011111 00111111 11011111
6	7.00	00011111 11111111 11111111 11011111 11111111 00111011 11111111 11011111
7	7.75	11011111 11111111 11011011 11011111 11011111 11111111 11111111 11111111
8	7.75	11111111 11011011 11111111 11111111 11011111 11011111 11011111 11111111
9	7.75	11011111 11111101 11011111 11111111 11011111 11111111 11111111 11011111
10	6.75	11011111 11011111 11111110 11111111 11111111 10011111 11001111 11010111

optimum, 11111111, and one nonglobal, local optimum, 00000000. Suppose we are interested in locating and maintaining both optima. Our GA must contend not only with selection noise and operator disruption, but also with selection pressure, which pushes the population toward the global optimum.

We run the same GA on $E3$ as we did for the previous two problems (with RWS, $p_c = .9$, and $p_m = .01$). At $n = 2$ and $n = 4$, the GA never locates either optimum. At $n = 8$, shown in Table 3.4, the GA never encounters the local optimum, but finds the global optimum in Generation 5, and maintains it through the end of the run. At $n = 16$, the GA finds the global optimum in Generation 1, loses it the next generation, rediscovers it in Generation 5, and keeps it until the end of the run; the GA never encounters the local optimum. At $n = 32$, the GA briefly discovers the local optimum in Generation 2, but permanently loses it the next generation; the initial population contains the global optimum, which the GA maintains through the end of the run. For all $n \geq 8$ that we test, the population converges in the end about the 11111111 point. For n of 32, 64, 128, 256, 512, 1024, and 2048, the initial population contains the global optimum, and the global optimum remains in the population through the end of the run. The local optimum is not so fortunate. At $n = 64$, the local optimum is never discovered. For n of 128, 256, 512, 1024, and 2048, the initial population contains between 1

and 10 copies of the local optimum, but in each case, the local optimum disappears from the population between Generations 1 and 5, and once absent, is never seen again.

When both global and nonglobal optima are of interest, we propose that selection pressure becomes an adversary (in addition to selection noise and operator disruption). Selection pressure is also a factor when less fit *subsolutions* are required for constructing one or more optima. For instance, if the problem is misleading, less fit schemata will be needed, in order to construct the global optimum. We further examine in Section 3.3 the proposition that selection pressure hinders the formation and maintenance of nonglobal optima.

3.2 Selection Noise

In the previous section, we conjectured that the high variance of roulette-wheel selection was the primary cause of the observed loss of solutions and subsolutions, especially in cases where solutions or subsolutions were identically fit. To further examine the variance of RWS, we must first isolate RWS from the rest of the GA, namely the crossover and mutation operators.

We conduct the following experiment using RWS only ($p_c = 0$, $p_m = 0$, and $n = 16$). Consider the one-bit function $E4$, in which $E4(0) = E4(1) = 1$. The initial population is uniform, with eight 0s and eight 1s. Figure 3.3 tracks the population distribution over time. With RWS isolated, the essential component of the GA's behavior on $E1$ and $E2$ repeats itself on $E4$. The population soon converges to all 1s. Note that it could just as easily have converged to all 0s; since the two solutions have the same fitness, the GA does not prefer one over the other. Over the course of 10 runs, the average time to full convergence is 21.9 generations. Six out of 10 runs converge to all 1s; 4 out of 10 runs, to all 0s.

Other GA researchers have noted similar behavior on a variety of problems. While many researchers blamed the vague notion of premature convergence, a few, starting with De Jong in his 1975 dissertation, were able to identify the real culprit — the variance of selection.

One might logically ask how long it takes a population of 0s and 1s to fully converge. If selection noise is the only factor at work, as on $E4$, convergence is linear in population size. The behavior of RWS is analogous in some ways to the *gambler's ruin*, in which a gambler with a certain amount of money places successive, fixed sized bets. The gambler plays until either going bankrupt or doubling his/her money, at which point the gambler leaves the table. Goldberg

27

Figure 3.3: RWS runs on E_4, with $n = 16$, $p_c = 0$, and $p_m = 0$. The starting distribution is uniform.

and Segrest (1987) use this analogy, along with Markov chain computations, to show that the expected time under RWS for either the '0' or '1' allele to take over a specified percentage of the population, is linear with respect to population size. They define the takeover point as the first generation in which either allele achieves or surpasses the specified percentage. The authors also show that when mutation is added, higher mutation GAs take longer to converge. They show that as mutation probability increases, takeover time grows from a linear function of population size to an exponential function.

Having identified the variance of RWS as a major problem, we now quantify that variance. Brindle (1981) computes the variance in the number of instances of solution i expected after application of RWS, where i is any element of the search space present in the population prior to application of RWS. This variance is $\sigma_i^2 = n p_s(i)(1 - p_s(i))$, where $p_s(i)$ is the selection probability for i. (We re-derive this variance in a later chapter, but for classes of solutions rather than single solutions.) She proves that RWS has higher variance than five other fitness-proportionate selection schemes, and demonstrates RWS's inferior performance on several test functions. Other authors (J. E. Baker, 1987; Booker, 1982; De Jong, 1975) also obtain worse empirical results with RWS than with lower variance selection schemes.

Figure 3.4: RWS runs on $E5$, with $n = 16$, $p_c = 0$, and $p_m = 0$. The starting distribution is uniform.

3.3 Selection Pressure

Having alleles of equal fitness is the best scenario when trying to maintain diversity. When alleles are of differing fitnesses, selection pressure becomes a factor, causing faster convergence (Goldberg & Deb, 1991). Goldberg and Segrest (1987) show that given two unequally fit alleles, the higher the fitness differential, the faster the expected convergence, and the less likely that the weaker allele wins in the long run. We illustrate their result, using the one-bit problem $E5$, in which $E5(0) = 1$ and $E5(1) = 2$. We again use RWS only ($p_c = 0$, $p_m = 0$, and $n = 16$), along with uniform initialization. Figure 3.4 shows the temporal population distribution.

The population soon converges to the global optimum, as it previously did on the more complex function $E3$. We conduct 10 additional runs to examine whether convergence is due to the expectation or the variance of selection. Unlike on $E4$, all 10 populations converge to all 1s. The average time to full convergence is only 4.4 generations, compared with 21.9 generations on $E4$. Barring the influence of some unknown force, the stronger force of selection pressure is primarily responsible for the loss of diversity on $E5$.

In the previous section, we demonstrated that selection noise (or variance) can cause a population to converge fairly rapidly to one of a set of identically fit solutions or subsolutions. In the current section, we have shown that selection pressure (or expectation) is an obstacle to the formation and maintenance of nonglobally optimal solutions and subsolutions. In the following section, we will consider a third factor — disruption due to genetic operators other than selection.

3.4   Operator Disruption

Operators allow GAs to explore the search space. However, operators typically have destructive as well as constructive effects. Ideally, if the GA has located an optimal solution or subsolution, the GA's operators should leave that solution or subsolution intact.

The mutation operator is typically applied in small doses, such as the .01 mutation probability we have been using. Researchers including De Jong (1975) and Brindle (1981) have found that higher rates of mutation push the GA's behavior towards that of random search. Subsequent studies have predominantly verified this discovery. Therefore, we assume that mutation occurs infrequently and hence causes negligible disruption. This assumption allows us to concentrate on disruption due to crossover, an operator typically applied in large doses.

Crossover disruption is minimal on some problems, but prominent on others. Chapter 2's re-derivation of the schema theorem shows that schemata with high defining lengths are likely to be disrupted by single-point crossover. However, this unary view is rather limited, since crossover is a binary operator. When a problem contains solutions or subsolutions of interest that differ greatly amongst themselves, one can expect disruption due to crossover to be at its highest: crossover must frequently mix complementary bits. Function $E2$ is one such function. We will examine crossover disruption on $E2$ in the next section, after separating out the issue of selection noise.

3.5   Previous Research

This section reviews proposed mechanisms for promoting genetic diversity that, by themselves, do not qualify as niching methods. We can classify previous diversification methods under three major categories: those which attempt to reduce selection noise, those which attempt to reduce

Table 3.5: Diversification mechanisms are categorized, according to whether they reduce selection noise, selection pressure, operator disruption, or a combination.

Category	Selection Noise	Selection Pressure	Operator Disruption
Direct Infusion			
Noise-Reduced Selection	X		
Multiple Sequential Runs	X		
Restricted Mating			X
Thermodynamic GAs		X	X
Adjusted Control Parameters	X	X	X
Multiple Parallel Runs	X	X	X

selection pressure, and those which attempt to reduce operator disruption. Diversification methods and their corresponding categories are summarized in Table 3.5. Some diversification schemes attempt a combination of the three approaches, and hence fall under more than one category.

The distinction between diversification and niching mechanisms is sometimes a fine one. It hinges on the utility of the diversity that the mechanism promotes. For example, many diversification schemes slow a GA's convergence, but are unable to stably maintain multiple solutions, even when those solutions are all globally optimal. Other schemes are somewhat successful at maintaining multiple global optima, but break down when trying to maintain multiple local optima. We further examine the distinction between diversification and niching methods in the final section of this chapter, and in Chapter 4.

Noise-reduced selection

Roulette-wheel selection is highly noisy, as we have previously demonstrated. Many researchers have suggested schemes with lower variance, but with the same or approximately the same expectation (as RWS). The alternative schemes lessen variance by introducing determinism, through techniques such as sampling without replacement. We briefly review two of these methods, *stochastic remainder selection with replacement* (Brindle, 1981) and *stochastic universal selection* (J. E. Baker, 1987). We choose these two methods because they have exactly the same expectation as RWS, but lower variance. The two methods deterministically allocate to each population element, its expected number of whole offspring; the methods probabilistically allocate expected fractional offspring. For instance, an individual whose fitness is 1.5 times the

average population fitness is guaranteed that one copy of itself advances to the crossover stage. The remainder probability of .5 determines the likelihood of other copies advancing.

Stochastic remainder selection with replacement (SRSWR; originally called "remainder stochastic sampling with replacement") allocates fractional offspring by invoking RWS on the remainder probabilities. The remainder probability r_s for element i is $r_s(i) = np_s(i) - \lfloor np_s(i) \rfloor$, where $p_s(i)$ is the selection probability for i, defined in Equation 2.2. Brindle (1981) computes the variance in the expected number of i elements selected by SRSWR. The variance is $r_s(i)(1 - r_s(i)/\sum_{i=1}^{n} r_i)$. She proves that the mean square error, over all population elements, for one generation of SRSWR is always less than or equal to the corresponding error for RWS.

J. E. Baker (1987) analyzes various fitness-proportionate selection methods, with respect to three criteria: efficiency, bias (distance from RWS in expected value), and spread ("the range of possible" offspring counts for an individual). Spread is a measure of selection noise. Baker states that although SRSWR is unbiased, its spread is "virtually unlimited", meaning that an element with above-average fitness could advance from 1 to n copies of itself to the crossover stage; an element with below-average fitness could advance from 0 to $n - 1$ copies. He unveils a selection scheme that is optimal with respect to the three criteria. This selection scheme is stochastic universal selection (SUS; originally called "stochastic universal sampling").

SUS works as follows. Recall that RWS simulates a roulette wheel with n slots, each of size $p_s(i)$ (for $i = 1$ to n), where the total circumference of the wheel is 1. SUS simulates the same wheel. However, instead of spinning the wheel n times, with a single pointer indicating the winner of each spin, SUS spins it only once, using n uniformly spaced pointers on the outside of the wheel.

SUS has zero bias, is extremely efficient, and has minimal spread. Individual i will always be selected between $\lfloor np_s(i) \rfloor$ times and $\lceil np_s(i) \rceil$ times, where $np_s(i)$ is the expected number of selections for i under fitness-proportionate selection, and the lower and upper brackets respectively denote the functions of rounding down and rounding up to the nearest integer. We employ SUS throughout this paper as our lowest noise, fitness-proportionate selection scheme, and employ RWS as our highest noise, fitness-proportionate selection scheme.

We illustrate the performance of SUS through 10 runs on $E4$, the one-bit function in which both solutions are of equal fitness. In all 10 runs, SUS maintains the initial distribution of eight 0s and eight 1s perpetually. (We stop each run after 100 generations.) We can prove

Table 3.6: The GA with SUS, $n = 8$, $p_c = .9$, and $p_m = .01$ runs on function *E1*. The generation number is g and the average population fitness is $\bar{f}$.

g	$\bar{f}$	The Population
0	3.25	10100000 11110010 11000100 11011100 10000100 10101100 11010000 01111111
1	6.75	01110010 11111111 01110100 11011010 01111111 01111111 01111110 11110001
2	12.75	11110001 11111111 11111111 01111111 11111111 01111011 01110011 11111111
3	17.50	11111111 11101111 11111111 11111111 11111111 01111111 11111111 11111111
4	20.00	11111111 11111111 11111111 11111111 11111111 11111111 11111111 11111111
5	17.50	11111011 11011111 11111111 11111111 11111111 11111111 11111111 11111111
6	17.50	11011111 11111111 11111111 11111111 11111111 11011111 11111111 11111111
7	17.50	11011111 11111111 11111011 11111111 11111111 11111111 11111111 11111111

that SUS, acting in isolation, will perpetually maintain a global optimum. Let j be a global optimum that is present in the population. Since global optima are at least of average fitness, the expected offspring count for j under SUS is greater than or equal to 1 ($n p_s(j) \geq 1$). Since SUS guarantees any individual at least the integer portion of its expected offspring count, j is guaranteed at least $\lfloor n p_s(j) \rfloor \geq 1$ offspring. Therefore, j survives indefinitely.

We test the GA with SUS on function *E1*, using $n = 8$, $p_c = .9$, and $p_m = .01$. (Recall that the GA with RWS failed at $n = 8$.) All parameters are the same as previously (under RWS); the initial population is also the same. This time, the GA successfully locates the global optimum and maintains it to convergence. The run is shown in Table 3.6.

So far, we have demonstrated that SUS successfully reduces selection noise. However, it is not a cure-all. For instance, with extremely small populations of size $n = 4$ and $n = 2$, SUS fails to converge to the global optimum on *E1*. This is not a disappointment, however, since such small populations are unlikely to contain very useful building blocks.

SUS also does not reduce selection pressure. On *E5*, with $n = 16$, $p_c = 0$, and $p_m = 0$, 10 out of 10 runs converge to the fitter solution, taking an average of 3.7 generations. The results for SUS are roughly equivalent to those for RWS.

Finally, SUS can be hampered by disruption due to crossover and mutation. On *E2*, with $p_c = .9$ and $p_m = .01$, SUS allows convergence to at least one of the two global optima for all $n \geq 8$. (RWS, under the same parameter settings and using the same initial populations, required $n \geq 32$ to converge to a global optimum.) SUS luckily finds and maintains both global

optima at $n = 32$ and $n = 64$, because these runs terminate early, after Generations 10 and 7, respectively. However, for population sizes of $128, 256, \ldots, 2048$, despite the initial population's containing copies of both global optima, the GA with SUS loses one of the global optima by the end of each run. (Runs take from 33 to 38 generations.) This loss is due to the disruption caused by crossover and mutation. Crossover is the main culprit, since the mutation rate is very low. Additional runs of the GA, with SUS, $n = 256$, and $p_m = .01$, verify the disruptive effects of crossover on $E2$. For $p_c < .2$, the GA generally converges to both global optima. For $p_c > .2$, the GA generally converges to only one of the two optima. Low crossover probabilities, however, are not practical for GAs, because they hinder the formation of good solutions and subsolutions, and they also allow selection to prematurely converge the GA, before crossover has performed sufficient mixing.

Adjusted control parameters

The simple GA gives the user control over population size (n), crossover probability (p_c), and mutation probability (p_m). Some studies utilize selection pressure as an additional control parameter (Goldberg, Deb, & Thierens, 1993; Thierens & Goldberg, 1993). (In fitness-proportionate selection, selection pressure can be controlled through the use of fitness scaling.) A number of studies seek the optimal set of control parameters for a collection of test functions (Bäck, 1992, 1993; De Jong, 1975; Grefenstette, 1986; Schaffer, Caruana, Eshelman, & Das, 1989). Some of these studies are frequently misinterpreted as calls for a single set of control parameters to solve all problems. More careful readers have come to the conclusion that parameter settings which are optimal for a particular set of test functions may hold little meaning for functions outside that set. Optimal parameter settings are undoubtedly problem dependent.

Ongoing research suggests that the proper setting of control parameters can probabilistically guarantee, on problems of bounded difficulty, that the GA converges to a single global optimum (Goldberg, Deb, & Clark, 1992; Goldberg, Deb, & Thierens, 1993; Thierens & Goldberg, 1993). This line of research does not rely on the reintroduction of diversity, but makes the initial population sufficiently large (and therefore diverse) before the GA commences. Nevertheless, the adjustment of control parameters has shown no promise to date for forming and maintaining multiple solutions or subsolutions (niching). The emphasis of this study is not on fine-tuning parameters, but rather on developing effective mechanisms for promoting niching in GAs.

In terms of the three mechanisms for promoting diversity, SUS can minimize selection noise, fitness scaling can reduce selection pressure, and lower p_c and p_m can reduce operator disruption. There is, however, a trade-off between convergence and diversity in the GA. If one takes diversity to the extreme, via an algorithm with $p_c = p_m = 0$ and no selection pressure, the GA will maintain the initial population indefinitely and perform no useful search. It is not beneficial to stunt crossover and mutation as search operators, or to slow indefinitely the speed of selection. A GA must allow beneficial search as well as controlled convergence.

Crossover and mutation have constructive as well as destructive effects. However, crossover does not introduce new alleles to the population. When a population starts to converge, crossover's effects diminish, since crossing two identical solutions yields the same two solutions. Mutation, on the other hand, does introduce new alleles. Goldberg and Segrest (1987) show that higher mutation rates significantly slow the GA's convergence. However, the random variation and increased disruption caused by higher mutation rates do not usually result in useful diversity. As mentioned earlier, even slight mutation rates move a GA towards random search, a very poor search method.

Direct infusion

Some techniques try to infuse diversity directly into the population, attacking the symptoms rather than the causes of premature convergence. Perhaps the most popular form of direct infusion is the use of higher mutation rates. Unfortunately, mutation alone, even at an optimal rate, is ineffective at maintaining useful diversity. First consider low mutation rates. Although they prevent the fixation of bit-positions, they do not slow convergence. The population still converges in the neighborhood of some optimum, then with high probability, repeatedly visits points within that neighborhood. High mutation rates, on the other hand, slow convergence and prevent the fixation of bit-positions. However, they do not allow meaningful convergence. As diversity due to mutation increases, so does the disruption of crucial building blocks. Performance worsens, as measured by the quality of final solutions.

One strategy that slightly increases diversity is *duplicate elimination*. Before a new element is inserted into the population, all other elements are checked to see if they match the new element. If one matches, the new element is mutated in one bit and the population is rechecked. This cycle continues until the new element is made unique. Duplicate elimination is used in

GAs such as *GENITOR* (Whitley & Kauth, 1988) that process only one to several population elements at a time. Although duplicate elimination ensures that population elements are at least one bit apart in Hamming space, all population elements are still likely to converge to the neighborhood of a single optimum. The additional overhead required by duplicate elimination is at least n^2 comparisons per generation, depending upon the number of mutations required to make each element unique.

A more general direct-infusion method is Mauldin's (1984) uniqueness-assurance method. Instead of checking whether each new string differs in at least one bit from all other strings, his algorithm checks whether each new string differs in at least k bits. If a new string differs in only $i < k$ bits from some population element, it is mutated in the matching bit-positions until it is k bits different. Once a newcomer is sufficiently different from all population elements, it is inserted into the population. Mauldin decreases k as the run progresses, making his method resemble a thermodynamic GA. (We will review thermodynamic GAs shortly.)

Mauldin finds his uniqueness strategy to be effective at promoting diverse alleles. It displays controlled convergence, and locates better final results on a set of test functions than do several competing GAs and hillclimbers. As in thermodynamic methods, increasing the initial k leads to better exploration and to better final results. However, the uniqueness strategy does not qualify as a niching method, because diverse alleles, by themselves, are insufficient to simultaneously concentrate the search upon multiple, meaningful regions of the search space.

Reinitialization: Multiple sequential runs

Instead of slowing a GA's convergence, one could encourage it to converge quickly and then start another one. The increased sampling introduced by multiple runs would reduce selection noise. Goldberg (1989d) suggests restarting GAs that have substantially converged, by reinitializing the population using both randomly generated individuals and the best individuals from the converged population. Since simple GAs arrive at a solution, then spend the remainder of their time floundering about that solution via mutation, reinitialization enables GAs to make productive use of excess computation time.

Reinitialization techniques reintroduce diversity either throughout the run or upon convergence. Reinitialization is somewhat similar to utilizing high rates of mutation; however, GAs

with reinitialization often do not employ ordinary mutation, but perform what can be viewed as periodic, mass mutation.

Given a fixed sized population, a GA with reinitialization can expect to find a better ultimate solution than the same GA without reinitialization. The reintroduction of lost alleles allows search to proceed from a wider variety of points. Since reinitialization techniques conduct search primarily within a neighborhood of the current best solution, they are not effective at forming or maintaining multiple solutions.

Micro-genetic algorithms are small-population GAs with reinitialization. Krishnakumar's (1989) micro-GA utilizes $n = 5$, $p_m = 0$, and $p_c = 1$, along with an elitist selection strategy that always advances the best string of the current population to the next generation. Selection holds four competitions between strings that are adjacent in the population array, advancing the fitter string in each competition. Krishnakumar compares his micro-GA to a simple GA with typical parameter settings ($n = 50$, $p_c = .6$, and $p_m = .001$). He reports faster and better results with the micro-GA on two simple stationary functions and on a real-world, engineering control problem. Micro-GAs have also been applied to the optimization of an "air-injected hydrocyclone" (Karr, 1991a), to the design of fuzzy logic controllers (Karr, 1991b), and to the solution of the k-queens problem (Dozier, Bowen, & Bahler, 1994).

The micro-GA can be viewed as a diversification method because it promotes diversity across runs. Over multiple runs, selection pressure rises and falls, but selection noise drops due to redundancy. The micro-GA does not qualify as a niching method, because after it finds one globally optimal solution, it perpetually passes that solution to future generations and future runs, effectively blocking the search for other optima. We will examine more promising techniques for the sequential location of optima, later in this thesis.

Mechanisms similar to the micro-GA have been proposed by Cobb and Grefenstette. Cobb's (1990) GA adaptively alternates between two modes of mutation, a standard mode with $p_m = .001$, and a *hypermutation* mode with $p_m = .5$. Cobb defines performance as the average, across all generations up to and including the current one, of the best population element of each generation. When performance increases or remains the same from one generation to the next, the standard mutation rate applies; when performance decreases, hypermutation emerges. The end effect is similar to reinitialization upon convergence, but slightly less abrupt. Grefenstette (1992) notes that hypermutation will fail to act should the environment become fitter

in a region away from the location to which the GA has converged. He suggests in place of hypermutation that a percentage of the population be reinitialized each generation. He calls this percentage the *replacement rate*, and calls his overall method *random immigrants*.

Cobb and Grefenstette (1993) compare triggered hypermutation, random immigrants, and the simple GA under high mutation, using two nonstationary test problems. Employing an experimentally determined set of optimal parameters for each algorithm, they find that the simple GA under high mutation (p_m ranges from .05 to .15) performs better than or equivalent to the other two strategies on functions that change incrementally over time. However, the simple GA underperforms the other two strategies on functions that oscillate drastically over time. The random immigrants method is generally best for tracking wildly changing environments. Its replacement rate ranges from .05 to .4 in the authors' experiments. Hypermutation probabilities range from .15 to .4. The three algorithms are similar to the micro-GA in terms of their selection pressure and selection noise. Due to higher than usual mutation and replacement rates, all three methods increase rather than decrease operator disruption. None of the three algorithms is capable of maintaining multiple, stable subpopulations within the same population.

The *CHC* genetic algorithm (Eshelman, 1991; Eshelman & Schaffer, 1991), upon convergence, uses the best string in the population as a template to reinitialize the entire population. To form each element of the new population, the best string is mutated in a large portion of its bits. (The authors suggest 35%.) If several successive reinitializations fail to yield an improvement, the population is completely (100%) randomly reinitialized. The authors set $p_m = 0$ prior to convergence.

Delta coding (Whitley, Mathias, & Fitzhorn, 1991) and *canonical delta folding* (Mathias & Whitley, 1993) are more sophisticated strategies that allow periodic reinitialization of the population. The strings of a reinitialized population are re-encoded as difference values (deltas), relative to the previous population's best solution. Re-encodings employ fewer bits per variable (usually one less) than the previous population's encoding. The goal is first to locate an interesting area of the search space, then to perform a reduced search about that area. In delta coding, strings decode to distances, in variable space, from the previous run's best solution; in canonical delta folding, strings are mapped directly in Hamming space. The authors compare GAs that use a delta strategy to GAs that do not. The GAs with delta strategies produce better final solutions on several test functions.

The delta strategies re-encode and reinitialize the population whenever it substantially converges. Convergence occurs when the Hamming distance between the best and worst strings in the population is less than or equal to 1. The re-encoding of strings shifts the encoding bias to new portions of the search space, thus alleviating encoding-related problems such as Hamming cliffs.

Whenever a re-encoded run converges to the same solution as the prior run, the number of bits in the encoding is increased. When a run converges to a worse solution than the prior run, search proceeds from the worse solution; this encourages escape from nonglobal optima.

Isolation and migration: Multiple parallel runs

Perhaps the most basic technique for locating multiple solutions is to perform multiple, independent runs of the simple GA, and to keep the best solution from each run off-line. In parallel, this technique is equivalent to Tanese's (1989a, 1989b) *partitioned GA*, in which each processor runs an independent GA on a local subpopulation, and subpopulations never interact. Since the serial and parallel versions are equivalent, we refer to both as partitioned GAs.

The partitioned GA is highly redundant. While this redundancy helps reduce selection noise, it also causes repeated exploration, from run to run, of certain regions of the search space. In multimodal function optimization, the partitioned GA will, with high likelihood, repeatedly converge to a few prominent optima. In standard function optimization, independent runs will blindly process the same subsolutions, since no coordination will exist among processes.

Consider a problem with b optima, and a GA that is guaranteed to converge to one of the optima. If each optimum is equally likely to be found, the partitioned GA is expected to find them all after $\sum_{i=1}^{b} 1/b$ runs (Beasley, Bull, & Martin, 1993). If the optima are not equally likely, the expected number of runs grows much higher, by an amount determined by the probabilities of finding individual optima.

Permitting a limited amount of communication among subpopulations improves the partitioned GA. In serial, this improved technique restarts a run using information from the best individuals of prior runs (like the reinitialization methods we previously discussed). In parallel, this technique models isolated subpopulations with migration. The serial and parallel versions are roughly equivalent.

Dependent parallel runs are similar to their independent counterparts. The $GAPE$ system (Cohoon, Hegde, Martin, & Richards, 1987; Cohoon, Martin, & Richards, 1991) runs P simple GAs in parallel, each assigned to a separate processor. Each simple GA utilizes a subpopulation of size n_s, and runs for g generations. (The authors choose g sufficiently large to allow all subpopulations to converge.) Upon convergence, each processor duplicates a number of randomly chosen individuals from its subpopulation, and packages them into k equally sized sets. One set is shipped to each of k neighboring processors. (Any topology may be used.) Every processor then reduces its surplus of individuals to n_s via fitness-based selection. The entire process, beginning with the simultaneous running of simple GAs, repeats for a predetermined number of *epochs.*

The authors compare GAPE to a simple GA, and record the best solution encountered by both algorithms. The simple GA employs a population of size $n = n_s$, and runs P times GAPE's total number of generations. Over five experiments, GAPE on the average finds better solutions and also finds the best overall solution. These results are consistent with prior sequential results, in which a GA that is reinitialized upon convergence outperforms a GA that is left converged. The authors also find that GAPE locates better ultimate solutions than the equivalent partitioned GA.

As a potential niching method, GAPE is subject to many of the same objections as are previously presented techniques. Efforts toward locating multiple solutions will be reproduced across processors. After a few epochs, all subpopulations will identically converge, due to exchange of solutions. The authors, in fact, search for only the global solution in the problems they solve. They provide no evidence that alternative solutions can be maintained.

Many parallel GAs provide communication between subpopulations prior to convergence. The objective is to introduce diversity from other subpopulations throughout the run, making each GA less likely to prematurely converge. Tanese (1987) examines the experimental performance of such a parallel GA. Every few generations, specified by the *exchange frequency*, an exchange takes place between each processor and one of its neighbors (the neighbor whose turn it is according to a cyclic schedule). Exchanges duplicate a predetermined percentage of each subpopulation, then ship the duplicated individuals to a neighboring processor. Duplication proceeds probabilistically, with probability of duplication increasing as fitness increases. Only individuals of better than subpopulation-average fitness are eligible for exchange. After a pro-

cessor receives its migrant individuals, it eliminates an equal number of individuals from its subpopulation. Only individuals of less than average fitness are eligible for elimination, with probability of elimination increasing as fitness decreases.

Tanese runs several simulations on a relatively easy problem, keeping the total population size constant but varying the number of processors. Population elements are evenly divided among the available processors. Nearly all runs converge to the global optimum, with multiple-processor versions of the GA yielding a near-linear speedup. Unfortunately, the author's experiment tells us nothing about the effects of migration on diversity.

GENITOR II (Whitley & Starkweather, 1990) is equivalent to Tanese's parallel GA, except that GENITOR (Whitley & Kauth, 1988) runs on each processor, and migration results in additional selection pressure. The k best strings from each source processor replace the k worst strings on each destination processor. GENITOR II is somewhat successful at delaying convergence.

Tanese (1989a, 1989b) develops another parallel GA, with lower migration-induced selection pressure. Migrants are not duplicate elements, but are produced by the standard process of selection, crossover, and mutation. Incoming immigrants replace randomly chosen subpopulation elements. The rest of the algorithm is identical to Tanese's (1987) parallel GA. Although most of Tanese's results are concerned with population sizing (Forrest & Mitchell, 1991), she also looks at the diversity of alleles. Tanese finds that the GA with migration is more successful at preserving allelic diversity than the simple GA. Reducing migration prolongs population diversity, up to a limit.

The idea behind migration is to exchange individuals at a rate where they contribute good building blocks to the receiving subpopulation, without taking it over. Tanese defines the *migration rate* as the percentage of each subpopulation that is exchanged. She finds that a migration rate as low as 10% every 5 to 500 generations delays but does not prevent takeover. All bit-positions in the combined population (over all processors) eventually become fixed at the same values. Tanese concludes that the migration rate has to be very low (about 1%) for any hope of maintaining stable, but differing subpopulations. This is consistent with one of Sewall Wright's calculations (Provine, 1986), which states that strict isolation — on the order of one migrant every two generations — is necessary to maintain variation of an allele. Subsequent population biology studies, summarized by Grosso (1985), indicate that the magic

number is actually smaller — about one migrant every four generations. As migration rates reach exceedingly low values, parallel GAs increasingly resemble the partitioned GA. They also increasingly inherit the partitioned GA's problems, such as redundant search.

So far we have examined techniques in which migrants replace the least fit elements, and techniques in which migrants replace randomly chosen elements. Another option is to replace the most similar elements (Pettey, Leuze, & Grefenstette, 1987). The specific type of *migration replacement* is likely to have little impact on diversity, since migration replacement will be overshadowed by selection pressure within each subpopulation.

One study of isolation and migration (Grosso, 1985) concludes that different subpopulations can successfully converge to different optima. However, Grosso's study incorporates diploidy and heterozygote advantage in addition to isolation; it is not apparent which mechanism or combination of mechanisms is responsible for the maintenance of diversity.

The multiple-run techniques we have so far reviewed work by dividing a population into subpopulations, either temporally or geographically. We now explore a finer grained subdivision that requires a geography. The geography is often composed of a landscape of parallel processors. These geographic methods, called *distributed GAs*, are characterized by a small number (often one) of elements per geographic region or processor. By their nature, these techniques enforce local, geography based, mating and competitive restrictions, although the effects of such restrictions fade as diffusion occurs.

Mühlenbein's (1991) parallel, distributed GA (PGA) places individuals on a two-dimensional grid, one individual per grid element or node. PGA works by hillclimbing to local minima, then hopping to others via crossover. Each individual performs local hillclimbing for a period, then crosses with one of its neighbors. A single resulting offspring then performs hillclimbing. If the improved offspring is fitter than its parent, it replaces the parent; otherwise the parent remains. The entire cycle then repeats.

Isolation of grid elements is an important component of PGA. Mühlenbein uses a population structure called a ladder, which requires $O(P)$ steps for a solution to propagate to all P nodes. A ladder is a circular strip in which each grid element is connected to four others. Mühlenbein contrasts this with a torus, which requires only $O(\sqrt{P})$ steps for propagation. (Schwehm, 1992, empirically examines convergence times for various distributed-GA architectures.) Predecessors, extensions, and further empirical tests of the PGA are described in a number of papers

(Gorges-Schleuter, 1989, 1991; Mühlenbein, 1989; Mühlenbein, Gorges-Schleuter, & Krämer, 1988; Mühlenbein, Schomisch, & Born, 1991a, 1991b). On some problems, the authors report a solution better than all previously published solutions. Gorges-Schleuter (1989) shows that the PGA substantially slows convergence compared to the simple GA, but that populations converge given several hundred to several thousand generations.

Another distributed GA is the *fine grained genetic algorithm* (FG) (Manderick & Spiessens, 1989; Spiessens & Manderick, 1991), which also assigns one population element to each processor. FG limits each element to selection and crossover within its neighborhood of processors. FG works as follows. After initialization, each processor replaces its population element with an element selected from one of the neighboring processors, according to the fitness distribution of the neighborhood. (A processor is considered a neighbor of itself.) Tournament selection, where k is both the neighborhood size and the tournament size, is one possible selection method. Each processor crosses its selected element with an element randomly chosen from the processor's neighborhood. The processor then forms its element for the next generation by picking one of the two offspring produced by crossover, and mutating that offspring.

FG is a clustering algorithm, in that small clusters of identical individuals form, shrink, grow, and eventually merge to form bigger clusters. Edges of the clusters combine to form hybrid clusters that explore new portions of the search space. As far as diversity is concerned, "a cluster with the same fitness as its surroundings will eventually disappear" (Spiessens & Manderick, 1991). Stable regions of the grid do not emerge, especially if fitnesses of neighboring clusters are even slightly different. The authors note, however, that diversity remains for a relatively long time.

Collins and Jefferson (1991) further explore the maintenance of stable subpopulations by distributed GAs. The authors employ toroidal, one and two-dimensional grids with one individual per node. Each node selects two individuals for reproduction by performing two random walks. All random walks are of fixed length, yielding an effective neighborhood size. A random walk starts at the current node, then visits a randomly chosen neighbor, then visits a randomly chosen neighbor of that neighbor, and so on. The best element encountered along a random walk is selected.

On a bimodal problem, both one and two-dimensional grid schemes locate both global optima and maintain them to the stopping point of 1000 generations. (The length of the

random walk is 1.) Given sufficient numbers of processing elements, a nearly 50/50 mixture of global optima results. The optima are separated by boundaries of low-fitness hybrids or *lethal* solutions. Optimal clusters remain stable for thousands of generations, but use thousands of population elements to do so (n ranges from 8,192 to 524,288). The authors run a simple GA for comparison, with the usual result that it converges to one or the other optimum. Collins and Jefferson find that one-dimensional grids are more stable than two-dimensional grids. They also note that grid methods maintain good mixtures of all alleles while the simple GA becomes fixed at all positions. This latter result is likely problem specific, since the two global optima of their problem are bitwise complements.

Davidor (1991b) presents another grid scheme that maintains diversity. His *ECO GA* employs a two-dimensional grid with opposite edges connected to form a torus. Each grid element has nine neighbors, including itself. The ECO GA is implemented serially, processing a few elements at a time, with replacement. Parallel versions are also possible.

To begin, the grid is randomly initialized, one population member per node. Each iteration, a node is chosen at random, and two individuals are selected from its neighborhood, based upon fitness. These two individuals mutate and then cross. The two resulting offspring each choose a grid element, at random, within the neighborhood of their parents. Each offspring competes with the individual currently occupying the chosen grid element. This represents an additional selection stage (survival or replacement selection), in which the survival probability of each competitor is proportional to its relative fitness.

On a test function with peaks of differing heights, ECO GA delays, but does not prevent convergence to the global optimum. (Davidor uses an unspecified population size, in the thousands. Full convergence requires over 5000 generations.) Davidor, Yamada, and Nakano (1993) present experiments using the ECO GA for job-shop scheduling. They find solutions to seven difficult problems, which they claim are better than the best previously published solutions.

In summary, geographic GAs provide a method of slowing convergence, with effectiveness depending on the number and arrangement of processing elements. Increasing the isolation or the number of processing elements typically slows convergence, as does decreasing the neighborhood size. Although the addition of geography slows the GA's convergence from logarithmic to at best linear, the GA is not stabilized. On multimodal functions with peaks of differing fitnesses, representatives of the peak with the highest fitness are expected to eventually take over

the entire population. On multimodal functions with peaks of identical fitness, representatives of one peak will eventually eliminate representatives of the others. However, a quasi-stable equilibrium may occur in cases where population size is hundreds to thousands of times the number of peaks. The diversity of alleles diminishes over time, although the rate of loss is slower than in the simple GA. Geographic GAs are successful at delaying convergence in the search for a single, global optimum, typically yielding better final results than the simple GA.

Thermodynamic genetic algorithms

Several authors suggest using an analog to temperature in the *simulated annealing* algorithm, in order to maintain diversity and control convergence in the GA. A summary of such techniques is contained in Mahfoud and Goldberg's (1995) study. Simulated annealing (Aarts & Korst, 1989; Kirpatrick, Gelatt, & Vecchi, 1983; Metropolis, Rosenbluth, Rosenbluth, Teller, & Teller, 1953) is a one-string-at-a-time, probabilistic hillclimbing technique that allows the user to regulate the degree of convergence via a temperature parameter. At high temperature, transitions to less fit states occur with high probability. As temperature falls, transitions to less fit states become less likely. In simulated annealing, both slower cooling and more iteration yield better end-solutions.

Goldberg's (1990) paper on Boltzmann tournament selection is the first to explicitly outline the close correspondence between simulated annealing and genetic algorithms. Boltzmann tournament selection proceeds by holding three-way tournaments each generation to fill population slots. Tournament winners are selected according to logistic acceptance and anti-acceptance probabilities. Goldberg gives both empirical and analytical evidence that the distribution of population members over time is nearly Boltzmann. Although his algorithm successfully reduces selection pressure, selection noise causes diverse solutions to gradually disappear from the population (Mahfoud, 1993, 1994a). At infinite temperature, given an initial population in which all n elements are unique, Boltzmann tournament selection, acting alone, maintains approximately $.4\sqrt{n}$ unique population elements to Generation 2000 (Mahfoud, 1994a). At lower temperatures, selection pressure drives additional solutions out of the population. Once crossover and mutation are added, operator disruption also becomes a problem.

Sirag and Weisser (1987) combine several genetic operators into a "unified thermodynamic operator" to solve ordering problems such as the travelling salesman. The unified operator is

applied to two selected parents to yield a single offspring. At high temperatures, the offspring differs dramatically from either parent; at low temperatures, the offspring greatly resembles one or both parents. While the authors' approach addresses the problem of operator disruption, it does nothing to reduce selection pressure or selection noise. It is likely that fitness-proportionate selection will heavily bias the population in favor of the best individuals, depriving it of lower fitness points from which to search.

Mahfoud and Goldberg (1992, 1995) introduce a parallel, genetic version of simulated annealing called *parallel recombinative simulated annealing*. The algorithm retains the convergence properties and guarantees of simulated annealing, while adding the implicit and explicit parallelism of genetic algorithms. Parallel recombinative simulated annealing repeatedly generates a new population from the current population as follows. First, all elements of the current population are randomly paired to form $n/2$ couples. Each couple generates two offspring via crossover and mutation. The two offspring compete against their parents for inclusion in the new population, and winners are selected via probabilistic Boltzmann trials. A gradually decreasing temperature parameter regulates selection pressure in the Boltzmann trials. As in simulated annealing and Boltzmann tournament selection, the temporal population distribution is approximately Boltzmann.

Parallel recombinative simulated annealing handles the problem of operator disruption through the use of *probabilistic elitism*. In short, newly generated solutions that are worse than current solutions are accepted only with a certain probability, determined by temperature. High rates of crossover and mutation can hence be used without excessive risk of losing good solutions, except at high temperatures.

At low temperatures, given a problem with multiple, identically fit solutions, simulated annealing, if successful, will fluctuate within the neighborhood of one of those solutions for long periods of time. On rare occasions, simulated annealing will transit to the neighborhood of another solution, after which it will fluctuate about the new solution for an extended period of time, until moved by the next rare transition. One might hope that parallel, genetic implementations of simulated annealing would be able to explore multiple solutions within a single population. However, even the best of these genetic methods follows the lead of simulated annealing, and at lower temperatures, explores multiple solutions temporally rather than simultaneously. It should be possible, however, to incorporate a niching method into thermo-

dynamic GAs such as parallel recombinative simulated annealing. One candidate is a niching method developed in Chapter 6.

Restricted mating

A *species*, according to the biological definition, is a class of organisms that are capable of interbreeding amongst themselves, but that typically do not breed with individuals outside their class (Cook, 1991). This notion of reproductive isolation has led some GA researchers to attempt to induce speciation by restricting mating among dissimilar individuals. Such an approach directly attacks the problem of crossover disruption: if only like individuals cross, disruption is kept to a minimum. However, biological mating restrictions may be a *consequence* of speciation, or a reinforcing mechanism, rather than a primary cause. The mating-restriction approach, unfortunately, does nothing to combat the pressure or the noise of selection. If selection is allowed to pit different species against each other, only one species will survive.

Booker (1982, 1985) tests a "restricted mating policy" in a classifier system; only classifiers that match the same message are allowed to cross. If not enough fully matching classifiers are available, partially matching classifiers are allowed to cross. Although Booker claims populations become partitioned into subpopulations, his additional use of a sharing scheme in selection raises the likelihood that it is the combination of sharing and mating restriction, or sharing alone, rather than mating restriction alone, that is responsible for the observed diversity. We will examine sharing methods in detail later in this thesis.

Booker (1982) and Goldberg (1989c) explore various approaches in which a mating tag is added to each individual. This tag must match another individual in some number of bit-positions before a cross is permitted. Many variations exist, including one-way matching, two-way matching, and partial matching. More advanced methods add a template to each individual, and matches must occur between tags and templates instead of between tags and individuals. Tags and templates evolve, along with the rest of a string. Perry (1984) attempts speciation using an evolving species tag, in combination with a collage of other techniques. His results are inconclusive. Deb applies evolving species tags and templates to restrict mating in multimodal function optimization (Deb, 1989; Deb & Goldberg, 1989). Two individuals cross if their tags and templates match both ways. If they match only one way, individuals cross with probability .5.

Todd and Miller (1991) concatenate an evolving, mating-preference tag to the genotype of each individual. The tag decodes to a real number on the interval $[0, 1]$, that indicates an individual's "preferred mating distance". A triangular function determines the partial probability that a given individual crosses with a second individual. If the normalized, phenotypic distance between two individuals is equal to the first individual's preferred mating distance, this probability is 1; as the distance between two individuals moves away from the preferred distance, this probability shrinks. (*Phenotypic distance* is usually defined as the Euclidean distance between two phenotypes.) Of course, the mating preference of the second individual must also be taken into account. Once the two individuals' partial probabilities are computed, the probability that they actually cross is the product of their individual, partial probabilities.

Todd and Miller employ the following selection scheme. They select, with probabilities proportional to fitness, two parents. If the parents successfully cross, the offspring advance to the next population. If not, the second parent is replaced by another individual from the population, selected proportional to fitness. If the first parent can not find a mate after five tries, both parents are replaced. Selection proceeds until the new population is full. The authors find that trying several times to find a match for the first parent maintains more diversity than replacing both parents if they are incompatible.

Todd and Miller initialize all population elements to be within a certain radius of each other in phenotypic and in mating-preference space. They perform simulations on a flat fitness function, at a variety of parameter settings, using small mutation rates. Over the course of a few hundred generations, several clusters of phenotypically similar individuals with like mating preferences evolve, merge with other clusters, and split into subclusters. At times, two (or more) neighboring clusters prefer mates from the other cluster over mates from their own cluster. On nonflat fitness functions, selection pressure and noise become dominant factors, and "species [do] not generally form on the different peaks".

Gorges-Schleuter (1992) compares the local mating strategies of several distributed GAs (Collins & Jefferson, 1991; Gorges-Schleuter, 1989; Manderick & Spiessens, 1989) without selection or mutation. The GAs process an idealized genotype having one gene and two alleles. She finds that clusters of like individuals form in both one-dimensional and two-dimensional architectures, but more quickly in one-dimensional architectures, due to their added isolation. She also finds that smaller neighborhood sizes are better than larger ones at forming clusters.

Mating restrictions, by themselves, can not preserve significant population diversity. Limiting mating to like individuals may result in the formation of like offspring; however, selection forces species to compete, resulting in the elimination of less fit species and the eradication, due to noisy selection, of all but one identically fit species. This is verified by considering two idealized mating-restriction schemes.

The first scheme is developed in this thesis. If two elements that are paired for crossover are of differing species, the cross is not permitted, and the individuals proceed to the mutation stage. Our implementation uses a phenotypic distance threshold of .1 to decide whether two elements belong to the same species.

Our GA with mating restriction employs random initialization, binary tournament selection, $n = 100$, $p_c = 1$, and $p_m = 0$. It runs on sinusoidal function *M1* of Chapter 5, which has five peaks, spaced at intervals of .2. Without mating restriction, the GA fully converges by Generation 40 to a single global optimum from the five possibilities. (All 100 individuals become identical.) With mating restriction, the GA fully converges by Generation 30 to a single global optimum; mating restriction actually accelerates convergence! This is due to the reduced number of crosses.

The second mating-restriction scheme is Spears' (1994) "simple subpopulation scheme". The algorithm employs a k-bit tag for each individual, and disallows mating between individuals with differing tags. Since mating is restricted in this way, tags are never modified by crossover. Tags are also never modified by mutation. The tag bits effectively divide the population into 2^k subpopulations, but with selection allowed across subpopulations. Selection is hence the only operator that can reach across subpopulations. On sinusoidal functions *M1* and *M3* of Chapter 5, each having five peaks of equal height, the algorithm maintains a few of the peaks. However, given peaks of differing heights, all subpopulations converge to the highest peak. Spears' results are consistent with our own.

Since idealized mating-restriction schemes are not effective at forming and maintaining multiple solutions, no mating-restriction scheme, in isolation, will be. This includes tag based, tag/template, and similarity-based mating-restriction methods. Cross-species competition will cause the loss of all but one species, unless some sort of competitive restriction is additionally enforced. (We will examine competitive restrictions in Chapter 4.)

Mating restrictions may be useful for improving the on-line performance of GAs, by preventing the formation of lethal solutions. Lethal solutions are the low-fitness offspring resulting from crosses between elements of different species. Two studies show that restricting mating to similar individuals produces more consistent results across multiple runs, and improves average population fitnesses (Deb, 1989; Deb & Goldberg, 1989).

The opposite approach, to restrict mating between individuals that are too similar, has also been tried. The idea is to prevent fixation within classes rather than between classes. Booker (1987) proposes crossover among *reduced surrogates* — nonmatching alleles of the strings being crossed. If the strings differ in more than one bit, crossover is guaranteed to produce offspring different from their parents.

Likewise, the CHC algorithm (Eshelman, 1991; Eshelman & Schaffer, 1991) seeks to maximize exploration through the combination of uniform crossover across reduced surrogates, and mating between dissimilar individuals only (whose Hamming distance exceeds a certain threshold). High disruption is prevented by an elitist selection method. Through striving to produce children different from their parents, CHC achieves limited success in delaying convergence. CHC takes, on the average, 2.2 times longer to converge than the equivalent algorithm without restricted mating, and yields fitter end results on several test functions. However, convergence is slowed only to a certain extent. Stable subpopulations do not form within the population.

It is an open question whether crosses between species are desirable for potential jumps to higher peaks in the search space. Both interspecies and intraspecies crosses may be beneficial, depending upon the algorithm, the problem, and the user's objectives. Perhaps a balance exists between mating like and unlike individuals. Later in this study, we will demonstrate the benefits and potential drawbacks of interspecies crosses.

3.6   A Formal Framework

The term *diversity* is frequently used by GA researchers, with rarely a definition, except in limited contexts. This section develops a formal framework for the study of diversity. It formally defines diversity, and examines prior measures of diversity from the perspective of the new framework. Chapter 4 later specializes this framework to apply to niching methods in the

context of multimodal function optimization. Chapter 5 incorporates the specialized framework into a modelling framework for niching GAs.

Diversity is a general term that describes variation or lack of similarity among a collection of objects. The diversity of a collection of objects is defined with respect to one or more properties of the objects under consideration, such as color or size. If the collection of objects changes from time to time, then the diversity of the collection may also change.

Measures of population diversity can be useful in the analysis and design of GAs. We define a general diversity measure Z for a population pop at timestep or generation t with respect to three things: a descriptive relation R that maps population elements to their relevant descriptive properties; a partition of R's range, designated by equivalence relation X; and a set of one or more goal distributions Φ. Note that pop is a *multiset* of population elements, meaning that it may contain duplicates. R, X, and Φ are detailed shortly.

Diversity (Z) is a function of the above five variables; Z takes on real values in the range 0 to MAX, with 0 indicating no diversity, higher values indicating higher levels of diversity, and MAX indicating maximal diversity. MAX depends upon R, X, and Φ, but can sometimes be treated as a constant. In mathematical form,

$$Z : \langle pop, t, R, X, \Phi \rangle \rightarrow [0, MAX] \quad . \tag{3.2}$$

The change in diversity of a population from time i to time j is calculated as

$$\Delta Z \;=\; Z(pop_j, j, R, X, \Phi) - Z(pop_i, i, R, X, \Phi) \;, \tag{3.3}$$

where pop_i and pop_j are the populations at timesteps i and j, respectively. Since R, X, and Φ do not vary with time, diversity can be written more simply, as a function of four variables: $Z(pop_t, R, X, \Phi)$.

A *descriptive element* is the result of applying relation R to one or more population elements. A *descriptive population* is the result of applying R to all population elements or combinations of population elements. Both the original population pop and the descriptive population $R(pop)$ are multisets, since they may contain duplicates. Relation R may be one to one, many to one, one to many, or many to many. Diversity is not a direct function of the population at time t, but a function of the descriptive population produced by applying R to pop_t. Hence Z can

Table 3.7: Examples are given of descriptive elements, descriptive relations, and descriptive populations. Each descriptive relation maps a single element of the search space to the specified range. Each descriptive population contains the m descriptive elements specified in the table.

Descriptive Element	*Range of R*	*Descriptive Population*	m
genotype	all l-bit strings	$\forall_{i \in [1, n]} \, s_i$	$m = n$
phenotype	all k-tuples of variables	$\forall_{i \in [1, n]} \, \langle x_{i1}, x_{i2}, \ldots, x_{ik} \rangle$	$m = n$
bit	$\{0, 1\}$	$\forall_{i \in [1, n]} \, \forall_{j \in [1, l]} \, s_{ij}$	$m = ln$
schema	all 3^l schemata	$\forall_{i \in [1, n]} \, \forall_{j \in [1, 2^l]} \, schema_j \supseteq s_i$	$m = n2^l$
three leftmost bits	all three-bit strings	$\forall_{i \in [1, n]} \, \langle s_{i1}, s_{i2}, s_{i3} \rangle$	$m = n$
bit and position	all $\langle bit, position \rangle$ tuples	$\forall_{i \in [1, n]} \, \forall_{j \in [1, l]} \, \langle s_{ij}, j \rangle$	$m = ln$
pair of strings	all tuples of l-bit strings	$\forall_{i \in [1, n]} \, \forall_{j \in [1, n]} \, \langle s_i, s_j \rangle$	$m = n^2$

alternatively be written as

$$Z(R(pop_t), X, \Phi) \ . \tag{3.4}$$

Let n_t be the number of population elements contained in pop_t. Since population size does not vary from timestep to timestep in traditional GAs, n_t is written omitting the t, as n. The descriptive population $R(pop_t)$ contains some number of elements m_t (m_t in most cases is also constant from timestep to timestep and can thus be written as m); m may be greater than, equal to, or less than n.

Table 3.7 presents examples of descriptive relations and descriptive populations. The table designates the n elements of pop as $s_1, s_2, \ldots, s_n$. Since the traditional GA's population elements are binary strings of length l, the jth bit of s_i is denoted as s_{ij}. The table assumes that each binary string encodes k variables, where x_{ij} represents the jth variable of s_i.

Recall that a mathematical relation is defined as a subset of the cross product of two sets, a domain M and a range G. In Table 3.7, the domain for all descriptive relations R is the set of all 2^l possible binary strings. The range G is the set of all possible descriptive elements. G may contain either binary strings, phenotypes, single bits, schemata, or a variety of other entities. Elements of $R \subseteq M \times G$ are ordered pairs of the form (a, b), where $a \in M$ and $b \in G$. The descriptive population $R(pop)$ contains those elements of G produced by applying R to each element s_i of pop.

We have so far defined three components of a diversity measure, the population *pop*, the current timestep t, and the descriptive relation R. The fourth component is a partition of G into c *equivalence classes* (often simply called *classes*). This partition is represented by the equivalence relation X. Every possible descriptive element is a member of exactly one equivalence class. Likewise, every element of the descriptive population is a member of exactly one class. We are interested in the frequency distribution of descriptive-population elements among classes. More specifically, we are interested in the *relative frequency distribution* (of proportions) that results from dividing all frequencies of occurrence by m.

We can further simplify Z. Instead of defining it as the function $Z(R(pop_t), X, \Phi)$, we can define it directly over the relative frequency distribution of $R(pop_t)$ under X:

$$Z(pop, t, R, X, \Phi) \;=\; Z(P = \langle \frac{I_0}{m}, \frac{I_1}{m}, \ldots, \frac{I_{c-1}}{m} \rangle, \Phi) \;\;, \tag{3.5}$$

where I_i is the number of elements of $R(pop_t)$ that are members of equivalence class i, and the c equivalence classes are numbered from 0 to $c - 1$. Note that distribution P is a function of *pop*, t, R, and X. Also note that since $\sum_{i=0}^{c-1} I_i/m = 1$, and since for all i, $0 \le I_i/m \le 1$, P can in many cases be treated as a probability distribution.

We have so far defined a general diversity measure as a function Z over distributions P and Φ. P is the relative frequency distribution of descriptive population elements among equivalence classes at time t. Φ is a nonempty set of goal distributions. Specifically, Φ contains all "fully diverse" relative frequency distributions of descriptive population elements among classes. The definitions that follow are applicable to situations in which both P and Φ vary with time. However, we will treat all distributions as if they were stationary. To complete the definition of a general measure of diversity, we must first define a difference measure D between two distributions P and Q.

Let $Q = \langle q_0, q_1, \ldots, q_{c-1} \rangle$ be a goal distribution. Let $P = \langle p_0, p_1, \ldots, p_{c-1} \rangle$ be the actual distribution. Assume the following properties hold:

$$\sum_{i=0}^{c-1} p_i = 1 \;\; \text{and} \;\; \forall_i \; 0 \le p_i \le 1 \;\; ; \tag{3.6}$$

$$\sum_{i=0}^{c-1} q_i = 1 \ \text{ and } \ \forall_i \ 0 \le q_i \le 1 \ . \tag{3.7}$$

$D(P; Q)$ is the *directed divergence* (Kapur & Kesavan, 1987) or distance, of distribution P from distribution Q. Where particular p_i and q_i are irrelevant, but only the absolute differences between corresponding p_i and q_i are relevant, D can be rewritten as $D(d_0, d_1, \ldots, d_{c-1})$, where $d_i = |p_i - q_i|$ for all classes i. We require D to possess certain properties to qualify as a distance measure upon which we can base a definition of diversity. Additional, desirable properties exist that we do not require; their presence, however, makes D a better diversity measure. The required and desirable properties are listed below. Note that properties 1,2,6, and 7 qualify D as a metric.

1. Required: *nonnegativity*: $\forall_{P,Q} \ D(P; Q) \ge 0$;

2. Required: *zero property*: $\forall_{P,Q} \ D(P; Q) = 0 \iff P = Q$, where $P = Q \iff p_0 = q_0$ and $p_1 = q_1$ and $\ldots$ and $p_{c-1} = q_{c-1}$;

3. Required: *continuity*: D must be defined at all legal values of P and Q, and must be continuous on all legal intervals for the p_i and the q_i (*legal* values are defined as those which meet the requirements of (3.6) and (3.7)) ;

4. Required: *strictly increasing*: $\forall_{i \in [0, c-1]} \ d_i \ge d_i'$ and $d_j > d_j'$ for some $j \in [0, c-1] \implies D(d_0, d_1, \ldots, d_{c-1}) > D(d_0', d_1', \ldots, d_{c-1}')$;

5. Required: *maximum value*: Given a fixed Q, D has a maximum value of $MAX_P[D(P; Q)]$, where MAX_P indicates the maximum over all legal distributions P ;

6. Desirable: *symmetry*: $\forall_{P,Q} \ D(P; Q) = D(Q; P)$;

7. Desirable: *triangle inequality*: $\forall_{P,Q,Y} \ D(P; Q) + D(Q; Y) \ge D(P; Y)$;

8. Desirable: *positional invariance*: For all $k_0, k_1, \ldots, k_{c-1} \in \{d_0, d_1, \ldots, d_{c-1}\}$ such that no two k map to the same d, $D(d_0, d_1, \ldots, d_{c-1}) = D(k_0, k_1, \ldots, k_{c-1})$.

Our general diversity measure will require the user to specify one or more goal distributions that represent "fully diverse" descriptive populations. In most cases, a single goal distribution Q will suffice; Q will often be the uniform distribution, in which case diversity will also be a measure of entropy.

Given multiple goals, we consider proximity to any one goal to be desirable. If Φ represents the set of k goals, $\{Q_0, Q_1, \ldots, Q_{k-1}\}$, then

$$D(P; \Phi) \; = \; \min[\; D(P; Q_0)\;,\; D(P; Q_1)\;,\; \ldots\;,\; D(P; Q_{k-1})\;] \;\;. \tag{3.8}$$

One could perhaps devise scenarios in which proximity to one goal distribution would be more beneficial than proximity to another. In such a case, one might wish to utilize a function other than minimum distance, such as minimum weighted distance, to select a single goal distribution for computing D. Other forms of (3.8) are also possible. One option is to give higher credit for proximity to many distributions (i.e., to favor better average proximities).

We now complete the definition of a general measure of population diversity. Given a fixed Φ, $MAX_P[D(P; \Phi)]$ indicates the maximum $D(P; \Phi)$ over all legal distributions P. Hence,

$$Z(pop, t, R, X, \Phi) \; = \; Z(P^*, \Phi) \; = \; MAX_P[D(P; \Phi)] - D(P^*; \Phi) \;\;, \tag{3.9}$$

where P^* is the particular distribution whose diversity is being measured. The distance measure used to compute $D(P; \Phi)$ must meet the first five required conditions outlined above. Values of Z close to 0 indicate little to no diversity, while values near $MAX_P[D(P; \Phi)]$ indicate full or nearly full diversity. The general definition of population diversity subsumes previous definitions and measures. Several examples follow that illustrate previous measures of diversity, from the perspective of our diversity framework. We revisit this diversity framework in Chapter 4, in the context of niching methods for multimodal function optimization.

Example 1: Shannon's entropy measure

Shannon's entropy is derived from the Kullback-Leibler measure of directed divergence, which assumes Φ contains a single goal distribution Q.

$$D(P; Q) \; = \; \sum_{i=0}^{c-1} p_i \ln \frac{p_i}{q_i} \;\;. \qquad \text{(Kullback-Leibler)} \tag{3.10}$$

All entropy measures are based on distance from uniformity. Hence, Q is the uniform distribution, in which $q_0 = q_1 = \ldots = q_{c-1} = 1/c$. Shannon's entropy measure can be derived as follows. Let $P^* = \langle p_0, p_1, \ldots, p_{c-1} \rangle$.

$$
\begin{aligned}
Z(P^*, Q) &= MAX_P[D(P;Q)] - D(P^*;Q) \\
&= \ln c - \sum_{i=0}^{c-1} p_i \ln \frac{p_i}{\frac{1}{c}} \\
&= \ln c - \sum_{i=0}^{c-1} (p_i \ln p_i + p_i \ln c) \\
&= \ln c - \left(\sum_{i=0}^{c-1} p_i \ln p_i \right) - \ln c \sum_{i=0}^{c-1} p_i \\
&= \ln c - \left(\sum_{i=0}^{c-1} p_i \ln p_i \right) - \ln c \\
&= -\sum_{i=0}^{c-1} p_i \ln p_i \quad . \qquad \text{(Shannon's entropy measure)} \qquad (3.11)
\end{aligned}
$$

Kullback-Leibler does not yield a general measure of diversity since it is undefined if any of the q_i are 0. However, when Q is uniform or when all q_i are specifically disallowed from taking on zero values, Kullback-Leibler meets Conditions 1–5 and 8. (It is not symmetric and does not satisfy the triangle inequality.) In such cases, a measure of diversity such as Shannon's entropy measure can be based upon it. Other entropy measures, based on other suitable measures of directed divergence, can also function as diversity measures.

Example 2: Distance metrics and associated entropies

The following equation defines a family of distance metrics or norms:

$$
D(P;Q) = \sqrt[k]{\sum_{i=0}^{c-1} |p_i - q_i|^k} \quad . \qquad (0 < k \leq \infty) \qquad (3.12)
$$

Equation 3.12 with $k = 1$ is known as *city-block distance*; with $k = 2$, *Euclidean distance*. Equation 3.12 satisfies Conditions 1–8 for all $k \in (0, \infty]$. The corresponding diversity measure follows from Equation 3.9, given a fixed Q. Uniform Q with $k = 2$ produces a Euclidean entropy measure.

Example 3: Mauldin's uniqueness threshold

To obtain Mauldin's (1984) uniqueness threshold, we let the descriptive population consist of all nonredundant pairs of genotypes: $\forall_{i \in [1,n],\ j \in [1,n],\ i \neq j} \langle s_i, s_j \rangle \in R(pop_t)$. The descriptive

population size is $m = n(n-1)$. Mauldin defines a uniqueness threshold $THRESHOLD = \lceil \frac{k(g-t)}{g} \rceil$ where g is the total number of trials to be run, t is the number of the current trial, and k is the *uniqueness* factor — the number of unique bits initially required. The partition X creates two equivalence classes, A and B; X is defined as follows: $\langle s_i, s_j \rangle \in A$ if the Hamming distance between s_i and s_j is greater than $THRESHOLD$; $\langle s_i, s_j \rangle \in B$, otherwise. Φ contains a single goal distribution $Q = \langle 1, 0 \rangle$ in which all descriptive elements are in Class A. Mauldin requires each string in the population to differ in Hamming distance from all other strings by more than $THRESHOLD$. Therefore, the descriptive population's distribution is required to match Q exactly. There is no concept of intermediate diversity.

$$Z(P^*, \langle 1,0 \rangle) = \begin{cases} 1 \, , & \text{if } P^* = \langle 1,0 \rangle \, ; \\ 0 \, , & \text{otherwise} \, . \end{cases} \tag{3.13}$$

Mauldin's uniqueness threshold is not useful as a general measure of diversity, because the distance measure upon which it is based violates the "strictly increasing" condition (Condition 4).

Example 4: Allele frequencies

Diversity measures based on allele frequencies are common in both genetic algorithms and biological genetics. Let j be an arbitrary locus from 1 to l. To measure diversity at the jth locus, define $R(pop)$ to contain the jth bit of each population element ($m = n$). X partitions the descriptive-element space into two classes, corresponding to '0' and '1'. The most common notion of "fully diverse" is a single goal containing 50% zeros and 50% ones: $\Phi = \{Q\} = \{\langle .5, .5 \rangle\}$. One possibility for D is city-block distance: $D(P; \langle .5, .5 \rangle) = 2|.5 - p_1|$, in which case

$$Z(P^*, Q) = MAX_P[D(P; Q)] - D(P^*; Q) = 1 - 2|.5 - p_1| \, . \tag{3.14}$$

Another possibility is $D(P; Q) = 4(.5 - p_1)^2$ (Collins & Jefferson, 1991), in which case

$$Z(P^*, Q) = MAX_P[D(P; Q)] - D(P^*; Q) = 1 - 4(.5 - p_1)^2 \, . \tag{3.15}$$

Bitwise diversity measures (3.14) and (3.15) are based upon difference metrics that meet all conditions (1–8) for a general measure of population diversity. Collins and Jefferson (1991) take

57

the diversity measure (3.15) at each bit-position, and use the average over all bit-positions as a combined diversity measure for the population.

De Jong (1975) uses the number of missing alleles as a measure of lost diversity. The number of alleles remaining in the population can serve as a thresholding measure of combined diversity, using either formula (3.14) or (3.15) to compute Z at each bit. This works as follows. The number of alleles remaining starts at l. For each bit-position of the string, if Z at that position is 0, the number of alleles remaining is decremented.

Louis and Rawlins (1993) employ the average Hamming distance of the population as a diversity measure. They calculate the Hamming distance between all nonredundant combinations of two strings — $m = n(n-1)/2$ total pairs — and take an average. The descriptive population of our framework is composed of all m such pairs. The partition X creates $l+1$ equivalence classes, one corresponding to each possible Hamming distance between two strings. Many maximally diverse goal distributions are possible, depending upon how the user chooses to define "maximally diverse". If random uniform populations are to be considered maximally diverse, then the authors suggest that the average Hamming distance of a fully diverse population is $l/2$. Let the index of each equivalence class in Q_i correspond to the Hamming distance it represents. Φ contains all distributions such that $\sum_{j=0}^{l} j q_j = l/2$. Euclidean distance makes an appropriate distance measure.

Combinations of diversity measures are useful if one is interested in examining the diversity of more than one aspect of the population. One can either combine the different measures, or keep them separate, looking at population diversity from a Pareto-optimality standpoint.

Example 5: Boltzmann distribution

Let the descriptive population contain all energies or negated fitnesses of the population elements ($m = n$). Let Q be the Boltzmann distribution at a particular temperature T. X divides the space of all possible energies into as many classes as there energies ($E_0, E_1, \ldots, E_{c-1}$ in the discrete case).

$$q_{E_i} = \frac{e^{-E_i/T}}{\sum_{j=0}^{c-1} e^{-E_j/T}} \quad . \tag{3.16}$$

P^* is the actual distribution of population energies. $D(P; Q)$ is any distance measure that meets Conditions 1–5. Since the Boltzmann distribution does not contain any zero propor-

tions, Kullback-Leibler can be used as a distance measure. Often in thermodynamic GAs, distributions are measured temporally. In such cases, t would not represent a single generation but a sequence of generations, and *pop* would actually be an enlarged population that contained all elements of all populations over the sequence of generations in question.

Example 6: Chi-square difference

For sufficiently large m, an expression for the difference between distributions can be based on the approximately chi-square statistic,

$$D(P;Q) \;=\; \sum_{i=0}^{c-1} \frac{(mp_i - mq_i)^2}{mq_i} \;=\; \left(\sum_{i=0}^{c-1} \frac{mp_i^2}{q_i} \right) - m \;\;. \qquad (3.17)$$

Disallowed is $q_i = 0$. The chi-square test will tell whether the p_i depart significantly from the q_i (statistically speaking).

Deb (1989) divides the population into classes based on peak membership. Any individual that lies on a hill leading up to a peak, and that has a fitness within ϵ of that peak's maximal fitness, is considered a representative of that peak. No population element may represent more than one peak, and some population elements represent no peaks (are members of a "none of the above" class). In terms of the current framework, $R(pop_t)$ consists of $m = n$ genotypes. X partitions R into classes based on peak membership. The goal distribution Q allocates individuals to peaks proportional to peak fitness, such that $q_i = f_i / \sum_{j=0}^{c-2} f_j$, where f_i is the height of Peak i. The none-of-the-above class, $c - 1$, has no expected members: $q_{c-1} = 0$. To compute the difference between ideal and actual distributions, Deb uses the "chi-square-like" distance measure,

$$D(P;Q) \;=\; \sqrt{\sum_{i=0}^{c-1} \left(\frac{mp_i - mq_i}{\sigma_i} \right)^2} \;\;, \qquad (3.18)$$

where $\sigma_i^2 = mq_i(1 - q_i)$ is the variance in the expected number of individuals allocated to the ith peak. The calculation of σ_i^2 assumes the expected number of individuals allocated to a peak is binomially distributed. A special case is $\sigma_{c-1}^2 = \sum_{i=0}^{c-2} \sigma_i^2$. The corresponding diversity measure follows from Equation 3.9, given a fixed Q.

3.7   Useful Diversity

GAs that were uniformly randomly initialized and immediately stopped would exhibit, by most definitions, near-maximal diversity at each locus. Earlier, we presented diversification methods that are capable of reducing all three criteria of this chapter — selection noise, selection pressure, and operator disruption — to arbitrarily low levels. Doing so, however, would result in the GA performing little beneficial search. Goldberg, in discussing those GAs which rely primarily on mutation or mutation-like mechanisms for diversification, has mentioned the concept of *useful diversity*. Goldberg and Richardson (1987) state that "maintaining diversity for its own sake is not the issue. Instead we need to maintain appropriate diversity — diversity that in some way helps cause (or has helped cause) good strings". More generally, diversity is useful if it helps in achieving some purpose or goal. We have formally defined utility in terms of the prior section's goal distributions. The set of one or more maximally diverse goal distributions may vary from one type of problem to the next.

In the GA literature, diversity is recognized as playing two potential roles: allowing exploration of more of the search space in order to generate a better, single solution; and allowing exploration for multiple solutions. These roles can be restated in terms of the formation and maintenance of diverse subsolutions on the way to a single solution, versus the formation and maintenance of diverse solutions. The two roles are not entirely independent. In searching more of the space for a single solution, the GA may encounter multiple solutions. Likewise, an algorithm for maintaining multiple solutions will undoubtedly be applicable, in some form, to the maintenance of diversity on the way to a single solution.

Most GA researchers interested in diversity have focused on the diversity of individual alleles. However, this microscopic perspective can be misleading. Biases in the search space such as optima at Hamming cliffs or such as complementary dual optima may allow diversity at all bit-positions, but typically do not maintain a variety of subsolutions. We regard mechanisms that strive to maintain allelic diversity as promoters of useful diversity, if they consistently lead to better final solutions at fixed n, on moderate to difficult problems, compared to their simple GA counterparts. Many of the methods of this chapter qualify as promoters of useful diversity. These methods generally fail, however, when the goal is to locate multiple, final solutions, especially when those solutions are of differing fitnesses.

In contrast, we expect a method that is capable of forming and maintaining multiple, final solutions, to also be effective at forming and maintaining multiple subsolutions on the way to a single, final solution. Even in cases where the final solutions or interim subsolutions are of differing fitnesses, we expect the method to find and maintain them. These are the ideal characteristics of a full-fledged *niching method*. The litmus test for a niching method, therefore, will be whether it possesses the capability to find multiple, final solutions within a reasonable amount of time, and to maintain them for an extended period of time. A niching method must be effective on problems in which the solutions have different fitnesses, as well as problems in which the solutions have identical fitnesses. We will assume that any niching method is capable, perhaps with slight modification, of forming and maintaining multiple, useful subsolutions, when the goal is to find only a single, final solution. For any niching method that needs slight modification to be applied to the formation and maintenance of subsolutions, we have at our disposal many of the diversification techniques of this chapter.

We focus our attention, for the remainder of this thesis, on the formation and maintenance of multiple solutions within a population. There are many possible sets of final solutions in which a GA user might be interested. Some users, for instance, might be interested in finding the k best optima of a multimodal function. Others might be interested in finding any k optima, as long as those optima are sufficiently dissimilar. We will define several desirable sets of final solutions in Chapter 4, using the goal distributions of the present chapter.

Chapter 4

Niching Methods

Niching methods are techniques that promote the formation and maintenance of stable sub-populations in the GA. Niching methods can be applied to the formation and maintenance of interim subsolutions on the way to a single, final solution. They are traditionally viewed, however, in the context of forming and maintaining multiple, final solutions. We restrict our attention to this latter context for the remainder of this thesis.

A niching method must be able to form and maintain multiple, diverse, final solutions, whether these solutions are of identical fitness or of varying fitness. A niching method must be able to maintain these solutions for an exponential to infinite time period, with respect to population size. (The simple GA and the diversification techniques of Chapter 3 have logarithmic to linear maintenance times.)

Chapter 3 showed that reducing selection pressure, selection noise, and operator disruption does not typically result in a niching GA. What is required is not just slower selection, noise-reduced selection, or less disruptive operators, but a new type of algorithm — one that promotes diversity along useful dimensions of diversity, while allowing other dimensions to converge. With respect to the search space, convergence may occur to some degree within local regions, but diversity must prevail across the most prominent or fittest regions. Niching methods alter the selection algorithm to provide selection pressure within, but not across regions of the search space. The selection pressure within individual regions can be substantial, and still preserve niching properties.

We will restrict our attention in much of the remainder of this thesis to the general domain of application, multimodal function optimization. We will also map covering problems in classification to multimodal function optimization problems, and will briefly review niching methods in multiobjective function optimization, nonstationary function optimization, and ecological simulation. Niching methods should be applicable without modification to many useful domains, and with slight modification to many more.

Given multiple optima and a limited capacity to locate optima, the best niching GAs will prefer the highest optima. Some niching GAs will also prefer optima that are far away from other, partially located optima. The best niching methods will be able to locate the highest peaks in the presence of a large number of lower peaks, as well as in the presence of deception. In addition, the best niching methods will not be overly selective; they will retain the ability to form and maintain nonglobal as well as global optima. We will focus on maintenance rather than formation of optima in our analysis; we will assume that since the simple GA can form good solutions, this capability will transfer to niching GAs. If this capability does not transfer to a particular algorithm, the algorithm will not be considered a niching method.

The remainder of this chapter first presents answers to common questions about niching in genetic algorithms. It then reviews previous research into niching mechanisms, and outlines some previously unresearched niching mechanisms as well. The unresearched mechanisms are derived from ecological genetic analogy. That analogy to natural methods should yield additional mechanisms should not be surprising, when one considers that such analogy has advanced GAs to their current level. The chapter next extends the diversity framework of Chapter 3 to multimodal function optimization, our general domain of application. Finally, it chooses the specific niching methods to be examined in the remainder of this thesis.

4.1   Frequently Asked Questions

Newcomers to the world of niching methods often ask three major questions. The most basic question is "Why do niching?" We have repeatedly motivated the need for niching methods in the first three chapters. Niching is required if one is interested in finding multiple solutions to a problem. Broad areas of application include multimodal function optimization, machine learning and classification, multiobjective function optimization, simulation of complex systems,

and tracking of nonstationary environments. Niching is also useful for finding better, single solutions to hard problems; the intermediate formation and maintenance of diverse subsolutions is often critical to the solution of hard problems. We apply niching methods later in this thesis to the optimization of multimodal functions and to the classification of data.

A second question that newcomers often ask is "Why *maintain* niches?" Is not the *location* of niches sufficient, if one keeps track of them off-line? Maintenance is the critical issue in niching, because one can never be sure whether a new solution in the population represents a new niche or a previously located niche. If the best solutions are kept off-line, many are likely to be instances of the same optimum. There is no pressure, except that exerted by the niching method, to spread solutions over multiple optima. Another reason for maintaining niches is that the extended presence of one solution or subsolution in the population might be critical to the location of another.

If one is interested in intermediate subsolutions rather than entire solutions, these subsolutions typically arise and then disappear in the span of a few generations. One can not be sure, without exhaustive enumeration each generation, which subsolutions are present in the population. Furthermore, in applications such as classification, one is not interested in individual population elements, but in the population as a whole, cooperatively acting as a solution.

A third question that newcomers often ask is "Why not locate multiple solutions sequentially, by iterating the GA?" We have already examined the poor results of naive iteration — repeated convergence to the same optima, and discrimination against less fit optima. We will examine in depth, later in this thesis, the effects of a more sophisticated scheme and the problems that arise under such a scheme.

4.2   Previous Research

Niching methods generally can be classified along two dimensions of behavior. The first dimension is space versus time. We have so far discussed niching methods as spatial algorithms — ones that form and maintain subpopulations within the space of a single population. However, effective temporal methods may also exist — ones that form and maintain subpopulations over time. In this study, spatial niching methods will often be called *parallel* niching methods, since they conceptually develop niches in parallel, within a single population. Temporal niching

methods will often be called *sequential* niching methods, since they conceptually develop niches sequentially, over time. Note that our designation of parallel versus sequential is independent of the number of physical processors employed.

The second dimension is niching within single environments versus niching due to multiple environments. We can view the overall environment as single or multiple, either spatially or temporally. Many niching methods are able to form and maintain stable niches within a single environment. (The environment in multimodal function optimization consists of a single fitness function.) These are the more general niching methods, because often they are also applicable to situations with multiple environments. Other niching methods require multiple environments, either spatially (multiple fitness functions) or temporally (nonstationary fitness functions). In some cases, the overall environment varies over both time and space.

The second dimension of behavior is roughly equivalent to the sympatric versus allopatric speciation dichotomy of population ecology, if one allows time as well as geography to act as a barrier. *Sympatric* speciation refers to the differentiation of species that coexist geographically, but that evolve to exploit different resources or ecological niches within the same environment. *Allopatric* speciation refers to the differentiation of species due to geographic isolation. Of course, different geographic locations in the natural world possess different environments. Each environment causes a different species to evolve that is well adapted to the environment.

This thesis represents the first study to exhaustively collect prior research on niching methods and to examine it in a unified setting. Table 4.1 summarizes the broad categories, developed in the remainder of this section, for niching methods. The table lists the location of each category of niching method, with respect to the two dimensions of behavior. The niching methods in each quadrant of the table are similar in the types of problems to which they are applicable. For instance, temporal, single-environment techniques should be most applicable to multimodal function optimization. Spatial, single-environment techniques should be well suited for multimodal function optimization, multiobjective function optimization, and classification. Spatial, multiple-environment techniques should be proficient at multiobjective function optimization. Finally, temporal, multiple-environment techniques should be well suited for adaptive simulation.

Table 4.1: Niching methods are classified along two dimensions of behavior — temporal versus spatial niching, and niching within single environments versus niching due to multiple environments.

	Single Environments	*Multiple Environments*
Temporal	Sequential Location	Overspecification Ecological GAs
Spatial	Heterozygote Advantage Crowding Restricted Competition Fitness Sharing	Ecological GAs Immune Systems

Sequential location of niches

We previously mentioned problems with the naive approach of running the GA multiple times to locate multiple niches. Beasley, Bull, and Martin (1993) present a more sophisticated strategy that they call *sequential niching*. Their algorithm runs a traditional GA multiple times and maintains the best solution of each run off-line. At the end of each run, their algorithm depresses the fitness function at all points within a certain radius of the best solution. This change to the fitness function discourages future runs from revisiting the same area.

Sikora and Shaw (1994) implement a sequential niching technique for generating classification rules. Upon convergence of the GA, their technique retains the best rule off-line and removes, from the examples set, all examples that the rule covers. The GA then restarts on the reduced problem. This cycle continues until all examples have been covered.

Of the single-environment niching methods, sequential niching is the only temporal method that has yet been developed. Some would argue, from a traditional standpoint, that sequential niching is not truly niching, since it does not form subpopulations within the same population. Others would argue that in single environments, temporal niching methods have the potential to be at least as effective as spatial niching methods. We will examine in depth, in Chapter 10, the differences between temporal and spatial niching methods for single environments, and whether they both achieve the same end objectives. Specifically, we will examine and test Beasley et al.'s sequential niching, and will use it as a base for comparing parallel niching methods.

Overspecification

A large percentage of the genome of higher organisms has no known function. This fact has prompted some GA researchers to examine the potential role of excess genetic material in diversification and niching. The resulting niching methods maintain excess genetic material in quantities sufficient to track a single optimum of an environment that changes over time. Overspecification, in combination with a changing environment, produces a temporal, multiple-environment niching method. Some researchers hope to adapt overspecification techniques to the location of multiple optima in stationary environments, but no one has yet demonstrated a technique with this potential. Given a stationary environment and sufficient time, all diversity disappears from both the expressed and the redundant genetic material.

Several genetic algorithms utilize some form of overspecification or redundancy in their representations. A biologically motivated approach is the use of multiple strands for the chromosome, otherwise known as *diploidy* (two strands) or *polyploidy* (many strands). Several early studies in the field of genetic algorithms employ diploid chromosomal structures. The driving motivation of most of these studies is to simulate biological genetic systems; problem solving is secondary. These early studies, reviewed by Smith (1988), show no advantage for diploidy in optimizing static functions. One exception is Grosso's (1985) diploid GA, which successfully maintains multiple peaks of a static function. However, the additional mechanisms of isolation, migration, and heterozygote advantage are undoubtedly at least partially responsible for his GA's success.

Goldberg and Smith (1987) propose that while GAs with diploidy may have no advantage in static environments, they should have a significant advantage in changing environments. The diploid chromosome should be able to store diversity that was useful in the past and that may again be useful in the future. After all, GAs were originally conceived by Holland (1975) as adaptive systems. Goldberg and Smith successfully apply diploid GAs to tracking a global optimum that alternates, at regular intervals, between two values (Goldberg & Smith, 1987; Smith, 1988; Smith & Goldberg, 1992). The population maintains, within the chromosomes of its members, both locations for the optimum. However, only one location at a time is expressed.

Expression occurs via a *dominance* map. Each locus of a diploid chromosome contains two alleles, but only one allele at a time is expressed, meaning that it is used to compute the

phenotype and subsequently the function value. If one type of allele (either '0' or '1' in binary encodings) is dominant at each locus, then when both types are present (the *heterozygous* case), the dominant one is expressed. When the locus contains two of the same allele (the *homozygous* case), that allele is expressed. The dominance map for each individual designates which type of allele is dominant at each locus. The dominance map may evolve, along with the rest of the genotype. Goldberg and Smith employ a compact triallelic encoding (Holland, 1992; Hollstien, 1971) of the diploid chromosome and its dominance map.

When the environment is constantly changing, previously useful solutions become stored in the unexpressed bits of the population's diploid chromosomes. Should the environment later reach a state similar to a previous state, stored solutions may return. Note that the population as a whole stores away useful solutions; these solutions must be reconstructed using pieces contained within various individuals. Currently unexpressed solutions are brought forward through evolution of the dominance map (via selection, crossover, and mutation).

Goldberg and Smith restrict their attention to a function that oscillates between two optima. It is not clear whether diploid chromosomes would be sufficient for tracking more than two optima, or whether polyploid chromosomes would be required. It is also not clear whether diploidy would be of help in tracking environments that move to novel states. Perhaps some degree of similarity would have to exist between the new environment and one of the old ones. Preliminary application to novel environments appears promising, as shown by Hillis (1990). (We review Hillis's work in the next section.) Preliminary application to dual-criteria optimization problems also appears promising (Kursawe, 1991). Given a stationary, single environment, however, most evidence points to the conclusion that a diploid GA's population will fully converge, although in time slower than the haploid GA's (Goldberg & Smith, 1987; F. Greene, 1994).

It is well known that redundancy, whether in computer science or in living organisms, protects against the loss of information. We have already examined the benefits of redundancy through diploidy. An alternative to diploidy would be to keep redundant, possibly conflicting information on an extra segment of the haploid chromosome. Goldberg (1989c) notes the possibility of using variable-length strings in combination with the duplication and deletion operators. To resolve conflicts, some kind of *intrachromosomal dominance* procedure would be necessary.

One GA that allows overspecification within haploid chromosomes is the *messy GA* (Deb, 1991; Goldberg, Deb, Kargupta, & Harik, 1993; Goldberg, Deb, & Korb, 1990; Goldberg, Korb, & Deb, 1989; Merkle & Lamont, 1993). The messy GA incorporates variable-length strings, and handles overspecification using positional precedence. That is, if a gene resides on more than one locus, the leftmost locus takes precedence. After a chromosome is altered by genetic operators, previously hidden genes may gain precedence. Goldberg, Korb, and Deb (1989) note that other precedence schemes are possible, such as adaptive precedence. No one has yet applied messy GAs to nonstationary functions.

The *structured GA* (Dasgupta & McGregor, 1992) redundantly encodes an individual's entire phenotype or parts of its phenotype, a prespecified number of times. Activation bits are responsible for activating and deactivating segments of encoding bits. Activation bits may be arranged hierarchically so that they activate and deactivate each other. Redundant bits are never simultaneously active. Dasgupta and McGregor restrict their attention to two-level hierarchies, consisting of activation bits and encoding bits. Using the same problem as Goldberg and Smith, the authors show that the structured GA is effective at tracking two alternating optima.

Ecological genetic algorithms

Some mechanisms induce niching by utilizing and often creating multiple environments. We call genetic algorithms that employ such niching mechanisms *ecological GAs*. Unlike the overspecification GAs of the previous section, ecological GAs interact to a greater extent with their environments, and sometimes modify their environments. The category of ecological GAs includes methods such as symbiotic and parasitic coevolution, resource modelling, fitness function decomposition, and full-scale ecological simulation.

The multiple environments required by an ecological GA may occur over time or over space (or both). An example of multiple, temporal environments is nonstationary fitness functions. Examples of multiple, spatial environments include multiobjective fitness functions and multiple fitness functions residing in different geographic regions.

An ecological GA that is similar to the diploid GAs of the prior section, but that modifies the fitness function over time, is Hillis's (1990) model of "co-evolving parasites". Hillis's GA simultaneously coevolves a population of sorters and a population of sets of test cases for the

sorters. The sorters and the sets define fitness functions for each other that evolve over time. The sorters are assigned fitnesses based on how well they perform on a set of test cases. The sets of 10–20 test cases are assigned fitnesses based on how poorly a sorter performs on them (hence the parasitism analogy). A grid is employed in which every grid element contains a sorter and a set of test cases. Both sorters and sets are represented by diploid structures, and are operated upon, in parallel, by two separate genetic algorithms.

Hillis states that increasingly complex test cases evolve in order to fool increasingly sophisticated sorters. Both populations remain diverse, even after several hundred thousand generations. We attribute this niching effect to the multiple, temporal environments created by the evolving fitness functions, in combination with the tracking ability of the diploid chromosomes. The geography is a secondary diversification mechanism. Hillis's simulation demonstrates the ability of diploid GAs to adapt to gradually changing environments. Adaptation of diploid GAs to drastically changing environments (with more than two states) remains an open area of research.

In Hillis's GA, population elements (sorters) coevolve with fitness functions (sets of test cases), but do not directly contribute to other population elements' fitnesses. A more general type of coevolution is one in which any population element can directly contribute to the fitness of any other. In some coevolutionary models, the entire fitness of an individual is determined through its interaction with others; the fitness function becomes spatially distributed, throughout the population, as well as temporally distributed, as other individuals evolve. Some coevolutionary models additionally make use of a geography.

One coevolutionary approach in which individuals derive fitness directly from each other is *Game World* (Adachi & Matsuo, 1991). Game World is a two-dimensional geographic system in which multiple strategies to a particular game move about and compete with their neighbors. The fitness of an individual is determined by the number, type, and location of other individuals. Adachi and Matsuo test 64 possible strategies to the multiplayer, iterated prisoner's dilemma. The Game World lattice initially contains randomly generated strategies, with some sites left blank (filled with a null strategy). The population is allowed to grow and shrink. One generation consists, to start, of each individual playing a series of games with its neighbors, and accumulating a score or fitness. Then selection, mutation, and, optionally, crossover occur within a local neighborhood. Various dynamic behaviors and quasi-steady states evolve. The

authors note that both competition and cooperation develop, and often two or more species (different bit strings or strategies) coexist in a quasi-equilibrium. Infrequently, the system shifts to another, possibly much different quasi-equilibrium.

Sannier and Goodman (1987) describe the *Asgard* system, in which an individual's fitness is implicit in its ability to utilize the food resources of a two-dimensional terrain. Individuals that are incapable of properly utilizing resources die. The environment varies over both space (terrain) and time in its distribution of food resources, but individuals that are geographically close tend to experience the same environment. The terrain is a 160 x 60, toroidal grid, divided into four quadrants of equal size. Each quadrant contains two "farms", surrounded by "desert". Food appears only within the farms. Each farm experiences two seasons — summer and winter. Seasons are timed so that it tends to be winter in one of a quadrant's farms while it is summer in the other. Food production within each farm depends upon consumption and land maintenance during the prior timestep, and also upon the current season. Consumption and land maintenance are determined solely by the number of individuals occupying a particular region. Potential food production reaches a high point during a region's summer season and a low point during its winter season.

Asgard works as follows. Individuals may replicate, or cross with others that are geographically close. Mutation and inversion operators apply to both types of reproduction. Offspring are initially placed near their parents on the grid. Each timestep, an individual burns one unit of its stored energy to maintain itself, and additional energy if it reproduces. Individuals with no energy left die. The overall population fluctuates in size from timestep to timestep.

Individuals consist of lists of instructions. Only two instructions exist: *Move-x* and *Food-x-y*, where x is one of eight neighboring locations, and y is the address of an instruction. The *Move* instruction moves an individual in one of eight directions. The *Food* instruction tests location x for food, and jumps to instruction y if the test is successful. An individual consumes the food at a location upon moving to that location.

Sannier and Goodman initialize randomly chosen locations of the grid with 1000 randomly generated elements. The authors describe Asgard's behavior in the initial generations as mostly random walks, and death for all individuals that do not locate and stay within a farm. After 1500 generations, eight distinct classes of individuals emerge. Some classes are all contained within the same farm and some are spread over multiple farms. Most individuals either repeatedly

test for food, or move in one direction until locating food. After 3000 generations, two classes of individuals evolve, which the authors dub *farmers* and *nomads*. Farmers tend to move in circular patterns about each farm. Nomads cycle between the two farms of a quadrant, encountering and remaining at each farm during its summer season: "they strike the left edge of a farm during the farm's 'spring' and move across it during its most fertile period, leaving the rightmost edge of the farm as its productivity wanes."

Sannier and Goodman test the farmers by removing all nomads from the terrain. In isolation, the farmers exhibit highly efficient group consumption during the winter, but consumption only half as efficient during the summer. In the presence of nomads, however, consumption is highly efficient during both seasons. Nomads in isolation do not survive, because they can not stay at a farm long enough to increase its efficiency.

Sannier and Goodman further report that before Generation 4000, crosses of farmers and nomads yield weak individuals. After Generation 4000, however, a "composite" genotype arises, and eventually takes over the entire population. The composite genotype's early success in locating food throws it into an infinite loop, effectively making it a farmer if it finds food soon after birth, or a nomad otherwise. The authors note that Asgard is rather sensitive to parameter settings, and that its application to more "useful" domains remains elusive.

Holland (1992) outlines a model of complex adaptive systems he calls *Echo*. Echo consists of a geography, divided into localities or sites. Each site produces various resources that may differ from timestep to timestep and from site to site. Individuals or *agents*, at each timestep, occupy a certain site, can consume the resources of that site, and can interact through combat, trading, or mating. Any interaction may result in the transfer of resources from one individual to another. For instance, the loser in a combat typically must turn over all resources to the winner. An Echo user can designate rules to regulate the system and to govern interactions between individuals. Echo does not enforce an explicit fitness function, but kills off individuals that use up all of their resources. Agents that better utilize environmental resources tend to survive.

Echo iterates the following loop at each site, and synchronizes the iteration across all sites. First, Echo randomly selects two agents at each site to interact via combat, trading, or mating. After interaction, the agents consume some of the site's resources. Agents must then pay a "maintenance cost"; if an agent can not afford the cost, it dies. If an agent has accumulated

enough resources of the right kinds, it replicates, possibly with mutation. (Offspring can be produced via replication or mating.) If an agent has not found any resources at the current site, the agent migrates to a neighboring site. Each site then replenishes its resources. Holland notes that complex behaviors, such as cooperative communities of specialist agents, emerge from the model.

Ecological simulations such as Sannier and Goodman's and Holland's are representative of numerous studies in the growing field of *artificial life* (Langton, 1989; Langton, Taylor, Farmer, & Rasmussen, 1992; Varela & Bourgine, 1992). Holland says that currently the main purpose of models such as Echo is to conduct "thought experiments".

One practical ecological GA is Husbands and Mill's (1991) parallel GA for job-shop scheduling. In job-shop scheduling, x jobs, composed of sequences of steps for manufacturing a component, are scheduled for execution on k machines. The jobs may each require up to k of the machines. Job-shop scheduling finds the sequence of jobs, for each machine, that minimizes one or more objectives such as time. The authors obtain good results simulating 2–4 jobs on up to 11 machines.

Husbands and Mill's GA employs $x + 1$ independent subpopulations whose elements derive fitnesses both locally and from interactions with other subpopulations. Each subpopulation element represents a plan for executing a job. Each plan specifies an ordering of steps. Each step involves one or more machines and their setups, and independent steps may execute in parallel.

Each subpopulation is assigned one of the x jobs, and therefore searches for a near optimal plan for executing that job. Internal GA representations of plans may differ from subpopulation to subpopulation. A unique subpopulation of arbitrators resolves conflicts between other subpopulations. An arbitrator is a precedence vector, specifying which plan gets precedence. Since representations may differ across subpopulations, no migration is allowed between subpopulations.

The overall algorithm consists of one GA running on each subpopulation. All subpopulations synchronize their fitness assignment. To start, each individual receives a preliminary fitness based upon the local criteria of its subpopulation. Next, each subpopulation ranks its elements. An individual's final fitness is a function of its preliminary fitness and of its ability to utilize shared resources when interacting with individuals from other subpopulations. Equally ranked

members across all subpopulations are tested together, with conflicts resolved by the equally ranked arbitrator. Such testing simulates the simultaneous execution of x plans. Arbitrators receive fitnesses according to how well they resolve conflicts. For instance, an arbitrator that gives precedence to a resource-hungry job will receive a low fitness.

We now turn our attention to methods that partition a static fitness function into multiple, spatial environments. Preliminary insight into such a partitioning comes from ecological niche theory. In one common ecological framework, an environment is composed of multiple *resources*. Types of resources define axes of a multidimensional space of possible environments (Perry, 1984; Shorrocks, 1979). A niche in this framework is any subspace of the overall environmental space. Niches may be either disjoint or overlapping.

Some ecologists conjecture that a one-to-one mapping exists between phenotypic variables and environmental resources (Shorrocks, 1979). For instance, the phenotypic variable, temperature tolerance, corresponds to the environmental resource, temperature. Similarly, beak depth in birds corresponds to the size of food in the environment. Phenotypic variables are thus an indication of an individual's ability to utilize an environmental resource. A problem with the one-to-one mapping approach is the subjectivity required in labelling axes. One must bias both phenotypic and environmental variables to bring them into a one-to-one correspondence.

The subspace of possible environments amenable to a particular species is called that species' *fundamental niche*. The portion of the real environment that the species actually occupies is called its *realized niche*. A species' realized niche may expand or contract, depending upon competitive forces and environmental fluctuations. The principle of competitive exclusion states that if two species occupy the same niche, one species will eventually disappear. To coexist, two species must occupy different niches. If two species occupy overlapping niches, either the two species will eventually differentiate or one will become extinct.

At first glance, the above knowledge of environments, niches, and resources does not appear to be useful in the standard GA. Environments, niches, and resources are implicitly coded into the GA's fitness function. The main difficulty is in extracting the resources from the fitness function. However, in cases where a fitness function is decomposable into independent criteria, each criterion is analogous to a separate resource.

An early incorporation of principles from ecological niche theory into the GA is Perry's (1984) dissertation. Perry responds to the lack of resources by inventing some. He introduces

objects called *external schemata* that are arbitrary specifications of fitter portions of the search space. The external schemata represent criteria used to determine the fitnesses of strings. Utilizing such artificially constructed fitness functions and a plethora of potential diversification mechanisms, Perry tries to induce speciation. His results are inconclusive.

A more direct approach (Elo, 1994) is the *dynamic division* algorithm, which forms an increasing number of subpopulations as it progresses, through splitting current subpopulations in half. Each subpopulation is independent: there is no migration or other interaction between existing subpopulations. Subpopulations that are sufficiently large and sufficiently diverse undergo division. During division, the fittest individual in the subpopulation serves as a seed for one resulting subpopulation, as does the fittest individual that is sufficiently distant from the first. The remaining individuals are assigned to the post-division subpopulation whose seed individual they better resemble. Dynamic division achieves an approximate subdivision of the fitness function, into as many categories as there are subpopulations. The algorithm requires an overall population size of approximately 100 times the number of peaks, to locate on the order of 100 peaks. However, this multiple appears to increase polynomially as the number of peaks increases. The author does not address problem difficulty along dimensions other than multimodality. Elketroussi and Fan (1993) define a more general algorithm that merges similar subpopulations in addition to splitting diverse subpopulations.

Two other GAs deserve mention. One performs *multiobjective* (or *multicriteria*) function optimization; the other, machine learning. In both GAs, the fitness function is divided into independent components, and different population elements optimize each of the components.

In multiobjective function optimization, the goal is to locate solutions that simultaneously optimize a set of objectives, without combining all objectives into a weighted sum or other function. The fitness of each individual is hence a vector, containing one vector element per objective function. The goal can be restated as finding individuals that are members of the Pareto-optimal set. A *Pareto-optimal* individual is one that is not dominated by any other individual. An individual *dominates* another if it is at least equally fit on all dimensions, and more fit on at least one dimension.

The first system (Schaffer, 1984, 1985) is called *VEGA* (short for "vector evaluated genetic algorithm"). VEGA works by dividing the population, each generation, into as many subpopulations as there are objectives, and allocating one objective to each subpopulation.

Each objective function assigns fitnesses to all of the individuals in its subpopulation. After fitness assignment, VEGA shuffles the entire population, and applies selection, crossover, and mutation. Since VEGA randomly assigns individuals to subpopulations each generation, an individual is likely to be evaluated with respect to different objective functions from one generation to the next. This will result in an averaging effect. "Analysis of VEGA shows that [its] effect is the same as if fitness were a linear combination of the attributes" (Richardson, Palmer, Liepins, & Hilliard, 1989). Despite the averaging effect, Schaffer reports some success in finding Pareto-optimal solutions.

The second system (D. P. Greene & Smith, 1993, 1994), called COGIN (short for "coverage-based genetic induction"), is for machine learning. COGIN assigns training examples to rules that cover them, with fitter rules receiving examples covered by more than one rule. Rules are ranked according to the number of training examples they cover, as well as the accuracy of the coverage. COGIN employs a selection process called *coverage-based filtration*, which repeatedly selects the best rule from the population, deletes the training examples it covers, and reranks the remaining rules. This continues until no more training examples remain. All unselected rules die, resulting in a variable sized population. Since fitness is based upon the number of training examples a rule covers, COGIN effectively creates multiple environments by partitioning the set of training examples among surviving rules. Each rule is adapted to a unique environment — the set of training examples that the rule covers. Although the set of training examples remains static, the partition boundaries vary from generation to generation.

Heterozygote advantage

We now focus upon spatial methods that form and maintain niches within a single, static environment. These methods are directly applicable to multimodal function optimization problems, as well as many other types of problems. Before exploring established methods, one theory of ecology that warrants mention is *heterozygote advantage*.

Recall that in diploid chromosomes, each locus contains a pair of alleles. Recall also that if the alleles do not match, an individual is said to be heterozygous at that locus. Under systems of complete dominance, the dominant allele will, in theory, eventually drive the recessive allele to extinction (Shorrocks, 1979). However, if heterozygotes are more fit than homozygotes (i.e., they have an advantage), the population will maintain both alleles in stable proportions.

Some biologists feel that heterozygote advantage is a major factor in the diversity of natural populations.

Grosso (1985) maintains diversity across separate peaks of the search space using diploidy, heterozygote advantage, subpopulations, and migration. It is not clear which combination of methods leads to the apparently successful niching, but a prime candidate would be the combination of diploidy and heterozygote advantage. Grosso implements heterozygote advantage by directly assigning different fitnesses to different combinations of alleles at a locus.

While it may seem that heterozygote advantage is not applicable to haploid populations, this is not the case. Haploid individuals become diploid at one point every generation — during crossover. We therefore propose the following mechanism for heterozygote advantage in haploid populations. Perform crossover prior to selection, using random pairings. Elevate the fitnesses of the two offspring by a function of the distance between their parents; this function should grow as distance increases. An alternative to directly elevating fitnesses would be to use the distance function as a secondary fitness criterion. The method of heterozygote advantage deserves further exploration in both haploid and diploid GA populations.

Crowding: Restricted replacement

In *generational GAs* such as the simple GA, the entire population is replaced every generation by offspring formed through crossover and mutation (except for elements lucky enough to pass through unaltered by selection, crossover, and mutation). *Steady-state GAs* (Syswerda, 1989, 1991), on the other hand, process only a few individuals a time, and insert the resulting offspring into the population. Some techniques strive to preserve diversity in this replacement step. Instead of practicing wholesale generational replacement, or replacement selection (replacing the worst), they attempt to replace population members in a way that maintains diversity. We call methods that insert new elements into a population by replacing similar elements, *crowding methods.*

De Jong (1975) presents an algorithm he calls the "crowding factor model". De Jong's crowding is a steady-state GA, since only a fraction of the population reproduces and dies each generation. Each newly generated population member replaces an existing member, preferably the most similar one. To accomplish an approximation of most-similar replacement, a small

77

sample is taken from the existing population, and the new member replaces the closest element of the sample.

Specifically, De Jong's crowding works as follows. A proportion of the population, specified by the *generation gap* (GG), is chosen via fitness-proportionate selection to undergo crossover and mutation. After crossover and mutation, $GG \times n$ individuals from the population are chosen to die (to be replaced by the new offspring). Each offspring finds the element it replaces by taking a random sample of CF individuals from the population, where CF is called the *crowding factor*. The offspring replaces the most similar individual from the sample. Similarity is defined using bitwise matching in Hamming space, otherwise known as *genotypic comparison*.

Crowding is inspired by a corresponding ecological phenomenon — the competition, among similar members of a natural population, for limited resources. That is, similar individuals in a natural population, often of the same species, tend to occupy the same environmental niches, and therefore must compete against each other for resources. Dissimilar individuals, often of differing species, tend to occupy different niches, so they typically do not compete for resources. When a niche has reached its carrying capacity, weaker members of that niche are crowded out of the population by stronger members. The end result is that new members of a particular species replace old members of that species. The overall number of members of each species does not change unless the environment changes.

Crowding does not model the method by which a population arrives at a stable mixture of species, but instead strives to maintain the diversity of the preexisting mixture. De Jong's original goal in designing crowding was to maintain the diversity of alleles in order to prevent premature convergence. The algorithm was successful to some extent at maintaining allelic diversity, but stochastic errors introduced by low CF and by other factors forced the algorithm to gradually drift toward fixation at each bit-position. De Jong's crowding has empirically been shown to be of limited use in multimodal function optimization (Deb, 1989; Deb & Goldberg, 1989), where the selective preservation of useful diversity is essential. As demonstrated in Chapter 6, however, this verdict is premature for crowding methods, in general.

Five years prior to De Jong's work, Cavicchio (1970) introduced several of what he called *preselection* schemes, and claimed that one of them was successful at preserving population diversity. Like De Jong's crowding, preselection schemes were intended to prevent premature convergence. Cavicchio stated that performing comparisons for each new offspring would yield

too expensive a replacement strategy. Instead, since an offspring is usually similar to its parents, an offspring could simply replace one of its parents. The validity of this approach is demonstrated in Chapter 6.

The best preselection scheme works as follows: if a child has higher fitness than the worse parent, it replaces that parent. Unfortunately, the effect of this scheme is not evident in Cavicchio's work, since many other factors, including self-modifying parameters and up-front selection, are also present. If one tries to optimize simple multimodal functions using preselection, one will generally not be successful at preserving representatives of more than one optimum.

Prior to 1992, preselection had largely remained untouched except in passing mention in a few studies, and De Jong's crowding had seen only sporadic application, mostly in classifier systems (Goldberg, 1983; Holland & Reitman, 1978). Slight variations of De Jong's crowding, with additional selection pressure in the replacement step, had been implemented in two studies. Stadnyk's (1987) closest-of-the-worst replacement strategy was one. This strategy selected CF candidates from the population inversely proportional to fitness, and replaced the candidate closest to the new element. Sedbrook, Wright, and Wright's (1991) closest-of-the-worst replacement strategy replaced the closest individual in the bottom third of the population. (The population was sorted in order of ascending fitness.) Unfortunately, no study had been able to modify crowding so that it was able to optimize multimodal functions or otherwise perform effective niching.

In 1992, through analysis of De Jong's crowding and Cavicchio's preselection, Mahfoud was able to determine why neither method was successful, and to develop design modifications that produced a successful crowding algorithm. The resulting algorithm, *deterministic crowding* (Mahfoud, 1992, 1994b), exhibited extensive niching capabilities when applied to both multimodal optimization problems and classification problems. Deterministic crowding has since been successfully applied to the optimization and design of statistical, chemical quality-control methods (Hatjimihail, 1993), as well as to the spin-glass problem (Pál, 1994). Chapters 6, 7, and 10 review the development of deterministic crowding, provide extensive analysis, and test the algorithm on several problems.

Deterministic crowding works as follows. First, it groups all population elements into $n/2$ pairs. Then it crosses all pairs and optionally mutates the offspring. Each offspring competes

in a tournament against one of the parents that produced it. Given a pair of parents and their two offspring, two sets of parent-child tournaments are possible: Parent 1 against Child 1 and Parent 2 against Child 2; or Parent 1 against Child 2 and Parent 2 against Child 1. The set of tournaments that forces the closest competitions is held, where closeness is the average distance between parent-child couples in a set. Closeness is computed according to some appropriate distance measure, preferably phenotypic distance.

Since the introduction of deterministic crowding, two new but similar crowding methods have been proposed and shown to be preliminarily successful at optimizing multimodal functions. The two methods employ worst-of-the-closest replacement strategies. Like deterministic crowding, they do no selection up front, but only upon replacement. A third method has also been proposed which is not a crowding method as it stands, but which can potentially be modified to emulate the behavior of deterministic crowding.

The first method (Cedeño & Vemuri, 1992) gives each population element exactly one trial as the first parent for crossover. The second parent is chosen from a random sample of MS population elements. The sample element that is phenotypically closest to the first parent becomes the second parent. After crossover and mutation, each of the two offspring is inserted into the population as follows. First, CF groups containing CS individuals apiece are randomly selected from the population. Second, the element of each group most similar to the waiting offspring advances to a selection pool. Third, the offspring replaces the worst element of the selection pool. Cedeño and Vemuri indicate that closest-of-the-worst replacement strategies (such as Stadnyk's, 1987) are too biased in favor of higher peaks. The authors do not indicate how often a parent gets replaced under their method, especially given the method's up-front mating restriction. They also do not indicate whether mating restriction is essential to their algorithm.

Cedeño and Vemuri test their crowding method, using population sizes of 200 and higher, on several two-dimensional, multi-peaked, function optimization problems. The authors use a 30-bit chromosome to encode each dimension. Other parameters are $p_c = .95$, $p_m = .06$, $MS = 4$, $CF = 5$, and $CS = 10$. On a function with two peaks of uniform height, their method maintains both peaks in roughly equal proportions. On functions with five peaks of nonuniform height, their method, after 10 generations, starts to transfer individuals from lower peaks to

higher peaks (a behavior of crowding methods that we later analyze in depth). After 100–150 generations, many of the smaller peaks stabilize, but others lose all representatives.

The second crowding method (Harik, 1994) adds elitism to its replacement strategy. The crowding method starts by choosing two elements at random from the population to undergo crossover and mutation. After production of two offspring, a random sample of CF individuals is taken from the population. Each offspring competes against the closest sample element. Winners advance to the next generation.

Harik runs his method on several multimodal functions having between 5 and 32 peaks. In each simulation, his algorithm maintains representatives at all peaks for 100 generations. While the distribution of population elements among peaks is roughly uniform to start, by Generation 100, some peaks lose a number of elements and others gain them. Harik calculates the algorithm's expected maintenance time to be exponential with respect to CF.

The third method is Culberson's (1992) $GIGA$ (short for "gene invariant genetic algorithm"). GIGA employs any type of crossover, but no mutation. Selection occurs only upon replacement, and mating is restricted to elements with similar fitnesses. GIGA selects two appropriate parents for crossover. The two parents undergo crossover s times, yielding s pairs of offspring. The best pair of offspring replaces the parents. If elitism is optionally invoked, the best offspring pair replaces the parent pair only if the offspring pair is better. (Unlike deterministic crowding, competition is pairwise.) The fitness of a pair of individuals is defined as the maximum fitness over both individuals.

Note that GIGA conserves all alleles. Since there is no mutation, and since offspring and their parents compete in pairs, all alleles remain in exactly the same proportions as were present in the initial population. GIGA's role is hence to shuffle bits between solutions in beneficial ways. The interesting thing about GIGA, from a niching perspective, is that it forms two subpopulations within the population. The first subpopulation is highly fit and converges about one optimal solution. The second subpopulation contains mostly complements of subsolutions from the first subpopulation: it contains all of the garbage that has been pushed out of the first subpopulation.

GIGA, as it stands, does not perform niching; it can not form more than two subpopulations, and one of these subpopulations is not very meaningful. However, with some modification, a hybrid of GIGA and deterministic crowding could potentially form and maintain multiple

subpopulations, as well as conserve all alleles. One potential hybrid is constructed as follows. First strip away from GIGA its mating restrictions and multiple crosses. Then invoke elitism.

Restricted competition

A slightly different approach to reaching stable subpopulations is to restrict competition among dissimilar individuals during selection. Borrowing from our earlier ecological discussion, species that coexist tend to occupy different environmental niches, so they typically do not compete for resources.

Goldberg, Deb, and Korb (1990) implement, in the messy GA, a competitive restriction that they call *thresholding*. Thresholding first picks a competitor at random from the population, without replacement, to undergo binary tournament selection. To pick the second competitor, thresholding draws elements at random from the population, without replacement, until it finds an element that possesses at least a prespecified threshold of genes in common with the first. (The two competitors may not be the same individual.) If no such second competitor is found after a fixed number of trials, the first competitor advances to the next stage uncrossed. Note that the alleles of the two competitors do not have to match; the corresponding loci just have to be present. (Recall that the messy GA allows underspecification of strings.) The authors recommend thresholding for solving problems of varying subfunction scale, not for locating multiple solutions.

A similar modification can be made to the standard GA under binary tournament selection. The basic idea is to disallow competition between individuals that are dissimilar, according to some dissimilarity criterion. Competition can be restricted in binary tournament selection as follows. A first competitor is randomly selected, with replacement, from the population. Potential second competitors are repeatedly selected at random from the population, also with replacement, until one is phenotypically (or alternatively, genotypically) within a specified distance θ of the first competitor. The competitor with higher fitness advances.

Many optional modifications are available, either alone or in combination. The first modification is to select the first competitor without replacement, giving each individual in the population exactly one trial as the first competitor. The second modification is to select candidates for the second competitor without replacement, but to place a limit on how many can be rejected. If the limit is reached, the first individual passes through without competing, as

Table 4.2: The GA with competitive restriction, $\theta = .1$, $n = 100$, $p_c = 1$, and $p_m = 0$ runs on function *M1*.

Generations	Peaks Maintained	Average Fitness
100	5	.89
200	5	.94
300	4	.97
400	3	.98

in the messy GA's thresholding technique. The third modification is to choose a fixed sized sample from the population, as many crowding techniques do, and to compete with the closest element from that sample.

We perform a test run of the GA with competitive restriction, but without the three optional modifications. The GA runs on function *M1* of Chapter 5, using a phenotypic θ of .1, $n = 100$, $p_c = 1$, and $p_m = 0$. The GA successfully locates all five peaks and maintains them to Generation 100, the stopping point. A second run maintains all five peaks to Generation 50, but loses one by Generation 100. A third run, of 400 generations, has mixed success, as displayed in Table 4.2. We hypothesize that solutions disappear from various peaks because crosses of good individuals form lethal individuals that dilute classes.

To avoid the formation of lethal individuals, we add mating restriction to the competitive-restriction GA. Basically, two elements must be within the distance θ of each other, or they are not allowed to cross. (If they do not cross, both proceed to the mutation stage.) This is the idealized mating-restriction scheme introduced in Chapter 3. With $n = 100$, $p_c = 1$, and $p_m = 0$, the population fully converges by Generation 24 to five different values, one corresponding to each peak. The least fit values are of fitness .8999; the most fit, .9932. Average population fitness is .9664.

The addition of mutation ($p_m = .01$) produces the results outlined in Table 4.3. Average population fitness is very high, and representatives of all peaks remain, even to Generation 1000. Mutation prevents full and final convergence of the population. Note that a few individuals with extremely low fitnesses are created by the mutation of high-order bits.

The above simulations show that although restricted-competition methods have seldom been attempted, they form a category of niching methods that deserves further consideration. Note that restricted-competition methods bear strong resemblance to restricted-replacement

Table 4.3: The GA with competitive restriction, mating restriction, and mutation runs on function *M1*. GA parameters are $\theta = .1$, $n = 100$, $p_c = 1$, and $p_m = .01$.

Generations	Peaks Maintained	Average Fitness	High Fitness	Low Fitness
100	5	.972	1.00	0.026
200	5	.913	1.00	0.000
300	5	.950	1.00	0.000
400	5	.952	1.00	0.003
500	5	.914	1.00	0.000
1000	5	.943	1.00	0.003

(crowding) methods. The difference is that restricted-competition methods perform selection up front, while successful restricted-replacement methods perform selection upon replacement. In that sense, both categories of algorithms restrict their selection method. However, GAs with replacement selection, such as CHC (Eshelman, 1991; Eshelman & Schaffer, 1991), tend to be elitist in nature, and exhibit markedly different properties than GAs with up-front selection. We expect this distinction to transfer to niching GAs, thereby producing two categories of algorithms with different behaviors. However, whether the behavior of restricted-competition GAs differs significantly from the behavior of restricted-replacement GAs remains to be seen. We leave this question to future research.

Fitness sharing

The use of fitness as a single, shared resource has enjoyed the most success, to date, of any niching method. In fact, prior to the commencement of this thesis, fitness sharing was the only niching technique with a successful track record in multimodal function optimization.

Holland (1975, 1992) discusses the concept of limiting the number of individuals occupying a niche, to the carrying capacity of that niche. He states that if a niche has associated with it a fixed *payoff* at every timestep, and if each individual occupying that niche is forced to equally share that payoff, then a stable situation arises when each niche contains a number of individuals proportional to its payoff. (Payoff in this context is equivalent to fitness in the GA context.) If some niches become overcrowded, it is to the advantage of individuals occupying those niches to seek out less crowded niches. Niching methods which utilize this concept of Holland's are called *sharing* methods. Sharing methods can be defined as algorithms that require similar

population elements to share fitnesses. To induce sharing, these algorithms typically alter the fitness of each population element based on its proximity to other population elements.

An early example of a sharing technique is in Booker's (1985) classifier system, in which similar classifiers are required to share payoff. Booker's scheme works as follows. Low-fitness classifiers are periodically deleted, supplying selection pressure. All classifiers that match a message pay tax, then share a "fixed-size tax rebate". Too many classifiers in a niche causes the decrease of average classifier fitness in that niche, "because they lose more [fitness] than they gain" in a transaction of tax followed by rebate; too few classifiers causes a gain of fitness in a transaction. Sharing the tax rebate results in a steady state in the number of classifiers each niche supports.

Horn, Goldberg, and Deb (1994) dub Booker's version of sharing, *LCS implicit fitness sharing*. They implement it in a simpler manner, by assigning fitness based on the number of positive examples a rule covers. For each positive example covered by more than one rule, the authors force the overlapping rules to evenly divide that example's payoff. Their simplified implementation assumes that all rules are of equal generality.

Explicit fitness sharing is introduced by Goldberg and Richardson (1987) as the "method of sharing functions". Their method, now known more simply as *sharing*, is directly applicable to multimodal function optimization. Sharing derates an individual's fitness by an amount related to the number of similar individuals in the population. Specifically, an individual's new *shared fitness*, f', is equal to its old fitness f divided by its *niche count*. An individual's niche count is a sum of *sharing function* (sh) values between itself and every individual in the population (including itself). The shared fitness of a population element i is

$$f'(i) \;=\; \frac{f(i)}{\sum_{j=1}^{n} sh(d(i,j))} \;\;. \tag{4.1}$$

The sharing function sh is a function of the distance d between two population elements; it returns a '1' if the elements are identical, a '0' if they exceed some threshold of dissimilarity, and an intermediate value for intermediate levels of dissimilarity. The threshold of dissimilarity is specified by a constant, σ_{share}; if the distance between two population elements is greater than or equal to σ_{share}, they do not affect each other's shared fitness. Most commonly-used

sharing functions are of the form,

$$sh(d) \;=\; \begin{cases} 1 - \left(\frac{d}{\sigma_{share}}\right)^{\alpha} \;, & \text{if } d < \sigma_{share} \;; \\ 0 \;, & \text{otherwise} \;. \end{cases} \qquad (4.2)$$

In the above equation, α is a constant (typically set to 1) used to regulate the shape of the sharing function. Both genotypic and phenotypic distance measures can be employed, the appropriate choice depending on the problem.

Goldberg and Richardson report stable clustering of population elements around the peaks of two sinusoidal functions. One function has five peaks of equal height, while the other has five peaks of differing heights. The number of individuals clustered around each peak is directly proportional to the height of the peak. Only when the expected number of individuals around a peak drops close to one, does the peak lose all representation in the population.

Further studies experiment with both genotypic and phenotypic sharing, using similar test problems (Deb, 1989; Deb & Goldberg, 1989). *Genotypic sharing* employs Hamming distance as its distance measure, while *Phenotypic sharing* employs Euclidean distance in decoded parameter space. Phenotypic sharing gives slightly better results, due to decreased noise in the decoded parameter space. Mating-restriction strategies are tested in combination with sharing. They successfully reduce the number of lethal crosses between individuals from different peaks.

Sharing can be implemented using any selection method, but the choice of method may either increase or decrease the stability of the algorithm. Stochastic remainder selection has been the most popular to date. Tournament selection is another possibility, but special provisions must be made to promote stability. Oei, Goldberg, and Chang (1991) propose a technique for combining sharing with binary tournament selection. Their technique calculates shared fitnesses with respect to the new population as it is being filled. This method, called *tournament selection with continuously updated sharing* is used by Goldberg, Deb, and Horn (1992), in conjunction with fitness scaling prior to sharing, to solve a massively multimodal, deceptive problem.

Yin and Germay (1993) propose that a clustering algorithm be implemented prior to sharing, in order to divide the population into niches. Each individual shares only with the individuals in its niche. Since each individual's shared fitness is not computed with respect to the entire population, the algorithm is expected to be faster than traditional sharing. Yin (1994) suc-

cessfully applies sharing with clustering to several test problems and to a real-world "load flow problem in electrical power systems".

Spears (1994) suggests a "simple subpopulation scheme" (SSS1) that incorporates a form of fitness sharing. SSS1, like its predecessor with only mating restriction, employs a k-bit tag for each individual, and disallows both mating between individuals with differing tags and mutation on tag bits. Sharing "divide[s] the fitness of each individual by the size of its subpopulation". Spears notes that SSS1 "maintain[s] stable subpopulations on all but the lowest one or two peaks" of functions possessing five peaks. Some subpopulations have individuals on more than one peak, creating competition within the subpopulation and encouraging migration to the highest peak.

Spears also defines an extension, SSS2, which adds a one-dimensional ring topology, and further restricts mating to neighbors with identical tags. (Each individual on a ring has two neighbors.) Spears notes that stable subpopulations consistently form about all but the lowest peak of the previous functions, and sometimes about the lowest peak as well. Spears also runs SSS2 on a two-dimensional function with six peaks. SSS2 finds five of the six peaks, using $n = 100$ and 16 subpopulations. It is questionable how much exploration is occurring in SSS2, because of its severe mating and topological restrictions, in addition to fitness sharing. The test problems are simple enough that initial populations can rapidly hillclimb the peaks, given an appropriately defined parallel hillclimber (see Chapter 10).

Spears finally suggests performing topological sharing — defining niches based on subpopulations that cluster together geographically. He notes that this would lessen or even eliminate the need for tag bits: a single tag bit could denote a change in subpopulation. This scheme would give the user less control over the number of subpopulations. Topological sharing remains an untested idea.

In machine learning, Smith and Valenzuela-Rendón (1989) try explicit fitness sharing in their stimulus-response classifier. They demonstrate, theoretically, the benefits of niching methods. Kargupta and Smith (1991) employ fitness sharing in evolving "polynomial networks". Packard (1990) utilizes a "diversity booster mechanism" that derates the fitnesses of population elements that are genotypically close to fitter population elements (effectively a sharing mechanism). Giordana, Saitta, and Zini (1994) use fitness sharing, along with isolated subpopulations and specialized operators, to learn disjunctive concepts.

In multiobjective function optimization, sharing has recently been shown to be highly effective. Fonseca and Fleming (1993) implement fitness sharing among individuals with similar objective function values. However, if one individual dominates another, the two do not share fitnesses. The authors assign a fitness to each individual based upon that individual's rank. The rank is 1 plus the number of population elements that dominate the individual. All non-dominated individuals receive a rank of 1. The authors note that if a niching method is good at maintaining diversity given a flat fitness function, then it should also perform well in multi-objective function optimization.

The *niched Pareto GA* (Horn & Nafpliotis, 1993; Horn, Nafpliotis, & Goldberg, 1994) works via tournament selection with continuously updated sharing. After picking two competitors at random from the population, the GA selects an additional sample of size k. If one competitor is dominated by any individual in the sample, but the other competitor is not, the nondominated competitor wins the tournament. If neither or both are dominated, sharing chooses a winner. The winner is the competitor with the smallest niche count, where niche counts are estimated by sampling the population. Choosing the winner with the smallest niche count spreads solutions across the Pareto-optimal front. The authors state that the sample size k allows a user control over selection pressure: smaller k values result in greater selection pressure.

The main drawback to using sharing is the additional time required to cycle through the population to compute shared fitnesses. Several authors have suggested calculating shared fitnesses from fixed sized samples of the population (Goldberg & Richardson, 1987; Oei, Goldberg, & Chang, 1991). Sharing with sampling has subsequently been applied, with success, to solving complex multimodal and multiobjective optimization problems (Goldberg, Deb, & Horn, 1992; Horn & Nafpliotis, 1993; Horn, Nafpliotis, & Goldberg, 1994). Clustering is another potential remedy, as Yin and Germay have shown. As far as GA time complexity is concerned, in real-world problems, a function evaluation requires much more time than a comparison; most GAs perform only $O(n)$ function evaluations each generation.

Immune system models

Somewhere between techniques that use multiple fitness functions, crowding schemes, and fitness-sharing methods, are pattern matchers based on models of the immune system. The immune system is able to differentiate foreign objects (*antigens*) from those native to an or-

ganism. To do so it forms *antibodies* whose job it is to recognize or match the antigens. So far, immune system models have been applied only to simple simulations in which a population of bit strings, called antibodies, evolves to match or cover a set of bit-string patterns, called antigens. The antigens define multiple, spatial environments that allocate increasing fitnesses to population elements as those population elements get closer to matching one or more antigens. The pattern-matching task in immune system models is similar to that in classifier systems, except no explicit code exists in antibodies to allow them to generalize. Research into immune system models is still in its early stages. No one has yet demonstrated a technique that allows them to optimize multimodal functions.

Stadnyk (1987) introduces one of the first immune-system pattern matchers. She tries two different fitness functions that allow credit for partial matches. The first function is the average number of matching bits across a set of patterns. The second function is a less generous average: it counts each matching bit only if the adjacent two bits also match. This is designed to prevent assignment of high fitnesses to rules likely to be disrupted by single-point crossover. The author employs a sample of k randomly chosen patterns to evaluate each population element.

Stadnyk's GA employs fitness-proportionate selection, crossover, mutation, and a replacement scheme. The author tries several replacement schemes that fail to induce niching, even on simple problems with two complementary patterns that occur in equal proportions in the pattern set. These schemes include replacing the worst element of the population, De Jong's crowding with small CF, and De Jong's crowding with $CF = n$. She finally settles upon a crowding scheme that replaces the closest of the worst population elements. That is, CF candidates are chosen from the population inversely proportional to fitness, and the one closest to the new element is replaced.

Stadnyk finds that stable subpopulations evolve around different patterns in the pattern set. Using $k = 3$, $CF = 10$, and $n = 50$, her GA successfully maintains subpopulations of strings that match a set of 10 patterns, in at least five bit-positions per pattern. The observed niching is not due to the crowding scheme in isolation, but to pattern sampling as well. If all patterns are used in the fitness computation, "then the entire recognizer population ends up matching only one pattern" (due to an implicit preference for generalists over specialists).

Forrest, Jarvonik, Smith, and Perelson (1993) present an extended immune system model that assigns fitnesses to the entire population through repeated application of the following

procedure. (They apply the procedure $3n$ times.) First, their algorithm randomly picks one antigen. Then it randomly chooses a sample of σ antibodies without replacement. The best-matching antibody wins the competition, and has its matching score added to its fitness. Ties are broken at random. The authors show that at very low σ, their algorithm exhibits a form of generalization. They also determine experimentally that their algorithm requires at least 15 population elements for each solution it maintains. Smith, Forrest, and Perelson (1993a, 1993b) give analytical results that roughly equate Forrest et al.'s immune system model to fitness sharing, and that verify the generalization capabilities of the model at low σ.

Note that Stadnyk starts to achieve successful niching by lowering the number of antigens used in computing the fitness of each population element to 3. Forrest et al. continue this trend by lowering the number all the way to 1. However, instead of sampling the antigens, they sample the antibodies.

The commonality between the two immune system models we have examined is that they both employ sampling, either of the antigens or the antibodies. Without sampling, both algorithms converge to a uniform population. That these sampling methods should work is not surprising if we look at them from the standpoint of emphasizing exceptional rather than average behavior. Showing the population only one or a few antigens at a time effectively varies the fitness function to emphasize specialization rather than generalization. In contrast, exhaustive application of every antigen to every antibody produces an averaging effect.

Note that Forrest et al.'s method is analogous to Schaffer's VEGA system, in which only one objective at a time (corresponding to an antigen) is shown to a subpopulation (corresponding to a sample of antibodies). However, Schaffer's VEGA differs in that all population elements in the sample are assigned their full fitnesses after this one showing; in Forrest et al.'s method, a single population element is assigned a partial fitness, and all population elements in the sample are likely to be sampled several more times.

4.3 A Formal Framework for Multimodal Landscapes

In Chapter 3, we presented a formal framework for the study of population diversity. We now specialize that framework to multimodal landscapes, the fitness terrains on which our niching

90

methods will run. In Chapter 5, we will incorporate this specialized framework into a modelling framework for niching GAs.

The optimization of multimodal landscapes or fitness functions involves issues such as the total number of optima in the search space, the quantity and quality of optima one wishes to locate, and the degree of problem difficulty due to isolated and misleading optima. We do not wish to perform all possible types of multimodal function optimization since, as explained later, the possible types are unlimited. Instead, we concentrate on solving a general type of multimodal function optimization problem that is likely to encompass most practical problems to which niching methods will be applied.

The general goal in multimodal function optimization is to find several solutions that are globally or locally optimal. Before we can be more specific, it is necessary to formally define optimization. Our treatment of optimization follows, to some extent, Aarts and Korst's (1989) book. Assume a search space S. Also assume an objective function f that assigns a real number to each element of S ($f : S \rightarrow \Re$, where $\Re$ represents the set of all real numbers). Assume without loss of generality that maximization with respect to f is the goal. Define a neighborhood generation function $N : S \rightarrow 2^S$ that assigns each element i of S a set of elements in S that are close to i in some sense (i need not be in its own neighborhood).

For an arbitrary $i \in S$, $N(i) \subseteq S$ is called the *neighborhood* of i; i is a *local maximum* if $f(i) \geq f(j)$ for all $j \in N(i)$; i is a *global maximum* if $f(i) \geq f(j)$ for all local maxima j. Call the number of local maxima c and the number of global maxima z.

Several types of multimodal function optimization problems should be of interest to people solving real-world problems. These include the following:

1. Find any $b < c$ maxima;

2. Find all c maxima;

3. Find at least the $b \leq c$ highest maxima;

4. Find any $b < z$ global maxima;

5. Find all z global maxima.

As an example, Type 5 problems are important in multiobjective function optimization. In multiobjective function optimization, Pareto-optimal solutions are typically global maxima, while dominated solutions are not.

Note that solving an equivalent Type 3 problem will yield a solution to any of the other four types. Therefore, we concentrate upon the most general task, finding the b best maxima (plus maybe a few extra) on a landscape with c maxima. Of course, there exist infinite types of multimodal function optimization problems — not just the five we list. There are less useful types of problems, such as locating the lowest maxima. There are problems with non-fitness-based criteria, such as finding maxima that are maximally spread across the search space, rather than in close proximity. One way to handle secondary criteria is to embed them into the fitness function. In general, there is no way to predetermine the non-fitness-based criteria a specific application may have. These considerations must be left to the applications programmer. We restrict our attention in this study to fitness-based multimodal function optimization, and to the most general of the five types of problems, Type 3.

We will utilize a definition of diversity that is useful for studying Type 3 multimodal function optimization problems. Recall the general definition of a diversity measure from Equation 3.9. For multimodal function optimization, we will work with descriptive populations, $R(pop_t)$, where R is a one-to-one mapping between individuals and either genotypes (bit strings) or phenotypes (k-tuples of variables). The partition X will divide either the genotypic or the phenotypic search space into equivalence classes, one class for every local maximum. We will use the Euclidean distance metric of Chapter 3 to measure distances between distributions. We need only specify a set of goal distributions for the Type 3 problem.

Assume c local optima — $o_0, o_1, \ldots, o_{c-1}$ — indexed in decreasing order of objective function value or fitness (i.e., o_0 has the highest fitness and o_{c-1} has the lowest fitness). Optima with identical fitnesses are given consecutive indices. Let $f(o_i)$ be the fitness of optimum o_i.

Several sets of goal distributions are possible. One could insist that each distribution Q_i be directly proportional to class fitness, as in fitness sharing. In this case Φ would consist of the following distributions:

$$Q_0 = \left\langle \quad \frac{f(o_0)}{\sum_{i=0}^{b-1} f(o_i)} \quad \cdots \quad \frac{f(o_{b-1})}{\sum_{i=0}^{b-1} f(o_i)}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$Q_1 = \left\langle \quad \frac{f(o_0)}{\sum_{i=0}^{b} f(o_i)} \quad \cdots \quad \frac{f(o_b)}{\sum_{i=0}^{b} f(o_i)}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$\vdots$$

$$Q_{c-b} = \left\langle \quad \frac{f(o_0)}{\sum_{i=0}^{c-1} f(o_i)} \quad \cdots \quad \frac{f(o_{c-1})}{\sum_{i=0}^{c-1} f(o_i)} \quad \right\rangle .$$

One could insist on a uniform distribution among the located peaks:

$$Q_0 = \left\langle \quad \frac{1}{b} \quad \cdots \quad \frac{1}{b}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$Q_1 = \left\langle \quad \frac{1}{b+1} \quad \cdots \quad \frac{1}{b+1}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$\vdots$$

$$Q_{c-b} = \left\langle \quad \frac{1}{c} \quad \cdots \quad \frac{1}{c} \quad \right\rangle .$$

We would like to be as inclusive as possible. Therefore, our set of goal distributions consists of all population distributions that allocate at least one element to each of the b optima of interest. This guarantees that all optima of interest will be returned by a niching method, if its population is in a maximally diverse state. Formally, we define Φ as containing, for all $q_{i,j} \geq 1/n$, all legal distributions of the following forms:

$$\left\langle \quad q_{0,0}, \quad q_{0,1} \quad \cdots \quad q_{0,b-1}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$\left\langle \quad q_{1,0}, \quad q_{1,1} \quad \cdots \quad q_{1,b}, \quad 0 \quad \cdots \quad 0 \quad \right\rangle ;$$

$$\vdots$$

$$\left\langle \quad q_{k,0}, \quad q_{k,1} \quad \cdots \quad q_{k,c-1} \quad \right\rangle ,$$

where k is an arbitrary identifier.

4.4   Further Research

In Chapter 3, we examined previous diversification methods that did not qualify as niching methods. In this chapter, we examined and classified various niching methods along two dimensions — space versus time, and single versus multiple environments. Since the main thrust of this thesis is multimodal function optimization, we will restrict our attention, in the remainder of this thesis, to spatial niching methods that operate within a single environment. One exception is the sequential niching technique of Beasley, Bull, and Martin (1993). We will utilize their technique as a simple benchmark for evaluating and comparing other niching methods.

Of the four categories of spatial, single-environment, niching methods — heterozygote advantage, crowding, restricted competition, and fitness sharing — we choose crowding and sharing for further analysis. We choose these categories because they are the most mature of the four categories. Although still only sparsely explored, crowding and sharing have had the most GA research devoted to them. Heterozygote advantage and restricted competition, while interesting and promising categories of methods, have had nearly zero research devoted to them.

We end this chapter by noting that there are no right or wrong choices. Had we chosen the other two categories for further examination, this thesis would have evolved along different, but still productive lines. There has been a tendency among GA researchers to seek out novel areas, rather than revisit previously researched areas, many of which have only superficially been investigated. Often, seemingly new directions, viewed from different perspectives, correspond to familiar scenery. We resist the aforementioned tendency in this study, in the hope that deep understanding of the two methods we examine will lead to a more complete understanding of niching methods for genetic algorithms.

Chapter 5

Models of Niching Methods

The framework used throughout this study for modelling niching methods builds upon Chapter 3's framework for diversity and Chapter 4's specialization of that framework. The present chapter first reviews previous methods for modelling GAs, then presents our framework for modelling niching methods and GAs that incorporate niching methods. Our framework is partly analytical and partly empirical. The analytical portion consists of several abstractions, definitions, and simplifying assumptions. It recommends particular equivalence-class models that simplify the niching methods under consideration. The empirical portion consists primarily of test problems. The test problems cover a range of difficulty levels, and vary along three dimensions of problem difficulty — multimodality, isolation, and misleadingness.

The predominant modelling technique in prior GA research has been Markov chains (Lial & Miller, 1989). As the next section explains, Markov chain models can be quite cumbersome. This thesis demonstrates analytical and empirical alternatives to Markov chain models.

Previously, little has been known about the properties and expected behavior of many of the niching methods we examine. Consequently, the primary purposes behind our modelling are to uncover intrinsic properties of niching methods and to fully understand their behavior. The models resulting from our framework should assist us in deriving optimal control parameters such as population size and crossover probability, and in selecting better algorithmic design alternatives. Through modelling ideal properties and behavior, we can hope to determine expected properties and behavior, such as the expected distribution of population elements across niches and the expected time to loss of desired solutions.

5.1 Previous Models

Previous models of GAs are for the most part highly complex, and are typically based on Markov chains. Several exact models exist that are equivalent in expected population trajectories to the GAs they model. Exact models of the GA typically are poorly suited for analysis, and have little predictive value. In most cases, one can derive more conclusions from analyzing the GA itself than from analyzing its exact model. Some researchers attempt to bypass the limitations of exact models by executing their models; they first pick a set of GA parameters and a test function, and then compute an expected end result for the GA. Unfortunately, such models grow at least exponentially in both execution time and memory requirements, as n and l increase. Therefore, even assuming that results using specific test functions can meaningfully be generalized, execution time for an exact model is typically orders of magnitude longer than for its corresponding GA. Again, better insights into the GA's expected behavior can often be obtained through multiple runs of the GA. We briefly cover previous GA models in this section, before presenting a new modelling framework for niching GAs, in subsequent sections.

Several Markov chain models of the simple GA have been constructed that incorporate fitness-proportionate selection, crossover, and mutation. These models assume population elements are fixed sized bit-strings of length l, that the population is finite, and that population size remains constant from generation to generation. Each possible population of n elements forms one state of a Markov chain. The models are complete and exact: they consist of complex expressions for the transition probabilities between all possible pairs of populations.

Nix and Vose (1992) give complete expressions for transition probabilities between the possible populations of a simple GA. They employ the general formula for a transition probability to describe the asymptotic behavior of a simple GA as population size increases. They find that as n increases, the GA converges asymptotically to the same attractor as would an infinite-population-size GA (assuming a single attractor for the population). Vose (1993) reinforces this relationship between large-population and infinite-population GAs. He describes the infinite-time behavior of a simple GA as visiting each attractor infinitely often, with preference for the attractor with the largest basin of attraction. Note that the space under which the authors define attractors is not the fitness landscape, but the space of all possible populations.

Nix and Vose (1992) calculate the number of possible populations, which is also the number of states in the GA's Markov chain. Rewritten in terms of a one-to-one mapping between binary strings and equivalence classes, the number of possible populations is

$$\frac{(n + c - 1)!}{n!(c - 1)!} \quad , \tag{5.1}$$

where $c = 2^l$ is the number of equivalence classes. Mahfoud (1993) gives an algorithm for enumerating all possible states of the Markov chain for a generalized, equivalence-class model, in which equivalence classes need not directly correspond to binary strings. The number of transitions in the Markov chain's transition matrix is the number of states, squared.

Nix and Vose's model tells us that finite-population combinatorics become overwhelming as n and l increase. For example, given a population containing only 10 strings of length 10, one would need to represent approximately 3.65×10^{23} states, and approximately 1.33×10^{47} transitions. Clearly, for any nontrivial set of GA parameters, one could not reasonably expect to calculate all transition probabilities or to process the resulting matrix.

Even if one collapses related states into equivalence classes, finite Markov chain models still require a prohibitive number of transitions for nontrivial population sizes, when $c > 2$. Mahfoud (1993) models Boltzmann tournament selection (selection only) at the equivalence-class level, using absorbing, finite Markov chains. His model computes expected absorption time for the selection procedure. As c and n become moderately large, the Markov chain model becomes unwieldy. Mahfoud is able to compute absorption times for population sizes of up to $n = 36$ with $c = 3$, and for up to six equivalence classes ($c = 6$) with $n = 6$.

De Jong, Spears, and Gordon (in press) employ Nix and Vose's model to determine expected convergence times for simple GAs with small populations, simple objective functions, and $l = 2$. They compute the expected time for the GA to first encounter an optimal solution, rather than for the full population to converge to some attractor. The authors find that the family of two-bit fitness functions they define (24 functions in all) can be partitioned into three categories, such that a GA is expected to exhibit identical behavior on all functions within a category. Future work will determine whether results observed at $l = 2$ extend to $l = 3$ and higher.

Infinite population-size models are also prevalent in GA research. Goldberg (1987) iterates the expected proportions of four competing order-2 schemata that form a partially deceptive

partition of the search space. He identifies the schema that is expected to take over the whole population, and computes its expected takeover time. Bridges and Goldberg (1987) present exact proportion-equations for string gains and string losses due to crossover and selection in the simple GA. Whitley (1993) extends their model, allowing it to execute for arbitrary specifications of GA parameters and objective functions. He states that his model can accommodate functions of up to 15 bits before becoming computationally impractical. Whitley further models alternative crossover operators and parallel GAs. Srinivas and Patnaik (1993) introduce an exact model of traditional GAs for objective functions of unitation. The time complexity for execution of their model is order-l^3, an improvement over the exponential execution time required by Whitley's model. However, their model applies only to functions of unitation.

Other studies (e.g., Goldberg & Deb, 1991; Goldberg & Bridges, 1990) also assume infinite population sizes and execute proportion-equations, in most cases to compute expected convergence time. The major shortcoming of most infinite population-size models is that by their nature, they can not simulate finite properties of the GA such as population size and genetic drift; they model only expectation. Diffusion-equation models are exceptions that approximate Markov chains (Harvey, 1993; Kargupta, 1993).

T. E. Davis and Principe (1993) develop a somewhat different GA model. They employ a nonstationary Markov chain in an attempt to develop a theory of convergence for GAs analogous to that for simulated annealing. The authors utilize a variable probability-of-mutation parameter as an analogue to simulated annealing's temperature parameter. They also experiment with limiting behavior as population size increases. Suzuki (1993) models an elitist GA using Markov chains, and analyzes convergence using the eigenvalues of the transition matrix.

Goldberg and Segrest (1987) present simplified, two-class models of fitness-proportionate selection, both with and without mutation. The authors use absorbing, finite Markov chains, and represent the two classes with the one-bit strings, '0' and '1'. We have previously discussed their results in Chapter 3. Horn (1993) extends Goldberg and Segrest's selection-only model to incorporate fitness sharing. He graphs expected absorption times for the two-class case.

In summary, models at a level of complexity similar to or higher than that of the GA typically do not provide a clearer perspective of the GA. Models must simplify the highly complex GA, and decompose it into subproblems if possible. The traditional GA has natural, quasi-separable subproblems — selection, crossover, and mutation. (Niching methods usually modify

the selection phase.) We present a modelling framework in the remainder of this chapter that allows the subdivision and simplification of the niching GA. The resulting models concentrate upon selection (with niching), the GA's major operator, and then add crossover, the GA's second most influential operator. The modelling framework allows successful abstraction of the most significant problem-solving behaviors of niching GAs.

5.2   Analytical Framework

The analytical portion of our modelling framework continues from Chapter 4's specialization (to multimodal landscapes) of Chapter 3's diversity framework. It consists of abstractions, definitions, and simplifying assumptions, from which we later construct specific models. These abstractions, definitions, and assumptions allow extraction of the major behaviors of each method. As demonstrated in later simulations, the predictive value of the models is not diminished by their assumptions. The models would in fact not be possible without these assumptions.

Our analytical framework represents a new approach to the modelling of genetic algorithms. The new approach is not entirely novel, however. It is based philosophically upon the problem decomposition approach discussed in the first two chapters. In addition, it borrows some of its explicit assumptions from implicit assumptions found in previous studies of niching GAs.

Equivalence classes

In Chapters 3 and 4, we defined a framework for the study of diversity and, more specifically, for the study of niching methods in multimodal function optimization. The framework was based on a partitioning of the search space into equivalence classes. In Chapter 4, we equated equivalence classes with maxima in the search space. We defined the ideal behavior of a niching method that performs multimodal function maximization, in terms of the maxima around which the niching method deposits population elements. Specifically, we stated that an effective niching method should locate a number of the best peaks; based upon this definition, we defined a set of goal distributions. This set of goals consisted of all population distributions that allocate at least one population element to each of b best optima.

One detail we omitted in Chapter 4 is exactly how one might determine the class to which each population element corresponds. Traditional schema analysis for the GA makes use of

schema partitions of the search space. That is, if one chooses an interesting set of fixed bit-positions, and is interested in how the GA assigns values to bits in those positions, then a natural partitioning is one in which every possible schema having those positions fixed, defines an equivalence class. If k bits are designated as fixed, the search space subdivides into 2^k schema-partitions or equivalence classes.

Several researchers, starting with Holland (1975), have suggested a more general framework in which partitions do not necessarily correspond to bit-positions. Arbitrary partitions can be considered, consisting of arbitrary equivalence classes that need not be schema based. These arbitrary equivalence classes are known as *generalized schemata* or *formae* (Radcliffe, 1991a, 1991b, 1993).

Since we are studying multimodal function optimization, we would like our classes to correspond to local optima in the search space. For certain well defined functions, local optima will in turn correspond to schema partitions, but we will not limit our models to such functions. In many cases, optima-based partitions will not correspond one-to-one with schema-based partitions. However, significant overlap may exist between any optima-based equivalence class and a corresponding schema.

If one were to visually examine a search space of three dimensions or fewer, one would intuitively know the peak to which each element belonged. Often, however, multiple pictorial representations exist of the same search space. These include Hamming cubes, graphs with phenotypic variables as axes, unitation-based representations, and others. The simple GA operates directly in Hamming space. However, bit-based or schema-based class definitions are not always the most natural ones for a given problem. In addition, a large number of uninteresting local optima may be present. In such a case, the question arises of whether to represent all the extraneous optima as classes, or perhaps lump them together, as Deb (1989) does, into an all-encompassing "none of the above" class. Another question arises when a point is not very close to any maximum — in the worst case the point is a local minimum that is equidistant from two or more maxima: to which class should we assign such a point? Furthermore, some research studies consider a point a member of a peak only if it is within some fitness of peak fitness, or if it is within a fixed number of moves from the peak, according to some neighborhood operator.

The first definition we make in our modelling framework is to decide what constitutes an equivalence class, and to which class each point in the search space belongs. In terms of Chapter 4's framework, we define an appropriate partition X of the search space, using the concepts of *attractors* and *basins of attraction*. We consider each local maximum to be an attractor, and assign to each class all points within its corresponding maximum's basin of attraction. Hence, each population element is considered a member of the class corresponding to the peak to which it is attracted, and all points in the search space are attracted to exactly one peak. We do not utilize none-of-the-above classes, and we assign degenerate points such as local minima through an appropriate tie-breaking procedure. (Local minima will, anyway, disappear rapidly in most GAs.)

We define a local maximum's basin of attraction in terms of hillclimbing under an appropriately defined neighborhood operator. A point is in the basin of attraction of a local maximum if it would hillclimb to that maximum using a given neighborhood operator and deterministic hillclimbing algorithm. We make the hillclimbing algorithm deterministic so that it enforces a tie-breaking procedure and so that the possibility of a point hillclimbing to two different peaks, given two tries, is zero. Our hillclimbing-based definition of equivalence class is consistent with the notion that GAs are global optimization methods that operate best when combined with local optimization methods such as hillclimbers. In practice, a hillclimber may be invoked on a GA's final population in order to force population elements to the nearest local optimum. We employ such a hillclimber in many experiments of this thesis.

The best neighborhood operator for a particular problem is problem dependent. Examples of neighborhoods are epsilon neighborhoods in multidimensional, real-valued variable spaces, and one-bit Hamming neighborhoods. As an example, given a neighborhood of $\pm\epsilon$ in the phenotype, where the phenotype contains only one variable, points which lie under a peak are in the basin of attraction of that peak. We assume an appropriate neighborhood operator for each problem the GA must solve. For problems in which a user-accessible phenotype exists, we assume a neighborhood operator that adds or subtracts ϵ from one of the problem's variables, where ϵ is the smallest increment encoded by the GA. Otherwise, we assume a genotypic neighborhood operator that flips one of the bits in the encoding.

We define a theoretical, deterministic, hillclimbing algorithm for our modelling framework. The hillclimbing algorithm finds the local optimum to which any point i in the search space

1. *Current Set* = {i}

2. *Old Set* = $\emptyset$

3. REPEAT

   (a) *Old Set* = *Old Set* $\bigcup$ *Current Set*

   (b) f_{max} = fitness of fittest element in *Current Set*

   (c) Remove all points from *Current Set* with fitness < f_{max}

   (d) IF *Current Set* contains one or more local optima
       - RETURN the first optimum in *Current Set*
       - TERMINATE algorithm

   (e) *Current Set* = exhaustively apply N to each element of *Current Set*

   (f) *Current Set* = *Current Set* − (*Current Set* $\bigcap$ *Old Set*)

Figure 5.1: Pseudocode is given for the theoretical, deterministic hillclimber of the modelling framework.

is attracted, given a neighborhood operator N. N must be deterministic in order for the hillclimber to be deterministic.

The hillclimbing algorithm, shown in Figure 5.1, works as follows. It maintains a *Current Set* of points, to which N is repeatedly applied. It also maintains an *Old Set* of points that have already been visited, in order to prevent repeated visits to the same points. The algorithm initializes *Current Set* to contain the single element i. Then it repeats the following cycle until it has found a local optimum. The cycle first discards all elements of *Current Set* that are not of maximal fitness. Then it applies N, exhaustively and in parallel, to all elements of *Current Set*, overwriting the old values in *Current Set* with the results from applying the neighborhood operator. Exhaustive application means that for each $j \in$ *Current Set*, all points in j's neighborhood must be generated. Finally, the algorithm eliminates elements of *Current Set* that have already been visited.

Just prior to termination, if *Current Set* contains more than one local optimum, the algorithm returns the first optimum in *Current Set*. The hillclimber maintains *Current Set* as an ordered set, and N deposits elements into *Current Set* in a deterministic manner. The algorithm is intended for situations in which all local optima have been identified beforehand. The purpose of the hillclimbing algorithm is not to locate optima, but to identify the single

optimum to which any point in the search space is attracted. Nevertheless, one does not need to know the local optima beforehand to execute the hillclimber. One can simply iterate it until no further improvement occurs.

Our definition of local optimum, from Chapter 4, includes all points on plateaus. Therefore, the hillclimber terminates upon encountering a plateau. If one wishes not to terminate on plateaus, one can redefine a local optimum using a requirement of strict increase. However, care must be taken to define the hillclimber's behavior in flat regions of the search space. (Some functions may not even possess optima.)

We index the c equivalence classes, as in prior chapters, from 0 to $c-1$. The c classes contain, at a fixed point of interest in time, $I_0, \ldots, I_{c-1}$ population elements, respectively ($\sum_{i=0}^{c-1} I_i = n$). A successful niching GA will match one of many goal distributions, as defined in the previous chapter. Relatively successful niching GAs will approximately match a goal distribution.

Note that the terms, equivalence class, class, niche, species, peak, subpopulation, maximum, and optimum are somewhat interchangeable in our framework. Of course, subtle distinctions do exist. *Equivalence class* is a mathematical term that refers to a component of the partitioning of a space by an equivalence relation; *class* is short for equivalence class. *Niche* is an ecological term that refers to the environmental factors which are favorable to a particular *species*. In our framework, a niche corresponds to a *peak* in the fitness landscape, and a species corresponds to the *subpopulation* of individuals occupying a peak. Finally, the terms *maximum* and *optimum* come from multimodal function optimization, our primary domain of application.

Representative fitnesses

We assume that each class i has a *representative fitness*, f_i, which we define to be the height of the corresponding peak; all elements of i will have identical fitness f_i. Our assumption becomes increasingly valid as runs progress towards equilibrium; from a maintenance standpoint, we are most interested in an algorithm's behavior at equilibrium. It would be possible to relax this assumption so that f_i is the mean fitness of Class i's elements, each class has a certain variance in the fitness of its elements, and fitnesses conform to some predefined distribution. This would require the modelling of an additional source of noise — internal class fitness variance. However, as demonstrated in this thesis, our assumption of identical fitnesses within each class buys a more powerful model, where power is measured by predictive rather than descriptive value.

We require that for all i, $f_i \geq 0$. This requirement of nonnegative fitnesses is already enforced by many GAs. Where objective functions can take on negative values, an appropriate scaling mechanism must transform negative objective function values into nonnegative fitnesses.

Desirable peaks

Some of our models assume we are interested in locating all of the c peaks. Such models correspond to the Type 2 optimization problem of Section 4.3. Other, more general models assume we wish to locate at least the highest b of c maxima. Such models correspond to the most general type of optimization problem defined in Section 4.3 — Type 3. We implement the latter models as follows. First, we assign a fitness threshold. Maxima with fitnesses below the threshold are considered *undesirable* or *extraneous*; maxima with fitnesses above the threshold, *desirable*. We let b be the number of desirable maxima, and c be the total number of maxima. The niching method's challenge is to locate the desirable peaks in the presence of possibly many more, undesirable peaks.

An additional challenge is deception, the combination of isolation and misleadingness. In terms of our modelling framework, deceptive problems can be equated to problems in which undesirable peaks have larger basins of attraction than desirable peaks, and also have fitnesses close to the threshold.

Maintenance of classes

Our models operate from the standpoint of class maintenance rather than class formation — of maintaining niches in a quasi-equilibrium in the population. Therefore, we assume that all desired classes are initially represented in the population. (At time or generation $t = 0$, for all desired classes i, $I_i > 0$.) Note that this is a requirement for our models and not for the actual niching GAs we will be running. If class formation is successful in a niching GA, then as that GA approaches equilibrium, the maintenance perspective becomes increasingly valid.

The process of locating a local maximum is addressed in the standard GA literature for unimodal function optimization. The population size required to locate a single peak is a signal-to-noise issue both within and across schemata (Goldberg, Deb, & Clark, 1992). The process of both locating and maintaining peaks can incorporate estimates based upon both signal-to-noise and class maintenance considerations.

We expect that for most multimodal optimization problems of interest, class maintenance requirements will override class formation (signal to noise) requirements. This has been the case in all experiments conducted to date. We utilize this rule of thumb due both to successful experimentation, and to the intuitive notion that the multimodality dimension of problem hardness will cause significantly more trouble for a niching GA than will the location of any single peak. If, as expected, class maintenance requirements subsume class formation requirements, a population size sufficient to maintain a number of classes will also be sufficient to form them.

Mutation removed

A key simplifying assumption of our framework is the removal of mutation, the GA's local neighborhood operator. We assume, as many authors have previously pointed out, that mutation in combination with selection has a hillclimbing effect. Since our framework incorporates a hillclimber that puts equivalence classes in one-to-one correspondence with local maxima, adding small rates of mutation to a GA should keep its behavior within the scope of our models.

In this thesis, we are primarily interested in the peaks under which niched-selection and crossover deposit points. The selection-plus-crossover combination has done its job if it deposits a desirable distribution of points among the peaks. We assume that selection-plus-mutation or some other type of genetic hillclimbing can move a point to the top of a peak, once the point has successfully been deposited.

Note that we are not ignoring mutation, just separating it out. Our models can hence concentrate on more germane aspects of multimodal function optimization — locating peaks via niched-selection and crossover — and ignore the variation about local points in the search space (typically variation *within* each equivalence class) induced by mutation. In fact, some of our later simulations incorporate mutation, with resulting behavior still within the scope of the corresponding model. In actual implementations of niching methods, the user will typically utilize some small level of mutation to improve performance.

Crossover

Our framework does not assume a particular variation of crossover. Instead, individual models, given a crossover operator and the class membership of both elements undergoing a cross, can assign probabilities that the offspring will be members of various classes.

One assumption that does not vary throughout this study is that crossing two elements of the same class will yield two offspring of that class. This assumption is generally accurate for the partitions of this study, and is universally accurate under standard crossover operators for certain other useful partitions, such as schema-based partitions. Radcliffe (1991a, 1991b, 1993) calls this assumption *respect*. He says that the presence of respect allows a GA to "converge upon good formae [equivalence classes]".

Preliminary models assume no crossover, in order to isolate selection. Such models, nevertheless, have predictive value for niching GAs with crossover (and mutation), on problems in which crossover's (and mutation's) disruptive effects are negligible. Such models also have predictive value for methods such as fitness sharing, whose restorative pressure compensates for minor disruptions. Advanced models incorporate crossover. Actual runs of GAs with crossover, unless otherwise specified, utilize single-point crossover.

Perfect discrimination

The perfect discrimination assumption states that a niching GA is able to determine, without error, the class memberships of all points in the search space. In addition, class membership alone determines the relative proximity of all points in the search space.

The perfect discrimination assumption goes by several names and has several consequences. The names are *perfect discrimination*, *perfect sharing* (Horn, 1993) (for fitness sharing), *nonoverlapping niches* (the partitioning of niches into equivalence classes already ensures they do not overlap), *perfect comparison*, and *noiseless comparison*. One consequence is that an element in a given class, regardless of the comparison or distance measure employed, is always closer to every element of its own class than to any element of another class. A second consequence is that a difference measure, without error, is able to determine whether two elements are members of the same class, and which of two elements is closer to a third element.

For models of fitness sharing, perfect discrimination means that all classes are fully distinguishable via the proper setting of σ_{share} or via a similar technique. In addition, elements of different classes do not contribute to each other's shared fitnesses. Let A and B be arbitrary equivalence classes under our framework, and let a, b, a_1, and a_2 be arbitrary elements of the solution space. If $a \in A$, $b \in B$, and $A \neq B$, then $sh(d(a,b)) = sh(d(b,a)) = 0$. (Recall that sh is the sharing function and d is the distance measure.) Elements within a class contribute

100% to each other's fitness (due to both perfect discrimination and our prior assumption that all elements of a class have identical fitness). Therefore, if $a_1 \in A$ and $a_2 \in A$, then $sh(d(a_1, a_2)) = sh(d(a_2, a_1)) = 1$. Note that if niches were allowed to overlap, neighboring niches would partially contribute to the fitnesses of each other's elements, and elements within a niche would not fully contribute to each other's fitnesses. Perfect discrimination is not a significant restriction, since sharing is most effective in cases of no overlap; as the amount of overlap grows, the effectiveness of sharing diminishes.

Our models do not require niches to be equidistant or spaced in any particular way. We expect our models to be applicable to both cases with prominent niche boundaries and cases with noisier boundaries between niches.

5.3   Empirical Framework

The empirical portion of our modelling framework consists of test problems in multimodal function optimization and classification, as well as default experimental parameters. The test problems allow models of niching GAs to be compared to the GAs themselves. Some test problems serve the purpose of verification: they are designed to very closely match the corresponding model, in order to test the model's validity. Other test problems probe the limits of the corresponding model to determine how well it applies to difficult, artificially constructed problems, and to representative, "average case" problems. Several test problems are used in the process of analysis and design. Finally, some test problems are included for completeness, since prior studies of GAs in multimodal function optimization have established these problems as preliminary benchmarks.

The test problems vary in difficulty along three dimensions — multimodality, isolation, and misleadingness. We consider the isolation and misleadingness dimensions together, as a single dimension of "deception". Table 5.1 summarizes each problem and its level of difficulty. The table lists the number of peaks a problem contains, broken down into desirable and undesirable peaks. It also lists whether the problem is deceptive, misleading, or nondeceptive, as well as whether the problem possesses a user-accessible phenotype.

Before examining the problems in detail, we list the default experimental parameters of our framework. Both phenotypic distance measures (Euclidean distances in variable space) and

Table 5.1: The test problems of the empirical framework are summarized. Highlighted problem characteristics include numbers and types of peaks, as well as degree of deception. Problems are rated according to relative difficulty. An appropriate distance measure is given for each problem.

Problem	Desirable Peaks	Extraneous Peaks	Deception	Difficulty	Distance Measure
M1	5	0	None	Low	Phenotypic
M2	5	0	None	Low	Phenotypic
M3	5	0	None	Low	Phenotypic
M4	5	0	None	Low	Phenotypic
M5	4	0	None	Low	Phenotypic
M6	25	0	None	Medium	Phenotypic
M7	32	5,153,600	Deceptive	High	Genotypic
M8	32	5,153,600	Misleading	High	Genotypic
M9	27	2,170	Deceptive	High	Genotypic
M10	2	0	None	Low	Phenotypic
M11	2	0	None	Low	Phenotypic
M12	4	0	None	Low	Phenotypic
M13	4	0	None	Low	Phenotypic
M14	4	0	None	Low	Phenotypic
MUX-6	4	5	None	Low	Either
PAR-5	16	0	None	Low	Either
PAR-8	128	0	None	Medium	Either
PAR-10	512	0	None	High	Either

genotypic distance measures (Hamming distances) are employed. Table 5.1 lists the appropriate distance measure for each test problem. Problems marked "Either" in Table 5.1 are served equally well by either a genotypic or a phenotypic measure; for such problems, genotypic distance is used by default. For simulations in which each class consists of exactly one element, phenotypic and genotypic measures have equivalent effects.

Initialization of GA populations is uniformly random, unless otherwise specified. For sharing, the default α is 1. (Recall that α is an exponent that determines the shape of the sharing function.) All variable encodings are binary (as opposed to Gray [Caruana & Schaffer, 1988] for instance) and tight (meaning that variables are concatenated — as opposed to loose or random [Goldberg, Deb, & Korb, 1990]).

Note that certain trivial problems are used throughout this thesis which are not mentioned in this chapter. These include one-bit problems and example problems. These trivial problems are easy to understand, and are described as they occur, in the accompanying text.

M1–M4: Sine functions

The first test problems *M1–M4* are one-dimensional, five-peaked, sinusoidal functions, and are shown in Figures 5.2–5.5. Similar functions were first used by Goldberg and Richardson in 1987, then in two subsequent studies (Deb, 1989; Deb & Goldberg, 1989). These four functions have subsequently been adopted as preliminary benchmarks for niching GAs (Beasley, Bull, & Martin, 1993). Despite the simplicity of these functions, most potential niching GAs have in the past had trouble locating and maintaining all five peaks. In fact, traditional GAs have been shown to rapidly converge upon only one of the five peaks (Goldberg & Richardson, 1987). We include these functions in our empirical framework for two reasons: to establish a preliminary litmus test for niching methods and to compare a few of our results with those of prior studies. The single variable x in *M1–M4* is restricted to the real-valued range $[0, 1]$ and is encoded using 30 bits. It is decoded through normalization of the resulting 30-bit, unsigned, binary integer, through division by $2^{30} - 1$.

The first function, *M1*, consists of equally spaced peaks of uniform height. Maxima are located at the x values of .1, .3, .5, .7, and .9. All maxima are of height 1.0. *M1* is defined below and displayed in Figure 5.2.

$$M1(x) \;=\; \sin^6(5\pi x) \;\;.$$
(5.2)

M2 consists of equally spaced peaks of nonuniform height. Maxima are located at x values of .1, .3, .5, .7, and .9. Maxima are of rounded height, 1.000, .917, .707, .459, and .250, respectively. *M2* is defined below and displayed in Figure 5.3.

$$M2(x) \;=\; e^{-2(\ln 2)\left(\frac{x-0.1}{0.8}\right)^2} \sin^6(5\pi x) \;\;.$$
(5.3)

M3 consists of unequally spaced peaks of uniform height. Maxima are located at x values, to three decimal places, of .080, .247, .451, .681, and .934. All maxima are of height 1.0. *M3* is defined below and displayed in Figure 5.4.

$$M3(x) \;=\; \sin^6(5\pi[x^{0.75} - .05]) \;\;.$$
(5.4)

Figure 5.2: Function *M1* is displayed.

Figure 5.3: Function *M2* is displayed.

Figure 5.4: Function *M3* is displayed.

Figure 5.5: Function *M4* is displayed.

Figure 5.6: Modified Himmelblau's function *M5* is displayed.

M4 consists of unequally spaced peaks of nonuniform height. Maxima are located at x values, to three decimal places, of .080, .247, .451, .681, and .934. Maxima are of rounded height, 1.000, .948, .770, .503, and .250, respectively. *M4* is defined below and displayed in Figure 5.5.

$$M4(x) = e^{-2(\ln 2)(\frac{x-0.08}{0.854})^2} \sin^6\left(5\pi[x^{0.75} - .05]\right) \ . \tag{5.5}$$

M5: Himmelblau's function

M5, shown in Figure 5.6, is the modified *Himmelblau's* function of Deb's (1989) study. Himmelblau's function is two dimensional, with four peaks of identical height. Like *M1–M4*, it has been used in prior studies of niching GAs. We include this function for consistency with prior studies, and as a preliminary multidimensional benchmark. The variables x and y of Himmelblau's function are both restricted to the real-valued range $[-6, 6]$ and are encoded using 15 bits per variable. The GA operates upon a 30-bit string, which it decodes by first splitting into halves — into components corresponding to x and y — and then transforming each resulting binary integer into a real number on the interval $[-6, 6]$.

Figure 5.7: Modified Shekel's Foxholes function *M6* is displayed.

Maxima in *M5* are located at the $\langle x, y \rangle$ coordinates of $\langle 3.000, 2.000 \rangle$, $\langle 3.584, -1.848 \rangle$, $\langle -3.779, -3.283 \rangle$, and $\langle -2.805, 3.131 \rangle$. All maxima are of height 1.0. *M5* is defined below.

$$M5(x, y) = \frac{2186 - (x^2 + y - 11)^2 - (x + y^2 - 7)^2}{2186} \quad .$$

(5.6)

Most prior studies of GAs in niching and multimodal function optimization do not move beyond simple problems such as *M1–M5*. This chapter, in most of the problems that follow, enters the territory of more complex — and more interesting — problems.

M6: Shekel's Foxholes

M6, shown in Figure 5.7, is the *Shekel's Foxholes* problem from De Jong's (1975) dissertation. *M6* is a two-dimensional function with 25 peaks. It has been used in prior studies, but mainly for the purpose of locating the single global optimum. *M6* is harder to solve than its predecessors,

because of its higher modality. It is not deceptive, however. The variables x and y of *M6* are both restricted to the real-valued range $[-65.536, 65.535]$ and are encoded using 17 bits per variable. The GA operates upon a 34-bit string, which it decodes by first splitting into x and y components, and then transforming each resulting binary integer into a real number on the interval $[-65.536, 65.535]$.

Maxima in *M6* are located at the $\langle x, y \rangle$ coordinates of $\langle 16i, 16j \rangle$, where i and j represent all integers in $[-2, 2]$. The 25 maxima are all of differing heights, ranging from 476.191 to 499.002. The global optimum occurs at $\langle -32, -32 \rangle$. The other optima form a staircase of spikes to the global optimum. *M6* is defined below.

$$M6(x, y) \;=\; 500 - \frac{1}{.002 + \sum_{i=0}^{24} \frac{1}{1+i+(x-a(i))^6+(y-b(i))^6}} \;, \qquad (5.7)$$

where $a(i) = 16[(i \bmod 5) - 2]$ and $b(i) = 16(\lfloor i/5 \rfloor - 2)$.

M7–M9: Massively multimodal / deceptive problems

M7 is the massively multimodal, deceptive function of Goldberg, Deb, and Horn (1992). Overall fitness is the sum of the fitnesses of five subfunctions. Each subfunction is a bimodal, deceptive function of unitation, as displayed in Figure 5.8. (Functions of unitation are explained in Section 3.1.) The GA operates upon a 30-bit string, formed from the concatenation of the 6-bit x values of each subfunction. In *M7*, the order of difficulty or deception is 6 bits. The total number of optima is 5,153,632, of which 32 are global and hence considered desirable. Subfunction global maxima are located at $x = 000000$ and $x = 111111$, and have fitness 1.0. Subfunction nonglobal maxima are located at the 20 x values containing exactly 3 ones, and have fitness .640576.

M8 is the same as *M7*, but exponentially scaled to create larger differentials between the fitnesses of global and nonglobal optima. The scaling function used is

$$M8 \;=\; 5 \left(\frac{M7}{5} \right)^{15} \;, \qquad (5.8)$$

where *M7* in the equation represents the value for *M7* over the entire 30-bit function. While *M8* is still massively multimodal and misleading, it is no longer deceptive.

114

Figure 5.8: *M7* is composed of five of these bimodal, deceptive subfunctions.

M9 is a *minimum distance* function (Horn & Goldberg, in press). Overall fitness is the sum of the fitnesses of three 8-bit subfunctions. The GA operates upon a 24-bit string, formed from the concatenation of the 8-bit x values of each subfunction. Each subfunction is maximally deceptive, and defined directly over Hamming space. Subfunction global optima occur at the arbitrarily chosen points, 00000000, 10001100, and 01001010. The fitness of any nonglobal, 8-bit substring is its Hamming distance from the nearest global substring. Global substrings receive a fitness of 10. *M9* has 2197 local optima, of which 27 are global and hence considered desirable. It is a difficult problem, because the further a substring is from all global substrings, the more fitness that substring receives, thus isolating the global optima, and placing most points in the search space on a misleading path towards a nonglobal optimum. *M9* is similar in structure to the one-dimensional, phenotypically defined, minimum distance function depicted in Figure 5.9.

M10–M14: Constructed problems

Problems *M10–M14*, displayed in Figures 5.10–5.14, are constructed so that classes correspond to both peaks and schemata. In *M10* and *M11*, the most significant (leftmost) bit of each

115

Figure 5.9: This one-dimensional, phenotypically defined, minimum-distance function is similar in structure to *M9*. Two isolated global optima occur at $x = .25$ and $x = .75$. Almost all points in the search space lead to three local optima at $x = 0$, $x = .5$, and $x = 1$.

population element indicates the schema or class to which it belongs. *M10* has two peaks of equal height. In *M11*, the peak corresponding to Class *A* is four times the height of the peak corresponding to Class *B*. In both *M10* and *M11*, chromosomes are eight bits in length, and decode from binary values to single-integer phenotypes ranging from $x = 0$ to $x = 255$. Chromosomes with a '0' in the leftmost position ($x \in [0, 127]$) are considered members of Class *A*, while chromosomes with a '1' in the leftmost position ($x \in [128, 255]$) are considered members of Class *B*.

M12–M14 contain four peaks each. In all three problems, the two most significant (leftmost) bits of the eight-bit chromosome determine class membership. Class *A* individuals have '00' in the most significant bits ($x \in [0, 63]$); Class *B*, '01' ($x \in [64, 127]$); Class *C*, '10' ($x \in [128, 191]$); and Class *D*, '11' ($x \in [192, 255]$). In *M12*, all four peaks are of identical height. In *M13*, Class *A* is twice as fit as the other three. In *M14*, Classes *A* and *D* are each twice the fitness of *B* and *C*.

Figure 5.10: Function *M10* is displayed.

Figure 5.11: Function *M11* is displayed.

Figure 5.12: Function *M12* is displayed.

Figure 5.13: Function *M13* is displayed.

Figure 5.14: Function *M14* is displayed.

Classification problems

The remaining problems are covering problems — idealized classification problems. Given a set of positive and negative training examples, the objective is to find a concept description that includes all of the positive examples, but none of the negative examples.

We present a new technique for solving classification problems, that is based upon the niching method. We will illustrate this new technique on the forthcoming boolean classification problems. Our technique can be extended to solve complex, real-world, classification and machine-learning problems. It maps classification problems to multimodal optimization problems by utilizing the entire population as a disjunctive-normal-form (DNF) concept description, in which each population element represents a disjunct. A DNF concept description is a disjunction of conjunctions, which takes the following form:

$$(A_1 \wedge A_2 \wedge \ldots \wedge A_a) \ \vee \ (B_1 \wedge B_2 \wedge \ldots \wedge B_b) \ \vee \ \ldots \ \implies \ C \ . \tag{5.9}$$

In the case of boolean concepts, A_i and B_i represent boolean variables or negated boolean variables, and C is either '0' or '1'.

119

The niching GA must form and maintain a set of optimal disjuncts. (In the above example, $(A_1 \wedge A_2 \wedge \ldots \wedge A_a)$ is a disjunct.) We assign fitnesses to individual disjuncts based on the number of positive examples they cover (POS), and the number of negative examples they cover (NEG). In general, disjuncts receive higher fitnesses for covering positive examples, and lower fitnesses for covering negative examples. Given NTX negative examples in the set of training examples, our fitness function is of the form,

$$f(POS, NEG) \;=\; \begin{cases} 1 + POS \;, & \text{if}\;\; NEG = 0 \;; \\ 1 - \frac{NEG}{NTX} \;, & \text{otherwise} \;. \end{cases} \qquad (5.10)$$

Such a fitness function is shown in Figure 5.15. Note that the fitness function in Figure 5.15 is not shown over the usual variable space (the variable space is typically of high dimensionality), but over the space defined by both the number of positive examples covered by the disjunct, and the number of negative examples covered by the disjunct.

Individuals in the population are composed of a concatenation of two-bit values that represent boolean variables. A '00' value corresponds to a boolean variable set to '0' or *off*; '11', to a boolean variable set to '1' or *on*; '01' and '10', to a boolean variable whose value does not matter (a wild-card or "don't care" symbol). A *repair* mechanism (Orvosh & Davis, 1994) is employed to effectively reduce the size of the search space. The repair mechanism flips all '10' alleles, upon fitness assignment, to '01' alleles.

We consider two types of boolean covering problems that are extensively used in the machine-learning literature — parity problems and multiplexer problems. Parity problems represent a category of boolean classification problems that are maximally hard in a sense: from an optimization point of view, a k-bit parity problem contains the largest number of optima of any k-bit boolean concept. Multiplexer problems represent average-case problems.

Our empirical framework includes three different odd-parity problems, which we call *PAR-5*, *PAR-8*, and *PAR-10*. The number in the problem name represents the number of boolean variables that the problem contains. In odd-parity problems, if an odd number of variables are on, the example is positive; if an even number of variables are on, the example is negative. Under our chosen DNF representation, the solution to a k-bit parity problem requires the formation and maintenance of 2^{k-1} disjuncts, because each positive example must be covered individually. (Any example left uncovered is considered negative.) Hence, *PAR-5* requires the

Figure 5.15: This function assigns fitness to a classification rule based upon the number of positive and negative training examples that the rule covers.

formation and maintenance of 16 disjuncts; *PAR-8*, 128; and *PAR-10*, 512. In *PAR-8* and *PAR-10*, the number of desirable solutions (that the GA must form and maintain) is greater than for any other problems in this thesis.

Our empirical framework also includes a 6-bit multiplexer problem, *MUX-6*. A multiplexer contains a number of address bits and a number of data bits. The address bits determine the data bit that is selected. *MUX-6* contains two address bits that specify which of four data bits is selected. For example, if both address bits are on, they form the binary address '11', which designates data-bit number 3. Subsequently, if Data-bit 3 is on, the example is positive; if Data-bit 3 is off, the example is negative.

Chapter 6

Crowding: Selection

A crowding method, according to our earlier definition, is a selection scheme that inserts a new element into the population by overwriting a similar element. This strategy of replacing similar population elements appears, intuitively, to be a solid foundation upon which an effective niching method can be built. However, 17 years after De Jong (1975) introduced crowding, no one had yet been able to get it to consistently maintain more than two peaks of a multimodal function, no matter how simple the function (Deb, 1989; Deb & Goldberg, 1989).

Prior to De Jong, Cavicchio (1970) had introduced several schemes in which children directly replaced the parents that produced them. Since parents are similar to their children, these so-called *preselection* schemes qualify as crowding methods, according to our definition. Again, intuitively it would seem that parental replacement techniques should be able to preserve representatives of multiple peaks. However, like De Jong's crowding, preselection schemes had never exhibited niching capabilities.

In this chapter, we repeat our prior analysis (Mahfoud, 1992) of De Jong's crowding and Cavicchio's preselection to determine why they are not effective niching methods, in order to design an improved crowding scheme that is an effective niching method. We seek to design an improved crowding method because it seems a plausible goal. Plausibility, as suggested earlier, is a partially intuitive notion. As additional motivation, De Jong's and Cavicchio's descriptions of crowding and preselection give the impression that these algorithms should be successful at preserving population diversity. A third motivating factor is that natural populations are effective at maintaining separate species through methods similar to crowding.

To determine why, in practice, that De Jong's crowding and Cavicchio's preselection are not successful, we postulate that one or more factors are at work, combining to create significant noise or variance in the selection process, ultimately leading to *genetic drift* — the loss of alternative solutions due to random fluctuation. As we saw in earlier chapters, even in the absence of selection pressure, a simple GA's population will eventually converge due to selection noise. It turns out that the main culprit in crowding methods is the stochastic error in the replacement of population members. We can combat genetic drift in crowding methods by introducing design alternatives that reduce or eliminate the replacement error.

When introducing potential design improvements to a GA, one should keep in mind that the design space can be highly nonlinear. In prior GA research, the result of this nonlinearity has often been that well intended design changes have had unforeseen consequences, leading to either poorer performance than the unmodified algorithm, or better behavior on one type of problem but worse behavior on others. When altering a design, it is important to pursue a clear goal in terms of desired algorithmic performance. Because the design space is nonlinear, it may be necessary to take steps backwards (i.e., accept poorer performance after intermediate changes) if such backward steps in some way lead towards the goal. Since the major premise underlying crowding methods is to replace similar elements, we define our goal in terms of the replacement process. Our primary criterion, the reduction of replacement error, guides us in making algorithmic changes. Sometimes this criterion takes us temporarily backward, and sometimes it takes us sideways, but in the end, forward progress produces a much improved crowding method.

The remainder of this chapter starts with De Jong's crowding, and through a series of tests and modifications, develops an effective form of crowding similar to Cavicchio's preselection. The resulting algorithm is named *deterministic crowding*. Under the modelling framework of the previous chapter, we employ an equivalence-class model for crowding that concentrates upon the replacement step. The model defines correct and incorrect replacements, allowing us to use the number of replacement errors as a performance and design criterion.

We include the test functions *M1* and *M2* in this chapter's model of crowding. These functions are complex enough to make prior crowding methods fail (Deb, 1989; Deb & Goldberg, 1989), but simple enough to allow full analysis and visualization of class distributions and stochastic errors. The functions are defined and displayed in Chapter 5.

6.1   Performance Criteria

In evaluating the performance of his genetic algorithms, De Jong (1975) operated under the assumption of unimodal function optimization. The ability to quickly adapt the population toward optimal regions of the search space was the primary goal. This resulted in the use of two performance measures: on-line and off-line performance. On-line performance is defined as the average fitness of all past and present population elements. Off-line performance is defined as the average fitness of a collection consisting of the fittest element of each past population and the fittest element of the present population. De Jong also employed a secondary performance measure for his crowding scheme, the number of bit-positions that had become fixed (i.e., were identical in each population element).

The above performance criteria are not of much use in our study of niching. Both on-line and off-line performance measures are overly concerned with speed of convergence, rather than quality of final results. Off-line performance is actually a second type of on-line performance measure, since current results are heavily degraded by past ones. In the current study, we are more concerned with maintaining optima and minimizing noise. Speed is a secondary consideration that, as later becomes apparent, need not be sacrificed. The number of fixed bit-positions, though not used as a criterion to develop algorithms, is nonetheless measured, and is discussed in a later section.

The two related performance criteria we use in this study are the number of peaks or classes an algorithm maintains and the number of replacement errors it makes. A peak is considered maintained if at least one population element (that has a fitness of at least 80% of the peak's height) exists in the basin of attraction of that peak. The basins are easy to visualize for *M1* and *M2*. They are represented by intervals of width .2 along the x axis at the base of each peak, and they are defined using a phenotypic neighborhood operator of size $\epsilon = 1/(2^{30} - 1)$, the minimum increment in our representation of x. *M1* and *M2* are simple enough so that the quality of final solutions will tend to be similar across algorithms. We are hence more interested in the quantity of diverse, but good, final solutions; this quantity corresponds to the number of peaks maintained, one of our performance criteria.

The number of replacement errors constitutes our other performance criterion. We define a replacement error as the replacement of a member of one class or peak by a member of another.

The primary goal of this chapter is to develop crowding methods that minimize the number of replacement errors and that, as a result, maintain the maximum number of peaks.

6.2 Algorithms and Results

Six crowding algorithms are tested. Algorithm 1 represents De Jong's crowding algorithm, with typical parameter settings. Algorithms 2–5 are cumulative attempts to eliminate unnecessary parameters and to reduce the number of replacement errors. Algorithm 6 is an improved version of preselection that approximates the desirable properties of De Jong's crowding. As we soon demonstrate, our algorithmic changes are definitely not linear in design space.

All algorithms are tested 100 times on both $M1$ and $M2$, using a population size of $n = 100$ and no mutation ($p_m = 0$). Populations are randomly initialized, and then the GA runs for $20,000$ function evaluations. We implement De Jong's crowding using roulette-wheel selection, and employ sampling without replacement to choose individuals for the replacement sample.

Figures at the end of the chapter illustrate performance of the six algorithms. Figures 6.1 and 6.2 show the distribution of population elements after a sample run of each algorithm on the test problems. Figure 6.3 shows the number of peaks each algorithm maintains over time, averaged over 100 runs. Figure 6.4 shows the cumulative number of replacement errors as each algorithm progresses, averaged over 100 runs. Figure 6.5 displays the number of diverse bit-positions each algorithm maintains over time, also averaged over 100 runs.

Algorithm 1: De Jong's crowding; Deb's parameters

Algorithm 1 (A1) is our starting point. It represents the variation of De Jong's crowding used by Deb (1989). Crossover probability is $p_c = .9$, generation gap is $GG = .1$, and crowding factor is $CF = 3$. (De Jong's crowding, along with its generation gap and crowding factor parameters, is described in Chapter 4.)

A1 maintains two peaks each run, on both test functions (see Figures 6.1–6.3). This is consistent with Deb's results, where crowding nearly always maintains exactly two peaks on every test function. Arriving at an explanation for this behavior is not overly difficult. A1 makes an average of 6600 replacement errors on $M1$ and an average of 5900 replacement errors on $M2$ (see Figure 6.4). Since crowding does a replacement for each function evaluation, nearly

one out of every three replacements is in error. This is an open invitation to genetic drift, so it is apparent why A1 maintains no more than two peaks.

With such high replacement error, should A1, instead, be maintaining only one peak? Let us examine, under selection alone, the mechanism by which an element is lost from one peak and gained on another. Suppose an element from peak A is awaiting insertion in the population. Crowding chooses CF candidates from the population and replaces the closest one. Assuming crowding makes no errors in comparison (an assumption of our modelling framework), an element of peak A will be replaced if it is among the CF. If an element of A is not among the CF, an element from some other peak will be replaced. As long as all peaks other than A together contain CF or more elements, these other peaks are vulnerable to loss of members. However, when peak A contains $n - 1$ of the population's elements, and peak B contains the final one, peak B will not lose that element, since for $CF \geq 2$, an element of A will be present among the CF. Note that crowding with $CF = 1$ is effectively a simple GA, in which all peaks but one will eventually lose all representatives.

Algorithm 2: De Jong's crowding; Steady state with full crossover

Algorithm 2 (A2) is a simplification of A1. The intent of A2 is to get rid of parameters whose fine-tuning detracts from the thorough analysis of crowding and its behavior. We eliminate crossover probability as a design parameter by raising p_c all the way to 1, the worst case (maximally disruptive) setting. An effective crowding method should be able to operate under full crossover.

Generation gaps that are too high result in the replacement of too much of the population, with behavior approaching that of a simple GA. In fact, crowding with $GG = 1$ is a simple GA. De Jong himself noted that crowding becomes less effective as GG increases. On the other hand, he avoided generation gaps less than .1, probably to allow a limited amount of parallelism. It seems a better design approach, however, to first worry about achieving niching and then about secondary considerations such as parallelism. Such a design approach previously led to the successful development of fitness sharing (Goldberg & Richardson, 1987). We consider parallelism later.

A2 is identical to A1, except that $p_c = 1$ and $GG = 1/(2n) = .02$. GG is set to the lowest possible value that still allows the binary operator, crossover, to function. Note that

126

the setting for *GG* effectively results in a steady-state crowding method. (A steady-state GA (Syswerda, 1989, 1991) processes few population elements — as few as one — each generation.) As expected, A2 produces results nearly identical to those produced by A1 (see Figures 6.1–6.5), except for a slightly greater number of replacement errors. The additional errors are due to the higher variance of the steady-state GA and to the increased disruption caused by additional crossover.

Algorithm 3: Phenotypic comparison

Having fixed most parameters, we can concentrate on reducing the number of replacement errors. In the analysis of A1, we assumed perfect comparison. This is equivalent to the perfect discrimination assumption of our modelling framework. Perfect comparison requires that the algorithm never make an error when deciding which of the *CF* elements is closest to the one being inserted. For the binary-encoded strings of this study, bitwise or genotypic comparison is far from ideal. Only the three most significant bits of each string are likely to be useful for distinguishing different peaks of *M1* and *M2*. Genotypic comparison is hence trying to detect roughly 3 bits of signal in the presence of approximately 27 bits of noise.

One solution is the phenotypic comparison measure outlined in the original study of fitness sharing (Goldberg & Richardson, 1987). Phenotypic comparison uses decoded variables rather than raw, binary strings. For *M1* and *M2*, the decoded variable is the real number x. Although Deb (1989) tried phenotypic sharing, he did not attempt phenotypic crowding. Another potential solution to noisy comparison (that we will not examine in this study) is genotypic comparison in combination with Gray codes. Gray codes (Caruana & Schaffer, 1988) ensure that variables which are one bit apart in binary integer space are also one bit apart in Hamming space.

Algorithm 3 (A3) is identical to A2, except that it employs phenotypic rather than genotypic comparison. As shown in Figures 6.1–6.3, A3 maintains two to three peaks, a marked improvement. The number of replacement errors also drops significantly (Figure 6.4).

Algorithm 4: Full sampling

A major source of error must still be eliminated: the sampling error due to low *CF*. De Jong noted that while higher crowding factors fix fewer bit-positions, they also degrade on-line and

off-line performance. In addition, higher crowding factors limit parallelism. Nevertheless, we are not yet concerned with speed or parallelism, but rather with maintaining maxima by reducing replacement error.

It is not difficult to calculate the CF required to perform correct replacement with some specified probability, if one makes a few assumptions regarding the fitness function and the initial distribution of population elements. Algorithm 4 (A4) strives to perform 100% of its replacements correctly: it examines every population element to find the closest match.

A4 is identical to A3, except that $CF = n$. Although this change adds an order of complexity to crowding (requiring $O(n^2)$ rather than $O(n)$ comparisons per generation), it produces a useful, intermediate algorithm. A4 is similar in complexity to fitness sharing (Goldberg & Richardson, 1987), which also cycles through the entire population (to compute each element's shared fitness). Goldberg and Richardson note that a version of sharing which estimates shared fitnesses using population sampling (as De Jong's crowding does) is a promising possibility. However, they first design sharing without sampling, in order to determine the ideal behavior of their algorithm in the absence of additional sources of noise. Although Deb (1989) compares genotypic crowding using sampling to phenotypic sharing without sampling, no one has yet compared equivalent versions of the two algorithms.

The percentage of replacement errors in A4 drops to nearly zero (about 0.1%). Errors are still made a fraction of the time, when the phenotypically closest element is near a boundary between peaks. A4 consistently maintains population elements at all five peaks of the search space (Figures 6.1–6.3). However, A4 also distributes population elements all across the search space, including inside the low-fitness valleys of $M1$ and $M2$. It is apparent that A4 does not exhibit any selection pressure — that it locates peaks due to the random, initial distribution of population elements. This absence of selection pressure can be explained as follows. While A4 chooses two good solutions for reproduction each generation, it also overwrites the most similar elements of the population, whether they are better, worse, or equally fit. Since A4 replaces the population element that most closely resembles the one being inserted, it is highly likely that A4 will replace a solution of nearly identical fitness. Therefore, even though highly fit elements are more frequently chosen by up-front selection, they are just as often replaced, resulting in a cancellation of the up-front selection pressure.

Runs of A4 on a one-dimensional, strictly increasing, linear function confirm that the algorithm exhibits little, if any, selection pressure: A4 maintains population elements at seemingly random positions up and down the line. A4's behavior is not surprising if one examines the method by which traditional crowding (with low CF) converges. Crowding selects elements for reproduction according to fitness. If the closest element in the population were consistently replaced, crowding would very likely replace an element similar in fitness, resulting in little or no improvement (as A4 demonstrates). Instead, if a certain percentage of the time, a dissimilar element were replaced in error, an improvement would be likely since the inserted element would probabilistically be of superior fitness. This surprising discovery leads to the conclusion that De Jong's crowding relies on replacement errors for its convergence.

Algorithm 5: Parental replacement

A curious thing happens when we raise CF to n. A high percentage of the time, one or both parents are replaced. On *M1* and *M2*, a parent is replaced 83% of the time. This suggests a method of approximating De Jong's crowding using a technique similar to Cavicchio's preselection: replace each parent with its closest offspring, making sure that the same offspring does not replace both parents. This method has the pleasant side effect of reducing algorithmic complexity to the level of the simple GA, since the previously required $O(n^2)$ comparisons reduce to $O(n)$. In addition, the CF parameter disappears, since a parent rather than an element of a sample is replaced.

Algorithm 5 (A5) is identical to A4, except that instead of sampling the whole population to find the phenotypically closest element, it samples only the two parents. Each pair of offspring is simultaneously inserted into the population as follows. There are two possible methods of replacing two parents with their two offspring: Offspring 1 replaces Parent 1 and Offspring 2 replaces Parent 2, or Offspring 1 replaces Parent 2 and Offspring 2 replaces Parent 1. The pair of replacements is employed that yields the smallest sum of absolute distances between an offspring and a replaced parent. Even though the end result is the same for either pair of replacements, the dichotomy allows flexibility in further modifications of crowding.

Although the number of replacement errors remains very low (0.7%), most solutions have gravitated toward the valleys rather than the peaks of the two functions. The algorithm actually exhibits very slight, reverse selection pressure, with cumulative effects. Somewhat surprising

at first, this behavior can be explained as follows. Recall that De Jong's crowding chooses two elements for reproduction according to fitness. While in A4, these chosen elements might or might not be replaced by their offspring, in A5 these elements are always replaced by their offspring. A5 is in effect killing off the best elements of the population, and replacing them by elements that, although similar, are more often worse than better. This results in the slight downward selection pressure displayed. A4 does not display such downward pressure, due to the small percentage of the time (about 17%) that a parent is not the closest element in the population. Neither A4 nor A5 displays any observable, upward selection pressure.

Algorithm 6: Deterministic crowding

What remains in our design process is to add selection pressure to either A4 or A5. We choose A5 because of its lower time complexity. Note that A5 resembles Cavicchio's preselection, except that parents are always replaced, regardless of offspring fitness, and a similarity measure is used to decide which of the two offspring replaces which parent. Results have shown that replacing parents arbitrarily or replacing only the worse parent (as preselection does) result in too many replacement errors to achieve niching.

Since the fitness-proportionate selection of A4 and A5 results in little to no favorable selection pressure, A6 dispenses with it. Instead, A6 pairs all population elements randomly each generation, without replacement. This means that each population member undergoes crossover once each generation. To add back selection pressure, A6 only replaces a parent if the competing offspring is better, as in preselection. A6 differs from preselection in that it processes two parents and two offspring at a time, does not utilize up-front selection, and employs a similarity measure to determine which offspring competes against which parent. Selection upon replacement (rather than up front) is known as *replacement selection* and has previously been employed in GAs (Eshelman, 1991; Whitley & Kauth, 1988).

A6 has the additional effect of adding the capacity for true parallelism, since all population elements can proceed simultaneously each generation. The algorithm is also simpler than its predecessors, since the generation gap (GG) parameter is no longer necessary. As shown in Figures 6.1–6.3, A6 successfully clusters solutions about all five peaks of both *M1* and *M2*. In Figure 6.1, it might appear that A6 is only maintaining a few elements on the central peak; however, upon closer examination we find that this is not the case: the central peak has a

130

significant share of the 100 population elements; they have just converged to two very close values at the peak's tip. As in A4 and A5, replacement errors are almost nonexistent (0.2%). We call A6 *deterministic crowding* (DC for short), because it makes several of the random processes in De Jong's crowding algorithm, deterministic.

It is illustrative to summarize the properties of deterministic crowding. The algorithm incorporates a similarity measure like that of De Jong's crowding, but replaces parents like preselection. Unlike De Jong's crowing, DC may utilize either genotypic or phenotypic difference measures, with phenotypic measures being preferable when they are available. On *M1* and *M2*, DC virtually eliminates replacement error and locates and maintains all five peaks. The algorithm is simple to implement, is faster than its predecessors (since it performs replacement selection via binary tournaments and does not sample the population for either selection or replacement), and is inherently parallel. Because of its preliminary success and high level of stability, DC is used as the representative crowding method in the remainder of this study. The full algorithm is shown in pseudocode form in Figure 7.1 of the next chapter.

6.3 Bitwise Diversity

The maintenance of bitwise diversity is a questionable goal. Let us examine an instance of why this is so. Figure 6.5 shows that the traditional crowding techniques, A1 and A2, converge at a few bit-positions; this is consistent with De Jong's results. A3, which uses phenotypic comparison, loses a large majority of its bits. With phenotypic comparison, difference in Hamming space is not important, but difference in decoded parameter space is. In our prior discussion of diversity, we stated that an algorithm which promotes useful diversity will allow a population to converge along all dimensions of diversity, except those dimensions related to solving the current problem. Losing bits while maintaining phenotypic spread is a prime example.

A4 regains a lot of bitwise diversity due to its near elimination of replacement error. A5 and A6 seem to be anomalies, since they never lose a single bit. Let us take a closer look at one point to which A6 has converged. For both *M1* and *M2*, the central optimum (at $x = .5$) resides on a Hamming cliff, with one neighboring solution containing 29 zeros, and another containing 29 ones. These neighbors are, in fact, full bitwise complements. It is not surprising that no bits

are lost, since the final population contains both neighboring solutions. This form of bitwise diversity does not appear to be very useful.

The primary argument for maintaining bitwise diversity is that such diversity will somehow prevent premature convergence. However, even with full bitwise diversity, premature convergence may still occur. What one would like is a mechanism for controlling convergence. One possibility is the addition of thermodynamic controls. (See Section 3.5 for a discussion of thermodynamic GAs.) Briefly, thermodynamic controls would augment deterministic crowding to allow offspring to replace better parents on occasion. The probability of replacement would be determined by a temperature parameter, equivalent to that of simulated annealing. The resulting algorithm would be a variation of *parallel recombinative simulated annealing* (Mahfoud & Goldberg, 1992, 1995).

6.4 Further Research

Since the introduction of deterministic crowding, two similar crowding methods (detailed in Chapter 4) have been proposed (Cedeño & Vemuri, 1992; Harik, 1994). These two methods have demonstrated preliminary success at optimizing simple multimodal functions. The designs of these new crowding methods have verified one of our design results for DC — that in order for crowding to be effective at niching, up-front selection must be eliminated, and replacement selection incorporated. We expect the distributional behaviors for members of the crowding family of niching GAs to be similar, although some crowding methods may maintain niches more stably than others. (Chapter 7 examines the distributional behavior of DC.) A comparative analysis of crowding algorithms is left to future research.

Two studies have recently employed deterministic crowding to solve applications problems. The first study (Hatjimihail, 1993) applies DC to designing and optimizing statistical, chemical quality-control procedures. Using DC, the author is able to obtain quality-control procedures equivalent to, and often better than, those procedures produced by current algorithms. The second study (Pál, 1994) applies DC and other GAs to the spin-glass problem. He shows that DC outperforms several simple GAs and that a modified version finds optimal or nearly optimal solutions. In both studies, DC is used to prolong GA convergence in order to locate better, single solutions, rather than multiple solutions. This supports our earlier conjecture, that

niching methods designed to maintain multiple solutions should also be effective at maintaining multiple subsolutions, on the way to a better, single solution.

Ongoing research examines the incorporation of deterministic crowding and other niching methods in machine learning systems. We demonstrate later in this thesis the rudiments of such a system, operating on boolean covering or classification problems. The extension of such a system has yielded promising initial results on real-world, machine-learning problems.

We make one final note regarding parental replacement methods. In parental replacement, the method of competition between children and parents is a critical step, whereby a small algorithmic change typically yields an enormous change in the algorithm's capability and behavior. For this reason, most previously suggested parental replacement techniques do not qualify as niching methods. As an example, Cavicchio's preselection throws away a randomly chosen offspring of the two produced from each cross. The retained offspring competes against both parents. Due to replacement errors, his algorithm is incapable of maintaining multiple peaks. Deterministic crowding, on the other hand, maintains multiple peaks by retaining both offspring, and forcing the pair of offspring-parent competitions that results in the lowest sum of distances.

Holding a *double acceptance/rejection* competition (Mahfoud & Goldberg, 1992, 1995) is an interesting possibility, in which either both children win the competition as a pair, and are accepted into the population of the next generation, or both parents win it. Culberson (1992) employs such a scheme in his GIGA system, where the fitness of a pair of elements is the maximum fitness over both individuals in the pair. The result of double acceptance/rejection, in the absence of up-front selection and mutation, is complete conservation of alleles; conservation of equivalence class representatives, however, is not guaranteed.

In the chapter that follows, we will further analyze the behavior of deterministic crowding, emphasizing the combination of selection and crossover. We will observe and explain several interesting behaviors that are characteristic of the crowding family of niching methods and to some extent, of all GAs that employ crossover.
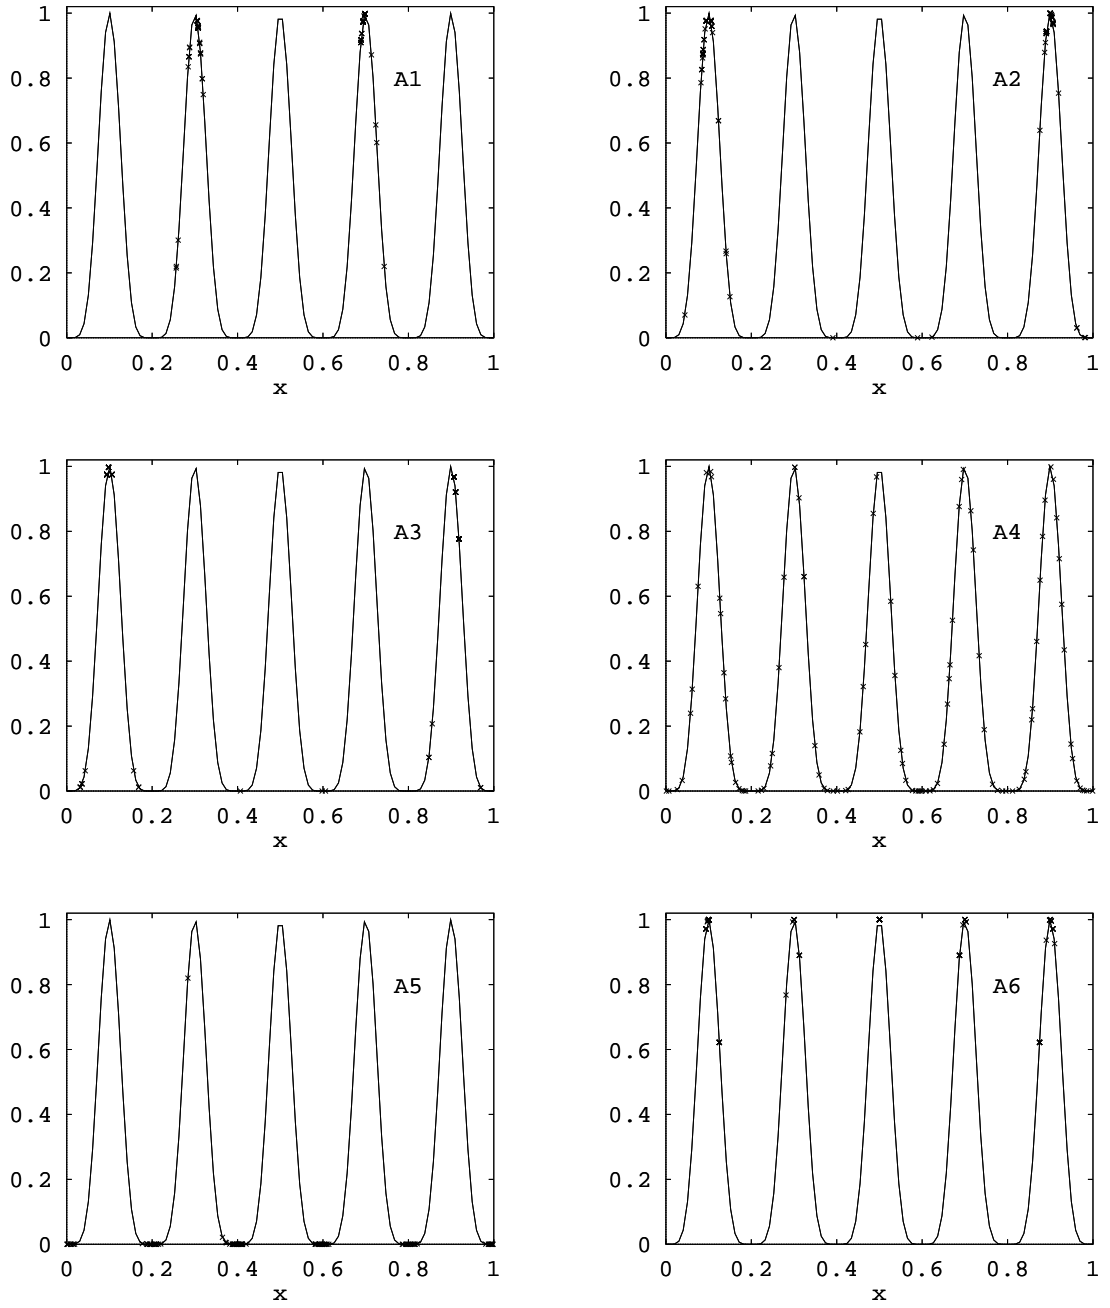
Figure 6.1: The final distribution of 100 population elements is shown for one run of each algorithm, A1–A6, on function *M1*. Each algorithm runs for 20,000 function evaluations.
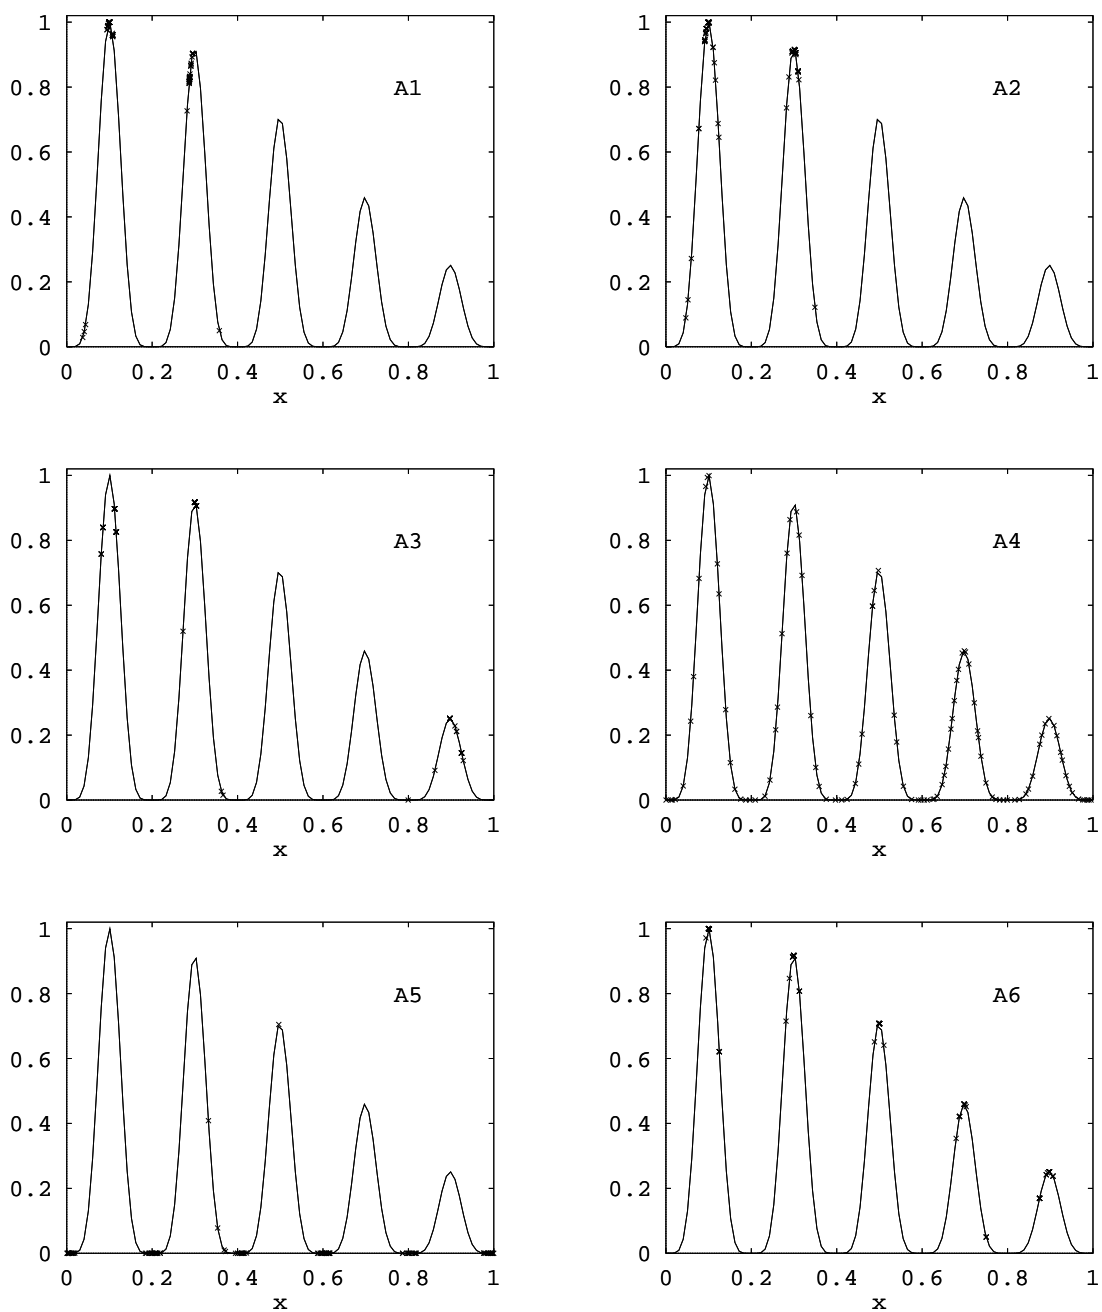
Figure 6.2: The final distribution of 100 population elements is shown for one run of each algorithm, A1–A6, on function *M2*. Each algorithm runs for 20,000 function evaluations.

Figure 6.3: The number of peaks maintained by each algorithm, A1–A6, is shown as a function of time, averaged over 100 runs. A peak is considered maintained if some population element exists in the peak's basin of attraction, whose fitness is at least 80% of the peak's height.

Figure 6.4: The total number of replacement errors for each algorithm, A1–A6, is shown as a function of time, averaged over 100 runs.
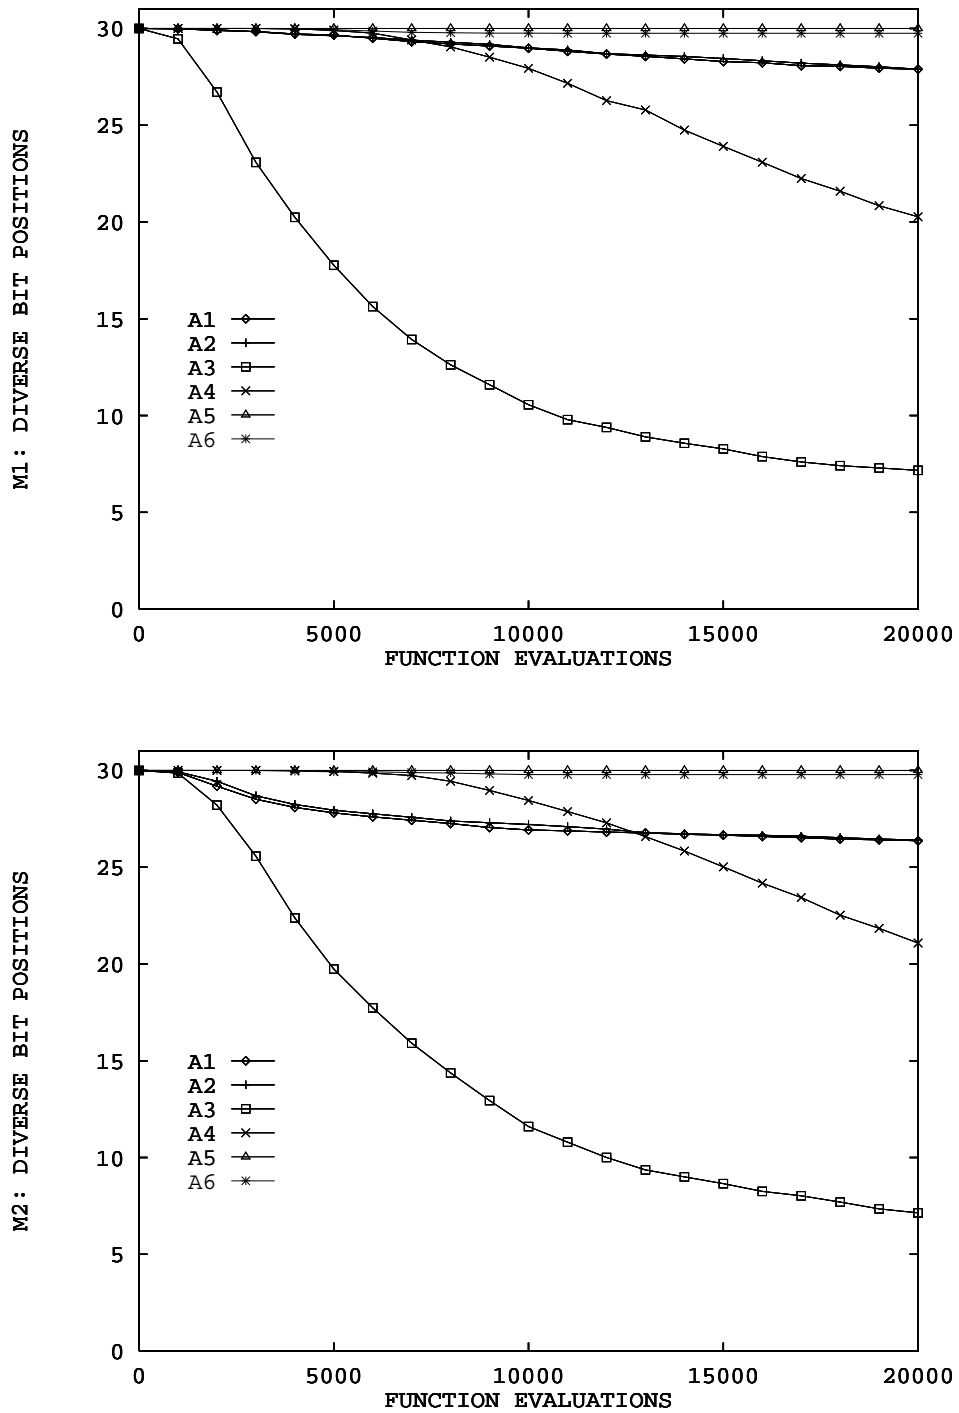
Figure 6.5: Bitwise diversity maintained by each algorithm, A1–A6, is shown as a function of time, averaged over 100 runs. A bit-position is considered diverse if both possible values occur at least once in the population.

Chapter 7

Crowding: Selection Plus Crossover

This chapter analyzes the deterministic crowding algorithm that was developed in the previous chapter. Specifically, it concentrates upon the distribution of population elements among classes, that arises from the combination of crossover and replacement selection. In fitness sharing, we know that the expected distribution is determined by representative class fitnesses: the proportion of the population expected in a class is that class's fitness divided by the sum of class fitnesses. In crowding methods, on the other hand, we currently known nothing about expected distributional properties. As we soon discover, crowding distributions are different from sharing distributions.

Recall from the previous chapter that deterministic crowding randomly pairs all population elements each generation, to yield $n/2$ pairs of parents, where n is the population size. Each such pair undergoes crossover, possibly followed by mutation or other genetic operators, to yield two children. Each of the two children competes with one of the two parents, according to similarity, for inclusion in the population. The deterministic crowding algorithm is presented in pseudocode form in Figure 7.1.

It turns out that crossover largely determines the behavior of deterministic crowding. This makes sense if one examines DC's behavior without crossover (and without mutation, which is absent in all models). Without crossover, DC simply advances both parents to the next generation, and the population never changes. Under crossover, however, a variety of interesting behaviors emerge. The combination of crossover and replacement selection produces interactions among niches, that we later isolate and explain. These interactions determine the

Deterministic Crowding

(REPEAT for g generations)

DO $n/2$ times:

1. Select two parents, p_1 and p_2, randomly, without replacement

2. Cross them, yielding c_1 and c_2

3. Apply mutation and possibly other operators, yielding c_1' and c_2'

4. IF $[Distance(p_1, c_1') + Distance(p_2, c_2')] \leq [Distance(p_1, c_2') + Distance(p_2, c_1')]$

   - IF $f(c_1') > f(p_1)$ replace p_1 with c_1'
   - IF $f(c_2') > f(p_2)$ replace p_2 with c_2'

   ELSE

   - IF $f(c_2') > f(p_1)$ replace p_1 with c_2'
   - IF $f(c_1') > f(p_2)$ replace p_2 with c_1'

Figure 7.1: Pseudocode is given for deterministic crowding.

distributional properties of DC and, we expect, of other crowding methods. This chapter's characterization of crossover-induced interactions among niches may additionally be applicable to other niching methods and GAs that incorporate crossover.

To simplify matters, we assume the crossover probability $p_c = 1$ throughout the models and runs of this chapter. Since DC is an elitist method, there is, in general, no reason to set p_c to values less than 1; crossover disruption is not a factor, since a child that is less fit than the parent with which it competes, will immediately be discarded.

7.1 Two Classes

We start by constructing the simplest possible model in which crossover can potentially change the existing distribution — a two-class model. We can predict the combined effect of crossover and replacement selection using this model. According to our modelling framework, crossing two elements of the same class yields two offspring of that class. Crossing two elements of differing classes, however, may yield two offspring from any of the c classes. We expect that most of the time, especially for $c = 2$, crossing elements of differing classes will yield one offspring from each class involved in the cross. Hence, we predict that the two-class case with crossover, like

the case of selection only, will ideally maintain the initial class distribution perpetually. (Unlike selection only, improvement will occur within each class.)

To illustrate this model behavior, let us consider the bimodal functions, *M10* and *M11*, displayed in Figures 5.10 and 5.11. *M10* and *M11* are constructed so that the most significant (leftmost) bit of each population element indicates the peak under which it lies, and defines the class (either A or B) to which it belongs.

We expect the following behavior under DC. If two elements of Class A get randomly paired, both offspring will also be of Class A, and the resulting tournament will advance two Class A elements to the next generation. The pairing of two Class B elements will similarly result in no net change to the distribution of the next generation. If an element of Class A gets paired with an element of Class B, one offspring will be from Class A, and the other, from Class B. According to the perfect discrimination assumption of our modelling framework, elements of the same class are always closer to each other than they are to elements of other classes. Therefore, the Class A offspring will compete against the Class A parent; the Class B offspring, against the Class B parent. The end result will be that one element of each class advances to the next generation — no net change in distribution. Since each element of the population receives exactly one trial per generation, the mean and variance for the number of population elements in A after one generation (n trials) are $\mu_A = I_A$ and $\sigma_A^2 = 0$, where I_i represents the number of elements in Class i at the beginning of the generation. Likewise, $\mu_B = I_B$ and $\sigma_B^2 = 0$.

We run DC on *M10* and *M11* for 100 generations, with $n = 32$, $p_c = 1$, and $p_m = 0$. We perform 100 runs on each function. DC, as expected, forms and maintains stable subpopulations atop both peaks. On each function, all 100 runs locate both global optima and maintain them to the end of the run. In fact, every population element typically moves to the top of one or the other peak by the end of a run. Although improvement occurs within classes, all 100 runs maintain the initial class distribution for 100 generations (zero variance). In some runs, random initialization gives Class B many more elements than Class A (or vice versa); however, crowding exhibits no pressure towards balance, even on *M11*, where A is four times as fit as B. Over the 100 runs, DC maintains a near uniform, average distribution of 15.79 Class A elements and 16.21 Class B elements, on both *M10* and *M11*. (The same 100 initial populations are used for both *M10* and *M11*.)

Crowding's behavior can hence be quite different from that of sharing, at least in the two-class case. The expected distribution is identical to the prior distribution, with no variance. No restoration is present, but no drift is either. Most notably, unlike sharing, the expected distribution is independent of fitness: a highly fit class does not gain elements from a lowly fit class. While lesser-fit classes may disappear in sharing (Deb, 1989), they appear to be in no danger in DC, at least in the two-class case. Stability is guaranteed, regardless of class-fitness differential. Even when *M11* is modified to have an increased peak fitness ratio of 256:1, the lesser-fit class maintains an average, over multiple runs, of roughly 50% of the population's elements.

One might ask what the distribution does depend upon, if not upon fitness. In the two-class case, since the steady-state distribution is expected to be identical to the initial distribution, and since the initial distribution is uniformly random with respect to each bit-position, we expect that the number of elements in a particular class i is proportional to the percentage of the search space in i. This percentage is proportional to the width, in the one-dimensional case, of the base of the peak corresponding to i. This width corresponds to the size of the peak's basin of attraction (defined in Chapter 5's modelling framework), under an appropriately chosen hillclimbing neighborhood operator.

One might also ask what crowding can accomplish if it never shifts elements between classes. The answer is partly that improvement does occur within each class. For instance, DC consistently locates and maintains both optima in *M10* and *M11*. In addition, multiclass problems, as we soon discover, allow beneficial interclass migrations and other interactions.

7.2   Multiple Classes

We now examine DC's behavior on problems with multiple classes, where crossover produces a variety of interactions among classes. Certain classes may gain an advantage over others, and ultimately take a larger share of the population. In our DC model, since all elements of a class have fitness identical to the corresponding peak's height, offspring from classes of higher fitness, when competing against parents from classes of lower fitness, will always win. We need only account for the possible offspring of interclass crossover.

Consider the four-peaked functions, *M12–M14*, of Figures 5.12–5.14, in which class membership is determined solely by the leftmost two bits of the eight-bit chromosome. Class A individuals have '00' in these bits; Class B, '01'; Class C, '10'; and Class D, '11'. Crossing A with B or C always yields one element from each class involved in the cross; after like competes against like, we obtain no net change in distribution. Likewise, crossing D with B or C yields no net change. The two remaining possibilities are of the most interest. With probability 6/7, single-point crossover will avoid the leftmost two bits, resulting in no net change. However, with probability 1/7, the crossover point will fall between these most significant bits, resulting in two offspring of classes different from their parents' classes. Such a cross between elements of A and D will yield two offspring, one from B and one from C; such a cross between B and C will yield one A and one D. In both special cases, pairing based on phenotypic proximity will result in two tournaments — one between A and B, and one between C and D — with the fitter class winning each tournament. The two cases that produce outside offspring are symmetric and equally probable, so that given a uniform starting distribution, no one class can expect to be generated by crossover, more frequently than any other.

Consider *M12*, in which all four peaks are of equal height. The DC algorithm prefers parents to their offspring in case of a tie. Therefore, no migration is expected between peaks. In an actual run with $n = 32$, $p_c = 1$, and $p_m = 0$, after a few generations the population reaches equilibrium, and no further migration occurs. This run is shown in Figure 7.2. Note the symmetry of the curves corresponding to the number of individuals in A and B, and also the symmetry of the curves for C and D. Prior to equilibrium, A can only gain an individual if it takes it from B, and vice versa. The combined number of individuals in C and D likewise remains constant.

As an example of analysis contributing to design, let us temporarily alter DC so that in case of a tie, it always advances the child over the parent. For functions with multiple global optima, such as *M12*, the migration discussed above will occur freely, even after equilibrium. The inevitable result of such migration will be genetic drift, though crowding will drift much more slowly than the simple GA. A preliminary version of deterministic crowding was, in fact, implemented in this way, and a small amount of drift did occur on *M12*. This observation of drift led to the slight design change requiring that offspring be strictly fitter than their parents in order to advance.

Figure 7.2: DC runs on *M12* with $n = 32$, $p_c = 1$, and $p_m = 0$. The top graph shows the elements atop the peaks they have climbed after 100 generations. The bottom graph tracks the number of elements in each class over the course of the run.

In *M13*, Class A is twice as fit as the other three classes. Our model predicts that when a cross between B and C yields one A and one D, A will win the subsequent tournament against B. Since such a cross occurs with fixed, positive probability, B will eventually contribute all of its elements to A. C and D, however, will not give each other any elements. Figure 7.3 shows a sample run of DC on *M13*, again with $n = 32$, $p_c = 1$, and $p_m = 0$. In the sample run, C and D reach equilibrium after 13 generations, at which point they no longer exchange elements. A and B, on the other hand, reach a quasi-equilibrium around Generation 32, when B has given most but not all of its elements to A. This quasi-equilibrium lasts to Generation 56, when A gains the rest of B's elements. We refer to Peak A as a *dominating* peak (or class), and to Peak B as a *dominated* peak.

We have just uncovered two important properties of crowding. First, delayed convergence allows a user to identify dominated classes by stopping the run upon preliminary convergence, rather than waiting for full convergence. Second, crowding realizes a couple of properties of sharing, albeit via different routes. We define a dominated class as one which can, with another's assistance, cross to form a fitter class, and which, due to close proximity to the fitter class, must transfer its elements to the fitter class. In sharing, peaks in close proximity to higher peaks may donate all of their elements to one or more higher peaks, if σ_{share} is too large to discriminate (Goldberg, Deb, & Horn, 1992). This is the first similar property. The second property is that crowding, like sharing, ultimately distributes elements to classes based partially on fitness.

One further property of crowding deserves mention — the *assist* — where an element of one class helps an element of another class to migrate, but does not necessarily go anywhere itself. On *M13*, migration from B to A occurs with the assistance of C, since for a single migration to occur, B and C must cross to produce A and D. We say that A dominates B with the assistance of C.

In *M14*, Class D is also made a dominating class, so that peaks A and D are each twice the height of B and C. We might at first expect B to eventually transfer all of its elements to A, and C to transfer all of its elements to D. However, according to our model, B and C must assist each other in migration. Whenever either B or C is depleted, the other becomes stable and no longer transfers any of its elements. Without the assist, the weaker class is no longer dominated. We verify this hypothesis by running DC on *M14*, once again with $n = 32$, $p_c = 1$, and $p_m = 0$. One of several similar runs is shown in Figure 7.4. After Class B contributes all

Figure 7.3: DC runs on *M13* with $n = 32$, $p_c = 1$, and $p_m = 0$. The top graph shows the elements atop the peaks they have climbed after 100 generations. The bottom graph tracks the number of elements in each class over the course of the run.

Figure 7.4: DC runs on *M14* with $n = 32$, $p_c = 1$, and $p_m = 0$. The top graph shows the elements atop the peaks they have climbed after 100 generations. The bottom graph tracks the number of elements in each class over the course of the run.

Table 7.1: For various functions, averages are given (over r runs) of the number of elements in each class after g generations. Population size is n. In all cases, $p_c = 1$ and $p_m = 0$.

	r	g	A	B	C	D	E	n
$M12$	100	100	7.80	7.99	7.95	8.96		32
$M13$	100	100	15.39	.40	8.90	7.31		32
$M14$	100	100	13.58	2.21	2.32	13.89		32
$M1$	100	200	23.33	15.76	21.13	16.62	23.16	100
$M2$	500	300	24.25	17.86	20.77	14.52	22.59	100

of its elements to Class A, Class C becomes stable, with four elements remaining. The four elements are then maintained perpetually. (We actually stop the run after 5000 generations.) In all runs that we perform on $M14$, once either B or C is depleted, the other class remains stable to the end of the run.

Recall our earlier conjecture for two classes, that the expected number of elements in a class is proportional to the width (in one dimension) of the base of its peak. For multiple classes, we revise this conjecture to the following: the expected number of elements in a class, at equilibrium, is proportional to the sum of the width of the base of its peak, and the widths of the bases of all peaks it dominates. Dominated peaks are expected to disappear, unless stabilized by the prior disappearance of their assisting peak(s). In a multidimensional search space, our revised hypothesis tells us that the expected proportion of Class i elements, in the long run, is equal to the sum of the percentage of the search space contained in i, and the percentage of the search space contained in all classes that i dominates. According to our modelling framework, the amount of search space contained in a class is equal to the corresponding peak's basin of attraction, under an appropriate neighborhood operator.

Table 7.1 illuminates the above conjecture regarding long-run class distributions. It displays, for $M12$–$M14$, the average number, over 100 runs of DC, of elements in each class after 100 generations. For $M12$, the distribution is very close to uniform, approximately 8 elements per class. On $M13$, Class A swallows all of B's elements, yielding an average, combined count of 15.79 elements for the two classes; Classes C and D are close to the expected count of 8 elements apiece. On $M14$, A and B combine for an average of 15.79 elements, while C and D average a combined total of 16.21; both combined totals compare favorably with the expected count of 16 per pair of classes.

Note on *M13*, that not all runs lose all Class B elements by Generation 100. On *M14*, over two elements, on the average, are maintained in both B and C to Generation 100. These two observations demonstrate the delayed-convergence property mentioned earlier. In fact, migration in *M14* markedly slows down as assisting peaks are depleted — some runs hold on for several hundred generations before losing all elements of either B or C. This is due to the decreasing probability of selecting both parents from declining classes.

We now test DC on two functions from the literature, to see whether it exhibits the same properties as on *M10–M14*. We examine DC on *M1* and *M2*, the test functions of the previous chapter. *M1* has five peaks of uniform height; *M2*, nonuniform height. Runs employ $n = 100$, $p_c = 1$, and $p_m = 0$.

Note in Table 7.1 that the average distribution for *M1* is not quite uniform, even though all peaks have the same width and are of the same height. A plausible explanation for this behavior is found by examining key hybrid crosses. Typically, interclass crosses do not produce fit offspring of outside classes. However, prior to convergence, crosses between B and C sometimes yield relatively fit elements of A, and crosses between C and D sometimes yield relatively fit elements of E. We therefore suggest that prior to full convergence, A dominates B, and E dominates D, in both cases with the assistance of C. If we add together the average number of elements for A and B, we obtain 39.09; for D and E, 39.78. Both sums are very close to 40, the expected combined value under domination. The partial symmetry between dominating and dominated classes, illustrated in the sample run of Figure 7.5, provides further evidence. Dominance and its resulting migration come to a full stop after 151 generations, and the distribution remains unchanged thereafter. (We continue the run to Generation 5000.)

On *M2*, which has peaks of differing fitnesses, we observe more of the type of behavior displayed on *M1*, plus additional interactions, as the less fit peaks D and E give up a few of their collective elements to A and B (see Table 7.1). We start to see a sharing-like phenomenon, in which the end distribution is partially related to fitness. In most of the 500 runs we perform, the distribution no longer changes after 200–300 generations.

The main purpose of monitoring an algorithm's behavior on simple problems is to isolate components of that behavior that will transfer to difficult problems and to applications problems. We next apply DC to *M7*, a massively multimodal, deceptive function, described in

149

Figure 7.5: DC runs on *M1* with $n = 100$, $p_c = 1$, and $p_m = 0$. The top graph shows the population elements after 200 generations. The bottom graph tracks the number of elements in each class over the course of the run. After 151 generations, the distribution no longer changes.

Chapter 5. $M7$ has 32 global optima, which we wish to locate, and over 5.15 million nonglobal optima, which we do not care to locate.

Employing a population of 1000 individuals, genotypic comparison, $p_c = 1$, and $p_m = 0$, DC successfully locates and maintains copies of all global optima, and ultimately gets rid of all nonglobal optima. Locating all global optima is not immediate, however; it first requires that nonglobal optima be found. In a sample run, by Generation 56, each global optimum gains at least one element, and one global optimum gains 53 elements. Full convergence occurs by Generation 167, at which point all elements belong to classes corresponding to global optima. Upon convergence, 2 of the 32 global classes contain only 4 elements (the smallest number), but one class contains 70 elements (the largest number). This result compares favorably with the results of Goldberg, Deb, and Horn (1992), who introduce $M7$ and present the only prior, published solution to it (prior to the preliminary writing of this thesis). The authors use sharing along with exponential scaling and a population size of 5000, to maintain roughly 40% of the population at global optima. With DC, we consistently solve $M7$ in one fifth the function evaluations.

Figure 7.6 illustrates, for $M7$, the crossing of ordinary population elements to form local optima, and local optima to form global optima. It shows the progress, from generation to generation, as additional global subfunctions are discovered by each population element. It also shows that intermediate optima disappear, as a diverse collection of global optima takes over the whole population.

Over multiple runs, no global class appears to have an overwhelming advantage over any other. For instance, over the course of 200 runs, with 200 generations per run and $n = 1000$, the lowest average number of elements for any of the 32 global classes is 21.93; the highest is 38.76. Nonglobal optima, however, eventually disappear. After 200 generations, the 1600 nonglobal optima having the highest fitness (with four out of five subfunctions globally converged) average a combined total of only .31 of a population element amongst them. Local optima with lower fitnesses average 0. We attribute the eventual disappearance of all nonglobal optima to their being dominated by global optima.

Why does DC slightly prefer some global optima over others, despite the fact that an equivalent amount of search space leads up to each? One possibility is that some global optima dominate more local optima than do other global optima. Perhaps more pertinent are the
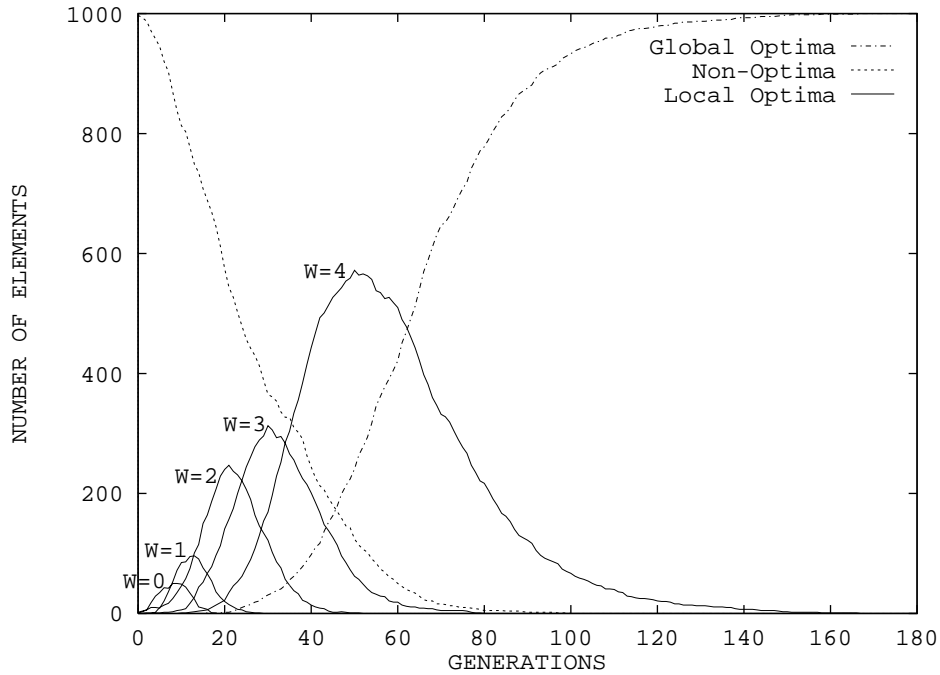
Figure 7.6: Various categories of optima are tracked for 180 generations of DC on $M7$, with $n = 1000$, $p_c = 1$, and $p_m = 0$. W is the number of globally converged subfunctions in a local optimum.

end biases of single-point crossover, which make the formation of certain global optima more difficult than others. Briefly, in single-point crossover, the end points of the strings serve as *de facto* crossover points. It is impossible to pass the central portion of a string from parent to offspring without also passing one of the end portions. Two-point crossover, on the other hand, does not possess any end biases.

We repeat the above experiment using two-point crossover in place of single-point. Over 200 runs, the lowest average number of elements for any of the 32 global classes is 28.01; the highest is 34.65. (Apportioning the 1000 population elements uniformly among global classes gives an average of 31.25 elements to each global class). The distribution among global classes is much closer to uniform using two-point crossover.

7.3 Discussion of Properties

We have formulated and tested the hypothesis that under DC, for all classes i, the expected long-run proportion of Class i elements is directly proportional to the sum of the size of i's

152

basin of attraction and the size of the basins of attraction of all peaks that i dominates. Basins of attraction are defined, whenever appropriate, in terms of phenotypic, epsilon-neighborhood hillclimbing. For $M7$, basins are defined in terms of single-bit hillclimbing in Hamming space. Relative peak fitnesses, proximity of peaks, and crossover biases help determine the dominance of peaks.

Without the power to migrate to higher peaks, DC would simply be a parallel hillclimber. However, the ability to jump from peak to peak allows DC to escape local optima, given the traditional definition of a local optimum. DC without mutation can be viewed as performing parallel, *crossover hillclimbing*. Crossover hillclimbing differs from traditional hillclimbing, because it employs the binary neighborhood operator of crossover, rather than a mutation-like, unary neighborhood operator. The crossover neighborhood of a particular pair of points may bounce either point all over the search space, when viewed from the vantage point of traditional, epsilon-neighborhood hillclimbing. The crossover neighborhood of a single point is defined relative to the other points in the search space with which it may cross. Therefore, we can attempt to predict the distribution resulting from crossover hillclimbing, in much the same way we have been — by examining the potential results of crossing every two classes.

We can reformulate our prior hypothesis in terms of crossover hillclimbing and *crossover basins of attraction*. In this new light, the single factor that determines the eventual distribution under deterministic crowding is the number of ordered pairs in $S \times S$ that potentially lead to a peak via crossover, where S is the original search space. We call this number of ordered pairs, a peak's crossover basin of attraction. Note that a point may be within more than one peak's crossover basin of attraction. However, it may be more likely to reach one peak than another. Jones (in press) provides a similar, but more detailed treatment of crossover landscapes and crossover hillclimbing. The concept of crossover hillclimbing should be a useful abstraction for future studies of crowding, as well as future studies of other niching methods, replacement selection methods, and genetic algorithms. It is of intermediate modelling complexity between the models of this study and the Markov chain models that we reviewed in Chapter 5.

The results of this chapter, in addition to shedding light on crossover interactions, crossover hillclimbing, and crowding distributions, also raise questions about adding mating restrictions to niching methods. Deb (1989) shows that mating restrictions yield a cleaner version of sharing, with better on-line performance. This is because mating restrictions prevent lethal

hybrid crosses between elements of different niches. However, mating restrictions also prevent beneficial hybrids from forming, essentially creating a trade-off between on-line and off-line performance. Ironically, we have found instances in this chapter, using the same five-peaked functions as Deb, where hybrid crosses lead to better offspring than their parents. On the massively multimodal and deceptive $M7$, hybrid crosses of local optima were essential to the formation of global optima. Had mating restrictions been additionally enforced, migration to higher peaks would have been prevented.

Chapter 8

Sharing: Selection

This chapter models fitness sharing in the absence of crossover. The models are analogous to a one-player game, in which the player chooses several "lucky numbers" between 1 and 6, inclusive. The player rolls a number of dice simultaneously, trying to match all of his or her lucky numbers. If successful, the player rolls again. Otherwise the game ends.

Recall from Chapter 4 that sharing is able to work in conjunction with any GA selection scheme. We model sharing with roulette-wheel selection, with the understanding that since RWS is the noisiest of commonly used fitness-proportionate selection schemes, our models will bound the behavior of sharing under noise-reduced, fitness-proportionate selection schemes such as stochastic remainder selection and stochastic universal selection. Specifically, we expect bounds on quantities such as drift time and population size, computed using RWS, to also serve as bounds for less noisy selection schemes. RWS not only gives us bounding behavior, but also offers the advantage of simplicity, making our modelling task a more promising venture.

This chapter first analyzes and illustrates sharing's distributional properties from the perspective of our modelling framework. It examines the likelihood that sharing loses important solutions, and derives closed-form expressions for the expected time to disappearance of a class. This time to disappearance or *drift time* is related to the population size n, number of classes c, and relative class fitnesses (which we designate using ratio r). We illustrate, both theoretically and empirically, the drift time of sharing.

The chapter next sets up models of fitness sharing in which all peaks in the fitness landscape are desirable. The first model assumes that all classes are of identical fitness. Subsequent models

allow classes of arbitrary fitness. The models yield lower bounds on required population size. We illustrate the utility of these population-size bounds on a set of test functions of varying difficulty. Afterwards, we relax the assumption that all peaks are desirable, in order to construct a model that handles both desirable and undesirable peaks. This new model results in a more general, population sizing expression.

8.1   Distributional Properties

This section examines class distributions resulting from sharing with roulette-wheel selection. It derives distributional properties for sharing with RWS, and illustrates the derived properties with several examples.

Roulette-wheel selection

Before examining sharing with roulette-wheel selection, we first look at RWS in isolation. In any given trial of RWS, the probability of selecting some element of an arbitrary class i, $P_s(i)$, is proportional to fitness:

$$P_s(i) \;=\; \frac{I_i f_i}{\sum_{j=0}^{c-1} I_j f_j} \;\;. \tag{8.1}$$

Note that the expression in Equation 8.1 for $P_s(i)$ is similar to that for $p_s(i)$ in Equation 2.2, except that i now designates a class instead of an individual, and P_s is that class's selection probability, while p_s is an individual's selection probability. Note also that $\sum_{i=0}^{c-1} P_s(i) = 1$.

Employing the formulas for the mean and variance of the binomial distribution for I_i, the expected number of population elements in Class i after one generation (n Bernoulli trials) is

$$\mu_i \;=\; n \, P_s(i) \;\;, \tag{8.2}$$

with variance,

$$\sigma_i^2 \;=\; n \, P_s(i) \, (1 - P_s(i)) \;\;. \tag{8.3}$$

The above results tell us a great deal about the behavior of a GA that runs RWS. The distribution at time $t + 1$ of population elements among classes depends both upon the distribution at time t and the relative fitnesses of the classes. Given multiple classes of identical

fitness, where one class has become more frequent than the others, there is no restoration — the expected number of each class in a given generation is the same as the actual number of each in the previous generation. (If all f_i are the same, Equations 8.1 and 8.2 tell us, respectively, that for all i, $P_s(i) = I_i/n$ and $\mu_i = I_i$.) This lack of restoration, in combination with the significant variance of RWS and finite population size, leads to the eventual disappearance of all but one class (Goldberg & Segrest, 1987). Given classes of nonuniform fitness, the most fit class is expected to take over the entire population: a rough estimate of takeover time can be obtained by iterating Equation 8.2 for the fittest class k until $\mu_k \geq n$.

Sharing with roulette-wheel selection

Let us now add fitness sharing to roulette-wheel selection. Shared fitness, denoted by f', for an arbitrary class i, is given by $f'_i = f_i/I_i \quad (I_i \neq 0)$. Substitution into Equation 8.1 yields

$$P_s(i) = \frac{I_i f'_i}{\sum_{j=0}^{c-1} I_j f'_j} = \frac{f_i}{\sum_{j=0}^{c-1} f_j} \quad . \tag{8.4}$$

The mean and variance for the number of population elements in Class i after one generation (n Bernoulli trials) are given in the following two equations:

$$\mu_i = n \ P_s(i) \ ; \tag{8.5}$$

$$\sigma_i^2 = n \ P_s(i) \ (1 - P_s(i)) \quad . \tag{8.6}$$

Two special cases of Equation 8.4 are of interest. For c classes of identical fitness ($f_0 = f_1 = \ldots = f_{c-1}$),

$$\forall_i \ P_s(i) = \frac{f_i}{\sum_{j=0}^{c-1} f_j} = \frac{1}{c} \quad . \tag{8.7}$$

For two classes ($c = 2$), A and B,

$$P_s(A) = \frac{f_A}{f_A + f_B} = \frac{r}{r+1} \quad , \tag{8.8}$$

where $r = f_A/f_B$.

For sharing with roulette-wheel selection, the expected distribution and its variance are independent of the starting distribution (since the I_i terms drop out of Equation 8.4); they

depend only upon relative class fitnesses. This explains the *restorative pressure* inherent in sharing methods, where classes that contain other than their share of population elements can expect to be restored to the proper number in one generation. This restorative pressure keeps drift from becoming an overwhelming factor, since unlike RWS alone, RWS with sharing does not give stochastic fluctuations the opportunity to accumulate over multiple generations. A related beneficial property of sharing, uncovered by the above model, is that no greater threat exists of a class becoming extinct when it has only one representative in the population than when it has many.

In the game we earlier outlined, a successful player picks up all the dice and rolls them again. The numbers that appear on the dice are independent of the numbers that appeared on the prior roll. In our sharing model, dice become population elements, rolls become generations, and numbers become classes. Successive generations are independent, allowing us to algebraically derive further properties of sharing such as drift time and population size. Because of this independence, sharing, unlike the simple GA, is not a Markov chain. (A Markov chain is a sequence of trials in which the outcome of a trial depends only upon the outcome of the previous trial.) Horn (1993) discovers this non-Markov property of sharing when the column entries of his transition matrix for "perfect sharing" turn out to be identical. Horn's observation, along with the results presented herein, suggest that Markov chains may be unnecessarily complex for modelling sharing.

Given classes of nonuniform fitness, the least fit class will not disappear if sharing is employed, unless its expected number of individuals is small enough to be overcome by noise. Deb (1989) notes that genotypic sharing sometimes loses peaks of relatively low fitness. He attributes this loss to the noisy discrimination of genotypic sharing: the setting of σ_{share} must take into account distances between peaks, as well as relative class fitnesses. Specifically, to maintain the lower of peaks i and j, the following σ_{share} is required:

$$\sigma_{share} \leq \frac{d(i,j)}{1 - r_{i,j}} \quad , \tag{8.9}$$

where $d(i,j)$ is the Hamming distance between the maxima of peaks i and j, and $r_{i,j} = \min(f_i/f_j, f_j/f_i)$. In our models, the perfect discrimination assumption allows us to ignore the setting of σ_{share}.

158

Note that Deb's calculation does not take population size into account. This suggests that another factor might also be at work in genotypic as well as phenotypic sharing: genetic drift due to the relative insignificance of lower peaks. Examples of this type of genetic drift follow.

Experimental verification

We now illustrate some of the previously derived distributional properties. Figures 8.1 and 8.2 show the expected distribution after one generation and its standard deviation, for RWS alone and sharing with RWS. Figure 8.1 is for two classes of equal fitness; Figure 8.2, two classes of unequal fitness. The two figures illustrate the previously stated concept, that for RWS, the expected distribution is a function of the prior distribution; with sharing added, however, the expected distribution becomes independent of the prior distribution.

Figures 8.3 and 8.4 compare sample runs of both algorithms to expected behavior. In particular, runs of RWS and sharing with RWS proceed for 100 generations, on the simple two-class optimization problems, $f_B = f_A = 1$ and $f_B = 4f_A = 4$. These are one-bit problems in which the '0' genotype corresponds to Class A; the '1', to Class B. The representative fitnesses of the two classes, A and B, are designated f_A and f_B, respectively. The initial class distribution in both sample runs is uniform. Other GA parameters are $n = 32$, $p_c = p_m = 0$, and $\sigma_{share} = 0.5$.

In Figure 8.3, where both classes are of equal fitness, under RWS, as expected, genetic drift allows Class A to eventually take over the whole population. (Class A is just as likely to disappear, allowing Class B to take over the entire population.) RWS's high variance and lack of restorative pressure cause the GA to wander aimlessly in both directions, until Class B loses all its elements. Sharing, since it incorporates RWS, also exhibits similar fluctuations, this time about the mean value of 16 elements per class. Sharing's restorative pressure, however, maintains both classes in equal proportions on the average, for the full 100 generations.

In Figure 8.4, selection pressure causes Class A, the weaker of the two, to disappear rapidly under RWS (by Generation 3). When sharing is added, Class A persists, with numbers proportional to its relative fitness, but also with some fluctuation about the expected niche size of 6.4. At times, the number of elements in Class A fluctuates dangerously close to zero. This suggests that sharing with RWS, although more stable than RWS alone, is not immune from genetic drift.

Figure 8.1: The mean and standard deviation of the expected proportion of elements in a class after a generation, are given for RWS alone, and sharing with RWS. Model parameters are $c = 2$ and $f_B = f_A$. The curves are calculated from Equations 8.2, 8.3, 8.5, and 8.6.

Figure 8.2: The mean and standard deviation of the expected proportion of elements in a class after a generation, are given for RWS alone, and sharing with RWS. Model parameters are $c = 2$ and $f_B = 4f_A$. The curves are calculated from Equations 8.2, 8.3, 8.5, and 8.6.

Figure 8.3: Sample runs of RWS, and sharing with RWS, are compared to the expectation and standard deviation for sharing with RWS, computed from Equations 8.5 and 8.6. Parameters are $n = 32$, $c = 2$, $f_B = f_A$, $p_c = p_m = 0$, and $\sigma_{share} = 0.5$.

Figure 8.4: Sample runs of RWS, and sharing with RWS, are compared to the expectation and standard deviation for sharing with RWS, computed from Equations 8.5 and 8.6. Parameters are $n = 32$, $c = 2$, $f_B = 4f_A$, $p_c = p_m = 0$, and $\sigma_{share} = 0.5$.

161

Figure 8.5: A sample run of sharing with RWS is compared to the expectation and standard deviation for sharing with RWS, computed from Equations 8.5 and 8.6. Parameters are $n = 16$, $c = 2$, $f_B = 4f_A$, $p_c = p_m = 0$, and $\sigma_{share} = 0.5$.

Figure 8.5 verifies our suspicion. On the one-bit, two-class problem $f_B = 4f_A = 4$, when population size is cut in half to $n = 16$, Class A's expected niche size also drops by half, down to 3.2. Sharing with RWS allows the stronger Class B to take over the entire population after the fatal fluctuation in Generation 12. (Again, $p_c = p_m = 0$ and $\sigma_{share} = .5$.) This proves the earlier conjecture regarding Deb's calculation, that given perfect discrimination between peaks, lower peaks may still disappear due to genetic drift. The next section examines the exact point at which a class can be expected to disappear and this point's relation to population size.

8.2   Genetic Drift

We have set the groundwork for computing the expected time to extinction of one of the c classes. To compute this drift time, we must first bound the probability of extinction occurring in a single generation. From this bound, since successive generations are independent in our sharing model, the calculation of expected extinction time is straightforward.

Loss probabilities

Equation 8.4 gives the probability $P_s(i)$ of selecting some element of an arbitrary class i in one trial of sharing with RWS (one spin of the roulette wheel using shared fitnesses). The probability of selecting some element outside Class i is $1 - P_s(i)$. After a full generation, the probability of losing all elements of Class i is simply the probability of selecting all n elements from outside Class i: $(1 - P_s(i))^n$. Since any of the c classes can potentially lose all its elements, the probability X of one or more classes disappearing after one generation is

$$X \leq \sum_{i=0}^{c-1} (1 - P_s(i))^n \; . \qquad (0 < X < 1) \qquad (8.10)$$

The probability of maintaining all classes for a generation is $1 - X$.

The above expression is an inequality because the probabilities of intersection of loss events are not subtracted off (e.g., for a three-class problem, the probability of losing both Class 1 and Class 2 would be included twice in the above summation). When the distribution of class selection probabilities does not stray too far from uniform, and when $n \gg c$, intersection terms are insignificant, and the inequality (8.10) can be treated as an equation. For example, consider the $c = 3$ case, where classes are of uniform fitness. We perform the full calculation, subtracting off intersections:

$$X \; = \; 3(1 - \frac{1}{3})^n - 3(\frac{1}{3})^n \; . \qquad (8.11)$$

Simplification yields

$$X \; = \; \frac{2^n - 1}{3^{n-1}} \; = \; \frac{2^n}{3^{n-1}} - \frac{1}{3^{n-1}} \; . \qquad (8.12)$$

This differs in only an exponentially decreasing increment from the bound derived from (8.10):

$$X \; \leq \; \frac{2^n}{3^{n-1}} \; . \qquad (8.13)$$

In fact, for $n = 32$, the bound is indistinguishable from the exact quantity up to the 15th decimal place.

In realistic optimization problems, we will not be interested in locating peaks that are of low fitness; hence we will have near uniformity. (Extraneous, lower-fitness peaks, including deceptive peaks, will be modelled later, separate from the c desired classes.) In addition,

maintaining c peaks typically requires n to be some multiple of c; hence $n \gg c$. We can tell if either condition is severely violated, because the summation in (8.10) will approach or exceed 1.0, leading to a useless probability bound.

For the special case where all classes have identical fitness ($\forall_i \ P_s(i) = 1/c$), we obtain

$$X \ \leq \ c(1 - \frac{1}{c})^n \ \ .$$

(8.14)

Note that the right-hand side of the above expression has two useful forms:

$$c(1 - \frac{1}{c})^n \ = \ \frac{(c-1)^n}{c^{n-1}} \ \ .$$

(8.15)

We can calculate the conditions under which (8.14) gives a useful estimate for X. Since all classes are of identical fitness, the near-uniformity condition is satisfied. That leaves the condition $n \gg c$. How much greater than c must n be? Starting with the inequality,

$$c(1 - \frac{1}{c})^n \ < \ 1 \ \ ,$$

(8.16)

with a little manipulation we obtain

$$n \ \ln(1 - \frac{1}{c}) \ < \ -\ln c \ \ .$$

(8.17)

Using approximation (A.9) of the appendix to simplify the above left-hand side, yields

$$n \ > \ c \ \ln c \ \ .$$

(8.18)

Therefore, (8.14) gives a useful bound for X for values of n greater than $c \ln c$, with increasing utility for larger n, relative to $c \ln c$.

For the special case of two classes ($c = 2$), A and B, the calculation of X is exact, since it is not possible to lose both classes:

$$X \ = \ \frac{r^n + 1}{(r+1)^n} \ \ ,$$

(8.19)

where $r = f_A/f_B$.

Loss distribution

Each generation, at least one class disappears from the population with probability X. This probability of disappearance remains constant from generation to generation, until a loss actually occurs. Therefore, the loss of one or more classes can be treated as a binomial event, occurring with probability of "success" X, with the stipulation that once a success occurs, X is no longer valid. Since we are not interested in repeated successes, but only in strings of failures followed by a success, this limitation will have no effect.

We now examine the time required to lose at least one class. Recall our prior assumption that all classes are initially represented in the population. This assumption tells us that the probability of losing one or more classes after exactly zero generations is 0; the probability of losing one or more classes after exactly one generation is X; after exactly two generations, $(1 - X)X$; after exactly g generations, $(1 - X)^{g-1}X$, which is simply the probability of $g - 1$ "failures" to lose a class, followed by a "success". Treating the number of generations L required to lose at least one class as a random variable, yields the *loss distribution*:

$$P(L = g) \; = \; (1 - X)^{g-1}X \;\; , \qquad (1 \leq g < \infty) \qquad (8.20)$$

where $P(L = g)$ denotes the probability that the generation at which a loss occurs is g. Figure 8.6 shows the loss distribution for two classes of equal fitness under a population of size $n = 8$. Note that the loss distribution is a *geometric distribution* (Freund & Walpole, 1980).

The expected number of generations required to lose at least one class is the mean of the geometric (loss) distribution,

$$\mu_L \; = \; \frac{1}{X} \;\; . \qquad (8.21)$$

This result is intuitively what one would expect, the inverse of the single generation loss probability. The derivation of Equation 8.21 is presented in the appendix. The variance in the expected number of generations to loss is the variance of the geometric (loss) distribution,

$$\sigma_L^2 \; = \; \frac{1 - X}{X^2} \; = \; \mu_L^2 - \mu_L \;\; . \qquad (8.22)$$

The derivation of Equation 8.22 is also presented in the appendix.

Figure 8.6: The loss distribution of Equation 8.20 is shown for $c = 2$, $n = 8$, and $f_B = f_A$.

Figure 8.7: The cumulative loss distribution of Equation 8.24 is shown for $c = 2$, $n = 8$, and $f_B = f_A$.

From summing values of the loss distribution from 1 to g generations, we obtain $P(L \leq g)$, the probability of losing, in g generations or fewer, all elements of one or more of the c classes. (All classes are tried simultaneously.)

$$P(L \leq g) \; = \; \sum_{i=1}^{g}(1-X)^{i-1}X \; = \; X\sum_{i=0}^{g-1}(1-X)^{i} \;\; . \tag{8.23}$$

X is independent of i, and $0 < X < 1$. This *cumulative loss distribution* is shown in Figure 8.7 for the case of $c = 2$, $n = 8$, and $f_B = f_A$. Using the formula for the partial sum of a geometric series (Equation A.5 of the appendix) yields

$$P(L \leq g) \; = \; X\frac{1-(1-X)^{g}}{1-(1-X)} \; = \; 1-(1-X)^{g} \;\; . \tag{8.24}$$

Thus the probability γ of maintaining all c classes for at least g generations is given by

$$\gamma \; = \; 1 - P(L \leq g) \; = \; (1-X)^{g} \;\; . \tag{8.25}$$

Taking the gth root of both sides allows us to solve for X, resulting in a generally applicable equation for sharing:

$$X \; = \; 1 - \gamma^{\frac{1}{g}} \;\; . \tag{8.26}$$

Drift time

Figure 8.8 shows, for $c = 2$, two fitness ratios, and various population sizes, the expectation and standard deviation in the number of generations to loss of a class. Curves above and below the mean curves indicate a range of one standard deviation from the mean. Note that the lower curve, $\mu_L - \sigma_L$, converges to .5 . This is because

$$\lim_{n \to \infty} \sqrt{\mu_L^2 - \mu_L} \; = \; \mu_L - .5 \;\; . \tag{8.27}$$

Even for very small n, $\sigma_L \approx \mu_L - .5$. The fact that one standard deviation in the downward direction takes us close to zero does not mean that typical runs, regardless of population size, will lose a class immediately: the standard deviation does not tell us everything about the loss distribution. Because of infinite potential upward deviation in the expected number of

Figure 8.8: $\mu_L \pm \sigma_L$, calculated from (8.21) and (8.22) is shown for increasing n, with $c = 2$ and $r = f_A/f_B$.

generations to loss, and limited potential downward deviation, we can decrease the probability of a run losing a class to as close to zero as we desire. This is accomplished by raising n, and will be detailed in the upcoming section on population sizing.

Figure 8.9 shows, for multiple classes of uniform fitness, a lower bound on the expected number of generations to loss of at least one class, as a function of c and n. The bound is obtained by substituting the right-hand side of (8.14) into Equation 8.21. It shows that drift time increases exponentially as a function of n but decreases at a similar pace as a function of c. This is evident if we employ a large-c approximation (Equation A.10 of the appendix):

$$\mu_L = \frac{1}{X} \approx \frac{1}{c(1 - \frac{1}{c})^n} \approx \frac{e^{n/c}}{c} \quad . \tag{8.28}$$

For the special case of $c = 2$, we can compare our drift-time expression with Horn's (1993) graph, drawn using a Markov chain model of sharing. Horn graphs the expected number of generations to loss of a class in the one-bit, two-class, "perfect sharing" case. We obtain an expression for this quantity by substituting the right-hand side of (8.19) into (8.21):

Figure 8.9: A lower bound on expected drift time, calculated from (8.14) and (8.21), is shown as a function of c and n, for multiple classes of uniform fitness.

$$\mu_L \;=\; \frac{(r+1)^n}{r^n + 1} \;\; ,$$

(8.29)

where r is the fitness ratio of the two classes. For two classes of equal fitness ($r = 1$), Equation 8.29 reduces to $\mu_L = 2^{n-1}$. Horn concludes from his graph that drift time is an exponential function of n. This is verified by Equation 8.29 as follows. Without loss of generality, assume that $0 < r \leq 1$. The denominator will be in the interval $(1, 2]$, but the numerator will grow exponentially at a rate of $(r+1)^n$. Therefore,

$$(1 + \epsilon)^n \;\; \leq \;\; \frac{(r+1)^n}{r^n + 1} \;\; \leq \;\; 2^{n-1} \;\; ,$$

(8.30)

where ϵ is an infinitesimal positive constant. Expression (8.30) states that growth is exponential, but not to the degree depicted in Horn's graph. This discrepancy in Horn's graph is due to a failure to divide by the dependent variable n (J. Horn, personal communication, June, 1993).

169

Table 8.1: Mean drift times $\bar{x}$ for 5000 runs of sharing with RWS are compared to expected drift times μ_L from the sharing model, for $c = 2$, and two different fitness ratios. A 95% confidence interval for μ_L is calculated from $\bar{x}$. Sample standard deviation is s, and the model's standard deviation is σ_L. For the actual GA, $p_c = p_m = 0$ and $\sigma_{share} = .5$. All drift-time statistics are in generations.

n	$\bar{x}$	μ_L	μ_L: 95% C.I.		s	σ_L
			$f_B = f_A$			
2	2.00	2.00	(1.96,	2.04)	1.42	1.41
4	8.04	8.00	(7.83,	8.25)	7.46	7.48
6	32.47	32.00	(31.60,	33.34)	32.02	31.50
8	129.18	128.00	(125.65,	132.71)	125.66	127.50
10	516.23	512.00	(502.05,	530.41)	512.43	511.50
			$f_B = 4f_A$			
2	1.46	1.47	(1.44,	1.48)	0.82	0.83
4	2.41	2.43	(2.36,	2.46)	1.82	1.86
6	3.80	3.81	(3.71,	3.89)	3.22	3.27
8	6.03	5.96	(5.88,	6.18)	5.42	5.44
10	9.47	9.31	(9.23,	9.71)	8.85	8.80
12	14.50	14.55	(14.11,	14.89)	14.07	14.04
14	23.01	22.74	(22.39,	23.63)	22.84	22.23
16	35.42	35.53	(34.45,	36.39)	34.58	35.03
18	55.83	55.51	(54.31,	57.35)	55.80	55.01
20	86.17	86.74	(83.78,	88.56)	86.59	86.24
22	134.89	135.53	(131.15,	138.63)	135.86	135.03
24	215.33	211.76	(209.47,	221.19)	217.80	211.26

Experimental verification

Table 8.1 compares the sharing-with-RWS model to statistics from 5000 runs on the one-bit, two-class functions defined earlier ($f_B = f_A$ and $f_B = 4f_A$). The table presents 95% confidence intervals for the expected drift time, given the average drift time over 5000 runs. All runs employ uniform initialization. For both functions, experimental results at various n with $p_c = p_m = 0$ and $\sigma_{share} = .5$ correspond nearly perfectly to derived results.

Compare Table 8.1's expected drift time for $f_B = 4f_A$ using $n = 16$, with the earlier runs from Figures 8.4 and 8.5. The table explains why the $n = 16$ run drifted but the $n = 32$ run did not: $n = 32$ requires an average of 1262 generations to drift, but we only ran it for 100; $n = 16$, on the other hand, requires only 35.53 generations on the average.

Table 8.2: For fixed $n = 32$, c varies, and the f_i are uniform. The estimator $\hat{\mu}_L$ (computed from (8.14) and (8.21)) for expected drift time in sharing is compared with the sample mean $\bar{x}$ from 5000 runs of sharing with RWS ($p_c = p_m = 0$; $\sigma_{share} = .5$). All drift-time statistics are in generations. Sample standard deviation is s.

c	$\bar{x}$	$\hat{\mu}_L$	% Error	s
12	1.74	1.35	22.41%	1.15
11	2.28	1.92	15.79%	1.68
10	3.27	2.91	11.01%	2.70
9	5.10	4.82	5.49%	4.61
8	9.19	8.97	2.39%	9.04
7	19.75	19.82	0.35%	18.96
6	57.93	56.97	1.66%	58.97
5	256.94	252.44	1.75%	252.90

The final simulations of this section test bounds (8.14) and (8.18). We fix n at 32 and vary c, employing starting distributions as close to uniform as possible. (Sometimes n is not evenly divisible by c.) Class fitnesses are uniform ($\forall_{i,j}$ $f_i = f_j$), and the GA (RWS plus sharing) employs a five-bit encoding in which each possible string is a class. However, only c of the 32 possible classes are placed into the initial population. Since crossover and mutation are turned off ($p_c = p_m = 0$), only those c classes will be present in subsequent populations. Using (8.18) to calculate at which c our bound (8.14) becomes useless, $13 \ln 13 = 33.34 > n$, but $12 \ln 12 = 29.82 < n$: the maximum useful c for $n = 32$ is 12. To verify, if we substitute $c = 13$ and $n = 32$ into the right hand side of (8.14), we obtain $X \leq 1.0036$, an illegally high probability. Substituting $c = 12$ and $n = 32$, on the other hand, yields $X \leq .7412$. We perform 5000 runs at each c and employ $\sigma_{share} = .5$ in the GA. Results are shown in Table 8.2.

As c decreases relative to n, accuracy of the computed bound increases. Note that even for the maximal $c = 12$, the bound is only off by approximately 22.41%. To solve real problems, we will be using a c to n ratio higher than that for even the $c = 5$ case in Table 8.2, which has only 1.75% error. Note that the computed sample mean contains slight sampling error which, as c decreases (to 7 or less), begins to overshadow the bounding error. We do not give drift times for $2 \leq c \leq 4$ because of prohibitive time requirements for 5000 runs. For $c = 2$, as noted earlier, the expected drift time calculated from Equation 8.29 would be exact. For $c = 3$, the calculated drift time would be nearly exact, indistinguishable from the exact value

up to the 10th significant digit. (An expression for the exact drift time can be obtained from Equations 8.12 and 8.21.) We expect that the $c = 4$ bound would also be extremely accurate.

8.3   Population Size: All Classes Desirable

The most effective method of combating genetic drift is to raise the population size to a level sufficient to protect the desired niches. Hence population sizing is a central issue in fitness sharing and, more generally, in the study of niching methods and traditional GAs. In fact, in most GAs, population size is the single, most critical parameter. This section uses the prior genetic drift results to derive bounds on minimum required population sizes for sharing, under the assumption that all classes are desirable. Specifically, this section derives, for three related models of sharing, bounds on the population size required to maintain, with probability γ, a fixed number of niches. The first model assumes all niches are equivalent with respect to fitness. The next two models allow niches to differ with respect to fitness. The models of this chapter do not consider the effects of crossover. These models still have predictive value for GAs with sharing and crossover (and mutation), however, on problems in which crossover and mutation are minimally disruptive. (Sharing's restorative pressure compensates for minor disruptions.) To verify the predictive value of these models, we run GAs with sharing, RWS, and even crossover and mutation, on seven test problems in optimization and classification, using population sizes derived from the models.

A sufficiently large population will ensure that the GA with sharing maintains, with confidence γ, all c desired classes for a sufficient number of generations g — that genetic drift and other disruptive forces do not eliminate any class. Substituting the right-hand side of Equation 8.26 into (8.10) yields the following general relationship for sharing under RWS:

$$1 - \gamma^{\frac{1}{g}} \ \leq \ \sum_{i=0}^{c-1}(1 - P_s(i))^n \ \ .\tag{8.31}$$

From (8.31), we can obtain bounds on n, for c classes of identical fitness, and for c classes of arbitrary fitness. It is important to keep in mind that n in the above equation represents the minimum population size required to maintain c classes for at least g generations with probability of at least γ. If we call the actual population size that the GA employs n', then

any $n' \geq n$ will be sufficient to maintain c classes for at least g generations with probability of at least γ. We call n the *minimum required population size* and n' the *actual population size*.

In practice, we will not be able to obtain an exact n except in the $c = 2$ case. However, we will be able to derive upper bounds on n that also serve as lower bounds on n'. The relationship among n, n', and the derived population-size bound is

$$n \ \leq \ \textit{population-size bound} \ \leq \ n' \ . \tag{8.32}$$

Classes of identical fitness

For c classes of identical fitness, for all classes i, $P_s(i) = 1/c$ (from Equation 8.7). Substituting $1/c$ for $P_s(i)$ in (8.31) and solving for n, yields

$$n \ \leq \ \frac{\ln \frac{1 - \gamma^{\frac{1}{g}}}{c}}{\ln \frac{c-1}{c}} \ \leq \ n' \ . \tag{8.33}$$

The middle term of the above expression is our desired population-size bound. We verify this as follows. Inequality (8.10) tells us that the more trials per generation the GA performs (the higher the actual population size), the smaller the probability X that one or more classes disappear after a generation. Equation 8.25 tells us that the smaller X is, the higher the probability γ that all c classes will be maintained for at least g generations. Therefore, a larger population will yield a higher γ. Hence, any $n' \geq n$ in (8.33) will maintain, with probability of at least γ, c classes of identical fitness for at least g generations.

For large c, the bound in (8.33) can be simplified, using approximation (A.9) of the appendix, as follows:

$$\frac{\ln \frac{1 - \gamma^{\frac{1}{g}}}{c}}{\ln \frac{c-1}{c}} \ \approx \ c \left[- \ln \left(\frac{1 - \gamma^{\frac{1}{g}}}{c} \right) \right] \ . \tag{8.34}$$

This simplification is useful for all but the smallest c. Either bound in (8.34) can also serve as a population sizing tool for c classes of nearly uniform fitness. In practice, the actual population size we employ will be one of the bounds in (8.34), rounded up to the nearest even integer.

For two classes of equal fitness, we can derive an exact expression for required population size, that corresponds to the bound in (8.33) with $c = 2$:

$$n = 1 - \frac{\ln(1 - \gamma^{\frac{1}{g}})}{\ln 2} \approx 1 - 1.4427 \ln(1 - \gamma^{\frac{1}{g}}) \quad . \tag{8.35}$$

Classes of arbitrary fitness: Derivation I

In the general case, inequality (8.31) can not be solved symbolically for n. However, one can obtain a lower bound for n', at or above which, one is guaranteed with confidence of at least γ, to maintain c classes for at least g generations. The representative fitness f_{min} of the least fit peak is at least some percentage, specified by r, of the representative fitness f_{max} of the globally optimal peak. The user need only specify this minimum-to-maximum fitness ratio of the desired peaks, $r = f_{min}/f_{max}$ $(0 < r \leq 1)$. For instance, $r = .8$ would model the maintenance of c peaks, where the fitnesses of all peaks are at least 80% of the fitness of the highest peak. This method for modelling class fitnesses allows population sizing in the presence of sketchy or partial information about the heights of peaks in the fitness landscape. The derivation of required population size proceeds as follows:

$$\forall_i \quad P_s(i) = \frac{f_i}{\sum_{j=0}^{c-1} f_j} = \frac{\frac{f_i}{f_{max}}}{\sum_{j=0}^{c-1} \frac{f_j}{f_{max}}} \quad ; \tag{8.36}$$

$$\forall_i \quad P_s(i) \geq \frac{\frac{f_{min}}{f_{max}}}{\sum_{j=0}^{c-1} \frac{f_{max}}{f_{max}}} \quad ; \tag{8.37}$$

$$\forall_i \quad P_s(i) \geq \frac{r}{c} \quad ; \tag{8.38}$$

$$\sum_{i=0}^{c-1}(1 - P_s(i))^n \leq \sum_{i=0}^{c-1}(1 - \frac{r}{c})^n \quad ; \tag{8.39}$$

$$\sum_{i=0}^{c-1}(1 - P_s(i))^n \leq c(1 - \frac{r}{c})^n \quad . \tag{8.40}$$

Substituting the left-hand side of (8.31) into inequality (8.40) yields

$$1 - \gamma^{\frac{1}{g}} \leq c(1 - \frac{r}{c})^n \quad . \tag{8.41}$$

Solving for n once again yields a lower bound for n':

$$n \le \frac{\ln \frac{1-\gamma^{\frac{1}{g}}}{c}}{\ln \frac{c-r}{c}} \le n' \quad . \tag{8.42}$$

For large c, the above bound can be simplified using approximation (A.9) as follows:

$$\frac{\ln \frac{1-\gamma^{\frac{1}{g}}}{c}}{\ln \frac{c-r}{c}} \approx \frac{c}{r} \left[-\ln \left(\frac{1-\gamma^{\frac{1}{g}}}{c} \right) \right] \quad . \tag{8.43}$$

Note the similarity between the derived population-size bounds in (8.34) and (8.43). In fact, substituting fitness ratio $r = 1$ into (8.43) yields (8.34).

Classes of arbitrary fitness: Derivation II

The fact that required experimental population sizes are in many cases much lower than those derived from (8.42) motivates the search for a tighter bound. We examine at which f_i the right-hand side of (8.10) is maximal (the worst case scenario). Maximizing the right-hand side of (8.10) is equivalent to maximizing the following sum, given f_{min} and f_{max}, under the constraints that for all classes $k \in [0, c-1]$, $f_{min} \le f_k \le f_{max}$:

$$\sum_{i=0}^{c-1} (1 - \frac{f_i}{\sum_{j=0}^{c-1} f_j})^n \quad . \tag{8.44}$$

We can prove that the maximum value of (8.44) is a tighter bound as follows. First, note that for all f_i and f_j in $[f_{min}, f_{max}]$,

$$X \le \sum_{i=0}^{c-1} (1 - \frac{\frac{f_i}{f_{max}}}{\sum_{j=0}^{c-1} \frac{f_j}{f_{max}}})^n \le (1 - \frac{\frac{f_{min}}{f_{max}}}{\sum_{j=0}^{c-1} \frac{f_{max}}{f_{max}}})^n \quad . \tag{8.45}$$

Therefore,

$$X \le \sum_{i=0}^{c-1} (1 - \frac{f_i}{\sum_{j=0}^{c-1} f_j})^n \le c(1 - \frac{r}{c})^n \quad , \tag{8.46}$$

proving that (8.44) is at least as close to X as the right-hand side of (8.41). It follows that (8.44) will also yield a tighter bound on n.

Figure 8.10: The population-size bound from (8.47) is shown as a function of c, at various r, for $\gamma = .95$ and $g = 100$.

Assuming c is an even number, we observe that the maximum for (8.44) occurs when half of the f_i are at f_{max} and half are at f_{min}. Hence, a better upper bound on n (and lower bound on n') is the solution to the inequality,

$$1 - \gamma^{\frac{1}{g}} \leq \frac{c}{2}\left[(1 - \frac{2r}{c(r+1)})^n + (1 - \frac{2}{c(r+1)})^n\right] \quad . \tag{8.47}$$

While one can not solve symbolically for n, one can easily obtain numeric solutions given γ, g, c, and r, via Newton's method. (Newton's method is summarized in the appendix.) As in Derivation I, given c classes of identical fitness ($r = 1$), (8.47) reduces to (8.33).

Figures 8.10–8.12 show the lower bound for n' as a function of c, r, and γ, respectively, with other parameters fixed. Note in Figure 8.10 that with r, γ, and g fixed, the bound for n' is a slightly superlinear function of c. We can derive the form of this relationship between n' and c from the right-hand side of (8.43) as follows:

$$\frac{c}{r}\left[-\ln\left(\frac{1 - \gamma^{\frac{1}{g}}}{c}\right)\right] \;=\; \frac{c}{r}\left[-\ln(1 - \gamma^{\frac{1}{g}}) + \ln c\right] \;=\; k_1 c[k_2 + \ln c] \quad , \tag{8.48}$$

176

Figure 8.11: The population-size bound from (8.47) is shown as a function of r, at various c, for $\gamma = .95$ and $g = 100$.

Figure 8.12: The population-size bound from (8.47) is shown as a function of γ, for $c = 32$, $r = 1$, and $g = 100$.

Table 8.3: This table displays the minimum percentage γ of runs we expect to maintain all c classes, versus the actual percentage from runs of the GA with sharing. Results are given for seven different test problems. For each problem, the actual population size employed n' comes from an appropriate population sizing formula, and $g = 100$.

Problem	γ	Experimental %	c	n'	Runs	p_c	p_m
C-2	95%	95%	2	12	5000	0.0	.00
C-32	75%	91%	32	300	100	0.0	.00
M1	90%	95%	5	44	100	0.9	.01
M2	90%	95%	5	94	100	0.9	.01
M6	95%	100%	25	266	100	1.0	.00
PAR-5	95%	99%	16	166	100	1.0	.00
PAR-10	95%	100%	512	7074	15	1.0	.00

where $k_1 = 1/r$ and $k_2 = -\ln(1 - \gamma^{\frac{1}{g}})$. Therefore,

$$n' \geq k_1 c(k_2 + \ln c) \ ,\tag{8.49}$$

where k_1 and k_2 are positive constants ($k_1 \geq 1$ and $k_2 > 0$). For $r = 1$, the form reduces to

$$n' \geq c(k_2 + \ln c) \ .\tag{8.50}$$

Experimental verification

We now examine the accuracy of the theoretical results on seven test problems of increasing difficulty. In all seven problems, the goal is to locate all peaks. The following parameters and conditions are enforced in all runs of sharing with RWS. Populations are randomly initialized; GAs run for a fixed number of generations $g = 100$; the sharing constant is $\alpha = 1$. Population sizes are derived from an appropriate model, indicated in the accompanying text. We perform a number of runs on each test function. Where indicated, we employ crossover and mutation, in addition to sharing with RWS. The results are summarized in Table 8.3.

To begin, we run a GA with sharing and RWS, 5000 times on the one-bit, two-class problem, $f_0 = f_1 = 1$ (f_i is the representative fitness of Class i). We call this problem C-2 ; C-2 is the simplest possible two-class problem. We start with this problem to verify the accuracy of the two-class population-sizing equation (8.35), in the absence of extraneous forces. We record the percentage of runs that maintain the two classes for the full 100 generations. Equation 8.35,

with $\gamma = .95$ and $g = 100$, tells us to use a population size of at least 11.93. We employ $n' = 12$. Recall that the population sizing equation for the two-class case is exact, rather than a bound. Out of 5000 runs, using $\sigma_{share} = .5$ (perfect discrimination) and $p_c = p_m = 0$, 4748 (95%) maintain both classes for the duration of the run. This obviously compares very favorably with the expected 95% .

Next, we test the multiclass population-sizing formula (8.34) on a five-bit, 32-class problem in which each genotype is a class, and for all classes i, $f_i = 1$. We call this problem *C-32*. This time we try a confidence level of $\gamma = .75$. Substitution into the large-c approximation of (8.34) yields a population-size bound of 298.18. Running the GA 100 times using $n' = 300$, $p_c = p_m = 0$, and $\sigma_{share} = .5$, 91% of the runs maintain all 32 classes to Generation 100.

Note that a significantly larger percentage (than 75%) of the runs maintain all 32 classes on *C-32*. The reason is that although (8.34) is exact for $c = 2$, it is a bound for $c > 2$, and additional confidence is purchased cheaply in terms of population size: from the right-hand side of (8.34), at least 90% confidence requires $n' = 332$; 95%, $n' = 354$; and 99%, $n' = 406$. This relationship between γ and the population-size bound is displayed in Figure 8.12.

The third problem *M1*, shown in Figure 5.2, is a five-peaked sine function, in which all peaks are of uniform height and are spaced evenly throughout the interval $[0, 1]$. Utilizing parameters similar to Deb's (1989), we run a GA with sharing, $p_c = .9$, $p_m = .01$, phenotypic comparison, and $\sigma_{share} = .1$, for 100 runs of 100 generations each. We once again employ population-sizing formula (8.34), since we expect disruption due to crossover and mutation, and the noise introduced by slightly overlapping niches, to be minimal. Substituting $\gamma = .9$, $g = 100$, and $c = 5$ into the right-hand side of (8.34), we obtain a population-size bound of 42.33. (Deb used a population size of 100.) Despite the addition of crossover and mutation, and the slightly imperfect comparison resulting from slightly overlapping niches, with $n' = 44$, 95% of the runs find all classes and maintain them for 100 generations. This is 5% better than the prespecified 90% confidence level.

The fourth problem *M2*, displayed in Figure 5.3, is a five-peaked sine function in which peaks are of nonuniform height, but are spaced evenly throughout the interval $[0, 1]$; the shortest peak is of height .25; the tallest, 1.0. Using the same parameters as we did on *M1* ($p_c = .9$, $p_m = .01$, $\sigma_{share} = .1$, phenotypic comparison; $c = 5$, $\gamma = .9$, $g = 100$), but with $r = .25$, the population-sizing formula (8.47) yields a recommended population size of 94. (Formula (8.42) yields a

recommended population size of 166; Deb used a population size of 100.) For 100 runs under a population of size 94, 95% of the runs find all classes and maintain them for 100 generations. This is once again 5% better than the prespecified 90% confidence level. Note that for both *M1* and *M2*, we are able to undercut Deb's population size of 100, with at least 90% confidence that our runs will converge. The population sizing formulas we employ explain Deb's success in locating and maintaining the five peaks using a population size of 100.

The fifth problem *M6*, displayed in Figure 5.7, is a two-dimensional problem with 25 peaks, ranging in height from 476 to 500. We run a GA with sharing, $p_c = 1$, $p_m = 0$, phenotypic comparison, and $\sigma_{share} = 11$; we perform 100 runs. Once again, we expect the disruption due to crossover and the noise due to overlapping niches to be minimal. We employ sizing formula (8.47), with $\gamma = .95$, $g = 100$, $c = 25$, and $r = .95$. This gives us a recommended population size of 266. (De Jong (1975) employs a population size of 50, but is searching for only one peak.) All 100 runs find and maintain the 25 peaks, once again surpassing the preset confidence level.

The last two problems, *PAR-5* and *PAR-10*, are classification problems or, more specifically, 5- and 10-bit parity problems. We map these to multimodal optimization problems as described in Chapter 5. The solution to *PAR-5* requires the formation and maintenance of 16 disjuncts ($c = 16$); the solution to *PAR-10* requires the formation and maintenance of 512 disjuncts ($c = 512$). Since all classes are of equivalent fitness and since we expect operator disruption and other sources of noise to be minimal, we employ the right-hand-side approximation of population-sizing formula (8.34) with $\gamma = .95$ and $g = 100$. According to (8.34), *PAR-5* requires a population size of at least 166; *PAR-10*, at least 7074.

We run sharing with RWS, $p_c = 1$, $p_m = 0$, genotypic comparison, and $\sigma_{share} = 3.5$. To get the GA to completely converge after application of sharing, a bitclimber (with a neighborhood size of one bit) runs after the GA has completed its 100 generations. Details regarding the bitclimber are in Chapter 10. On *PAR-5*, 99 out of 100 runs converge to the correct final concept (a population containing the desired 16 disjuncts). On *PAR-10*, all 15 runs converge to the correct final concept (a population containing the desired 512 disjuncts). On both parity problems, the number of runs converging again exceeds the preset confidence level.

8.4 Population Size: Desirable and Undesirable Classes

In many problems that one encounters, it is costly or impractical to find all peaks. In the worst case, given a massively multimodal problem, enumerating the local optima may be of the same order of complexity as enumerating the entire search space (Horn & Goldberg, in press). When one does not wish to locate all peaks, one is almost always interested in locating the highest peaks. Locating the highest peaks is equivalent to the Type 3 optimization problem that we described in Chapter 4. This section concentrates upon the solution of this general class of problems. Like earlier, we assume the existence of c local optima. However, instead of wanting to locate all maxima, we are interested in locating at least the b highest maxima, where $b \leq c$. Assuming no tiebreaking is necessary, we call the b highest maxima *desirable* peaks and the $c - b$ lowest maxima *undesirable* or *extraneous* peaks. (If tiebreaking is necessary, some peaks will have the option of being either desirable or undesirable.)

Let us index the b desirable classes from 0 to $b - 1$, and the $c - b$ undesirable classes from b to $c - 1$. We can rewrite Equation 8.4 to separate the summation of desirable and undesirable classes:

$$P_s(i) \; = \; \frac{f_i}{\sum_{j=0}^{b-1} f_j + \sum_{j=b}^{c-1} f_j} \quad . \tag{8.51}$$

Since we do not care if any of the extraneous classes disappear, we can rewrite expression (8.10) as a summation over just the desirable classes. Hence, the probability X of one or more desirable classes disappearing after one generation, becomes

$$X \; \leq \; \sum_{i=0}^{b-1} (1 - P_s(i))^n \quad . \qquad (0 < X < 1) \tag{8.52}$$

The loss distribution and cumulative loss distribution remain the same as previously calculated in Equations 8.20 and 8.24. Therefore, Equation 8.26 also still holds. Substituting the right-hand side of Equation 8.26 for X in (8.52) yields the following relationship for sharing under RWS:

$$1 - \gamma^{\frac{1}{g}} \; \leq \; \sum_{i=0}^{b-1} (1 - P_s(i))^n \quad . \tag{8.53}$$

From (8.53), we can once again obtain bounds on n. We proceed immediately to the most general case, c classes of arbitrary fitness. We relate the minimum population size required to

maintain at least the b best classes for at least g generations with probability of at least γ, to n and n'.

In order to determine the f_i at which the right-hand side of (8.53) is maximal, we first make a few definitions. As before, let f_{max} be the fitness of the global maximum or maxima. Let f_{min} be the fitness of the least fit, desirable class or classes. Let f_{extra} be the fitness of the most fit, extraneous class or classes. The following ordering holds: $f_{extra} \le f_{min} \le f_{max}$. We define the following ratios: $r = f_{min}/f_{max}$ and $r' = f_{extra}/f_{max}$. Note that $r' \le r$. We would like to maximize the right-hand side of (8.53), which corresponds to the following sum:

$$\sum_{i=0}^{b-1}(1 - \frac{f_i}{\sum_{j=0}^{b-1}f_j + \sum_{j=b}^{c-1}f_j})^n \quad . \tag{8.54}$$

The above sum is maximized when the expression $\sum_{j=b}^{c-1}f_j$ is maximal, which occurs when for all extraneous classes j, $f_j = f_{extra}$. Letting s be the number of extraneous classes $(s = c - b)$, the maximization of the above sum reduces to the maximization of

$$\sum_{i=0}^{b-1}(1 - \frac{f_i}{\sum_{j=0}^{b-1}f_j + sf_{extra}})^n \quad . \tag{8.55}$$

Assuming b is an even number, the maximum for (8.55) occurs when half of the f_i are at f_{max} and half are at f_{min}. Substituting half f_{max} and half f_{min} yields the following expression for the maximum:

$$\frac{b}{2}\left[(1 - \frac{f_{min}}{\frac{b}{2}(f_{max} + f_{min}) + sf_{extra}})^n + (1 - \frac{f_{max}}{\frac{b}{2}(f_{max} + f_{min}) + sf_{extra}})^n\right] \quad . \tag{8.56}$$

Simplification yields a lower bound on n', the solution to the inequality,

$$1 - \gamma^{\frac{1}{g}} \le \frac{b}{2}\left[(1 - \frac{2r}{b(r+1) + 2r's})^n + (1 - \frac{2}{b(r+1) + 2r's})^n\right] \quad . \tag{8.57}$$

Given γ, g, b, s, r, and r', one can easily obtain numeric bounds for n and n' via Newton's method. To make matters simpler, one could substitute for both r and r' an intermediate value that denotes a threshold between desirable fitness ratios and undesirable fitness ratios. When all classes are desirable $(b = c, s = 0, \text{ and } r' = r)$, (8.57) reduces to (8.47).

We test population sizing formula (8.57) on a constructed eight-bit problem. In the problem we construct, each of the 256 individual solutions is in a separate class ($c = 256$). We designate four solutions — 00111111, 01111111, 10111111, and 11111111 — as desirable, and assign them fitnesses of 4. We designate the other 252 solutions as undesirable and assign them fitnesses of 1. Formula (8.57), with $b = 4$, $s = 252$, $r = 1$, $r' = .25$, $g = 100$, and $\gamma = .90$, recommends a population size of 550. This is a considerable savings over the population size of 7488 recommended by prior formula (8.47) that assumes we want to maintain all 256 classes. We perform 90 runs of a GA with sharing, RWS, and genotypic comparison, employing $n' = 550$, $p_c = 1$, $p_m = 0$, $\sigma_{share} = .5$, and $\alpha = 1$. 100% of the runs locate all four desired optima and maintain them to Generation 100. Population sizing according to our model of desirable and undesirable classes exceeds the prespecified confidence level of 90%.

Chapter 9

Sharing: Selection Plus Crossover

The sharing models developed in the previous chapter bound the behavior of a full GA with sharing on problems in which crossover and other operators are not overly disruptive. The current chapter incorporates crossover into the modelling. The incorporation of crossover provides a template for modelling various sources of noise, such as mutation.

Recall that $P_s(i)$ is the probability of selecting some element of an arbitrary class i, in one trial of sharing with RWS. Let j be an arbitrary class. In two consecutive trials of sharing with RWS, the probability of selecting some element of i, followed by some element of j (in that order) is $P_s(i)P_s(j)$. We use the following notation to indicate the probability of producing one Class k and one Class l element by crossing a Class i element (the first parent) with a Class j element (the second parent):

$$p(i \times j \Rightarrow \{k, l\}) \ . \tag{9.1}$$

Given an arbitrary class A, the probability of producing either one or two Class A elements by crossing a Class i element (first parent) with a Class j element (second parent) is

$$p(i \times j \Rightarrow \{A, A\}) + \sum_{k=0, k \neq A}^{c-1} p(i \times j \Rightarrow \{A, k\}) \ . \tag{9.2}$$

The probability of not producing any Class A elements in such a cross is

$$Y(i, j, A) \ = \ 1 - p(i \times j \Rightarrow \{A, A\}) - \sum_{k=0, k \neq A}^{c-1} p(i \times j \Rightarrow \{A, k\}) \ . \tag{9.3}$$

The probability of selecting an element of Class i followed by an element of Class j (in that order), and after applying crossover with probability p_c, obtaining two elements from outside Class A, is

$$P_s(i)P_s(j)\Big[\; p_c \, Y(i,j,A) + (1 - p_c)(1 - P_s(A))^2 \; \Big] \quad . \tag{9.4}$$

After the combined operation of selecting two elements (via sharing with RWS) and crossing them with probability p_c, the probability of not producing any Class A elements is

$$\sum_{i=0}^{c-1}\sum_{j=0}^{c-1}\Big[\; p_c \, P_s(i)P_s(j) \, Y(i,j,A) + (1 - p_c)P_s(i)P_s(j)(1 - P_s(A))^2 \; \Big] \quad . \tag{9.5}$$

Simplification yields

$$(1 - p_c)(1 - P_s(A))^2 + p_c \sum_{i=0}^{c-1}\sum_{j=0}^{c-1} P_s(i)P_s(j) \, Y(i,j,A) \quad . \tag{9.6}$$

Substituting back the expression for $Y(i,j,A)$, the probability of not having any Class A elements in the population after one generation ($n/2$ of the operations in (9.6)) is

$$\Bigg[\quad (1 - p_c)(1 - P_s(A))^2 + p_c$$
$$- \quad p_c \sum_{i=0}^{c-1}\sum_{j=0}^{c-1} P_s(i)P_s(j) \, p(i \times j \Rightarrow \{A, A\})$$
$$- \quad p_c \sum_{i=0}^{c-1}\sum_{j=0}^{c-1} P_s(i)P_s(j) \sum_{k=0, k\neq A}^{c-1} p(i \times j \Rightarrow \{A, k\}) \quad \Bigg]^{n/2} \quad . \tag{9.7}$$

Since any of the c classes can potentially lose all its elements, the probability X of one or more classes disappearing after one generation is

$$X \; \leq \; \sum_{m=0}^{c-1} \Bigg[\quad (1 - p_c)(1 - P_s(m))^2 + p_c$$
$$- \quad p_c \sum_{i=0}^{c-1}\sum_{j=0}^{c-1} P_s(i)P_s(j) \, p(i \times j \Rightarrow \{m, m\})$$
$$- \quad p_c \sum_{i=0}^{c-1}\sum_{j=0}^{c-1} \sum_{k=0, k\neq m}^{c-1} P_s(i)P_s(j) \, p(i \times j \Rightarrow \{m, k\}) \quad \Bigg]^{n/2} \quad . \tag{9.8}$$

The probability of maintaining all classes for a generation is $1 - X$. The above inequality takes into account gains as well as losses of class elements, yielding an extremely detailed model of the GA with sharing and crossover.

Expression (9.8) can be specialized and simplified to produce models of greater utility. Consider all possible crosses from the perspective of an arbitrary class A. We partition all such crosses into the following nine cases:

1. $A \times A \Rightarrow \{A,A\}$,

2. $A \times non\text{-}A \Rightarrow \{A,A\}$,

3. $non\text{-}A \times non\text{-}A \Rightarrow \{A,A\}$,

4. $A \times A \Rightarrow \{A,non\text{-}A\}$,

5. $A \times non\text{-}A \Rightarrow \{A,non\text{-}A\}$,

6. $non\text{-}A \times non\text{-}A \Rightarrow \{A,non\text{-}A\}$,

7. $A \times A \Rightarrow \{non\text{-}A,non\text{-}A\}$,

8. $A \times non\text{-}A \Rightarrow \{non\text{-}A,non\text{-}A\}$, and

9. $non\text{-}A \times non\text{-}A \Rightarrow \{non\text{-}A,non\text{-}A\}$.

We will now eliminate less significant terms from expression (9.8). Note that the inequality still holds after elimination of these terms. Our modelling framework tells us that crossing two elements from the same class always produces two elements of that class. Therefore, the right-hand side of Case 1 has probability 1, given the left-hand side, and the right-hand sides of Cases 4 and 7 have probability 0, given the left-hand sides. Cases 2, 3, and 6 involve gaining elements of Class A. We will assume such gains are negligible (leaving Cases 5, 8, and 9 as the only remaining cases of interest). Therefore, $p(i \times j \Rightarrow \{m, m\}) = 1$ if $i = m$ and $j = m$; 0, otherwise. In addition, $p(i \times j \Rightarrow \{m, k\}) = 0$ if $i \neq m$ and $j \neq m$. We rewrite (9.8) as

$$X \leq \sum_{m=0}^{c-1} \Bigg[\quad (1 - p_c)(1 - P_s(m))^2 + p_c - p_c(P_s(m))^2$$
$$- 2p_c \sum_{j=0, j \neq m}^{c-1} \sum_{k=0, k \neq m}^{c-1} P_s(m)P_s(j)\, p(cross(m, j) \Rightarrow \{m, k\}) \quad \Bigg]^{n/2} . \quad (9.9)$$

We use the notation $cross(m, j)$ in place of $m \times j$ to indicate that order does not matter; our crossover operator is symmetric.

We now examine the probability of losing an element of Class m via crossover with an element of a different class (Case 8). For all classes $a \neq m$, in a cross between elements of a and m, the probability of producing one Class m element is greater than or equal to the probability that crossover does not disrupt m. In other words, disruptions happen more frequently than actual losses, because losses are caused only by disruptions, but not every disruption leads to a loss. (A class may be regained after a disruptive cross.)

$$\forall_{a \neq m} \sum_{k=0, k \neq m}^{c-1} p(cross(m, a) \Rightarrow \{m, k\}) \geq (1 - p_d(m)) , \tag{9.10}$$

where $p_d(m)$ is the probability that crossover disrupts m. The above bound yields a further simplification:

$$X \leq \sum_{m=0}^{c-1} \Big[\quad (1 - p_c)(1 - P_s(m))^2 + p_c(1 - (P_s(m))^2)$$
$$- 2p_c(1 - p_d(m))P_s(m)(1 - P_s(m)) \quad \Big]^{n/2} . \tag{9.11}$$

With $p_c = 0$ (no crossover) or $\forall_m \ p_d(m) = 0$ (disruptionless crossover), (9.11) simplifies to the previously derived expression (8.10) for selection alone. Note that under full crossover ($p_c = 1$), the first term inside the above summation drops out. Under full disruption ($\forall_m \ p_d(m) = 1$), the third term drops out; consequently, in order for a class to survive to the next generation, either at least one cross must occur involving only elements of that class, or the class must pass through at least once without crossover. For schema-based class definitions, $p_d(m) = \delta(m)/(l - 1)$ for single-point crossover, and $p_d(m) = 1 - .5^{o(m)}$ for uniform crossover, where δ is the defining length and o is the order of a schema, as defined in Chapter 2.

The previously derived Equation 8.26, $X = 1 - \gamma^{\frac{1}{g}}$, also holds for sharing with crossover. Therefore,

$$1 - \gamma^{\frac{1}{g}} \leq \sum_{m=0}^{c-1} \Big[\quad (1 - p_c)(1 - P_s(m))^2 + p_c(1 - (P_s(m))^2)$$
$$- 2p_c(1 - p_d(m))P_s(m)(1 - P_s(m)) \quad \Big]^{n/2} . \tag{9.12}$$

From (9.12), we can once again obtain bounds on n and n'. We give the derivations in the following sections for two cases. In the first case, we assume that all classes are desirable and that we have c classes of identical fitness. The second, more general case, models both desirable and undesirable classes of arbitrary fitness. In the remainder of this chapter, we assume that disruption probabilities for all classes can be bounded by a maximal disruption probability p_d. The inequality (9.12) still holds if we substitute p_d for all $p_d(m)$.

9.1 All Classes Desirable: Identical Fitnesses

For c classes of identical fitness,

$$
1 - \gamma^{\frac{1}{g}} \ \leq \ c \left[\ (1 - p_c)(1 - \frac{1}{c})^2 + p_c(1 - \frac{1}{c^2}) - 2p_c(1 - p_d)\frac{1}{c}(1 - \frac{1}{c}) \ \right]^{n/2} \ . \tag{9.13}
$$

Solving for n gives an overestimate for the minimum population size required to maintain, with probability γ, c classes of identical fitness for at least g generations:

$$
n \ \leq \ \frac{2 \ln \frac{1 - \gamma^{\frac{1}{g}}}{c}}{\ln \left[\ (1 - p_c)(1 - \frac{1}{c})^2 + p_c(1 - \frac{1}{c^2}) - 2p_c(1 - p_d)\frac{1}{c}(1 - \frac{1}{c}) \ \right]} \ \leq \ n' \ . \tag{9.14}
$$

Figure 9.1 shows bounded population size as a function of c, at various p_d, with other parameters fixed ($\gamma = .95$, $g = 100$, $p_c = 1$). Note that even with full crossover, the slightly superlinear relationship is maintained between the population-size bound and c. Note also that population size at first increases slowly as the probability of disruption increases. However, the rate of increase quickens as p_d approaches 1. It is illuminating to compare the bounds on X (calculated from (9.11)) at two extremes, $p_d = 0$ and $p_d = 1$, for c classes of identical fitness, with $p_c = 1$. At $p_d = 0$,

$$
X \ \leq \ c(1 - \frac{1}{c})^n \ . \tag{9.15}
$$

At $p_d = 1$,

$$
X \ \leq \ c(1 - \frac{1}{c^2})^{n/2} \ . \tag{9.16}
$$

Figure 9.1: The population-size bound of (9.14) is shown as a function of c, at various disruption probabilities p_d, with $\gamma = .95$, $g = 100$, and $p_c = 1$, for classes of identical fitness.

9.2 Desirable and Undesirable Classes: Arbitrary Fitnesses

Our starting point is the summation in (9.12), over b desirable classes, with p_d substituted for $p_d(m)$:

$$
1 - \gamma^{\frac{1}{g}} \leq \sum_{m=0}^{b-1} \Big[\quad (1 - p_c)(1 - P_s(m))^2 + p_c(1 - (P_s(m))^2)
$$
$$
- \quad 2p_c(1 - p_d)P_s(m)(1 - P_s(m)) \quad \Big]^{n/2} \quad , \tag{9.17}
$$

where, as in the previous chapter,

$$
P_s(m) = \frac{f_m}{\sum_{j=0}^{b-1} f_j + \sum_{j=b}^{c-1} f_j} \quad . \tag{9.18}
$$

As before, the b desirable classes are indexed from 0 to $b - 1$; the $c - b$ undesirable classes, from b to $c - 1$.

We maximize the right-hand side of (9.17) over the $P_s(m)$, as we did in the previous chapter. The maximal value occurs when all $P_s(m)$ are minimal. Like last time, maximization of the

189

right-hand side of (9.17) reduces to the simultaneous minimization, for all desired m, of

$$\frac{f_m}{\sum_{j=0}^{b-1} f_j + s f_{extra}} \quad , \tag{9.19}$$

where s is the number of extraneous peaks and f_{extra} is the fitness of the fittest extraneous peak. This minimum occurs when half of the f_m for the b desired peaks are at f_{min} and half are at f_{max}. For the half at f_{min},

$$U = P_s(min) = \frac{f_{min}}{\frac{b}{2}(f_{max} + f_{min}) + s f_{extra}} = \frac{2r}{b(r+1) + 2r's} \quad , \tag{9.20}$$

where $r = f_{min}/f_{max}$, and $r' = f_{extra}/f_{max}$, as in the previous chapter. For the half at f_{max},

$$V = P_s(max) = \frac{f_{max}}{\frac{b}{2}(f_{max} + f_{min}) + s f_{extra}} = \frac{2}{b(r+1) + 2r's} \quad . \tag{9.21}$$

A lower bound on n' comes from the solution to the inequality,

$$1 - \gamma^{\frac{1}{g}} \leq \frac{b}{2} \Big[(1 - p_c)(1 - U)^2 + p_c(1 - U^2) - 2p_c(1 - p_d)U(1 - U) \Big]^{n/2} +$$
$$\frac{b}{2} \Big[(1 - p_c)(1 - V)^2 + p_c(1 - V^2) - 2p_c(1 - p_d)V(1 - V) \Big]^{n/2} \quad . \tag{9.22}$$

Given γ, g, b, s, p_c, p_d, r, and r', one can obtain numeric bounds for n and n' via Newton's method. Again, one could substitute an intermediate threshold r'' for both r and r'.

Figure 9.2 illustrates the effect of varying p_c when disruption is maximal ($p_d = 1$). In this worst case scenario, crossover is always destructive, never constructing any desirable solutions. Population-size bounds are derived from expression (9.22), using $\gamma = .95$, $g = 100$, $r'' = 1$, $b = 32$, and $s = 1000$. As the figure shows, given high disruption, required population size increases exponentially as p_c approaches 1. A crossover probability of $p_c = .6$ requires a population size of roughly $28,500$; $p_c = .7$, $38,000$; $p_c = .8$, $57,000$; $p_c = .9$, $113,000$; and $p_c = 1$, $23,500,000$. The model tells us to steer away from crossover probabilities close to 1. Of course, one should not make p_c too low, to avoid stifling exploration and the location of niches.

Figure 9.2: The population-size bound of (9.22) is shown as a function of p_c, for $\gamma = .95$, $g = 100$, $p_d = 1$, $r'' = 1$, $b = 32$, and $s = 1000$.

9.3 Further Research

There are numerous promising extensions to the research presented in this chapter. We briefly mention some of the possibilities.

We have modelled the major forces in GAs that incorporate fitness sharing: selection, sharing, and crossover. In order to apply a sharing model to problem solving, one must have a method of setting that model's parameters. Although the models of Chapter 8 require just a few parameters, the most general model from the current chapter, corresponding to expression (9.22), requires eight parameters: γ, g, b, s, p_c, p_d, r, and r'. The parameter γ is easy to set: the user merely specifies a minimum level of confidence, such as 95%, that he/she would like to have in the values produced by the model. The number of generations g is also easy to set. GA theory tells us that the minimum number of generations should be of the order of some multiple of $\log n$ (or alternatively, $\log l$, where l is the string length), where selection pressure determines the base of the logarithm. The user can set g to some arbitrary value greater than the number of generations the GA will run. We have been using $g = 100$. As shown in the next

chapter's empirical results, GAs with sharing typically require nowhere near 100 generations to finish running. Furthermore, our models are not sensitive to increases in g.

A user can set the parameter b by specifying the minimum number of sufficiently diverse solutions the GA is to return. Of course, if the fitness landscape does not contain b peaks, sharing will be unable to find b peaks. A user can set r and r' by substituting a parameter r'' for both r and r', where $r' \leq r'' \leq r$; r'' indicates a desired quality of solution. The user can specify that the GA return at least b diverse solutions, all of which are within $100 \times r''$ percent of the quality or fitness of the global optimum or optima.

The modelling parameter p_d can be set in several ways. One can assume minimal disruption and set it to zero, as we did in the models of Chapter 8. Alternatively, one can assume the worst and set it to 1, in which case p_c should be set to a value less than 1 to allow at least some solutions to pass through from generation to generation, undisrupted. The effect of intermediate values of p_d is shown in Figure 9.1.

The user should set p_c to whatever value of p_c is to be used in the GA. We recommend a value of less than 1 for real-world problems, because if good solutions and subsolutions are susceptible to disruption, we would like many of them to be passed through untouched from generation to generation. Values of p_c between .6 and .9 are most common in the GA literature. Figure 9.2 gives some insight into why these are popular values.

The modelling parameter that might give a user the most trouble is s, the number of extraneous peaks in the landscape. One approach to setting this parameter, given an unknown fitness landscape, is to attempt overestimation. That is, one can simply choose a large number for s. Problem-specific information can be applied whenever available. If none is immediately available, one approach to obtaining some is dynamic estimation. From prior runs or from early generations, one might be able to develop some idea of the number of peaks in the search space. Another approach would be to run hillclimbers from every point within several small, randomly chosen regions, and to count the number of unique optima returned. From this sample, one could form a rough estimate of the modality of the search space. Methods for estimating the number of peaks in the fitness landscape deserve further research.

One can model minor forces in GAs that incorporate sharing, in much the same way we model crossover — as potential disrupters or sources of noise. Mutation, for instance, can be modelled similar to the way it is incorporated in the schema theorem: the probability that

mutation disrupts schema H is $(1 - p_m)^{o(H)}$. Whether or not class definitions correspond to schema definitions, we expect that mutation will disrupt a class very rarely, assuming that the usual low rates of mutation are employed. Higher mutation rates can be incorporated into a model if one can estimate the probability of mutation disrupting each individual class.

If sampling is employed in the computation of shared fitnesses, it would be useful to model the additional noise introduced by sampling, in order to determine the minimum required sample size. Goldberg and Richardson (1987) suggest that to compute an individual's niche count (sum of sharing function values between it and all population elements), k samples of the population be taken, where $k \ll n$. They call the mean sharing function value for the sample μ_i, and set niche count equal to $(n - 1)\mu_i + 1$. The '1' subtracted and added is for the individual itself, since all individuals share fully with themselves. We do not employ sampling in this thesis.

Another useful extension of this study would be to model class formation as well as maintenance. We have already modelled the effects of higher crossover rates. Considering the constructive as well as the destructive effects of crossover would perhaps yield an expression for the optimal crossover probability in a GA with sharing. As mentioned earlier, niche formation is an issue of signal-to-noise within each niche (Goldberg, Deb, & Clark, 1992). It has been the case in practice that interniche population sizing considerations (niche maintenance) override intraniche considerations (niche formation). This has allowed a population that is sized to maintain niches, more than enough elements to form the niches. However, this may not be the case for some problems. A potential combined method is to size populations using both approaches, and to employ the maximum recommended population size, perhaps adjusted upwards by a constant to account for interactions.

We have so far used only RWS in our models of sharing, with the understanding that computed bounds will also apply to sharing under any of the more stable, fitness-proportionate selection schemes. (We in fact employ SUS rather than RWS in Chapter 10's runs of a GA with sharing.) One could attempt to locate tighter bounds by modelling individual fitness-proportionate selection schemes that are less noisy than RWS. It is not clear how well our models apply to sharing under ranking and tournament selection schemes. Selection methods such as tournament selection have not usually been combined with sharing, to avoid stability

problems. Oei, Goldberg, and Chang (1991), however, propose a new tournament-sharing scheme that is more stable than the naive combination.

It is possible to allow flexibility in the definition of a class's representative fitness. Deb (1989), for instance, represents the mean and the variance of fitnesses, among the individuals clustered atop each peak. We have chosen an alternative approach in order to preserve the independence of successive generations. However, a hybrid model may be possible in which successive generations are still independent, but class fitnesses are fuzzy rather than definite, and are modelled as sources of disruption or noise.

One final possibility is to model overlapping niches. Horn (1993) did this in his Markov chain model, but limited his attention to one-bit problems. We have chosen not to model overlapping niches because our models are already predictive in cases of minor overlap, because the effectiveness of sharing diminishes as overlap increases, and to preserve the independence of successive generations. Again, we could attempt to model overlap as a source of disruption, in order to preserve independence.

Chapter 10

Parallel Versus Sequential Niching

We have analyzed crowding and sharing methods in depth in the previous four chapters, and through modelling, have gained considerable knowledge regarding better parameter settings and improved algorithmic designs. We previously designated crowding and sharing to be parallel niching methods, since they conceptually form and maintain niches simultaneously within a single population (regardless of the number of processors employed). Earlier, we also discussed sequential niching (SN) methods — techniques that locate multiple niches temporally. In this chapter, we test crowding and sharing on several problems of increasing difficulty. We compare their behaviors to those of both a parallel hillclimbing method we develop and a sequential niching method. Parallel hillclimbing and sequential niching serve as benchmarks for examining the power of parallel niching methods. We illustrate the strengths and the limitations of all four niching methods.

Our parallel hillclimber starts with a randomly generated initial population, and ideally forces each element to converge to its nearest attractor. Attractors are defined with respect to some neighborhood operator that is appropriate to the problem being solved. The neighborhood operator may be either phenotypic (operating over variable space) or genotypic (operating over Hamming or bit space).

The parallel hillclimbing algorithm is similar in functionality to binary search. When operating in phenotypic space, the hillclimber starts with a large neighborhood size. Each population element hillclimbs until it can no longer improve. The hillclimber then cuts its neighborhood size in half, and each population element again hillclimbs until it can not further improve. The

hillclimber repeats the previous step until it has hillclimbed using a neighborhood size of ϵ, where ϵ is the smallest increment encoded by the GA. The phenotypic neighborhood operator simply consists of either an addition or a subtraction of the neighborhood size from one of the problem's variables. Figure 10.1 gives an algorithmic description of the phenotypic hillclimber.

The initial neighborhood size requires some care in setting. Too small a neighborhood will delay the hillclimber's convergence. For example, given a unimodal function of a single binary-encoded variable, a random starting point, and an initial neighborhood size of ϵ, convergence time is, on the average, exponential in encoding length. Too large a neighborhood, on the other hand, will cause points to leave their local neighborhoods and skip from local optimum to local optimum. Unfortunately, parallel hillclimbing has no mechanism for allocating individuals to optima. Therefore, to locate multiple optima it is preferable that each point remain within the basin of attraction (of the optimum) in which it currently resides. This strategy prevents repetitive location of the same few optima. For genotypic hillclimbing, also called *bitclimbing*, we set all neighborhood sizes to one bit.

Both genotypic and phenotypic hillclimbers implement versions of *next-ascent* hillclimbing (Mühlenbein, 1991, 1992). Starting with a randomly chosen variable, they cycle through the variables, trying perturbations on each one. For the phenotypic hillclimber, a perturbation is either an upward or a downward change in the value of a variable. For the genotypic hillclimber, a perturbation is the flipping of a bit-variable. The hillclimbers take each improvement they find. They terminate, for each neighborhood size, after they have completed a full cycle through the variables without improvement.

The parallel hillclimber is an important base for comparison, because if a niching method is merely shuffling points and then converging to the nearest attractor, there is no point in pursuing that niching method — parallel hillclimbing is superior. We instead seek niching methods that intelligently decide which maxima to pursue, and that tend to prefer higher maxima over lower maxima. Note that the hillclimbing method we have chosen may not be the most efficient of all possible hillclimbers. However, it has predictable behavior, very close to that of the theoretical, deterministic hillclimber of our modelling framework. The parallel hillclimber will tend to climb to the nearest local optimum in whose basin of attraction it currently lies, a desirable property for a hillclimber that strives to locate multiple optima. Also to its merit, our hillclimber is general purpose, requiring exactly the same problem information as the GA.

196

Parallel Hillclimbing (Phenotypic)

1. Initialize *Neighborhood Size*

2. WHILE *Neighborhood Size* $\geq \epsilon$

    (a) FOR each population element

        • Randomly pick a starting variable

        • *Change* = TRUE

        • WHILE *Change*

            – *Change* = FALSE

            – FOR each variable

                ∗ IF adding neighborhood size to current variable yields improved fitness

                    · Perform the addition

                    · *Change* = TRUE

                ∗ ELSE IF subtracting neighborhood size from current variable yields improved fitness

                    · Perform the subtraction

                    · *Change* = TRUE

    (b) *Neighborhood Size = Neighborhood Size / 2*

Figure 10.1: Pseudocode is given for the phenotypic variation of parallel hillclimbing.

The sequential niching method we consider is that of Beasley, Bull, and Martin (1993). This is the most promising sequential niching method proposed to date. It works by iterating a simple GA, and maintaining the best solution of each run off-line. The authors call the multiple runs that sequential niching performs to solve a single problem, a *sequence*. To avoid converging to the same area of the search space multiple times, whenever their algorithm locates a solution, it depresses the fitness landscape at all points within some radius of that solution. This *niche radius* plays a role in SN similar to that of σ_{share} in sharing. In fact, Beasley et al. suggest that SN is a sequentialization of fitness sharing. We examine this possibility later in the chapter.

Like the parallel hillclimber, sequential niching is an important base to which parallel niching methods can be compared. If a parallel niching method can offer no advantage, there is no point in wasting time with it — one would be better off locating solutions sequentially. As we soon demonstrate, the advantages of parallel niching methods go beyond mere aesthetics or consistency with biological parallelism: there are many things parallel niching can accomplish that sequential niching can not. Beasley et al. mention three potential advantages of sequential

niching. The first is simplicity: SN is conceptually a simple add-on to existing optimization methods. The second is the ability to work with smaller populations, since the goal during each run of a sequence is to locate only one peak. The third and most important is speed, and is partially a by-product of the second. We find that the latter two potential advantages never materialize. In fact, many disadvantages quickly become apparent. These include:

- Loss, through deration, of optimal solutions and their building blocks;

- Repeated search of depressed portions of the space;

- Repeated convergence to the same solutions;

- Loss of cooperative population properties, including cooperative problem solving and niche maintenance along the way to a single solution;

- Slower runtime, even on serial machines.

10.1 Methodology

We attempt to keep corresponding parameters the same across algorithms whenever possible. Our secondary emphasis is choosing parameters that are best for each of the four algorithms.

SN and sharing run under stochastic universal selection. We employ SUS because it is the least noisy of commonly used, unbiased, fitness-proportionate selection methods. DC and parallel hillclimbing have built-in selection mechanisms. None of the selection methods utilize fitness scaling.

As our modelling framework suggests, GAs are global optimization methods that operate best when combined with local optimization methods such as hillclimbers. Therefore, after each GA terminates, including GAs internal to SN sequences, hillclimbing is invoked upon that GA's final population. This allows the three GA-based niching methods to always locate actual optima rather than near optima. We use the parallel hillclimbing algorithm as our post-GA hillclimber. (There is no need, of course, to call the parallel hillclimber to optimize its own final population.) For sequential niching (and for the other three niching methods as well), hillclimbing occurs on the original, non-derated, fitness landscape. This is one of Beasley et al.'s suggested improvements to their algorithm.

We employ full crossover ($p_c = 1$) and no mutation ($p_m = 0$) in all three genetic algorithms. Without mutation, runs converge faster. Hillclimbing upon termination, using the parallel hillclimbing algorithm, substitutes for mutation. This substitution of hillclimbing for mutation highlights better the merits and drawbacks of the four algorithms as well as the differences between algorithms. The absence of mutation also gives more consistent (less noisy) results from run to run of a GA.

All four niching algorithms return all "unique" elements in the final population (post-hillclimbing) as solutions. Uniqueness is determined by applying the same distance measure as the niching method uses. If a new solution is greater than some threshold (a distance of .001 in this study) from all previously located solutions, the new solution is added to the list of final solutions. This allows dual solutions that occur on Hamming cliffs to be treated as a single solution. The final population for SN consists of the combination of final populations (post-hillclimbing) of all runs in a sequence.

The problems we consider are *M1–M9*, *MUX-6*, and *PAR-8*, of Chapter 5. We run each algorithm — parallel hillclimbing, SN, sharing, and DC — 10 times on each problem, employing the same 10 random number seeds (and initial populations) for each algorithm. The runs internal to an SN sequence all use the same random number seed (and initial population). This guarantees that the discovery of a new solution is due to the deration of the fitness landscape and not to random factors. This also provides a convenient stopping point for the SN algorithm at insufficient population sizes. Once a sequence goes one run without generating a new solution, it will do so perpetually. (We derate the area about each returned solution only once.) There is no need to do more runs in a sequence than there are peaks; if no new solution is returned during a particular run, the sequence is stopped.

We could alternatively restart the runs of an SN sequence with new, randomly generated populations each time. This would require that SN perform additional runs until either it finds all desired solutions or exceeds a total function-evaluations limit. This alternative would bring into question whether the lowest population size that succeeds is optimal. One could imagine a population of size $n = 2$ restarting thousands of times, until it begins with a population element close enough to a desired optimum to converge to that optimum. A higher population size, on the other hand, might more easily find the same solution in much faster time. Our methodology implies that the lowest n to locate all peaks is most likely the optimal population size for the

particular algorithm and problem under consideration (because the number of generations a GA runs tends to increase with population size).

We terminate each run after it effectively stops improving. This stopping criterion applies to full runs of parallel niching methods as well as to internal runs of SN sequences. (Parallel hillclimbing is an exception; it runs until no further improvement occurs.) We determine the point of little further improvement in a fashion similar to Beasley et al., by employing a halting window of five generations. Call the current generation t_0, and the prior four generations, t_{-1}, t_{-2}, t_{-3}, and t_{-4}. If the average fitness of the population at t_0 is not more than some increment *inc* greater than the average fitness of the population at t_{-4}, the run halts. We use *inc* = .001 on all problems except *M6*, where we use *inc* = .1. (Although Beasley et al. use *inc* = 0, we employ nonzero *inc* parameters because the hillclimber is responsible for the end convergence of our algorithms.) Runs of an SN sequence terminate when the average *derated* fitness of the population stagnates. Sharing and DC terminate when the average *raw* fitness stagnates.

The stopping criterion for the overall SN algorithm is not easy to set, at least when the algorithm is successful. We give SN ideal behavior by stopping successful sequences immediately after they have located all desirable peaks. Unsuccessful sequences are stopped after some run fails to yield a new solution, or after they exceed a maximum number of function evaluations.

For all four algorithms, we set a number of function evaluations, above which we give up hope that an algorithm will be able to solve a problem at reasonable population sizes. This number is 1.5 million GA function evaluations or, alternatively, 2 million combined function evaluations for the GA plus the parallel hillclimber.

We employ the following niche radius for each problem. This value serves as the niche radius for sequential niching, σ_{share} for sharing, and the initial neighborhood size (lowered to the nearest power of two) for parallel hillclimbing, both on its own and when tacked on to the end of a GA. *M1–M4* use a phenotypic niche radius of .1; *M5*, 4.24 (phenotypic); *M6*, 8.0 (phenotypic); *M7* and *M8*, 5.5 (genotypic); *M9*, 2.5 (genotypic); *MUX-6* and *PAR-8*, 3.5 (genotypic). We choose these values to allow discrimination among desired peaks.

We utilize sharing function (4.2), but with $\alpha = 2$. Beasley et al. recommend $\alpha = 2$ rather than $\alpha = 1$, in order to produce lower false optima in the derated fitness landscape. Like Beasley et al., we set SN's derating function to 1 minus the equivalent sharing function.

For SN, we make an additional modification, to its benefit. For consistency with the other niching methods, SN returns all unique, previously unlocated elements of each final population (post-hillclimbing) as solutions. In case SUS successfully locates and maintains multiple solutions, we allow SN to use them all. We then derate about each solution, as if multiple runs had been performed. Note that Beasley et al. return only the single best solution of the population. We have found that returning only the best solution yields worse results.

We employ a single performance criterion that is common to all algorithms. Assume that all algorithms, given a sufficiently large population, will be able to solve all problems that we consider. This is not an unrealistic assumption, since the probability of all desired solutions appearing in the initial population approaches one as population size approaches infinity. Our performance criterion is the total number of function evaluations at the minimum population size at which an algorithm returns all desired solutions. (We do not penalize or reward an algorithm for returning extraneous solutions.) Over multiple trials, the fewer the average number of function evaluations, the better an algorithm is for solving a particular problem. This performance criterion is consistent with our previous, theoretical, population-sizing models, in which we bound the minimum population size required for a run to converge properly. In the current chapter, we instead compute empirical bounds.

We start with the minimum population size that is both a power of two and large enough to locate all peaks ($n = 2$ for SN). We double n until the algorithm locates and maintains to termination all desired peaks. That is, we redo all algorithmic trials at higher and higher n, up to the limit on the number of function evaluations, until one trial correctly converges. We use only population sizes that are powers of two, in order to avoid massive numbers of trials and fine-grained distinctions between nearly optimal population sizes. For SN, like Beasley et al., we employ a constant population size across all runs in a sequence. An alternative performance criterion would be to hold population size constant and measure quality of final solutions. However, we expect that different algorithms would have different optimal population sizes.

10.2   Results

Tables 10.1–10.3 summarize all results and allow comparison of the four algorithms on problems of varying difficulty. The tables compare the number of GA function evaluations (without

hillclimbing function evaluations added) for the three GAs. They then compare the total number of function evaluations (GAs plus hillclimbers) for all four algorithms. The best average results on each problem are shown in boldface.

We group test problems based on difficulty. The first of three groups, shown in Table 10.1, consists of the easiest problems, *M1–M5* and *MUX-6*. Each of these problems has fewer than 10 peaks and has negligible isolation or misleadingness (hence no deception). The second group, shown in Table 10.2, consists of two problems of intermediate difficulty, *M6* and *PAR-8*. These problems have a moderate number of peaks — *M6* has 25 and *PAR-8* has 128 — and have negligible isolation or misleadingness (hence no deception). The third group of problems, shown in Table 10.3, consists of the three hardest problems, *M7–M9*. Each of these problems has between several thousand and several million peaks, and also displays both isolation and misleadingness. (*M7* and *M9* are deceptive.)

In Tables 10.1–10.3, average subpopulation size is computed by dividing average population size by the number of desirable peaks (subpopulations). For sequential niching, average population size is first computed across all sequences and all runs within a sequence. This average population size is then divided by the number of desirable peaks to yield average subpopulation size. For SN, $\bar{g}$ indicates the average, across all sequences, of the combined number of generations for all runs within a sequence. 95% confidence intervals for the mean number of function evaluations are computed according to the t-distribution.

Table 10.1 shows that on the easiest problems, parallel hillclimbing is the overall winner, with deterministic crowding a close second. Parallel hillclimbing is fastest on four of the six easy test functions, and DC is fastest on the other two. Second place goes twice to hillclimbing, twice to DC, and twice to sharing. Third place goes four times to sharing, once to DC, and once to SN. SN takes last place all but one time: it edges out DC on *M5*.

These results confirm prior results of GA researchers who show that hillclimbing algorithms can outperform GAs on easy test functions, such as the five functions of De Jong's (1975) test suite. (Note that DC, a crossover hillclimber, is a close second and sometimes defeats the parallel hillclimber.) There is hence no point in doing performance comparisons involving GAs, exclusively using easy problems. We are more interested in how performance scales up to harder problems. Doing comparisons on a range of problems from very easy to very hard tells a lot about the merits of an algorithm and the scalability of its performance.

202

Table 10.1: Performances of the four algorithms, parallel hillclimbing (HC), sequential niching (SN), fitness sharing (SH), and deterministic crowding (DC), are given on the six easiest test functions. Statistics are taken over 10 runs. Average subpopulation size is $\bar{n}$; average number of generations is $\bar{g}$. The mean number of function evaluations (μ), its standard deviation (σ), and a 95% confidence interval for the mean are given for each GA alone, and for each combination of GA and hillclimber. The best average results on each problem are shown in boldface.

Method	$\bar{n}$	$\bar{g}$	GA Function Evaluations			Total Function Evaluations		
			μ	σ	μ: 95% C.I.	μ	σ	μ: 95% C.I.
					M1			
HC	2.72					**1017**	284	(813, 1220)
SN	3.68	46.40	738	359	(481, 994)	4112	1478	(3055, 5169)
SH	5.76	8.00	**264**	147	(159, 369)	2431	1262	(1529, 3333)
DC	2.40	28.00	380	270	(187, 573)	1246	536	(863, 1629)
					M2			
HC	2.72					**1021**	292	(812, 1230)
SN	4.64	75.60	1770	1047	(1022, 2519)	8632	4216	(5618, 11647)
SH	8.96	8.70	442	203	(297, 587)	3827	1436	(2801, 4854)
DC	2.40	27.40	**372**	239	(201, 543)	1264	537	(879, 1648)
					M3			
HC	3.04					1150	728	(629, 1671)
SN	5.92	26.80	719	435	(408, 1030)	4375	1709	(3153, 5598)
SH	6.08	8.70	294	133	(199, 390)	2579	1182	(1733, 3424)
DC	2.08	20.30	**262**	290	(55, 470)	**1013**	529	(635, 1391)
					M4			
HC	3.04					1140	713	(630, 1649)
SN	5.12	72.40	2445	2572	(606, 4284)	10231	8789	(3947, 16515)
SH	6.72	9.20	352	201	(209, 495)	2892	1047	(2143, 3640)
DC	2.08	17.00	**210**	185	(77, 342)	**975**	447	(655, 1295)
					M5			
HC	2.50					**901**	479	(558, 1243)
SN	1.30	32.30	180	77	(125, 235)	1456	433	(1146, 1766)
SH	2.80	8.00	**103**	57	(62, 144)	1111	411	(817, 1405)
DC	5.60	25.60	603	275	(406, 800)	2459	941	(1787, 3132)
					MUX-6			
HC	10.40					**1257**	998	(544, 1971)
SN	6.40	140.70	4423	3879	(1650, 7197)	9439	7367	(4172, 14706)
SH	13.60	8.90	**534**	186	(401, 667)	1931	555	(1537, 2328)
DC	12.00	44.30	2816	3287	(466, 5166)	3654	3750	(972, 6335)

We give results without hillclimbing to show how long it takes a GA to get close to the final desired solutions (anywhere within their hillclimbing basins of attraction), without actually having to pinpoint them. When hillclimbing function evaluations are not counted, sharing and deterministic crowding share three victories apiece, with sharing having the lowest overall number of function evaluations. Second place goes to sharing three times, crowding twice, and SN once. Third and final place goes to SN five times and crowding once. (There is no fourth place, since we do not compare the hillclimber to the GAs in terms of GA function evaluations.)

It is somewhat surprising that SN performs so poorly in comparison with the other algorithms on the six easiest test functions, especially since Beasley et al. employ *M1–M5* in their study. In addition, our SN results without hillclimbing are in all cases faster than Beasley et al.'s on *M1–M5*. (This comparison is valid, because their convergence results are approximate, varying between 76% and 100% correct convergence.) Ignoring hillclimbing, SN requires roughly 3–8 times as many function evaluations as sharing, on five of the six easy test functions, and roughly 1.7 times as many on the other. With hillclimbing, SN requires roughly 4–9 times as many function evaluations as parallel hillclimbing on five out of six functions, and roughly 1.6 times as many on the other.

A plausible explanation for SN's behavior is that once it has squashed several peaks in the fitness landscape, locating the final peak is harder because that peak is isolated. One observation about SN on the five-peaked functions *M1–M4* is that once its population grows large enough to locate one peak, it has grown large enough to locate multiple peaks. Many of the runs on *M1–M4* locate no peaks with $n = 8$, but 2–3 peaks at a time with $n = 16$.

All of the algorithms, including SN, find all desired optima on all six test functions in under 11,000 function evaluations. Therefore, all four algorithms are general-purpose enough to handle the easiest problems. We already know that the GA, although a capable optimizer of easy functions, may not be the optimal algorithm for any particular easy test function. A good algorithm, however, must scale up to solving harder problems.

We now examine the performances of the four algorithms on two functions of intermediate complexity, *M6* and *PAR-8*. Table 10.2 tells us that sharing is the clear winner on both functions, with or without hillclimbing. Sharing performs only 14% to 44% of the function evaluations of its closest competitor. On *M6*, hillclimbing comes in second and SN comes in third. DC fails to locate and maintain the 25 optima, even given 1.5 million GA function

Table 10.2: Performances of the four algorithms, parallel hillclimbing (HC), sequential niching (SN), fitness sharing (SH), and deterministic crowding (DC), are given on the two functions of intermediate difficulty. Statistics are taken over 10 runs. Average subpopulation size is $\bar{n}$; average number of generations is $\bar{g}$. The mean number of function evaluations (μ), its standard deviation (σ), and a 95% confidence interval for the mean are given for each GA alone, and for each combination of GA and hillclimber. The best average results on each problem are shown in boldface.

| Method | $\bar{n}$ | $\bar{g}$ | *GA Function Evaluations* | | | *Total Function Evaluations* | | |
			μ	σ	μ: 95% C.I.	μ	σ	μ: 95% C.I.
					M6			
HC	12.29					29017	10221	(21709, 36326)
SN	3.58	146.30	12202	10755	(4512, 19892)	46657	28680	(26150, 67163)
SH	5.12	11.80	**1638**	329	(1403, 1874)	**12910**	340	(12667, 13153)
DC			$> 1.5 \times 10^6$					
					PAR-8			
HC	48.08					202387	71122	(151535, 253239)
SN	19.20	36.40	100557	72359	(48820, 152294)	263666	139844	(163677, 363654)
SH	9.60	12.60	**17203**	8400	(11197, 23209)	**54402**	21280	(39187, 69617)
DC	11.20	87.40	125850	44106	(94314, 157385)	149022	52635	(111388, 186657)

evaluations. However, it consistently locates and maintains the single global optimum. On *PAR-8*, DC comes in second, taking roughly three times as many total function evaluations as sharing. Hillclimbing comes in third, taking about four times as many as sharing. SN comes in fourth, taking about five times as many as sharing. Without hillclimbing, SN edges out DC for second, using 80% of DC's function evaluations.

The reason for DC's peculiar performance on *M6* is that the function is not multimodal in crossover-hillclimbing space. In other words, DC uses the nonglobal optima as stepping stones to the global optimum: they are all on the crossover path to the global optimum. However, DC converges relatively slowly to the global optimum and, in any particular run, locates many nonglobal optima along the way to the global optimum. This raises the possibility that one could locate some or all of the 24 nonglobal optima by stopping the run early. It is interesting that the crossover interactions we observed in Chapter 7 on constructed problems, now emerge on a test problem. Specifically, the global optimum is dominating nearby local optima that are in turn dominating other local optima, and so on, creating a cascading effect.

The method by which SN solves *PAR-8* is of interest: SN locates all 128 disjuncts in a single run. Its success is therefore due to SUS's stability and not to iterating the GA. Again, once

Table 10.3: Performances of the four algorithms, parallel hillclimbing (HC), sequential niching (SN), fitness sharing (SH), and deterministic crowding (DC), are given on the three functions of greatest difficulty. Statistics are taken over 10 runs. Average subpopulation size is $\bar{n}$; average number of generations is $\bar{g}$. The mean number of function evaluations (μ), its standard deviation (σ), and a 95% confidence interval for the mean are given in thousands (indicated by the letter K) for each GA alone, and for each combination of GA and hillclimber. The best average results on each problem are shown in boldface.

Method	$\bar{n}$	$\bar{g}$	GA Function Evaluations			Total Function Evaluations		
			μ	σ	μ: 95% C.I.	μ	σ	μ: 95% C.I.
M7								
HC						> 2000K		
SN			> 1500K					
SH			> 1500K					
DC	20.80	119.80	**81K**	31K	(58K, 103K)	**101K**	39K	(73K, 129K)
M8								
HC						> 2000K		
SN			> 1500K					
SH	19.20	19.20	**13K**	5K	(9K, 16K)	**38K**	14K	(28K, 49K)
DC	22.40	134.40	98K	39K	(70K, 126K)	119K	47K	(86K, 153K)
M9								
HC						> 2000K		
SN			> 1500K					
SH						> 2000K		
DC	136.53	337.80	**1253K**	646K	(791K, 1715K)	**1342K**	691K	(847K, 1836K)

n is high enough to find all desired solutions, it is high enough to find them in one run of a parallel niching method. Unfortunately, we have seen in Chapter 3 that this ability of SUS to maintain multiple solutions of identical fitness is not consistent from problem to problem, and evaporates when the solutions have differing fitnesses.

SN's location of multiple peaks in a single run points to a potential, small improvement in its design. Instead of SUS, one could employ a faster selection scheme such as tournament selection, which is virtually guaranteed to return only one solution per run. This would make SN's actual behavior more like its expected behavior, but would not correct its fundamental problems.

On the three hardest problems (see Table 10.3), DC is the only method to solve all three in the allotted number of function evaluations. Sharing solves *M8* in fewer evaluations than DC, but reaches the function-evaluations limit on *M7* and *M9*. Sequential niching and parallel

hillclimbing fail on all three problems to find the required optima in 1.5 million GA function evaluations and 2 million overall function evaluations, respectively.

The results for sharing on $M7$ and $M8$ are consistent with those reported by Goldberg, Deb, and Horn (1992), where sharing is unable to solve the unscaled, massively multimodal, deceptive problem, but is able to solve the scaled version. The trouble is that sharing without scaling is too generous in allocating individuals to the millions of extraneous peaks, while sharing with scaling can minimize the heights of those peaks. The authors estimate that a population size of over 3.46 million would be necessary to solve the unscaled version using sharing. Plugging the number of extraneous peaks and other appropriate parameters ($\gamma = .9$, $g = 200$, $b = 32$, $s = 5,153,600$, $p_c = .9$, $p_d = .5$, $r = 1$, $r' = .928$) into our most general population sizing equation (9.22) gives a higher required population size of over 95 million. Sharing encounters problems with $M9$ for the same reasons it has trouble with $M7$ — too many extraneous peaks of too high a fitness. For instance, with $n = 32,768$, sharing converges after 19 generations, performing $655,360$ GA function evaluations and $2,227,120$ total function evaluations in the process. It returns 2009 total optima, of which 5 are global.

Parallel hillclimbing fails on all three problems for obvious reasons: the problems have become too complex for hillclimbing. The misleading or deceptive attractors, which correspond to extraneous peaks, draw most population elements that are not exact instances of global optima.

On $M7$ and $M8$, SN has the choice of trying to squash millions of extraneous peaks and then trying to converge to 32 remaining needles (in a huge haystack), or trying to locate the global optima one or a few at a time. Derating millions of undesirable optima is not a very appealing option. With sufficiently large populations, SN's first few runs locate several global optima. However, after SN derates these and other local optima, it has a harder time locating any more global optima. This problem becomes progressively worse until the algorithm fails. SN runs into similar problems on $M9$.

We make some final observations from Tables 10.1–10.3. Sharing shows the greatest stability of the tested algorithms, typically exhibiting lower standard deviation in the average number of function evaluations to convergence. Sharing also typically runs the fewest generations. Deterministic crowding succeeds with the smallest subpopulations on the average.

10.3 Discussion of Results

We can draw the following general conclusions about algorithmic performance versus problem hardness from the runs we have performed:

- Parallel hillclimbing is best for the easiest problems. It may also work in a reasonable time frame on problems of intermediate complexity. However, it fails on problems of high complexity.

- Sequential niching is weak on easy problems, and is unable to solve harder problems. In general, parallel hillclimbing is a better method that is also parallel. Parallel hillclimbing outperforms sequential niching because the parallel hillclimber does not destroy the fitness landscape.

- Sharing generally works on problems of all levels of complexity. However, it runs into trouble on problems in which many extraneous peaks exist that are of similar fitness to the desired peaks. It may be able to overcome this difficulty, however, through the intelligent application of fitness scaling.

- Deterministic crowding is generally good for problems of all levels of complexity. However, it may ultimately lose lower optima that lie on a crossover path to higher optima. Nevertheless, it only loses lower optima in favor of higher optima. Deterministic crowding tells us that crossover hillclimbing can solve problems that are much more difficult than those solvable by traditional, mutation-based hillclimbing.

We have found that parallel niching methods outperform sequential niching methods. Furthermore, sequential niching does not achieve a sequentialization of fitness sharing, and it yields few of the benefits that its authors claim. In general, parallel niching methods offer the following advantages over sequential niching methods:

- Parallel niching methods can easily be implemented on parallel machines. Sequential niching methods, by their nature, can not.

- Parallel niching methods should be faster than SN methods and give better results (contrary to Beasley et al.'s assertions and analysis), even when run on a single processor.

- Parallel niching GAs can be applied to maintain internal diversity along the way to a single solution. SN can not.

- SN is likely to locate the same solutions repeatedly, despite its deration of peaks.

- SN creates false optima in the derated fitness landscape that are in close proximity to peaks that were previously located. SN can also offset an optimum's location as a consequence of deration. Beasley et al. suggest hillclimbing in the original fitness landscape to overcome both problems. However, in both cases, hillclimbing stands a good chance of rediscovering a previously derated peak.

- For classification and machine-learning problems as well as simulation problems, parallel niching allows the whole population to cooperatively act as a solution. Sequential niching does not. For instance, one might not be interested in global optima at any point in time but in the state of the population as a whole — the optimal population rather than the optimal population element. SN achieves no cooperation within the population, except, arguably, for a weak form of temporal cooperation. However, such temporal cooperation is static — the size of reserved niches can never grow or shrink.

- SN's deration of optima may delete other optima of interest within the deration neighborhood. This is to some extent also a problem in sharing. However, deration neighborhoods are dynamic in sharing, while they irreparably alter the fitness landscape in SN. Even worse for SN, solutions that have been derated — whether local optima, global optima, or near optima — might take with them important building blocks for locating other solutions. On hard problems such as *M7–M9*, eliminating one global optimum hinders the location of others.

- In SN, as optima are derated, the remaining optima become increasingly difficult to locate. Derated regions, containing mostly plateaus and small ridges, occupy a greater and greater percentage of the space — and SN must repeatedly search through these derated portions of the space. One can observe this repeated search, even on simple functions. As a sequence progresses, the remaining optima become like multiple needles in a haystack. After only a few optima have been derated, the required population size for SN can easily exceed that for a parallel niching method capable of locating all niches within its single

population. Beasley et al.'s claim that SN requires a population of only $1/c$ of the size of a parallel niching method's population does not hold, except possibly for locating the first peak. In contrast to SN, parallel niching methods spread the population out across the entire search space, allowing cooperative subpopulations to locate the niches.

We have also found that parallel niching GAs outperform parallel hillclimbers, on all but the easiest functions. When many extraneous attractors are present or extraneous attractors with large basins are present in the search space, parallel hillclimbing will in probability converge to several of these peaks. Parallel niching GAs, on the other hand, have the power to escape these attractors and to converge to the desired solutions.

Chapter 11

Conclusion

To conclude this study of niching methods for genetic algorithms, we first summarize its contents and highlight its contributions. We then present avenues for continuation or extension of the research presented in this thesis. Finally, we draw conclusions from the modelling and experimentation conducted herein.

11.1 Summary

This section presents a summary of the thesis, and highlights its contributions to research in genetic algorithms. The thesis began by motivating both the study of niching methods and our adopted methodology for modelling niching methods. Our models embody a decomposition perspective. They emphasize the extraction of the most important components of a complex system, and the study of those components in order of importance. Decomposition is effective whether the components of a system are truly separable or only quasi-separable. Our specific approach has been to first study, in isolation, the most important component of niching genetic algorithms — selection with niching. Afterward, we have added the second most important component, crossover, in order to study the combination of niched selection and crossover. Minor factors such as mutation have been abstracted out of the models, but could have been modelled as sources of noise.

This study has reviewed genetic algorithms, and has illustrated the significance of its research within the overall spectrum of genetic algorithm research. This thesis is most related to prior work on modelling, analyzing, and designing genetic algorithms, through use of an

algorithmic decomposition approach. It is also related to prior genetic algorithm research on problem hardness, convergence proofs, and methods for niching and population diversification.

Since niching methods have historically been related to methods for promoting population diversity, this thesis has presented a comprehensive treatment of the topic of diversity. The thesis first describes the motivation for studying population diversity in genetic algorithms. It next illustrates simple problems on which traditional genetic algorithms fail because of a lack of diversity. It proceeds to isolate three major culprits in the loss of diversity — selection noise, selection pressure, and operator disruption. It then reviews previous research into the maintenance of diversity, bringing together and categorizing prior diversification mechanisms that, by themselves, do not qualify as niching methods. Such mechanisms include noise-reduced selection; control parameter adjustment; direct infusion of diversity; reinitialization and multiple sequential runs; isolation, migration, and parallel genetic algorithms; thermodynamic genetic algorithms; and mating restrictions.

Prior to this study, diversity was a term used very loosely by GA researchers. It was viewed as a method of combatting the equally vague notion of premature convergence. We have presented a formal framework for the study of population diversity and have examined prior diversity measures from the perspective of that framework. We have applied a specialization of the diversity framework to niching methods that perform multimodal function optimization. A very general type of multimodal function optimization problem emerges from the specialization of the diversity framework — to find at least the b highest maxima. The resulting goal distributions are more general than the goal distributions of prior studies, which typically are biased to conform to the algorithm under consideration. Biased goal distributions include fitness-proportionate distributions, uniform distributions, and Boltzmann distributions. In contrast, our goal distributions cover any population that is likely to be a good solution to a problem. This thesis goes one step further, incorporating the specialized diversity framework into a framework for modelling niching methods. The frameworks of this thesis should benefit future research into diversity and niching methods.

An important contribution of this study has been the unification of prior genetic algorithm research on diversification methods and niching methods. Previous research is spread over more than 100 studies, and often does not distinguish between mechanisms for diversity and niching. We have made a distinction between the two types of algorithms, based upon the concept of

useful diversity. We have required bona fide niching methods to be able to form and maintain stable subpopulations of differing quality or fitness as well as stable subpopulations of identical quality or fitness.

Insights that were previously not apparent from analyzing diversification and niching methods individually or in small groups, have become clear through analysis of prior methods *en masse*. Through the collection and interrelation of prior niching methods, we have gained a substantial amount of knowledge about broad categories of niching genetic algorithms. This thesis provides a foundation for the further study of niching methods, and a first reference for researchers to consult when studying mechanisms for diversity and niching in genetic algorithms.

Niching methods can be classified along two dimensions of behavior: spatial versus temporal niching, and niching within single environments versus niching due to multiple environments. Within each resulting quadrant, broad categories of niching methods exist. Categorization allows insight into niching method design space, enabling the identification of unexplored regions of design space, and illuminating meaningless distinctions among niching methods that are actually very close in design space. The broad categories of niching methods are sequential niching; overspecification; ecological genetic algorithms; heterozygote advantage; crowding or restricted replacement; restricted competition; fitness sharing; and immune system models. We have chosen crowding and sharing for detailed examination, both of which represent spatial, single-environment methods. We have also examined, to a lesser extent, sequential niching. Prior to this study, of all potential niching methods, crowding and sharing were the best known, yet were still only sparsely explored. Of the remaining niching methods, we have highlighted categories that are the most promising, and others that are likely to be less fruitful.

This thesis has presented a comprehensive framework for the modelling of niching methods, and has used this framework throughout to build models of specific niching methods — primarily models of niched selection alone and niched selection with crossover. The framework places equivalence classes into a one-to-one correspondence with peaks in the fitness landscape, where peaks are defined by a hillclimber operating under an appropriate neighborhood operator. Individual models typically emerge directly from the modelling framework, requiring little to no specialization for the particular niching method being modelled. Other researchers should be able to use the theoretical and empirical aspects of our modelling framework to construct further models for the analysis and design of niching methods. In connection with introducing

our modelling framework, we have reviewed prior modelling techniques and have discussed why they are inadequate for our purposes.

Our modelling framework contains a number of test problems of varying difficulty, that are intended for use in analysis, design, and empirical testing. One type of problem that was previously difficult for genetic algorithms, but can now be attacked in a straightforward fashion through use of a niching method, is the classification problem. Niching methods allow genetic algorithms to solve classification problems without the addition of complex machinery. Our mapping of classification problems to multimodal optimization problems demonstrates how niching methods can extend from optimization to classification. Our classification methodology is one of the first successful, yet straightforward extensions of genetic algorithms to classification problems. This thesis has demonstrated the new methodology on both difficult and average-case boolean classification problems. The methodology extends to real-world problems with a few modifications.

The first category of niching methods that this thesis investigates in depth is crowding methods. We have analyzed several crowding and preselection schemes under our modelling framework, and have determined, in terms of replacement errors, why these schemes fail to perform effective niching. Proposed design modifications targeted the reduction of replacement error. Through a series of modifications, tests, and analyses, we succeeded in designing a new crowding mechanism, deterministic crowding, that successfully met our design criteria. We have demonstrated the benefit of pursuing a single design goal, whether it takes us backward or forward in design space. For crowding, the goal of reducing replacement error ultimately leads to the maintenance of additional peaks. An alternative goal, maintaining bitwise diversity, does not typically lead to the maintenance of more peaks.

Our extensive analysis of the combination of replacement selection and crossover in deterministic crowding has led to the discovery of meaningful interactions among dominating, dominated, and assisting classes, that help determine algorithmic behavior. Through controlled testing, we have found that deterministic crowding allocates to a class or peak a number of elements that is proportional to the sum of the size of that peak's basin of attraction, and the sizes of the basins of attraction of all peaks it dominates. (Basins are defined, in this case, from the vantage point of traditional hillclimbing.) Factors determining the dominance of peaks include relative peak fitnesses, proximity of peaks, and crossover biases.

Alternatively, the distribution of population elements among classes in deterministic crowding can be viewed as depending solely upon the sizes of the basins of attraction of individual peaks in a transformed crossover space. This abstraction yields two new concepts — crossover hillclimbing and crossover basins of attraction — that are useful both for the analysis of algorithms and the characterization of fitness landscapes. Crossover hillclimbing allows migration to higher peaks in traditional hillclimbing space, thus yielding an algorithm more powerful than hillclimbing. The use of mating restrictions becomes questionable, since they potentially limit this migratory power.

The second category of niching methods that this thesis investigates in depth is sharing methods. Properties of fitness sharing are modelled, including distribution, drift time, population size, crossover disruption, and crossover probability. Less important properties, such as mutation, can be modelled as sources of noise. Closed-form expressions are derived for the time to disappearance of a class, and for population size. Derivations are verified using several test functions. Modelled cases include those in which the goal is to locate all peaks, and those in which a boundary exists between peaks of interest and extraneous peaks.

This study has derived lower bounds on required population sizes for sharing with roulette-wheel selection, both with and without crossover. Specifically, it has derived several expressions, under various assumptions, for the minimum population size required to maintain a number of classes, with a certain confidence, for at least a given number of generations. Because roulette-wheel selection has the highest variance among commonly used fitness-proportionate selection methods, the derived bounds also apply to sharing under more stable, fitness-proportionate selection schemes such as stochastic remainder selection and stochastic universal selection.

We have demonstrated, on several test problems, that the number of runs of sharing that form and maintain all classes, consistently surpasses the confidence level provided to the sharing model. Models have been predictive, even on problems with small disruptive forces or low levels of noise, introduced by crossover, mutation, and overlapping niches. We have developed a technique for providing to a sharing model, a practical assessment of the significant characteristics of a problem. This assessment includes a relative fitness threshold that estimates the boundary between desirable and undesirable classes. It also includes the minimum number of peaks that the user is interested in locating, plus an estimate of the number of extraneous peaks in the fitness landscape.

Besides yielding lower bounds on disappearance time and required population size, the sharing models also allow insight into the setting of parameters such as crossover probability. Furthermore, the methodology for incorporating crossover into the sharing model can potentially be emulated to incorporate mutation and other sources of disruption or noise. We have derived from one of the models a general formula that interrelates crossover probability and population size.

This study has examined the behaviors of the methods it has modelled, on three sets of test problems of increasing difficulty. It has compared the parallel niching genetic algorithms, crowding and sharing, to the more simplistic niching methods, sequential niching and parallel hillclimbing. (The term, parallel niching, describes methods that conceptually form multiple niches simultaneously within a single population — regardless of the number of physical processors employed.) The parallel niching genetic algorithms are significantly more powerful than the more elementary methods. This additional power is especially evident on harder problems.

Through empirical testing, we have been able to determine the strengths and the weaknesses of all four niching methods. Parallel hillclimbing, although not a match for the parallel niching genetic algorithms, is an effective local optimizer. When tacked onto the end of a parallel niching genetic algorithm, it allows the genetic algorithm to converge from anywhere within the basin of attraction of a local optimum, to the actual local optimum.

Overall, modelling and analysis has led to a better understanding of results from prior studies. For instance, we now understand why De Jong's crowding consistently maintains exactly two niches. We also understand better the behavior of sharing on the massively multimodal and deceptive functions of a previous study (Goldberg, Deb, & Horn, 1992).

Most importantly, modelling and analysis has led to improved algorithmic designs in our study. The first such case is deterministic crowding, developed via a model of crowding constructed from our framework. Deterministic crowding is the first crowding algorithm to successfully perform niching. Its design comes 20 years after the first crowding method was introduced. A second case is the derivation of optimal population sizes, and insight into optimal crossover probabilities, for fitness sharing. A third design improvement, for sequential niching, is the utilization of a fast convergence scheme such as tournament selection. Finally, a general design improvement for all niching genetic algorithms is the running of a hillclimber on the niching genetic algorithm's final population.

11.2 Future Research

Future research on niching methods for genetic algorithms promises extended genetic algorithms that are proficient at performing a variety of complex tasks. Niching genetic algorithms have previously enjoyed preliminary success in multimodal function optimization and multiobjective function optimization. Niching genetic algorithms should emerge in the near future that are also proficient in general-purpose classification and machine learning. In addition, niching genetic algorithms should ultimately enjoy a variety of applications to the simulation of complex and adaptive systems.

The research foundation presented in this thesis, along with extensions of this research, should facilitate the application of niching genetic algorithms to the aforementioned areas of interest. Many logical and promising extensions exist to the research presented in this thesis. We highlight several of the possibilities below.

First of all, the separation of diversification mechanisms from niching methods, and the further categorization of niching methods, have pointed out yet unexplored regions of niching-method design space. Just within the quadrant of spatial, single-environment techniques lie two categories of niching methods that await exploration — heterozygote advantage and restricted competition.

Within the crowding category of niching methods are several alternatives to and extensions of deterministic crowding. A comparative analysis of successful crowding techniques would be a beneficial area of future research. One aim of such a study could be to verify that distributional behaviors are similar for members of the crowding family.

Further extensions of crowding methods are possible that incorporate techniques borrowed from elsewhere. We have previously mentioned the possibility of adding a convergence control parameter similar to that of simulated annealing (Mahfoud & Goldberg, 1995). We have also mentioned a potential hybrid of deterministic crowding and GIGA (Culberson, 1992), which would assure complete conservation of alleles. Ultimately, one could strive to design a generalized crowding algorithm that encompassed all successful crowding methods. A further generalized algorithm might also subsume algorithms such as immune system models. Whether or not generalized algorithms proved useful in actual runs, they would most probably be beneficial for modelling the niching methods they encompassed.

The concept of crossover hillclimbing should be a useful abstraction for future studies of crowding, as well as future studies of crossover, replacement selection methods, niching methods in general, and genetic algorithms. Models based on crossover hillclimbing will be of intermediate complexity between the models of this study and the Markov chain models reviewed in Chapter 5. Since deterministic crowding is a crossover hillclimber, it could be used to detect the crossover-hillclimbing landscape. The further study of deterministic crowding should lead to a better understanding of crossover in any genetic algorithm.

Mating restrictions offer a trade-off between on-line and off-line performance. In this study, we have primarily been concerned with the off-line performance of niching methods. Therefore, mating restrictions have been too limiting for our purposes. For applications requiring excellent on-line performance, one can easily modify both crowding and sharing to incorporate mating restrictions. The most straightforward modification is to employ a mating-restriction threshold and to disallow mating between population elements that differ by more than the threshold. Adding mating restrictions to deterministic crowding may help preserve dominated peaks on problems such as Shekel's Foxholes.

For sharing, the most useful research path would be to develop practical methods for estimating some of the modelling parameters, especially the disruption probability and the number of extraneous peaks in the search space. We have detailed, in Chapter 9, potential methods for setting all modelling parameters. The disruption probability can be set to either 0 or 1 if one wishes to assume either minimal disruption or worst case disruption. Alternatively, it can be set to some intermediate, estimated bounding value. The most difficult parameter to set, the number of extraneous peaks in the fitness landscape, can be set via overestimation. Given a completely unknown fitness landscape, one could estimate the number of peaks using prior runs or from early generations. Another approach would be to run hillclimbers from every point within several small, randomly chosen regions, and to count the number of unique optima returned. From this sample, one could roughly estimate the modality of the search space.

A second potential research path for sharing would be to model additional sources of noise in the same way we modelled crossover. High levels of mutation could be modelled as they are in the schema theorem — as sources of disruption. Fuzzy boundaries between niches could also be modelled as sources of disruption. Population sampling to compute shared fitnesses is another potential source of noise that could be modelled in the same way. Modelling sampling

would hopefully tell us the minimum sample size required to emulate sharing's behavior when it samples the full population. A final source of noise that may also be beneficial to model, perhaps as another source of disruption, is internal class fitness variance.

A third research path for sharing would be to investigate alternative algorithmic designs. Our empirical simulations revealed that sharing stops improving a population after only a few generations. A possible remedy would be to put a gradually increasing selection pressure on sharing. Our empirical simulations also verified a prior result of Goldberg, Deb, and Horn (1992). The authors found that exponential scaling can be a useful addition to sharing. Exponential scaling enables sharing to overcome its tendency to overallocate population elements to extraneous peaks that sit in close proximity to desirable peaks. An exponential scaling parameter could be made a permanent part of the sharing algorithm; the parameter's adjustment would determine the number of peaks that sharing located. Although scaling is capable of adjusting the fitness threshold between desirable and undesirable peaks to arbitrary levels, care must be taken to avoid squashing good building blocks or desirable peaks in the process.

Several researchers have complained about the rigidity of sharing's thresholding parameter which sets the boundaries between niches. Clustering is one alternative to thresholding that has been suggested (Yin, 1994; Yin & Germay, 1993). However, parameters similar to sharing's thresholding parameter must be set for the clustering method. Adaptive niche-sizing techniques are a critical area of research for the extension of sharing methods and other niching schemes. However, these future, adaptive techniques should be designed to be relatively insensitive to changes in their own parameters.

A final possibility for sharing would be to locate tighter bounds on drift time and population size by modelling individual, fitness-proportionate selection schemes that are less noisy than roulette-wheel selection. Furthermore, sharing under non-fitness-proportionate selection schemes such as ranking and tournament selection has received only a little attention (Oei, Goldberg, & Chang, 1991), and deserves additional research.

From the modelling perspective, extending individual models or portions of the modelling framework offers a wide array of opportunities. For instance, an enhanced modelling framework could include test functions that more stringently isolated individual dimensions of problem hardness. In addition, the search for good test functions of varying complexity is an area of research that deserves further consideration.

Another possibility on the modelling side would be to model niche formation as well as maintenance. A comparison of required population sizes under formation-based population sizing and maintenance-based population sizing would help determine how the two types of population sizing could be integrated. A simple formula for the combination might be sufficient, such as the maximum of the two recommended population sizes, perhaps times a constant. As mentioned earlier, we have found that interniche population sizing considerations (niche maintenance) consistently override intraniche considerations (niche formation), allowing a population that is sized to maintain niches, more than enough elements to form the niches. However, assuming the existence of practical problems in which the opposite is the case (intraniche population sizing overrides interniche population sizing), an integrated methodology would be useful.

Since our models are of niche maintenance, we have concentrated upon the destructive rather than the constructive effects of crossover. Considering constructive effects as well would potentially lead to expressions for optimal crossover probabilities for sharing and for genetic algorithms in general. We suggest extensive future experiments with sharing, on constructed problems, using crossover probabilities between .5 and 1, to determine the trade-offs between crossover probability, population size, and desired performance.

In certain applications such as the simulation of ecologies or economies, one may be interested in maintaining multiple individuals within each niche. Our models have assumed that one individual within each desired niche is sufficient. Our models could be extended, however, starting with the goal distributions of Chapter 4, to require a minimum number of elements in each desired niche.

Our models and experiments have demonstrated the benefit of running local optimization methods (hillclimbers) on the final populations produced by global optimization methods (niching genetic algorithms). Local methods faster than our parallel hillclimber are undoubtedly possible, and would not even need to be hillclimbers. Several numerical methods exist (Press, Teukolsky, Vetterling, & Flannery, 1992) that do not require the computation of derivatives. One must be careful, however, that the local optimization method employed does not exhibit global optimization properties, which can lead to the loss of lower maxima that the niching method has located.

Most niching methods would benefit from additional research into appropriate distance measures. Where possible, we have employed phenotypic distances in our niching methods.

The use of Gray codes (Caruana & Schaffer, 1988) in combination with genotypic distances is an untested possibility. For classification and machine learning problems, there is no need to stick to naive phenotypic or genotypic comparison. One could instead base distances upon the number of common training examples covered.

From the applications perspective, niching methods are ready to be utilized in a variety of settings. Future research should extend our classification methodology to real-world classification and machine learning problems. In fact, extension of the classification system presented in this thesis has already yielded promising initial results on real-world problems. One difference between solving real-world problems and boolean covering problems is that the fitness function in real-world problems must be more tolerant of covering negative examples. A second difference is that the genetic algorithm will most probably have to learn multiple rather than single concepts, with some form of conflict resolution required when different classifications overlap. A third difference is the problem representation; an appropriate representation must be chosen for each problem's variables. Finally, available data must typically be divided into training and testing sets, in order to prevent overfitting to the training data.

A second promising area of application for niching methods is multiobjective function optimization. When viewed from the perspective of multimodal function optimization, the task of multiobjective function optimization is similar to the formation and maintenance of multiple solutions that lie on a plateau. Therefore, the study of individual niching methods on flat fitness functions will point to their applicability to multiobjective function optimization.

A third promising area of application is standard function optimization under difficult fitness landscapes. As Hatjimihail (1993) and Pál (1994) have demonstrated, niching methods (deterministic crowding in their studies) designed to maintain multiple solutions are also effective at maintaining multiple subsolutions, on the way to a better, single solution.

11.3   Conclusions

We can draw multiple conclusions from this study of niching methods for genetic algorithms. Starting with the subject of population diversity, three factors play a part in the loss of diversity — selection pressure, selection noise, and operator disruption. While methods exist that will mitigate the loss of diversity due to all three factors, the result will not necessarily be a niching

method. In fact, many previous methods promoted by their authors as being capable of niching do not qualify as niching methods under our definition: they are incapable of maintaining solutions about multiple peaks of the fitness landscape, especially when those peaks are of differing heights. One such category of diversification methods that do not qualify as niching methods is parallel, geographic genetic algorithms.

An organized categorization is possible for both diversification and niching mechanisms. For niching mechanisms, several distinct categories exist, all of which can be described along two dimensions of behavior — space versus time, and single versus multiple environments. Within each category, expected behaviors of member niching methods are similar. Promising niching methods from all categories await further exploration; most likely, a few promising niching methods remain undiscovered.

The decomposition approach to modelling and design is a fruitful approach, even when the system under study is not truly decomposable. We recommend that the designer first study the most important piece or the most important pieces (but individually) of the decomposition. Afterward, the designer can combine primary pieces or combine primary and secondary pieces, and so on. If the major pieces to a great extent determine the behavior of the system, minor pieces may never have to be studied. In design, it is important to study a simple domain of application first, and then proceed to more complex domains. Such an approach yields additional benefit when the chosen, first domain of application contains fundamental characteristics present in other domains. For example, the lessons we learned in multimodal function optimization extended to tasks such as classification and maintaining diversity in standard function optimization.

The definitions, abstractions, and simplifying assumptions of our modelling framework form a general, yet effective set of tools for modelling niching methods. In many cases, models follow directly from the framework, with little additional setup required. Constructed models are useful for designing experiments; understanding, explaining, and predicting algorithmic behavior; bounding control parameters; improving algorithmic designs; assessing the impact of design alterations; and even discovering better niching methods. We expect that ours or a similar framework could effectively be applied not only to niching methods, but to the modelling of other genetic algorithms and complex systems.

Previous studies have assumed ideal distributions that are fitness based, uniform, or even Boltzmann, in order to coincide with the distributions produced by the particular method under examination. It is better, however, to let ideal distributions be a function of the solutions one is trying to obtain rather than a function of the algorithm. A very general set of goal distributions, from our modelling framework, is the set of all distributions in which every solution of interest is allocated at least one population element.

One of the underlying hypotheses of this research has been that through modelling an algorithm and through analyzing an algorithm under one or more models, a person is likely to discover ways to improve the algorithm. In this thesis, our goal has been to model, analyze, and improve the designs of niching genetic algorithms. We have verified our underlying hypothesis through the design, via modelling and analysis, of deterministic crowding.

It is extremely important to have a well motivated goal in mind when making design changes, and to pursue that goal single-mindedly, whether the algorithmic performance resulting from individual alterations progresses or regresses. (The design space for complex systems is rarely linear.) For crowding, the reduction of replacement error is a well motivated goal, while the maintenance of bitwise diversity is not. Reducing replacement error ultimately leads to the maintenance of multiple peaks, while maintaining bitwise diversity typically does not.

We have found that crowding methods can make highly effective nichers, in some cases solving problems that elude other niching methods. One of the critical ingredients in the design of an effective crowding mechanism is the substitution of replacement selection for up-front selection. One should be aware that minor algorithmic changes can have dramatic effects on the proficiency of not only crowding methods, but other genetic algorithms. For instance, in parental replacement techniques, the method of competition between children and parents is critical: most design choices will result in an algorithm that is not a niching method.

In crowding methods, the distribution of population elements among peaks depends upon each peak's crossover basin of attraction. Alternatively, the number of elements ultimately expected at a peak is directly proportional to the sum of the size of that peak's hillclimbing basin of attraction and the size of the hillclimbing basins of attraction of all peaks that it dominates. Deterministic crowding is a useful tool not only for performing niching, but also for analyzing the crossover landscape.

Through analysis of crossover interactions, we have determined that rigid mating restrictions would in most cases be a detrimental addition to deterministic crowding. While mating restrictions limit migration from peak to peak in deterministic crowding and other genetic algorithms, they potentially present a trade-off between on-line and off-line performance.

In fitness sharing, practical lower bounds exist for population sizes and drift times. It should also be feasible to derive bounds on crossover probability in the future. The models of sharing that we have developed are ready for immediate use in sizing populations. They should also be able to provide some guidance in setting crossover probabilities for sharing and possibly other genetic algorithms.

In sharing, the minimum population size required to solve a problem with fixed certainty is in most cases a slightly superlinear function of the number of desired classes. The expected time to extinction or drift time of a class is an exponential function of population size, with relative class fitnesses determining the base of the exponential. However, drift time decreases rapidly as the number of classes increases. The drift time for a class is a useful value that tells a user how long desired subpopulations are expected to persist within the population. A desirable drift time should be significantly longer than the time required for genetic operators such as selection, crossover, and mutation, to perform their tasks. In general, we prefer niching methods that have exponential to infinite drift times with respect to population size. Crossover probabilities less than 1 are better for sharing, with lower probabilities required for increasingly disruptive problems. However, some minimum crossover probability is necessary to maintain sufficient exploration of the search space.

Difficult problems provide a testing platform for distinguishing between niching methods that are only proficient when faced with easy problems, and niching methods that are more likely to thrive in real-world situations. The parallel niching methods we have designed and the models we have constructed fare well not only on simple problems, but on difficult problems. The behavior of niching genetic algorithms on test functions of varying difficulty corresponds well to their prior modelled behavior. When moving from constructed problems to difficult test problems, the same underlying forces are at work.

For empirical testing, it is important to choose test problems across a range of difficulty levels. Furthermore, it is vital to include functions that are difficult for the methods under consideration to solve. Many of the widely studied functions in the genetic algorithm literature

on optimization have been found easy to solve using standard optimization methods and simple hillclimbers (L. Davis, 1991a). Likewise, many of the widely studied functions in the machine learning literature have been found relatively easy to learn with simple classification rules (Holte, 1993). With rare exception, studies of niching methods have used only a few of the simpler functions of this thesis.

We have been able to differentiate the behavior of parallel niching genetic algorithms from sequential niching genetic algorithms and parallel hillclimbers. While all four niching methods are applicable to simple problems, only the parallel niching genetic algorithms — crowding and sharing — perform well on the harder problems. Parallel niching genetic algorithms add value because they can solve problems that parallel hillclimbing and sequential niching can not, given a reasonable time frame.

Although parallel hillclimbing, by itself, can not solve the hardest problems, it makes a good local optimizer for use in conjunction with other niching methods. Fast local optimization techniques (that are not themselves capable of global optimization) are beneficial augmentations to parallel niching genetic algorithms, that force those genetic algorithms to converge to exact local optima. A genetic algorithm need only get into the hillclimbing basins of attraction of good local optima, and the local optimization technique will do the rest.

The four algorithms we tested exhibit the following overall performances when faced with problems at various levels of difficulty. Parallel hillclimbing is best for the easiest problems and has some success on problems of intermediate difficulty, before failing on problems of high difficulty. Sequential niching is generally the weakest algorithm at all levels of problem difficulty, consistently underperforming parallel hillclimbing. Sharing and deterministic crowding generally fare well on problems at all levels of difficulty. However, sharing reaches its limit on problems containing many extraneous peaks that are of fitness similar to the desirable peaks. Deterministic crowding experiences problems maintaining lower optima that lie on a crossover path to higher optima.

Parallel niching methods consistently outperform sequential niching methods, offering many advantages such as faster runtime, ability to run on parallel processors, cooperative problem-solving capacity, and sufficient flexibility to maintain internal diversity along the way to a single solution. Sequential niching, on the other hand, destroys critical building blocks and solutions in the fitness landscape, leading to duplication of effort both in rediscovering previous

solutions and in re-searching derated portions of the fitness landscape. Parallel niching genetic algorithms also outperform parallel hillclimbers, on all but the easiest test problems. When either significant numbers of extraneous attractors or extraneous attractors with large basins are present in the search space, parallel hillclimbing, in probability, will converge to several of these attractors. Parallel niching genetic algorithms, on the other hand, will have the power to escape these attractors and to converge to the desired solutions.

We recommend that niching methods be applied, with only minor augmentations, to classification problems. We have demonstrated the beginnings of a full-fledged system for classification and machine learning, based upon the niching method. Previous implementations of classifier systems have glossed over the niching method, often making it either a minor or an implicit player. However, poor performance of the classifier system could often be at least partially attributed to an ineffective form of niching. We suggest that future implementations of classification systems make the niching method the central player.

Those who wish to solve problems using niching methods now have a choice between two types of algorithms, crowding and sharing, whose behavioral properties we have uncovered and presented in detail. The other option is to venture into less travelled territory to attempt to develop a niching method from one of the other promising categories we have outlined. Of course, specific applications of niching methods may find it useful to tailor certain algorithmic design parameters to the problems under consideration. Whether one is exploring an entire category of niching methods, or tailoring the design of a niching method to a particular problem, the framework we have presented for modelling niching methods should be of benefit. Niching methods await extensive application in broad areas such as multimodal function optimization, classification and machine learning, multiobjective function optimization, and simulation of complex and adaptive systems.

Appendix

Mean of the Loss Distribution

Let L be a random variable representing the generation in which a loss occurs. The mean μ_L of the geometric (loss) distribution is the expected value of L, given by

$$\mu_L \;=\; \sum_L L \cdot P(L) \;=\; \sum_{i=1}^{\infty} iX(1-X)^{i-1} \;=\; \frac{X}{1-X} \sum_{i=1}^{\infty} i(1-X)^i \quad . \tag{A.1}$$

By Equation A.7,

$$\mu_L \;=\; \frac{X}{1-X} \frac{1-X}{(1-(1-X))^2} \;=\; \frac{1}{X} \quad . \tag{A.2}$$

Variance of the Loss Distribution

The variance σ_L^2 of the geometric (loss) distribution is the second moment about the mean of the random variable L:

$$
\begin{aligned}
\sigma_L^2 \;&=\; \sum_L (L - \mu_L)^2 \cdot P(L) \tag{A.3}\\
&=\; \sum_{i=1}^{\infty} (i - \frac{1}{X})^2 \cdot X(1-X)^{i-1}\\
&=\; \frac{X}{1-X} \sum_{i=1}^{\infty} (i^2 - \frac{2i}{X} + \frac{1}{X^2})(1-X)^i\\
&=\; \frac{X}{1-X} \left[\left(\sum_{i=1}^{\infty} i^2(1-X)^i \right) - \frac{2}{X} \left(\sum_{i=1}^{\infty} i(1-X)^i \right) + \frac{1}{X^2} \left(\sum_{i=1}^{\infty} (1-X)^i \right) \right] \quad .
\end{aligned}
$$

Summing the three infinite series, according to Equations A.6–A.8, yields

$$
\begin{aligned}
\sigma_L^2 &= \frac{X}{1-X}\left[\frac{(1-X)(1+1-X)}{(1-(1-X))^3} - \frac{2}{X}\frac{1-X}{(1-(1-X))^2} + \frac{1}{X^2}\left(\frac{1}{1-(1-X)} - 1\right)\right] \\
&= \frac{1-X}{X^2} \quad .
\end{aligned}
\tag{A.4}
$$

Sums of Series

$$
\sum_{i=0}^{k} x^i = \frac{1-x^{k+1}}{1-x} \qquad (0 < x < 1)
\tag{A.5}
$$

$$
\sum_{i=0}^{\infty} x^i = \frac{1}{1-x} \qquad (0 < x < 1)
\tag{A.6}
$$

$$
\sum_{i=0}^{\infty} i x^i = \frac{x}{(1-x)^2} \qquad (0 \le x < 1)
\tag{A.7}
$$

$$
\sum_{i=0}^{\infty} i^2 x^i = \frac{x(1+x)}{(1-x)^3} \qquad (0 \le x < 1)
\tag{A.8}
$$

Approximations

$$
\lim_{\epsilon \to 0} \ln(1-\epsilon) = -\epsilon \qquad (-\ln(1-\epsilon) > \epsilon)
\tag{A.9}
$$

$$
\lim_{b \to \infty} (1 - \frac{a}{b})^c = e^{-\frac{ac}{b}}
\tag{A.10}
$$

Newton's Method

To find an x such that $f(x) = 0$:

1. Choose an initial estimate x_0 .

2. Iterate the equation, $x_{i+1} = x_i - f(x_i)/f'(x_i)$.

References

Aarts, E., & Korst, J. (1989). *Simulated annealing and Boltzmann machines: A stochastic approach to combinatorial optimization and neural computing.* Chichester: John Wiley & Sons.

Ackley, D. H. (1987). An empirical study of bit vector function optimization. In L. Davis (Ed.), *Genetic algorithms and simulated annealing* (pp. 170–204). London: Pitman.

Adachi, N., & Matsuo, K. (1991). Ecological dynamics under different selection rules in distributed and iterated prisoner's dilemma game. *Lecture Notes in Computer Science: Parallel Problem Solving from Nature, 496*, 388–394.

Bäck, T. (1992). The interaction of mutation rate, selection, and self-adaptation within a genetic algorithm. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 85–94). Amsterdam: Elsevier.

Bäck, T. (1993). Optimal mutation rates in genetic search. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 2–8.

Bäck, T. (1994). Selective pressure in evolutionary algorithms: A characterization of selection mechanisms. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 57–62.

Baker, E. (1993). Evolving line drawings. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 627.

Baker, J. E. (1987). Reducing bias and inefficiency in the selection algorithm. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 14–21.

Bauer, R. J., Jr. (1994). *Genetic algorithms and investment strategies.* New York: John Wiley & Sons.

Beasley, D., Bull, D. R., & Martin, R. R. (1993). A sequential niche technique for multimodal function optimization. *Evolutionary Computation, 1*(2), 101–125.

Bethke, A. D. (1980). Genetic algorithms as function optimizers (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 41*(9), 3503B. (University Microfilms No. 8106101)

Booker, L. B. (1982). Intelligent behavior as an adaptation to the task environment (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 43*(2), 469B. (University Microfilms No. 8214966)

Booker, L. B. (1985). Improving the performance of genetic algorithms in classifier systems. *Proceedings of an International Conference on Genetic Algorithms and Their Applications,* 80–92.

Booker, L. B. (1987). Improving search in genetic algorithms. In L. Davis (Ed.), *Genetic algorithms and simulated annealing* (pp. 61–73). London: Pitman.

Bridges, C. L., & Goldberg, D. E. (1987). An analysis of reproduction and crossover in a binary-coded genetic algorithm. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms,* 9–13.

Brindle, A. (1981). *Genetic algorithms for function optimization.* Unpublished doctoral dissertation, University of Alberta, Edmonton.

Caldwell, C., & Johnston, V. S. (1991). Tracking a criminal suspect through "face-space" with a genetic algorithm. *Proceedings of the Fourth International Conference on Genetic Algorithms,* 416–421.

Caruana, R. A., & Schaffer, J. D. (1988). Representation and hidden bias: Gray vs. binary coding for genetic algorithms. *Proceedings of the Fifth International Conference on Machine Learning,* 153–161.

Cavicchio, D. J., Jr. (1970). *Adaptive search using simulated evolution.* Unpublished doctoral dissertation, University of Michigan, Ann Arbor.

Cedeño, W., & Vemuri, V. (1992). Dynamic multimodal function optimization using genetic algorithms. *Proceedings of the XVIII Latin-American Informatics Conference*, 292–301.

Cobb, H. G. (1990). *An investigation into the use of hypermutation as an adaptive operator in genetic algorithms having continuous, time-dependent nonstationary environments* (NRL Memorandum Report No. 6760). Washington, DC: Naval Research Laboratory.

Cobb, H. G., & Grefenstette, J. J. (1993). Genetic algorithms for tracking changing environments. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 523–530.

Cohoon, J. P., Hegde, S. U., Martin, W. N., & Richards, D. (1987). Punctuated equilibria: A parallel genetic algorithm. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 148–154.

Cohoon, J. P., Martin, W. N., & Richards, D. S. (1991). A multi-population genetic algorithm for solving the K-partition problem on hyper-cubes. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 244–248.

Collins, R. J., & Jefferson, D. R. (1991). Selection in massively parallel genetic algorithms. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 249–256.

Cook, L. M. (1991). *Genetic and ecological diversity: The sport of nature.* London: Chapman & Hall.

Culberson, J. C. (1992). *Genetic invariance: A new paradigm for genetic algorithm design* (Technical Report No. TR92-02). Edmonton: University of Alberta, Department of Computing Science.

Das, R., & Whitley, D. (1991). The only challenging problems are deceptive: Global search by solving order-1 hyperplanes. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 166–173.

Dasgupta, D., & McGregor, D. R. (1992). Nonstationary function optimization using the structured genetic algorithm. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 145–154). Amsterdam: Elsevier.

Davidor, Y. (1991a). Epistasis variance: A viewpoint on GA-hardness. In G. J. E. Rawlins (Ed.), *Foundations of genetic algorithms* (pp. 23–35). San Mateo: Morgan Kaufmann.

Davidor, Y. (1991b). A naturally occurring niche & species phenomenon: The model and first results. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 257–263.

Davidor, Y., Yamada, T., & Nakano, R. (1993). The ECOlogical framework II: Improving GA performance at virtually zero cost. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 171–176.

Davis, L. (1991a). Bit-climbing, representational bias, and test suite design. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 18–23.

Davis, L. (Ed.). (1991b). *Handbook of genetic algorithms*. New York: Van Nostrand Reinhold.

Davis, L. (1994). Genetic algorithms for optimization: Three case studies. In J. M. Zurada, R. J. Marks II, & C. J. Robinson (Eds.), *Computational intelligence: Imitating life* (pp. 416–426). New York: IEEE Press.

Davis, T. E., & Principe, J. C. (1993). A Markov chain framework for the simple genetic algorithm. *Evolutionary Computation 1*(3), 269–288.

Deb, K. (1989). *Genetic algorithms in multimodal function optimization* (Masters thesis and TCGA Report No. 89002). Tuscaloosa: University of Alabama, The Clearinghouse for Genetic Algorithms.

Deb, K. (1991). Binary and floating-point function optimization using messy genetic algorithms (Doctoral dissertation, University of Alabama). *Dissertation Abstracts International, 52*(5), 2658B.

Deb, K., & Goldberg, D. E. (1989). An investigation of niche and species formation in genetic function optimization. *Proceedings of the Third International Conference on Genetic Algorithms*, 42–50.

Deb, K., & Goldberg, D. E. (1993). Analyzing deception in trap functions. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 93–108). San Mateo: Morgan Kaufmann.

Deb, K., & Goldberg, D. E. (1994). Sufficient conditions for deceptive and easy binary functions. *Annals of Mathematics and Artificial Intelligence, 10*(4), 385–408.

Deb, K., Horn, J., & Goldberg, D. E. (1993). Multimodal deceptive functions. *Complex Systems, 7*(2), 131–153.

De Jong, K. A. (1975). An analysis of the behavior of a class of genetic adaptive systems (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 36*(10), 5140B. (University Microfilms No. 76-9381)

De Jong, K. A., Spears, W. M., & Gordon, D. F. (in press). Using Markov chains to analyze GAFOs. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 3*. San Mateo: Morgan Kaufmann.

Dozier, G., Bowen, J., & Bahler, D. (1994). Solving small and large scale constraint satisfaction problems using a heuristic-based microgenetic algorithm. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 306–311.

Elketroussi, M., & Fan. D. (1993). GADELO: A multi-population genetic algorithm based on dynamic exploration of local optima. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 633.

Elo, S. (1994). A parallel genetic algorithm on the CM-2 for multi-modal optimization. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 818–822.

Eshelman, L. J. (1991). The CHC adaptive search algorithm: How to have safe search when engaging in nontraditional genetic recombination. In G. J. E. Rawlins (Ed.), *Foundations of genetic algorithms* (pp. 265–283). San Mateo: Morgan Kaufmann.

Eshelman, L. J., & Schaffer, J. D. (1991). Preventing premature convergence in genetic algorithms by preventing incest. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 115–122.

Eshelman, L. J., & Schaffer, J. D. (1993). Real-coded genetic algorithms and interval-schemata. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 187–202). San Mateo: Morgan Kaufmann.

Fonseca, C. M., & Fleming, P. J. (1993). Genetic algorithms for multiobjective optimization: Formulation, discussion and generalization. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 416–423.

Forrest, S. (Ed.). (1993). *Proceedings of the Fifth International Conference on Genetic Algorithms.* San Mateo: Morgan Kaufmann.

Forrest, S., Jarvonik, B., Smith, R. E., & Perelson, A. S. (1993). Using genetic algorithms to explore pattern recognition in the immune system. *Evolutionary Computation, 1*(3), 191–211.

Forrest, S., & Mitchell, M. (1991). The performance of genetic algorithms on Walsh polynomials: Some anomalous results and their explanation. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 182–189.

Forrest, S., & Mitchell, M. (1993). Relative building-block fitness and the building-block hypothesis. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 109–126). San Mateo: Morgan Kaufmann.

Freund, J. E., & Walpole, R. E. (1980). *Mathematical statistics* (3rd ed.). Englewood Cliffs, NJ: Prentice-Hall.

Giordana, A., Saitta, L., & Zini, F. (1994). Learning disjunctive concepts with distributed genetic algorithms. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 115–119.

Goldberg, D. E. (1983). Computer-aided gas pipeline operation using genetic algorithms and rule learning (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 44*(10), 3174B. (University Microfilms No. 8402282)

Goldberg, D. E. (1987). Simple genetic algorithms and the minimal, deceptive problem. In L. Davis (Ed.), *Genetic algorithms and simulated annealing* (pp. 74–88). London: Pitman.

Goldberg, D. E. (1989a). Genetic algorithms and Walsh functions: Part I, a gentle introduction. *Complex Systems, 3*, 129–152.

Goldberg, D. E. (1989b). Genetic algorithms and Walsh functions: Part II, deception and its analysis. *Complex Systems, 3*, 153–171.

Goldberg, D. E. (1989c). *Genetic algorithms in search, optimization, and machine learning.* Reading, MA: Addison-Wesley.

Goldberg, D. E. (1989d). Sizing populations for serial and parallel genetic algorithms. *Proceedings of the Third International Conference on Genetic Algorithms*, 70–79.

Goldberg, D. E. (1990). A note on Boltzmann tournament selection for genetic algorithms and population-oriented simulated annealing. *Complex Systems, 4*, 445–460.

Goldberg, D. E. (1991a). Real-coded genetic algorithms, virtual alphabets, and blocking. *Complex Systems, 5*, 139–167.

Goldberg, D. E. (1991b). The theory of virtual alphabets. *Lecture Notes in Computer Science: Parallel Problem Solving from Nature, 496*, 13–22.

Goldberg, D. E. (1992). Construction of high-order deceptive functions using low-order Walsh coefficients. *Annals of Mathematics and Artificial Intelligence, 5*(1), 35–48.

Goldberg, D. E. (1993a). Making genetic algorithms fly: A lesson from the Wright Brothers. *Advanced Technology for Developers, 2*, 1–8.

Goldberg, D. E. (1993b). The Wright Brothers, genetic algorithms, and the design of complex systems. *Proceedings of the Symposium on Neural-Networks; Alliances and Perspectives in Senri*, 1–7.

Goldberg, D. E. (1993c). A Wright-Brothers theory of genetic-algorithm flight. *Journal of the Institute of Systems, Control, and Information Engineers, 37*(8), 450–458.

Goldberg, D. E., & Bridges, C. L. (1990). An analysis of a reordering operator on a GA-hard problem. *Biological Cybernetics, 62*, 397–405.

Goldberg, D. E., & Deb, K. (1991). A comparative analysis of selection schemes used in genetic algorithms. In G. J. E. Rawlins (Ed.), *Foundations of Genetic Algorithms* (pp. 69–93). San Mateo: Morgan Kaufmann.

Goldberg, D. E., Deb, K., & Clark, J. H. (1992). Genetic algorithms, noise, and the sizing of populations. *Complex Systems, 6*, 333–362.

Goldberg, D. E., Deb, K., & Horn, J. (1992). Massive multimodality, deception, and genetic algorithms. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 37–46). Amsterdam: Elsevier.

Goldberg, D. E., Deb, K., Kargupta, H., & Harik, G. (1993). Rapid, accurate optimization of difficult problems using fast messy genetic algorithms. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 56–64.

Goldberg, D. E., Deb, K., & Korb, B. (1990). Messy genetic algorithms revisited: Studies in mixed size and scale. *Complex Systems, 4*, 415–444.

Goldberg, D. E., Deb, K., & Thierens, D. (1993). Toward a better understanding of mixing in genetic algorithms. *Journal of the Society of Instrument and Control Engineers, 32*(1), 10–16.

Goldberg, D. E., Korb, B., & Deb, K. (1989). Messy genetic algorithms: Motivation, analysis, and first results. *Complex Systems, 3*, 493–530.

Goldberg, D. E., & Richardson, J. (1987). Genetic algorithms with sharing for multimodal function optimization. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 41–49.

Goldberg, D. E., & Rudnick, M. (1991). Genetic algorithms and the variance of fitness. *Complex Systems, 5*, 265–278.

Goldberg, D. E., & Segrest, P. (1987). Finite Markov chain analysis of genetic algorithms. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 1–8.

Goldberg, D. E., & Smith, R. E. (1987). Nonstationary function optimization using genetic algorithms with dominance and diploidy. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 59–68.

Gorges-Schleuter, M. (1989). ASPARAGOS an asynchronous parallel genetic optimization strategy. *Proceedings of the Third International Conference on Genetic Algorithms*, 422–427.

Gorges-Schleuter, M. (1991). Explicit parallelism of genetic algorithms through population structures. *Lecture Notes in Computer Science: Parallel Problem Solving from Nature, 496*, 150–159.

Gorges-Schleuter, M. (1992). Comparison of local mating strategies in massively parallel genetic algorithms. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 553–562). Amsterdam: Elsevier.

Greene, D. P., & Smith, S. F. (1993). Competition-based induction of decision models from examples. *Machine Learning, 13*(2/3), 229–257.

Greene, D. P., & Smith, S. F. (1994). Using coverage as a model building constraint in learning classifier systems. *Evolutionary Computation, 2*(1), 67–91.

Greene, F. (1994). A method for utilizing diploid/dominance in genetic search. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 439–444.

Grefenstette, J. J. (1986). Optimization of control parameters for genetic algorithms. *IEEE Transactions on Systems, Man, and Cybernetics, 16*(1), 122–128.

Grefenstette, J. J. (1992). Genetic algorithms for changing environments. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 137–144). Amsterdam: Elsevier.

Grefenstette, J. J. (1993). Deception considered harmful. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 75–91). San Mateo: Morgan Kaufmann.

Grosso, P. B. (1985). *Computer simulation of genetic adaptation: Parallel subcomponent interaction in a multilocus model* (Doctoral dissertation, University of Michigan). (University Microfilms No. 8520908)

Harik, G. (1994). *Finding multiple solutions in problems of bounded difficulty* (IlliGAL Report No. 94002). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Harvey, I. (1993). The puzzle of the persistent question marks: A case study of genetic drift. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 15–22.

Hatjimihail, A. T. (1993). Genetic algorithms-based design and optimization of statistical quality-control procedures. *Clinical Chemistry, 39*(9), 1972–1978.

Hillis, W. D. (1990). Co-evolving parasites improve simulated evolution as an optimization procedure. *Physica D, 42*, 228–234.

Holland, J. H. (1975). *Adaptation in natural and artificial systems.* Ann Arbor: University of Michigan Press.

Holland, J. H. (1992). *Adaptation in natural and artificial systems.* Cambridge, MA: MIT Press.

Holland, J. H., & Reitman, J. S. (1978). Cognitive systems based on adaptive algorithms. In D. A. Waterman & F. Hayes-Roth (Eds.), *Pattern-directed inference systems* (pp. 313–329). New York: Academic Press.

Hollstien, R. B. (1971). Artificial genetic adaptation in computer control systems (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 32*(3), 1510B. (University Microfilms No. 71-23773)

Holte, R. C. (1993). Very simple classification rules perform well on most commonly used datasets. *Machine Learning, 11*, 63–90.

Homaifar, A., Guan, S., & Liepins, G. E. (1993). A new approach on the traveling sales-man problem by genetic algorithms. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 460–466.

Homaifar, A., Qi, X., & Fost, J. (1991). Analysis and design of a general GA deceptive problem. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 196–203.

Horn, J. (1993). Finite Markov chain analysis of genetic algorithms with niching. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 110–117.

Horn, J., & Goldberg, D. E. (in press). Genetic algorithm difficulty and the modality of fitness landscapes. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 3*. San Mateo: Morgan Kaufmann.

Horn, J., Goldberg, D. E., & Deb, K. (1994). Implicit niching in a learning classifier system: Nature's way. *Evolutionary Computation, 2*(1), 37–66.

Horn, J., & Nafpliotis, N. (1993). *Multiobjective optimization using the niched Pareto genetic algorithm* (IlliGAL Report No. 93005). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Horn, J., Nafpliotis, N., & Goldberg, D. E. (1994). A niched Pareto genetic algorithm for multiobjective optimization. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 82–87.

Horner, A., & Goldberg, D. E. (1991). Genetic algorithms and computer-assisted music com-position. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 437–441.

Husbands, P., & Mill, F. (1991). Simulated co-evolution as the mechanism for emergent planning and scheduling. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 264–270.

Jones, T. (in press). Crossover, macromutation, and population-based search. *Proceedings of the Sixth International Conference on Genetic Algorithms*.

Kapur, J. N., & Kesavan, H. K. (1987). *The generalized maximum entropy principle (with applications)*. Sandford Educational Press.

Kargupta, H. (1993). Drift, diffusion and Boltzmann distribution in simple genetic algorithm. *Proceedings of the Workshop on Physics and Computation*, 137–145. Los Alamitos, CA: IEEE Computer Society Press.

Kargupta, H., Deb, K., & Goldberg, D. E. (1992). Ordering genetic algorithms and deception. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 47–56). Amsterdam: Elsevier.

Kargupta, H., & Goldberg, D. E. (1994). *Decision making in genetic algorithms: A signal-to-noise perspective* (IlliGAL Report No. 94004). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Kargupta, H., & Smith, R. E. (1991). System identification with evolving polynomial networks. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 370–376.

Karr, C. L. (1991a). Air-injected hydrocyclone optimization via genetic algorithm. In L. Davis (Ed.), *Handbook of genetic algorithms* (pp. 222–236). New York: Van Nostrand Reinhold.

Karr, C. L. (1991b). Design of an adaptive fuzzy logic controller using a genetic algorithm. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 450–457.

Kinnear, K. E., Jr. (1994). Fitness landscapes and difficulty in genetic programming. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 142–147.

Kirpatrick, S., Gelatt, C. D., Jr., & Vecchi, M. P. (1983). Optimization by simulated annealing. *Science, 220*(4598), 671–680.

Koza, J. R. (1992). *Genetic programming: On the programming of computers by means of natural selection*. Cambridge, MA: MIT Press.

Krishnakumar, K. (1989). Micro-genetic algorithms for stationary and non-stationary function optimization. *SPIE Proceedings: Intelligent Control and Adaptive Systems, 1196*, 289–296.

Kursawe, F. (1991). A variant of evolution strategies for vector optimization. *Lecture Notes in Computer Science: Parallel Problem Solving from Nature, 496*, 193–197.

Laine, P., & Kuuskankare, M. (1994). Genetic algorithms in musical style oriented generation. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 858–862.

Langton, C. G. (Ed.). (1989). *Artificial life.* Redwood City, CA: Addison-Wesley.

Langton, C. G., Taylor, C., Farmer, J. D., & Rasmussen, S. (Eds.). (1992). *Artificial life II.* Redwood City, CA: Addison-Wesley.

Lial, M. L., & Miller, C.D. (1989). *Finite mathematics* (4th ed.). Glenview, IL: Scott, Foresman.

Liepins, G. E., & Vose, M. D. (1990). Representational issues in genetic optimization. *Journal of Experimental and Theoretical Artificial Intelligence, 2*, 101–115.

Liepins, G. E., & Vose, M. D. (1991). Deceptiveness and genetic algorithm dynamics. In G. J. E. Rawlins (Ed.), *Foundations of genetic algorithms* (pp. 36–50). San Mateo: Morgan Kaufmann.

Louis, S. J., & Rawlins, G. J. E. (1993). Syntactic analysis of convergence in genetic algorithms. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 141–151). San Mateo: Morgan Kaufmann.

Mahfoud, S. W. (1992). Crowding and preselection revisited. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 27–36). Amsterdam: Elsevier.

Mahfoud, S. W. (1993). Finite Markov chain models of an alternative selection strategy for the genetic algorithm. *Complex Systems, 7*(2), 155–170.

Mahfoud, S. W. (1994a). *An analysis of Boltzmann tournament selection: Part II: An experimental analysis of Boltzmann tournament selection* (IlliGAL Report No. 94007). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Mahfoud, S. W. (1994b). Crossover interactions among niches. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 188–193.

Mahfoud, S. W., & Goldberg, D. E. (1992). A genetic algorithm for parallel simulated annealing. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 301–310). Amsterdam: Elsevier.

Mahfoud, S. W., & Goldberg, D. E. (1995). Parallel recombinative simulated annealing: A genetic algorithm. *Parallel Computing, 21*, 1–28.

Mahfoud, S. W., & Mani, G. (in press). Genetic algorithms for predicting individual stock performance. *Proceedings of the Third International Conference on Artificial Intelligence Applications on Wall Street*.

Manderick, B., & Spiessens, P. (1989). Fine-grained parallel genetic algorithms. *Proceedings of the Third International Conference on Genetic Algorithms*, 428–433.

Mason, A. J. (1991). Partition coefficients, static deception and deceptive problems for non-binary alphabets. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 210–214.

Mathias, K., & Whitley, D. (1993). Remapping hyperspace during genetic search: Canonical delta folding. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 167–186). San Mateo: Morgan Kaufmann.

Mauldin, M. L. (1984) Maintaining diversity in genetic search. *Proceedings of the National Conference on Artificial Intelligence*, 247–250.

Mc Intyre, R. A. (1994). Bach in a box: The evolution of four part baroque harmony using the genetic algorithm. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 852–857.

Merkle, L. D., & Lamont, G. B. (1993). Comparison of parallel messy genetic algorithm data distribution strategies. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 191–198.

Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H., & Teller, E. (1953). Equation of state calculations by fast computing machines. *The Journal of Chemical Physics, 21*(6), 1087–1092.

Mühlenbein, H. (1989). Parallel genetic algorithms, population genetics and combinatorial optimization. *Proceedings of the Third International Conference on Genetic Algorithms*, 416–421.

Mühlenbein, H. (1991). Evolution in time and space — The parallel genetic algorithm. In G. J. E. Rawlins (Ed.), *Foundations of Genetic Algorithms* (pp. 316–337). San Mateo: Morgan Kaufmann.

Mühlenbein, H. (1992). How genetic algorithms really work I: Mutation and hillclimbing. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 15–25). Amsterdam: Elsevier.

Mühlenbein, H., Gorges-Schleuter, M., & Krämer, O. (1988). Evolution algorithms in combinatorial optimization. *Parallel Computing, 7*, 65–85.

Mühlenbein, H., Schomisch, M., & Born, J. (1991a). The parallel genetic algorithm as function optimizer. *Parallel Computing, 17*, 619–632.

Mühlenbein, H., Schomisch, M., & Born, J. (1991b). The parallel genetic algorithm as function optimizer. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 271–278.

Nix, A. E., & Vose, M. D. (1992). Modeling genetic algorithms with Markov chains. *Annals of Mathematics and Artificial Intelligence, 5*(1), 79–88.

Oei, C. K., Goldberg, D. E., & Chang, S. J. (1991). *Tournament selection, niching, and the preservation of diversity* (IlliGAL Report No. 91011). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Orvosh, D., & Davis, L. (1994). Using a genetic algorithm to optimize problems with feasibility constraints. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 548–553.

Packard, N. H. (1990). A genetic learning algorithm for the analysis of complex data. *Complex Systems, 4*(5), 543–572.

Pál, K. F. (1994). Selection schemes with spatial isolation for genetic optimization. *Lecture Notes in Computer Science: Parallel Problem Solving from Nature — PPSN III, 866,* 170–179.

Perry, Z. A. (1984). Experimental study of speciation in ecological niche theory using genetic algorithms (Doctoral dissertation, University of Michigan). *Dissertation Abstracts International, 45* (12), 3870B. (University Microfilms No. 8502912)

Pettey, C. B., Leuze, M. R., & Grefenstette, J. J. (1987). A parallel genetic algorithm. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms,* 155–161.

Press, W. H., Teukolsky, S. A., Vetterling, W. T., & Flannery, B. P. (1992). *Numerical recipes in C: The art of scientific computing* (2nd ed.). Port Chester, NY: Cambridge University Press.

Provine, W. P. (1986). *Sewall Wright and evolutionary biology.* Chicago: University of Chicago Press.

Radcliffe, N. J. (1991a). Equivalence class analysis of genetic algorithms. *Complex Systems, 5*(2), 183–205.

Radcliffe, N. J. (1991b). Forma analysis and random respectful recombination. *Proceedings of the Fourth International Conference on Genetic Algorithms,* 222–229.

Radcliffe, N. J. (1993). Genetic set recombination. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 203–219). San Mateo: Morgan Kaufmann.

Richardson, J. T., Palmer, M. R., Liepins, G., & Hilliard, M. (1989). Some guidelines for genetic algorithms with penalty functions. *Proceedings of the Third International Conference on Genetic Algorithms,* 191–197.

Romeo, F., & Sangiovanni-Vincentelli, A. (1991). A theoretical framework for simulated annealing. *Algorithmica, 6,* 302–345.

Rudnick, M., & Goldberg, D. E. (1991). *Signal, noise, and genetic algorithms* (IlliGAL Report No. 91005). Urbana: University of Illinois, Illinois Genetic Algorithms Laboratory.

Sannier, A. V., II, & Goodman, E. D. (1987). Genetic learning procedures in distributed environments. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 162–169.

Schaffer, J. D. (1984). *Some experiments in machine learning using vector evaluated genetic algorithms.* Unpublished doctoral dissertation, Vanderbilt University, Nashville.

Schaffer, J. D. (1985). Multiple objective optimization with vector evaluated genetic algorithms. *Proceedings of an International Conference on Genetic Algorithms and Their Applications*, 93–100.

Schaffer, J. D., Caruana, R. A., Eshelman, L. J., & Das, R. (1989). A study of control parameters affecting online performance of genetic algorithms for function optimization. *Proceedings of the Third International Conference on Genetic Algorithms*, 51–60.

Schwehm, M. (1992). Implementation of genetic algorithms on various interconnection networks. In M. Valero, E. Oñate, M. Jane, J. L. Larriba, & B. Suárez (Eds.), *Parallel computing and transputer applications* (pp. 195–203). Barcelona: IOS Press.

Sedbrook, T. A., Wright, H., & Wright, R. (1991). Application of a genetic classifier for patient triage. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 334–338.

Shorrocks, B. (1979). *The genesis of diversity.* Baltimore: University Park Press.

Sikora, R., & Shaw, M. J. (1994). A double-layered learning approach to acquiring rules for classification: Integrating genetic algorithms with similarity-based learning. *ORSA Journal on Computing, 6*(2), 174–187.

Sirag, D. J., & Weisser, P. T. (1987). Toward a unified thermodynamic genetic operator. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 116–122.

Smith, R. E. (1988). *An investigation of diploid genetic algorithms for adaptive search of nonstationary functions* (Masters thesis and TCGA Report No. 88001). Tuscaloosa: University of Alabama, The Clearinghouse for Genetic Algorithms.

Smith, R. E. (1993). Adaptively resizing populations: An algorithm and analysis. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 653.

Smith, R. E., Forrest, S., & Perelson, A. S. (1993a). Population diversity in an immune system model: Implications for genetic search. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 153–165). San Mateo: Morgan Kaufmann.

Smith, R. E., Forrest, S., & Perelson, A. S. (1993b). Searching for diverse, cooperative populations with genetic algorithms. *Evolutionary Computation, 1*(2), 127–149.

Smith, R. E., & Goldberg, D. E. (1992). Diploidy and dominance in artificial genetic search. *Complex Systems, 6*, 251–285.

Smith, R. E., & Valenzuela-Rendón, M. (1989). A study of rule set development in a learning classifier system. *Proceedings of the Third International Conference on Genetic Algorithms*, 340–346.

Spears, W. M. (1994). Simple subpopulation schemes. *Proceedings of the Third Annual Conference on Evolutionary Programming*, 296–307.

Spiessens, P., & Manderick, B. (1991). A massively parallel genetic algorithm. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 279–286.

Srinivas, M., & Patnaik, L. M. (1993). Binomially distributed populations for modelling GAs. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 138–145.

Stadnyk, I. (1987). Schema recombination in a pattern recognition problem. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 27–35.

Suzuki, J. (1993). A Markov chain analysis on a genetic algorithm. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 146–153.

Syswerda, G. (1989). Uniform crossover in genetic algorithms. *Proceedings of the Third International Conference on Genetic Algorithms*, 2–9.

Syswerda, G. (1991). A study of reproduction in generational and steady-state genetic algorithms. In G. J. E. Rawlins (Ed.), *Foundations of genetic algorithms* (pp. 94–101). San Mateo: Morgan Kaufmann.

Tanese, R. (1987). Parallel genetic algorithm for a hypercube. *Genetic algorithms and their applications: Proceedings of the Second International Conference on Genetic Algorithms*, 177–183.

Tanese, R. (1989a). Distributed genetic algorithms. *Proceedings of the Third International Conference on Genetic Algorithms*, 434–439.

Tanese, R. (1989b). *Distributed genetic algorithms for function optimization.* Unpublished doctoral dissertation, University of Michigan, Ann Arbor.

Thierens, D., & Goldberg, D. E. (1993). Mixing in genetic algorithms. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 38–45.

Todd, P. M., & Miller, G. F. (1991). On the sympatric origin of species: Mercurial mating in the quicksilver model. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 547–554.

Varela F. J., & Bourgine, P. (Eds.). (1992). *Toward a practice of autonomous systems: Proceedings of the First European Conference on Artificial Life.* Cambridge, MA: MIT Press.

Vose, M. D. (1993). Modeling simple genetic algorithms. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 63–73). San Mateo: Morgan Kaufmann.

Whitley, L. D. (1991). Fundamental principles of deception in genetic search. In G. J. E. Rawlins (Ed.), *Foundations of genetic algorithms* (pp. 221–241). San Mateo: Morgan Kaufmann.

Whitley, D. (1993). An executable model of a simple genetic algorithm. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 2* (pp. 45–62). San Mateo: Morgan Kaufmann.

Whitley, D., & Hanson, T. (1989). Optimizing neural networks using faster, more accurate genetic search. *Proceedings of the Third International Conference on Genetic Algorithms*, 391–396.

Whitley, D., & Kauth, J. (1988). GENITOR: A different genetic algorithm. *Proceedings of the Rocky Mountain Conference on Artificial Intelligence*, 118–130.

Whitley, D., Mathias, K., & Fitzhorn, P. (1991). Delta coding: An iterative search strategy for genetic algorithms. *Proceedings of the Fourth International Conference on Genetic Algorithms*, 77–84.

Whitley, D., & Starkweather, T. (1990). GENITOR II: A distributed genetic algorithm. *Journal of Experimental and Theoretical Artificial Intelligence, 2*, 189–214.

Whitley, D., Starkweather, T., & Bogart, C. (1990). Genetic Algorithms and neural networks: Optimizing connections and connectivity. *Parallel Computing, 14*, 347–361.

Yin, X. (1994). *Investigations on the application of genetic algorithms to the load flow problem in electrical power systems*. Unpublished doctoral dissertation, Université Catholique de Louvain, Belgium.

Yin, X., & Germay, N. (1993). A fast genetic algorithm with sharing scheme using cluster analysis methods in multimodal function optimization. In R. F. Albrecht, C. R. Reeves, & N. C. Steele (Eds.), *Artificial neural nets and genetic algorithms: Proceedings of the international conference in Innsbruck* (pp. 450–457). Berlin: Springer-Verlag.

Vita

Samir W. Mahfoud

3665 East Bay Drive #204-429

Largo, FL 34641

E-mail: sam@lbs.com

Education

Ph.D. Computer Science: University of Illinois at Urbana-Champaign
May 1995   GPA: 4.9/5.0

M.S. Computer Science: University of Wisconsin-Madison
May 1987   GPA: 3.7/4.0

B.S. Mathematics: Murray State University, Murray, KY
May 1985   GPA: 3.9/4.0   Graduated Summa Cum Laude
Board of Regents Scholarship, Math Department Scholarships, Dean's List 6/6 Semesters

Professional Experience

11/94 – Present: **Computer Research Scientist**, **LBS Capital Management**, Clearwater, FL. In charge of designing and building an investment system that integrates genetic algorithms, neural networks, and other artificial-intelligence and computer technologies.

7/87 – 7/89: **Software Engineer**, **Xerox Corporation**, Rochester, NY. In charge of software in digital imaging, control, and parallel processing for one of the first multifunctional office machines ever built. This machine was a copier, a laser printer, a fax machine, a scanner, and a networked image-editing workstation, combined into a single product. The project commenced in early 1987; by July, 1989, 30 prototype machines had been built.

Assistantships

8/91 – 11/94: **Research Assistant, Illinois Genetic Algorithms Laboratory**, Urbana, IL. Conducted research on the topics of niching methods in genetic algorithms, and combining genetic algorithms with simulated annealing.

8/89 – 5/90: **Research Assistant, Construction Engineering Research Lab, U.S. Army Corps of Engineers**, Champaign, IL. Worked on a project to develop an expert system and English-language database interface for managing construction projects.

8/85 – 5/87: **Teaching Assistant, Department of Computer Sciences, University of Wisconsin-Madison**. Assisted in an introductory artificial intelligence course. Also helped students at office hours and as a site consultant, with Pascal, Fortran, assembly language, the Unix system, and the Macintosh. Graded programs, assignments, and exams for two sections.

Publications

Mahfoud, S. W. (1992). Crowding and preselection revisited. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 27–36). Amsterdam: Elsevier.

Mahfoud, S. W. (1993). Finite Markov chain models of an alternative selection strategy for the genetic algorithm. *Complex Systems, 7*(2), 155–170.

Mahfoud, S. W. (1993). Simple analytical models of genetic algorithms for multimodal function optimization. *Proceedings of the Fifth International Conference on Genetic Algorithms*, 643.

Mahfoud, S. W. (1994). Crossover interactions among niches. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 188–193.

Mahfoud, S. W. (1994). Genetic drift in sharing methods. *Proceedings of the First IEEE Conference on Evolutionary Computation, IEEE World Congress on Computational Intelligence*, 67–72.

Mahfoud, S. W. (in press). A comparison of parallel and sequential niching methods. *Proceedings of the Sixth International Conference on Genetic Algorithms*.

Mahfoud, S. W. (in press). Niching methods for genetic algorithms (Doctoral dissertation, University of Illinois at Urbana-Champaign). *Dissertation Abstracts International*.

Mahfoud, S. W. (in press). Population size and genetic drift in fitness sharing. In L. D. Whitley (Ed.), *Foundations of genetic algorithms, 3*. San Mateo: Morgan Kaufmann.

Mahfoud, S. W., & Deb K. (1993, February). Structural design by parallel recombinative simulated annealing. In O. O. Storaasli (Chair), *Second symposium on parallel computational methods for large-scale structural analysis and design*. Symposium held in Norfolk, VA.

Mahfoud, S. W., & Goldberg, D. E. (1992). A genetic algorithm for parallel simulated annealing. In R. Männer & B. Manderick (Eds.), *Parallel problem solving from nature, 2* (pp. 301–310). Amsterdam: Elsevier.

Mahfoud, S. W., & Goldberg, D. E. (1995). Parallel recombinative simulated annealing: A genetic algorithm. *Parallel Computing, 21*, 1–28.

Mahfoud, S. W., & Mani, G. (in press). Genetic algorithms for predicting individual stock performance. *Proceedings of the Third International Conference on Artificial Intelligence Applications on Wall Street.*

Mani, G., Quah, K. K., Mahfoud, S., & Barr, D. (1995). An analysis of neural-network forecasts from a large-scale, real-world stock selection system. *Proceedings of the IEEE/IAFE 1995 Conference on Computational Intelligence for Financial Engineering (CIFEr)*, 72–78. Piscataway, NJ: IEEE Press.

Other Activities

Review Board, *Applied Intelligence* journal.

Program Committee, *Sixth International Conference on Genetic Algorithms.*

Organizer and Co-chair, *Workshop on Niching Methods*, held July 20, 1993, during the *Fifth International Conference on Genetic Algorithms*, in Urbana, Illinois. The workshop attracted 60 participants.

Member, Association for Computing Machinery (ACM).

Have designed and implemented large-scale intelligent systems in areas such as machine learning, chess playing, parallel image processing, time-series forecasting, database management, and optimization.